

STAT 216 Coursepack



Fall 2024
Montana State University

Melinda Yager
Jade Schmidt
Stacey Hancock

This resource was developed by Melinda Yager, Jade Schmidt, and Stacey Hancock in 2021 to accompany the online textbook: Hancock, S., Carnegie, N., Meyer, E., Schmidt, J., and Yager, M. (2021). *Montana State Introductory Statistics with R*. Montana State University. <https://mtstateintrostats.github.io/IntroStatTextbook/>.

This resource is released under a Creative Commons BY-NC-SA 4.0 license unless otherwise noted.

Contents

Preface	1
1 Basics of Data	2
1.1 Activity 1: Intro to Data	2
1.2 Reading Guides	4
1.3 Video Notes: Intro to data	8
1.4 Guided Activity 1: Intro to Data Analysis	12
2 Study Design	15
2.1 Reading Guide: Sampling, Study Design, and Scope of Inference	15
2.2 Video Notes: Study Design	19
2.3 Guided Activity 2: American Indian Address	24
2.4 Guided Activity 3: American Indian Address (continued)	29
2.5 Activity 2: Study Design	34
3 Exploring Categorical and Quantitative Data	40
3.1 Reading Guide: Introduction to R, Categorical Variables, and a Single Quantitative Variable . .	40
3.2 Video Notes: Exploratory Data Analysis of Categorical and Quantitative Variables	46
3.3 Guided Activity 4: Summarizing Categorical Data	63
3.4 Guided Activity 5: Summarizing Quantitative Variables	71
3.5 Module 3 Lab: EDA Review	77
4 Exploring Multivariable Data	84
4.1 Reading Guide: Two Quantitative Variables and Multivariable Concepts	84
4.2 Video Notes: Regression and Correlation	91
4.3 Guided Activity 6: IPEDS (continued)	99
4.4 Activity 3: Moneyball — Linear Regression	106
4.5 Module 4 Lab: Penguins	110
5 Inference for a Single Categorical Variable: Simulation-based Methods	114
5.1 Reading Guide: Categorical Inference	114
5.2 Video Notes: Inference for One Categorical Variable using Simulation-based Methods	121
5.3 Guided Activity 7: Sleep Study — Simulation-based Inference Methods	130
5.4 Activity 4: Helper-Hinderer Part 1 — Simulation-based Hypothesis Test	136
5.5 Activity 5: Helper-Hinderer (continued)	141
5.6 Activity 6: Helper-Hinderer — Simulation-based Confidence Interval	145
6 Inference for a Single Categorical Variable: Theory-based Methods + Errors and Power	150
6.1 Reading Guide: Categorical Inference	150
6.2 Lecture Notes: Inference for One Categorical Variable using Theory-based Methods	162
6.3 Guided Activity 8: Cashless Society	171
6.4 Activity 7: Handedness of Male Boxers	177
6.5 Activity 8: What does confidence mean?	183
6.6 Activity 9: Errors and Power	187
6.7 Module 6 Lab: Mixed Breed Dogs in the U.S.	191
7 Inference for Two Categorical Variables: Simulation-based Methods	195
7.1 Reading Guide: Hypothesis Testing for a Difference in Proportions	195
7.2 Video Notes: Inference for Two Categorical Variables using Simulation-based Methods	202
7.3 Guided Activity 9: The Good Samaritan	208

7.4	Activity 10: Poisonous Mushrooms	214
8	Inference for Two Categorical Variables: Theory-based Methods	218
8.1	Reading Guide: Theory-based Inference for a Difference in Proportions	218
8.2	Video Notes: Theoretical Inference for Two Categorical Variables	223
8.3	Guided Activity 10: Winter Sports Helmet Use and Head Injuries — Theory-based Methods . . .	230
8.4	Activity 11: Diabetes	236
8.5	Module 8 Lab: Is yawning contagious?	240
9	Inference for Inference for a Single Mean or Paired Mean Difference	244
9.1	Reading Guide: Inference for a Single Mean or Paired Mean Difference	244
9.2	Video Notes: Inference for Paired Data	255
9.3	Guided Activity 11: Color Interference	263
9.4	Guided Activity 12: Snakes	270
9.5	Activity 12: Swearing	275
10	Inference for a Quantitative Response with Independent Samples	280
10.1	Reading Guide: Inference for a Difference in Two Means	280
10.2	Video Notes: Inference for Independent Samples	287
10.3	Guided Activity 13: Does behavior impact performance?	296
10.4	Guided Activity 14: Moon Phases and Virtual Reality	301
10.5	Activity 13: Dinosaurs	305
10.6	Module 10 Lab: Trustworthiness	309
11	Inference for Two Quantitative Variables	314
11.1	Reading Guide: Inference for Slope and Correlation	314
11.2	Video Notes: Inference for Two Quantitative Variables	320
11.3	Guided Activity 15: Prediction of Crocodilian Body Size	332
11.4	Guided Activity 16: Golf Driving Distance	338
11.5	Module 11 Lab: Big Mac Index	344
12	Probability	349
12.1	Reading Guide: Special Topics	349
12.2	Video Notes: Probability	352
12.3	Guided Activity 17: Probability Studies	356
12.4	Activity 14: What’s the probability?	359
13	Semester Review	362
13.1	Group Final Exam Review	362
13.2	Golden Ticket to Descriptive and Inferential Statistical Methods	375
	References	377

Preface

This coursepack accompanies the textbook for STAT 216: Montana State Introductory Statistics with R, which can be found at <https://mtstateintrostats.github.io/IntroStatTextbook/>. The syllabus for the course (including the course calendar), data sets, and links to D2L Brightspace, Gradescope, and the MSU RStudio server can be found on the course webpage: <https://math.montana.edu/courses/s216/>. Other notes and review materials are linked in D2L.

Each of the activities in this workbook is designed to target specific learning outcomes of the course, giving you practice with important statistical concepts in a group setting with instructor guidance. In addition to the in-class activities for the course, reading guides are provided to aid in taking notes while you complete the required readings, as well as video notes to be completed while watching the required videos. Bring this workbook with you to class each class period, and take notes in the workbook as you would your own notes. A well-written completed workbook will provide an optimal study guide for exams!

All activities and labs in this coursepack will be completed during class time. Parts of each lab will be turned in on Gradescope. To aid in your understanding, read through the introduction for each activity before attending class each day.

STAT 216 is a 3-credit in-person course. In our experience, it takes six to nine hours per week outside of class to achieve a good grade in this class. By “good” we mean at least a C because a grade of D or below does not count toward fulfilling degree requirements. Many of you set your goals higher than just getting a C, and we fully support that. You need roughly nine hours per week to review past activities, read feedback on previous assignments, complete current assignments, and prepare for the next day’s class. A typical week in the life of a STAT 216 student looks like:

- *Prior to class meeting:*
 - Read assigned sections of the textbook, using the provided reading guides to take notes on the material.
 - Watch the provided videos, taking notes in the coursepack.
 - Read through the introduction to the day’s in-class activity.
 - Read through the week’s homework assignment and note any questions you may have on the content.
- *During class meeting:*
 - Work through the guided activity, in-class activity or weekly lab with your classmates and instructor, taking detailed notes on your answers to each question in the activity.
- *After class meeting:*
 - Complete any parts of the activity you did not complete in class.
 - Review the activity solutions in the Math and Stat Center, and take notes on key points.
 - Complete any remaining assigned readings for the week.
 - Complete the week’s homework assignment.

Basics of Data

1.1 Activity 1: Intro to Data

1.1.1 Learning outcomes

- Creating a data set

1.1.2 Terminology review

Statistics is the study of how best to collect, analyze, and draw conclusions from data. This week in class you will be introduced to the following terms:

- Observational units or cases
- Variables: categorical or quantitative

For more on these concepts, read Chapter 1 in the textbook.

1.1.3 General information on the Coursepack

Information is provided throughout each activity and lab to guide students through that day's activity or lab. Be sure to read ALL the material provided at the beginning of the activity and between each question. At the end of each activity is a section called *Take-home messages* that contains key points from the day's activity. Use these to review the day's activity and make sure you have a full understanding of that material.

1.1.4 Steps of the statistical investigation process

As we move through the semester we will work through the six steps of the statistical investigation process.

1. Ask a research question.
2. Design a study and collect data. *Modules 1–2*
3. Summarize and visualize the data. *Modules 3–4*
4. Use statistical analysis methods to draw inferences from the data. *Modules 5 - 11*
5. Communicate the results and answer the research question. *Weeks 6–14*
6. Revisit and look forward.

Today we will focus on the first two steps.

Step 1: The first step of any statistical investigation is to *ask a research question*. As stated in the textbook, “with the rise of data science, however, we might not start with a research question, and instead start with a data set.” Today we will create a data set by collecting responses on students in class.

Step 2: To answer any research question, we must *design a study and collect data*. Our study will consist of answers from each student. Your responses will become our observed data that we will explore.

Observational units or **cases** are the subjects data are collected on. In a spreadsheet of the data set, each row will represent a single observational unit.

1. Open the Google Form linked in D2L and fill in the responses for the following questions. When creating a data set for use in R it is important to use single words or an underscore between words. Each outcome must be written the same way each time. Make sure to use all lowercase letters to create this data set to have consistency between responses. Do not give units of measure for numerical values within the data set. For **Residency** use `in_state` or `out_state` as the two outcomes.
 - Major: what is your declared major?
 - Residency: do you have in-state or out-of-state residency?
 - Num_Credits: how many credits are you taking this semester?
 - Dominant_hand: are you left or right-handed?
 - Hand_span: what is the width of your dominant hand from the tip of your thumb to the tip of your pinky with your hand spread out measured in cm?
 - Grip_dominant: what is the grip strength measured in lbs for your dominant hand?
 - Grip_nondominant: what is the grip strength measured in lbs for your non-dominant hand?

1.1.5 Take-home messages

1. When creating a data set, each row will represent a single observational unit or case. Each column represents a variable collected. It is important to write each variable as a single word or use an underscore between words.
2. Make sure to be consistent with writing each outcome in the data set as R is case sensitive. All outcomes must be written exactly the same way.

1.1.6 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered, and to write down the names and contact information of your teammates.

1.2 Reading Guides

Reading guides are designed to be completed while reading the required sections in the course textbook (<https://mtstateintrostats.github.io/IntroStatTextbook/index.html>) to aid students in taking notes. These reading guides are not turned in in class but will be useful in understanding key concepts each week. Solutions to the reading guides will be posted on D2L.

1.2.1 Module 1 Reading Guide: Basics of Data

Textbook Sections 1.1: Case study and 1.2: Data basics

Vocabulary

Data:

Summary statistic:

Case/Observational unit:

Variable:

Quantitative variable:

Discrete variables:

Examples of discrete variables using the **County** data:

Continuous variables:

Examples of continuous variables using the **County** data:

Example of a number which is NOT a numerical variable:

Categorical variable:

Ordinal variable:

Example of an ordinal variable using the **County** data:

Nominal variable:

Examples of nominal variables using the **County** data:

Note: Ordinal and nominal variables will be treated the same in this course. We recommend taking more statistics courses in the future to learn better methods of analysis for ordinal variables.

Data frame:

Scatterplot:

Each point represents:

Positive association:

Negative association:

Association or Dependent variables:

Independent variables:

Explanatory variable:

Response variable:

Observational study:

Experiment:

Placebo:

Notes

Big Idea: Variability is inevitable! We would not expect to get *exactly* 50 heads in 100 coin flips. The statistical question then is whether any differences found in data are due to random variability, or if something else is going on.

The larger the difference, the **less we believe the difference was due to chance**.

In a data frame, rows correspond to _____

and columns correspond to _____.

How many types of variables are discussed? Explain the differences between them and give an example of each.

True or False: A pair of variables can be both associated AND independent.

True or False: Given a pair of variables, one will always be the explanatory variable and one the response variable.

True or False: If a study does have an explanatory and a response variable, that means changes in the explanatory variable must **cause** changes in the response variable.

True or False: Observational studies can show a naturally occurring association between variables.

Example: Section 1.1 — Case study: Using stents to prevent strokes

1. What is the principle question the researchers hope to answer? (We call this the **research question**.)
2. When creating two groups to compare, do the groups have to be the same size (same number of people in each)?
3. What are the cases or observational units in this study?
4. Is there a clear explanatory and response variable? If so, name the variable in each role and determine the type of variable (discrete, continuous, nominal, or ordinal).
5. What is the purpose of the control group?
6. Is this an example of an observational study or an experiment? How do you know?
7. Consider Tables 1.1 and 1.2. Which table is more helpful in answering the research question? Justify your answer.
8. Describe in words what is shown in Figure 1.1. Specifically, compare the proportion of patients who had a stroke between the treatment and control groups after 30 days as well as after 365 days.

9. Given the notion that the larger the difference between the two groups (for a given sample size), the less believable it is that the difference was due to chance, which measurement period (30 days or 365 days) provide stronger evidence that there is an association between stents and strokes, or that the differences are not due to random chance?
10. This study reported finding evidence that stents *increase* the risk of stroke. Does this conclusion apply to all patients and all stents?
11. This study reported finding evidence that stents *increase* the risk of stroke. This conclusion implies a causal link between stents and an increased risk of stroke. Is that conclusion valid? Justify your answer.

1.3 Video Notes: Intro to data

Read through Sections 1.1 – 1.3 in the course textbook and watch the course videos prior to coming to class on Friday. Fill in the following questions to aid in your understanding of the material. Many of the following questions are asked on the video quiz on Gradescope.

1.3.1 Course Videos

- 1.2.1and1.2.2
- 1.2.3to1.2.4

Data basics: Video 1.2.1and1.2.2

Data: _____ used to answer research questions

Observational unit or case: the people or things we _____ data from; represents the _____ in each data set

Variable: characteristics measured on each _____.

Types of variables

- Categorical variable:
 - Ordinal: levels of the variable have a natural ordering
Examples: ‘Scale’ questions, years of schooling completed
 - Nominal: levels of the variable do not have a natural ordering
Examples: hair color, eye color, zipcode
- Quantitative variable:
 - Continuous variables: value can be any value within a range.
Examples: percentage of students who are nursing majors
 - average hours of exercise per week
 - distance or time (measured with enough precision)
 - Discrete variables: can only be specific values, with jumps between
Examples: SAT score
 - number of car accidents

Example: The Bureau of Transportation Statistics (“Bureau of Transportation Statistics” 2019) collects data on all forms of public transportation. The data set seen here includes several variables collect on flights departing on a random sample of 150 US airports in December of 2019.

```
airport <- read.csv("data/airport_delay.csv")
glimpse(airport)
#> Rows: 150
#> Columns: 19
#> $ airport      <chr> "ABI", "ABY", "ACV", "ACY", "ADQ", "AEX", "ALB", "~
#> $ city          <chr> "Abilene", "Albany", "Arcata/Eureka", "Atlantic Ci~
#> $ state         <chr> " TX", " GA", " CA", " NJ", " AK", " LA", " NY", "~
#> $ airport_name  <chr> " Abilene Regional", " Southwest Georgia Regional"~
#> $ hub           <chr> "no", "no", "no", "no", "no", "no", "no", "n~
#> $ international <chr> "no", "no", "no", "yes", "no", "yes", "yes", "yes"~
#> $ elevation_1000 <dbl> 1.7906, 0.1932, 0.2223, 0.0748, 0.0787, 0.0881, 0.~
#> $ latitude      <dbl> 32.4, 31.5, 41.0, 39.5, 57.7, 31.3, 42.7, 35.2, 45~
#> $ longitude     <dbl> -99.7, -81.2, -124.1, -74.6, -152.5, -92.5, -73.8, ~
#> $ arr_flights   <int> 195, 81, 215, 293, 54, 282, 943, 410, 53, 32314, 6~
#> $ perc_delay15  <dbl> 16.410256, 13.580247, 23.255814, 15.358362, 12.962~
#> $ perc_cancelled <dbl> 0.5128205, 0.0000000, 4.1860465, 0.6825939, 14.814~
#> $ perc_diverted <dbl> 0.00000000, 0.00000000, 2.32558139, 0.68259386, 0.~
#> $ arr_delay     <int> 1563, 1244, 4763, 2905, 329, 1293, 15127, 9705, 25~
#> $ carrier_delay <int> 459, 890, 1613, 476, 180, 302, 5627, 2253, 439, 10~
#> $ weather_delay <int> 21, 43, 549, 124, 1, 58, 2346, 168, 1236, 13331, 2~
#> $ nas_delay     <int> 257, 39, 154, 771, 51, 112, 2096, 616, 746, 45674, ~
#> $ security_delay <int> 0, 0, 0, 25, 0, 0, 44, 0, 0, 375, 0, 83, 0, 23, 0, ~
#> $ late_aircraft_delay <int> 826, 272, 2447, 1509, 97, 821, 5014, 6668, 108, 10~
```

- What are the observational units?
- Identify which variables are categorical.
- Identify which variables are quantitative.

Exploratory data analysis (EDA)

Summary statistic: a single number which _____ an entire data set

- Also called the point estimate.

Examples:

proportion of people who had a stroke

mean (or average) age

- The summary statistic and type of plot used depends on the type (categorical or quantitative) of variable(s)!

Roles of variables: 1.2.3to1.2.4

Explanatory variable: predictor variable

- The variable researchers think *may be* _____ the other variable.
- In an experiment, what the researchers _____ or _____.
- The groups that we are comparing from the data set.

Response variable:

- The variable researchers think *may be* _____ by the other variable.
- Always simply _____ or _____; never controlled by researchers.

Examples:

Can you predict a criminal's height based on the footprint left at the scene of a crime?

- Identify the explanatory variable:
- Identify the response variable:

Does marking an item on sale (even without changing the price) increase the number of units sold per day, on average?

- Identify the explanatory variable:
- Identify the response variable:

In the Physician's Health Study ("Physician's Health Study," n.d.), male physicians participated in a study to determine whether taking a daily low-dose aspirin reduced the risk of heart attacks. The male physicians were randomly assigned to the treatment groups. After five years, 104 of the 11,037 male physicians taking a daily low-dose aspirin had experienced a heart attack while 189 of the 11,034 male physicians taking a placebo had experienced a heart attack.

- Identify the explanatory variable:
- Identify the response variable:

Relationships between variables

- Association: the _____ between variables create a pattern; knowing something about one variable tells us about the other.
 - Positive association: as one variable _____, the other tends to _____ also.
 - Negative association: as one variable _____, the other tends to _____.
- Independent: no clear pattern can be seen between the _____.

1.3.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. What is the explanatory variable in the Male Physicians study?
2. What is the response variable in the Male Physicians study?

1.4 Guided Activity 1: Intro to Data Analysis

1.4.1 Learning outcomes

- Identify observational units, variables, and variable types in a statistical study.
- Creating a data set
- Creating a plot of the data
- Calculating a summary statistic of the data

1.4.2 Terminology review

Statistics is the study of how best to collect, analyze, and draw conclusions from data. This week in class you will be introduced to the following terms:

- Observational units or cases
- Variables: categorical or quantitative

For more on these concepts, read Chapter 1 in the textbook.

Further analysis of class data set

1. What are the observational units or cases for the data collected in class on day 1?
2. How many observations are reported in the data set? This is the **sample size**.
3. The header for each column in the data set describes each variable measured on the observational unit. For each column of data, fill in the following table identifying the type of each variable, and if the variable is categorical whether the variables is binary and if the variable is quantitative the units of measure used.

Column	Type of Variable	Binary?	Units?
Major			
Residency			
Num Credits			
Dominant hand			
Hand Span			
Grip strength dominant hand			
Grip strength non-dominant hand			

4. Review the completed data set with your table. Remember that when creating a data set for use in R it is important to use single words or an underscore between words. Each outcome must be written the same way each time to have consistency between responses. Do not include units of measure in the data set when reporting numerical values. Write down some issues found with the created class data set.

Summary Statistics

5. Brainstorm with your table on ways to summarize the variable **Residency** both visually and mathematically. Write down your tables ideas.
6. Your instructor will use RStudio (a statistical software package) to create a graph and numerical summaries of the variable **Residency**. Draw the graph created here.
7. Your instructor will use RStudio to find the count for each level of the variable **Residency**. Calculate the proportion of Stat 216 students that have in-state residency.

8. Brainstorm with your table on ways to summarize the variable `Grip_Strength_Dominant` both visually and mathematically. Write down your table's ideas.

9. Your instructor will use RStudio to create a graph and numerical summaries of the variable `Grip_Strength_Dominant`. Draw the graph created here.

1.4.3 Take-home messages

1. There are two types of variables: categorical (groups) and quantitative (numerical measures).
2. We will learn more about summarizing variable later in the semester. Categorical variables are summarized by calculating a proportion from the data and quantitative variables are summarized by finding the mean and the standard deviation.

1.4.4 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered, and to write down the names and contact information of your teammates.

Study Design

2.1 Reading Guide: Sampling, Study Design, and Scope of Inference

Textbook Chapter 2: Study Design

Section 2.1: Sampling principles and strategies

Vocabulary

(Target) Population:

Sample:

Anecdotal evidence:

Bias:

Selection bias:

Non-response bias:

Response bias:

Convenience sample:

Simple Random Sample:

Non-response rate:

Representative:

Notes

Ideally, how should we sample cases from our target population? Using what sampling method?

Notes on types of sampling bias

- Someone must first be *chosen* to be in a study and refuse to participate in order to have **non-response bias**.
- There must be a valid reason for someone to lie or be untruthful to justify saying **response bias** is present. Yes, anyone could lie at any time to any question. Response bias is when those lies are predictable and systematic based on outside influences.

True or False: Convenience sampling tends to result in non-response bias.

True or False: Volunteer sampling tends to result in response bias.

True or False: Random sampling helps to resolve selection bias, but has no impact on non-response or response bias.

Section 2.2 & 2.3: Study Design

Reminders from Section 1.2

Explanatory variable: The variable researchers think *may be* affecting the other variable. What the researchers control/assign in an experiment. If comparing groups, the explanatory variable puts the observational units into groups.

Response variable: The variable researchers think *may be* influenced by the other variable. This variable is always observed, never controlled or assigned.

Vocabulary

Observational study:

Observational data:

Prospective study:

Retrospective study:

Confounding variable:

Experiment:

Randomized experiment:

Blocking:

Treatment group:

Control group:

Placebo:

Placebo effect:

Blinding:

Scope of inference:

Generalizability:

Causation:

Notes

What are the four principles of a well-designed randomized experiment?

Section 2.4: Scope of inference

Fill in the appropriate scope of inference for each study design.

	Study Design	
Selection of Cases	Randomized experiment	Observational study
Random sample (and no other sampling bias)		
Non-random sample (or other sampling bias)		

True or False: Observational studies can show an association between two variables, but cannot determine a causal relationship.

True or False: In order for an experiment to be valid, a placebo must be used.

True or False: If random sampling of the target population is used, and no other types of bias are suspected, results from the sample can be generalized to the entire target population.

True or False: If random sampling of the target population is used, and no other types of bias are suspected, results from the sample can be inferred as a causal relationship between the explanatory and response variables.

2.2 Video Notes: Study Design

Read sections 2.1 - 2.5 in the course textbook. Use the following videos to complete the lecture notes for Module 2.

2.2.1 Course Videos

- 2.1
- 2.2to2.4

Sampling Methods: Video 2.1

The method used to collect data will impact

- Target population: all _____ or _____ of interest
- Sample: _____ or _____ from which data is collected

Example: Many high schools moved to partial or fully online schooling in Spring of 2020. Did students who graduated in 2020 tend to have a lower GPA during freshman year of college than the previous class of college freshmen? A nationally representative sample of 1000 college students who were freshmen in AY19-20 and 1000 college students who were freshmen in AY20-21 was taken to answer this question.

- What is the target population?
- What is the sample?

Good vs. bad sampling

GOAL: to have a sample that is _____ of the _____ on the variable(s) of interest

- Unbiased sample methods:

Simple random sample

- Biased sampling method:

Types of Sampling Bias

- Selection bias:

Example of Selection Bias: Newspaper article from 1936 reported that Landon won the presidential election over Roosevelt based on a poll of 10 million voters. Roosevelt was the actual winner. What was wrong with this poll? Poll was completed using a telephone survey and not all people in 1936 had a telephone. Only a certain subset of the population owned a telephone so this subset was over-represented in the telephone survey. The results of the study, showing that Landon would win, did not represent the target population of all US voters.

- Non-response bias:

- To calculate the non-response rate:

$$\frac{\text{number of people who do not respond}}{\text{total number of people selected for the sample}} \times 100\%$$

- For non-response bias to occur must first select people to participate and then they choose not to.

Example of Non-response bias: A company randomly selects buyers to complete a review of an online purchase but some choose not to respond.

- Response bias:

Example of Response Bias: Police officer pulls you over and asks if you have been drinking. Expect people to say no, whether they have been drinking or not.

- Need to be able to predict how people will respond.

Words of caution:

- Convenience samples: gathering data for those who are easily accessible; online polls

Selection bias?

Non-response bias?

Response bias?

- Random sampling reduces _____ bias, but has no impact on _____ or _____ bias.

Video Example

A radio talk show asks people to phone in their views on whether the United States should pay off its debt to the United Nations.

- Selection?
- Non-response?
- Response?

The Wall Street Journal plans to make a prediction for the US presidential election based on a survey of its readers and plans to follow-up to ensure everyone responds.

- Selection?
- Non-response?
- Response?

A police detective interested in determining the extent of drug use by high school students, randomly selects a sample of high school students and interviews each one about any illegal drug use by the student during the past year.

- Selection?
- Non-response?
- Response?

Observational studies, experiments, and scope of inference: Video 2.2to2.4

- Review
 - Explanatory variable: the variable researchers think *may be* effecting the other variable.
 - Response variable: the variable researchers think *may be* influenced by the other variable.
- Confounding variable:
 - associated with both the explanatory and the response variable
 - explains the association shown by the data

Example:

Study design

- Observational study:
- Experiment:

Principles of experimental design

- Control: hold other differences constant across groups
- Randomization: randomized experiment
- Replication: large sample size or repeat of study
- Blocking: group based on certain characteristics

Example: It is well known that humans have more difficulty differentiating between faces of people from different races than people within their own race. A 2018 study published in the *Journal of Experimental Psychology* (Levin 2000): Human Perception and Performance investigated a similar phenomenon with gender. In the study, volunteers were shown several pictures of strangers. Half the volunteers were randomly assigned to rate the attractiveness of the individuals pictured. The other half were told to rate the distinctiveness of the faces seen. Both groups were then shown a slideshow of faces (some that had been rated in the first part of the study, some that were new to the volunteer) and asked to determine if each face was old or new. Researchers found people were better able to recognize faces of their own gender when asked to rate the distinctiveness of the faces, compared to when asked to rate the attractiveness of the faces.

- What is the study design?

Example: In the Physician's Health Study ("Physician's Health Study," n.d.), male physicians participated in a study to determine whether taking a daily low-dose aspirin reduced the risk of heart attacks. The male physicians were randomly assigned to the treatment groups. After five years, 104 of the 11,037 male physicians taking a daily low-dose aspirin had experienced a heart attack while 189 of the 11,034 male physicians taking a placebo had experienced a heart attack.

- What is the study design?
- Assuming these data provide evidence that the low-dose aspirin group had a lower rate of heart attacks than the placebo group, is it valid for the researchers to conclude the lower rate of heart attacks was caused by the daily low-dose aspirin regimen?

Scope of Inference

1. How was the sample selected?
 - Random sample with no sampling bias:
 - Non-random sample with sampling bias:

2. What is the study design?

- Randomized experiment:
- Observational study:

Scope of Inference Table:

Scope of Inference: If evidence of an association is found in our sample, what can be concluded?

	Study Type		
Selection of cases	Randomized experiment	Observational study	
Random sample (and no other sampling bias)	Causal relationship, and can generalize results to population.	Cannot conclude causal relationship, but can generalize results to population.	➡ Inferences to population can be made
No random sample (or other sampling bias)	Causal relationship, but cannot generalize results to a population.	Cannot conclude causal relationship, and cannot generalize results to a population.	➡ Can only generalize to those similar to the sample due to potential sampling bias

↓
 Can draw cause-and-
effect conclusions

↓
 Can only discuss association
due to potential confounding
variables

Example: It is well known that humans have more difficulty differentiating between faces of people from different races than people within their own race. A 2018 study published in the Journal of Experimental Psychology (Levin 2000): Human Perception and Performance investigated a similar phenomenon with gender. In the study, volunteers were shown several pictures of strangers. Half the volunteers were randomly assigned to rate the attractiveness of the individuals pictured. The other half were told to rate the distinctiveness of the faces seen. Both groups were then shown a slideshow of faces (some that had been rated in the first part of the study, some that were new to the volunteer) and asked to determine if each face was old or new. Researchers found people were better able to recognize faces of their own gender when asked to rate the distinctiveness of the faces, compared to when asked to rate the attractiveness of the faces.

- What is the scope of inference for this study?

2.2.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. Purpose of random assignment:
2. Purpose of random selection:

2.3 Guided Activity 2: American Indian Address

2.3.1 Learning outcomes

- Explain why a sampling method is unbiased or biased.
- Identify biased sampling methods.
- Explain the purpose of random selection and its effect on scope of inference.

2.3.2 Terminology review

In this activity, we will examine unbiased and biased methods of sampling. Some terms covered in this activity are:

- Random sample
- Unbiased vs biased methods of selection
- Generalization

To review these concepts, see Chapter 2 in the textbook.

2.3.3 Class Preparation

Prior to the next class, complete questions 1–3.

2.3.4 American Indian Address

For this activity, you will read a speech given by Jim Becenti, a member of the Navajo American Indian tribe, who spoke about the employment problems his people faced at an Office of Indian Affairs meeting in Phoenix, Arizona, on January 30, 1947 (Moquin and Van Doren 1973). His speech is below:

It is hard for us to go outside the reservation where we meet strangers. I have been off the reservation ever since I was sixteen. Today I am sorry I quit the Santa Fe [Railroad]. I worked for them in 1912–13. You are enjoying life, liberty, and happiness on the soil the American Indian had, so it is your responsibility to give us a hand, brother. Take us out of distress. I have never been to vocational school. I have very little education. I look at the white man who is a skilled laborer. When I was a young man I worked for a man in Gallup as a carpenter's helper. He treated me as his own brother. I used his tools. Then he took his tools and gave me a list of tools I should buy and I started carpentering just from what I had seen. We have no alphabetical language.

We see things with our eyes and can always remember it. I urge that we help my people to progress in skilled labor as well as common labor. The hope of my people is to change our ways and means in certain directions, so they can help you someday as taxpayers. If not, as you are going now, you will be burdened the rest of your life. The hope of my people is that you will continue to help so that we will be all over the United States and have a hand with you, and give us a brotherly hand so we will be happy as you are. Our reservation is awful small. We did not know the capacity of the range until the white man come and say "you raise too much sheep, got to go somewhere else," resulting in reduction to a skeleton where the Indians can't make a living on it. For eighty years we have been confused by the general public, and what is the condition of the Navajo today? Starvation! We are starving for education. Education is the main thing and the only thing that is going to make us able to compete with you great men here talking to us.

By eye selection

1. Circle ten words in Jim Becenti's speech which are a representative sample of the length of words in the entire text. Describe your method for selecting this sample.
2. Fill in the table below with your selected words from the previous question and the length of each word (number of letters/digits in the word):

Observation	Word	Length
1		
2		
3		
4		
5		
6		
7		
8		
9		
10		

3. Calculate the mean (average) word length in your selected sample. Is this value a parameter or a statistic?

2.3.5 Class Activity

1. Report your mean word length in the Google sheet. Your instructor will create a visualization of the distribution of results generated by your class. Draw a picture of the plot here. Include a descriptive x -axis label. Report the mean and standard deviation of the data.

The plot created in the previous question is a sampling distribution of statistics. This sampling distribution plots the mean word length from many samples taken from the population of words.

2. Based on the plot and summary statistics of sample mean word lengths, what is your best guess for the average word length of the population of all 359 words in the speech?
3. The true mean word length of the population of all 359 words in the speech is 3.95 letters. Is this value a parameter or a statistic?

Where does the value of 3.95 fall in the plot given? Near the center of the distribution? In the tails of the distribution?

4. If the class samples were truly representative of the population of words, what proportion of sample means in the sampling distribution would you expect to be below 3.95?
5. Using the graph in Q1, estimate the proportion of students' computed sample means that were lower than the true mean of 3.95 letters?
6. Based on your answers to questions 4 and 5, would you say the sampling method used by the class is biased or unbiased? Justify your answer.
7. If the sampling method is biased, what type of sampling bias (selection, response, non-response) is present? What is the direction of the bias, i.e., does the method tend to overestimate or underestimate the population mean word length?
8. Should we use results from our "by eye" samples to make a statement about the word length in the population of words in Becenti's address? Why or why not?

Types of bias

9. To determine if the proportion of out-of-state undergraduate students at Montana State University has increased in the last 10 years, a statistics instructor sent an email survey to 500 randomly selected current undergraduate students. One of the questions on the survey asked whether they had in-state or out-of-state residency. She only received 378 responses.

Sample size:

Observational units sampled:

Target population:

Justify why there is non-response bias in this study.

10. A television station is interested in predicting whether or not local voters will pass a referendum to legalize marijuana for adult. The TV station asks its viewers to phone in and indicate whether they are in favor or opposed to the referendum. Of the 2241 viewers who phoned in, forty-five percent were opposed to legalizing marijuana.

Sample size:

Observational units sampled:

Target population:

Justify why there is selection bias in this study.

11. To gauge the interest of Bozeman City Voters in a new swimming pool, a local organization stood outside of the Bogart Pool in Bozeman, MT, during open hours. One of the questions they asked was, “Since the Bogart Pool is in such bad repair, don’t you agree that the city should fund a new pool?”

Sample size:

Observational units sampled:

Target population:

Justify why there is response bias in this study.

Justify why there is selection bias in this study.

2.3.6 Take-home messages

1. There are three types of bias to be aware of when designing a sampling method: selection bias, non-response bias, and response bias.
2. When we use a biased method of selection, we will over or underestimate the parameter.
3. To see if a method is biased, we compare the distribution of the estimates to the true value. We want our estimate to be on target or unbiased. When using unbiased methods of selection, the mean of the distribution matches or is very similar to our true parameter.
4. If the sampling method is biased, inferences made about the population based on a sample estimate will not be valid.

2.3.7 Additional notes

Use this space to summarize your thoughts and take additional notes on today’s activity and material covered.

2.4 Guided Activity 3: American Indian Address (continued)

2.4.1 Learning outcomes

- Explain the purpose of random selection and its effect on scope of inference.
- Select a simple random sample from a finite population using a random number generator.
- Explain why a sampling method is unbiased or biased.
- Explain the effect of sample size on sampling variability.
- Explore a sampling distribution of statistics.

2.4.2 Terminology review

In today's activity, we will examine unbiased and biased methods of sampling. Some terms covered in this activity are:

- Random sample
- Unbiased vs biased methods of selection
- Generalization

To review these concepts, see Section 2.1 in the textbook.

Random selection

Today we will return to the American Indian Address introduced in the guided activity. Suppose instead of attempting to select a representative sample by eye (which did not work), each student used a random number generator to select a simple random sample of 10 words. A **simple random sample** relies on a random mechanism to choose a sample, without replacement, from the population, such that every sample of size 10 is equally likely to be chosen.

To use a random number generator to select a simple random sample, you first need a numbered list of all the words in the population, called a **sampling frame**. You can then generate 10 random numbers from the numbers 1 to 359 (the number of words in the population), and the chosen random numbers correspond to the chosen words in your sample.

1. Use the random number generator at <https://istats.shinyapps.io/RandomNumbers/> to select a simple random sample from the population of all 359 words in the speech.
- Set “Choose Minimum” to 1 and “Choose Maximum” to 359 to represent the 359 words in the population (the sampling frame).
 - Set “How many numbers do you want to generate?” to 10 and ensure the “No” option is selected under “Sample with Replacement?”
 - Click “Generate”.

Fill in the table below with the random numbers selected and use the Becenti.csv data file found on D2L to determine each number's corresponding word and word length (number of letters/digits in the word):

Observation	Number	Length
1		
2		
3		
4		
5		
6		
7		
8		
9		
10		

2. Calculate the mean word length in your selected sample in question 1. Is this value a parameter or a statistic?
3. Report your mean word length in the Google sheet. Your instructor will create a visualization of the distribution of results generated by your class. Draw a picture of the plot here. Include a descriptive x -axis label. Report the mean and standard deviation of the data.

4. Where does the value 3.95, the true mean word length, fall in the distribution given? Near the center of the distribution? In the tails of the distribution? Circle this value on the provided distribution.

5. How does the plot given in this activity compare to the plot generated in the guided activity?

Is the shape similar?

Is the range (smallest to largest values) similar?

Is the mean of the distribution similar?

Why didn't everyone get the same sample mean?

One set of randomly generated sample mean word lengths from a single class may not be large enough to visualize the distribution results. Let's have a computer generate 1,000 sample mean word lengths for us.

- Navigate to the “One Variable with Sampling” Rossman/Chance web applet: <http://www.rossmanchance.com/applets/2021/sampling/OneSample.html?population=gettysburg>.
 - Click “Clear” below the text box containing data from the Gettysburg address to delete that data set.
 - Download the Becenti.csv file from D2L and open the spreadsheet on your computer.
 - Copy and paste the population of word lengths (column C) into the applet from the data set provided making sure to include the header. Click “Use Data”. Verify that the mean for the data set is 3.953 with a sample size of 359. If these are not the values you got, check with your instructor for help with copying in the data set correctly.
 - Click the check-box for “Show Sampling Options”
 - Select 1000 for “Number of samples” and select 10 for the “Sample size”.
 - Click “Draw Samples”.
6. The plot labeled “Statistics” displays the 1,000 randomly generated sample mean word lengths (the sampling distribution). Sketch this plot below. Include a descriptive x -axis label and be sure to write down the provided mean and SD (standard deviation) of the distribution.

7. What is the center value (mean) of the distribution created in question 6?

8. Explain why the sampling method of using a random number generator to generate a sample is a “better” method than choosing 10 words “by eye”.
9. Is random selection an unbiased method of selection? Explain your answer. Be sure to reference your plot from question 6.

Effect of sample size

We will now consider the impact of sample size.

10. First, consider if each student had selected 20 words, instead of 10, by eye. Do you think this would make the plot from the out-of-class activity centered on 3.95 (the true mean word length)? Explain your answer.
11. Now we will select 20 words instead of 10 words at random.
 - In the “One Variable with Sampling” Rossman/Chance web applet(<http://www.rossmanchance.com/applets/2021/sampling/OneSample.html?population=gettysburg.>), change the Sample size to 20.
 - Click “Draw Samples”.

The plot labeled “Statistics” displays the 1,000 randomly generated sample mean word lengths. Sketch this plot below. Include a descriptive x -axis label and be sure to write down the provided mean and SD (standard deviation) of the distribution.

12. Compare the distribution created in question 11 to the one created in question 6.

Is the shape similar?

Is the range (smallest to largest values) similar?

Is the mean of the distribution similar?

13. Compare the values of the standard deviation of the plots in question 11 and in question 6. Which plot shows the smallest standard deviation?
14. Using the evidence from your simulations, answer the following research questions:
Does changing the sample size impact whether the sample estimates are unbiased? Explain your answer.

Does changing the sample size impact the variability (spread) of sample estimates? Explain your answer
15. What is the purpose of random selection of a sample from the population?

2.4.3 Take-home messages

1. Random selection is an unbiased method of selection.
2. To determine if a sampling method is biased or unbiased, we compare the distribution of the estimates to the true value. We want our estimate to be on target or unbiased. When using unbiased methods of selection, the mean of the distribution matches or is very similar to our true parameter.
3. Random selection eliminates selection bias. However, random selection will not eliminate response or non-response bias.
4. The larger the sample size, the more similar (less variable) the statistics will be from different samples.
5. Sample size has no impact on whether a *sampling method* is biased or not. Taking a larger sample using a biased method will still result in a sample that is not representative of the population.

2.4.4 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

2.5 Activity 2: Study Design

2.5.1 Learning outcomes

- Explain the purpose of random assignment and its effect on scope of inference.
- Identify whether a study design is observational or an experiment.
- Identify confounding variables in observational studies and explain why they are confounding.

2.5.2 Terminology review

In this activity, we will examine different study designs, confounding variables, and how to determine the scope of inference for a study. Some terms covered in this activity are:

- Scope of inference
- Explanatory variable
- Response variable
- Confounding variable
- Experiment
- Observational study

To review these concepts, see Sections 2.2 through 2.5 in the textbook.

2.5.3 Atrial fibrillation

Atrial fibrillation is an irregular and often elevated heart rate. In some people, atrial fibrillation will come and go on its own, but others will experience this condition on a permanent basis. When atrial fibrillation is constant, medications are required to stabilize the patient's heart rate and to help prevent blood clots from forming. Pharmaceutical scientists at a large pharmaceutical company believe they have developed a new medication that effectively stabilizes heart rates in people with permanent atrial fibrillation. They set out to conduct a trial study to investigate the new drug. The scientists will need to compare the proportion of patients whose heart rate is stabilized between two groups of subjects, one of whom is given a placebo and the other given the new medication.

1. Identify the explanatory and response variable in this trial study.

Explanatory variable:

Response variable:

Suppose 24 subjects with permanent atrial fibrillation have volunteered to participate in this study. There are 16 subjects that self-identified as male and 8 subjects that self-identified as female.

2. One way to separate into two groups would be to give all the males the placebo and all the females the new drug. Explain why this is not a reasonable strategy.

3. Could the scientists fix the problem with the strategy presented in question 2 by creating equal sized groups by putting 4 males and 8 females into the drug group and the remaining 12 males in the placebo group? Explain your answer.

4. A third strategy would be to **block** on sex. In this type of study, the scientists would assign 4 females and 8 males to each group. Using this strategy, out of the 12 individuals in each group what **proportion** are males?

5. Assume the scientists used the strategy in question 4, but they put the four tallest females and eight tallest males into the drug group and the remaining subjects into the placebo group. They found that the proportion of patients whose heart rate stabilized is higher in the drug group than the placebo group.

Could that difference be due to the sex of the subjects? Explain your answer.

Could it be due to other variables? Explain your answer.

While the strategy presented in question 5 controlled for the sex of the subject, there are more potential **confounding variables** in the study. A confounding variable is a variable that is *both*

1. associated with the explanatory variable, *and*
2. associated with the response variable.

When both these conditions are met, if we observe an association between the explanatory variable and the response variable in the data, we cannot be sure if this association is due to the explanatory variable or the confounding variable—the explanatory and confounding variables are “confounded.”

Random assignment means that subjects in a study have an equally likely chance of receiving any of the available treatments.

6. You will now investigate how randomly assigning subjects impacts a study's scope of inference.

- Navigate to the “Randomizing Subjects” applet under the “Other Applets” heading at: <http://www.rossmanchance.com/ISIapplets.html>. This applet lists the sex and height of each of the 24 subjects. Click “Show Graphs” to see a bar chart showing the sex of each subject. Currently, the applet is showing the strategy outlined in question 3.
- Click “Randomize”.

In this random assignment, what proportion of males are in group 1 (the placebo group)?

What proportion of males are in group 2 (the drug group)?

What is the difference in proportion of males between the two groups (placebo - drug)?


7. Notice the difference in the two proportions is shown as a dot in the plot at the bottom of the web page. Un-check the box for Animate above “Randomize” and click “Randomize” again. Did you get the same difference in proportion of males between the placebo and drug groups?

8. Change “Replications” to 998 (for 1000 total). Click “Randomize” again. Sketch the plot of the distribution of difference in proportions from each of the 1000 random assignments here. Be sure to include a descriptive x -axis label.

9. Does random assignment *always* balance the placebo and drug groups based on the sex of the participants? Does random assignment *tend* to make the placebo and drug groups *roughly* the same with respect to the distribution of sex? Use your plot from question 8 to justify your answers.

10. Change the drop-down menu below Group 2 from “sex” to “height”. The applet now calculates the average height in the placebo and drug groups for each of the 1000 random assignments. The dot plot displays the distribution of the difference in mean heights (placebo - drug) for each random assignment. Based on this dot plot, is height distributed equally, on average, between the two groups? Explain how you know.

The diagram below summarizes these ideas about confounding variables and random assignment. When a confounding variable is present (such as sex or height), and an association is found in a study, it is impossible to discern what caused the change in the response variable. Is the change the result of the explanatory variable or the confounding variable? However, if all confounding variables are *balanced* across the treatment groups, then only the explanatory variable differs between the groups and thus *must have caused* the change seen in the response variable.



11. What is the purpose of random assignment of the subjects in a study to the explanatory variable groups? Cross out the arrow in the figure above that is eliminated by random assignment.


12. Suppose in this study on atrial fibrillation, the scientists did randomly assign groups and found that the drug group has a higher proportion of subjects whose heart rates stabilized than the placebo group. Can the scientists conclude the new drug *caused* the increased chance of stabilization? Explain your answer.


13. Is the sample of subjects a simple random sample or a convenience sample?

14. Both the sampling method and the study design will help to determine the *scope of inference* for a study: To *whom* can we generalize, and can we conclude *causation or only association*? Use your answers to question 12 and 13 and the table on the next page to determine the scope of inference of this trial study described in question 12.

Scope of Inference: If evidence of an association is found in our sample, what can be concluded?

Selection of cases	Study Type		
	Randomized experiment	Observational study	
Random sample (and no other sampling bias)	Causal relationship, and can generalize results to population.	Cannot conclude causal relationship, but can generalize results to population.	→ Inferences to population can be made
No random sample (or other sampling bias)	Causal relationship, but cannot generalize results to a population.	Cannot conclude causal relationship, and cannot generalize results to a population.	→ Can only generalize to those similar to the sample due to potential sampling bias


 Can draw cause-and-
effect conclusions


 Can only discuss association
due to potential confounding
variables

2.5.4 Scope of Inference

The two main study designs we will cover are **observational studies** and **experiments**. In observational studies, researchers have no influence over which subjects are in each group being compared (though they can control other variables in the study). An experiment is defined by assignment of the treatment groups of the *explanatory variable*, typically via random assignment.

For the next exercises identify the study design (observational study or experiment), the sampling method, and the scope of inference.

15. The pharmaceutical company Moderna Therapeutics, working in conjunction with the National Institutes of Health, conducted Phase 3 clinical trials of a vaccine for COVID-19 in the Fall of 2021. US clinical research sites enrolled 30,000 volunteers without COVID-19 to participate. Participants were randomly assigned to receive either the candidate vaccine or a saline placebo. They were then followed to assess whether or not they developed COVID-19. The trial was double-blind, so neither the investigators nor the participants knew who was assigned to which group.

Study design:

Sampling method:

Scope of inference:

16. In another study, a local health department randomly selected 1000 US adults without COVID-19 to participate in a health survey. Each participant was assessed at the beginning of the study and then followed for one year. They were interested to see which participants elected to receive a vaccination for COVID-19 and whether any participants developed COVID-19.

Study design:

Sampling method:

Scope of inference:

2.5.5 Take-home messages

1. The study design (observational study vs, experiment) determines if we can draw causal inferences or not. If an association is detected, a randomized experiment allows us to conclude that there is a causal (cause-and-effect) relationship between the explanatory and response variable. Observational studies have potential confounding variables within the study that prevent us from inferring a causal relationship between the variables studied.
2. Confounding variables are variables not included in the study that are related to both the explanatory and the response variables. When there are potential confounding variables in the study we cannot draw causal inferences.
3. Random assignment balances confounding variables across treatment groups. This eliminates any possible confounding variables by breaking the connections between the explanatory variable and the potential confounding variables.
4. Observational studies will always carry the possibility of confounding variables. Randomized experiments, which use random assignment, will have no confounding variables.

2.5.6 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

Exploring Categorical and Quantitative Data

3.1 Reading Guide: Introduction to R, Categorical Variables, and a Single Quantitative Variable

Textbook Chapter 3 Applications: Data

Notes

R is case sensitive, meaning it reads `data` differently from `Data`. If you get an error message, check that your capitalization is correct.

R does not like spaces or special characters. This means the column and row headers in the data set should not have spaces, periods, commas, etc. Instead of titling the variable `column header`, use `column_header` or `ColumnHeader`.

Tidy data: Data frames should have

1 row per _____,

1 column per _____.

We highly recommend completing the R/RStudio tutorials in section 3 to help understand R better.

We will not expect you to be able to write full code independently for this course. For Stat 216, you will need to understand types (categorical or quantitative) and roles (explanatory or response) of variables, as well as the structure of data, in order to fill in a few blanks in provided code to graph or analyze data.

Functions

State what these introductory functions do in R:

`glimpse(data_set_name)`

`head(data_set_name)`

`data_set_name$variable_name`

`<-`

`%>%`

Textbook Chapter 4: Exploring categorical data

Vocabulary

Frequency table:

Relative frequency table:

Contingency or two-way table:

Association (between two variables):

Unconditional proportion:

Conditional proportion:

Row proportions:

Column proportions:

Statistic:

Sample proportion:

Notation:

Parameter:

Population proportion:

Notation:

Bar plot:

Segmented bar plot:

Mosaic plot:

Simpson's Paradox:

Notes

In a contingency table, which variable (explanatory or response) generally will make the columns of the table? Which variable will make the rows of the table?

In a segmented bar plot, the bars represent the levels of which variable? The segments represent the levels of which variable?

What type of plot(s) are appropriate to display a single categorical variable?

What type of plot(s) are appropriate to display two categorical variables?

What is the difference between a standardized segmented bar plot and a mosaic plot?

True or false: Pie charts are generally highly recommended ways to graphically display categorical data.

True or false: Two categorical variables are associated if the conditional proportions of a particular outcome (typically of the response variable) differ across levels of the other variable (typically the explanatory variable).

True or false: When a segmented bar plot has segments that sum to 1 (or 100%), the segment heights correspond to the proportions conditioned on the **segment**.

Review of Simpson's Paradox

Based on the segmented bar plot in Figure 4.6, which race of defendant was more likely to have the death penalty invoked?

Based on the segmented bar plot in Figure 4.7 and Table 4.9, which race of defendant was more likely to have the death penalty invoked when the victim was Caucasian?

Based on the segmented bar plot in Figure 4.7 and Table 4.9, which race of defendant was more likely to have the death penalty invoked when the victim was African American?

The direction of the relationship between the _____ and _____ variables is **reversed** when accounting for a _____ variable.

Textbook Chapter 5: Exploring quantitative data

Type of Plots

Scatterplot:

Dot plot:

Histogram:

Density plot:

Box plot:

Vocabulary

Four characteristics of a scatterplot:

Form:

Strength:

Direction:

Unusual observations or outliers:

Distribution (of a variable):

Four characteristics of the distribution of one quantitative variable:

Center:

Variability:

Shape:

Outliers:

Point estimate:

Histogram:

Data density:

Tail:

Skew:

Symmetric:

Modality:

Density plot:

Deviation:

Variance:

Standard deviation:

Boxplot:

Five number summary:

Median:

X^{th} percentile:

Interquartile range (IQR):

Robust statistics:

Notes

What type of plot(s) are appropriate for displaying one quantitative variable?

What type of plot(s) are appropriate for displaying two quantitative variables?

What type of plot(s) are appropriate for displaying one quantitative variable and one categorical variable?

What are the two ways to measure the ‘center’ of a distribution? Which one is considered robust to skew/outliers?

What are the three ways to measure the ‘variability’ of a distribution? Which one is considered robust to skew/outliers?

How are variance and standard deviation related?

Fill in the following table with the appropriate notation.

Summary Measure	Parameter	Statistic
Mean		
Variance		
Standard deviation		

How are outliers denoted on a box plot? How can you mathematically determine if a data set has outliers?

Summarizing Chapters 4 and 5

Look at the table of vocabulary terms in the final section of each chapter. If there are any you do not know, be sure to review the appropriate section of your text.

Notes

Statistics summarize _____.

Parameters summarize _____.

Fill in the following table with the appropriate notation for each summary measure.

Summary measure	Statistic	Parameter
Sample size		
Proportion (used to summarize one categorical variable)		
Mean (used to summarize one quantitative variable)		
Correlation (used to summarize two quantitative variables)		
Regression line slope (used to summarize two quantitative variables)		

Data visualization summary

Fill in the following table to help associate type of plot for each of several scenarios.

	Appropriate plot(s)
One categorical variable (categorical response, no explanatory)	
One quantitative variable (quantitative response, no explanatory)	
Two categorical variables (categorical response, categorical explanatory)	
One of each (quantitative response, categorical explanatory)	
Two quantitative variables (quantitative response, quantitative explanatory)	

3.2 Video Notes: Exploratory Data Analysis of Categorical and Quantitative Variables

Read Chapters 3, 4, and 5 in the course textbook. Use the following videos to complete the lecture notes for Module 3.

3.2.1 Course Videos

- 4.1
- 4.2
- 5.2to5.4
- 5.5to5.6
- 5.7

Summarizing categorical data - Video 4.1

- A _____ is calculated on data from a sample
- The parameter of interest is what we want to know from the population.
- Includes:
 - Population word (true, long-run, population)
 - Summary measure (depends on the type of data)
 - Context
 - * Observational units
 - * Variable(s)

Categorical data can be numerically summarized by calculating a _____ from the data set.

Notation used for the population proportion:

- Single categorical variable:
- Two categorical variables:
 - Subscripts represent the _____ variable groups

Notation used for the sample proportion:

- Single categorical variable:
- Two categorical variables

Categorical data can be reported in a _____ table, which plots counts or a _____ frequency table, which plots the proportion.

When we have two categorical variables we report the data in a _____ or two-way table with the _____ variable on the columns and the _____ variable on the rows.

Example from the Video: Gallatin Valley is the fastest growing county in Montana. You'll often hear Bozeman residents complaining about the 'out-of-staters' moving in. A local real estate agent recorded data on a random sample of 100 home sales over the last year at her company and noted where the buyers were moving from as well as the age of the person or average age of a couple buying a home. The variable age was binned into two categories, "Under30" and "Over30." Additionally, the variable, state the buyers were moving from, was created as a binary variable, "Out" for a location out of state and "In" for a location in state.

The following code reads in the data set, `moving_to_mt` and names the object `moving`.

```
moving <- read.csv("data/moving_to_mt.csv")
```

The R function `glimpse` was used to give the following output.

```
glimpse(moving)
```

```
#> Rows: 100
#> Columns: 4
#> $ From      <chr> "CA", "CA", "CA", "CA", "CA", "CA", "CA", "CA", "CA", "CA", ~
#> $ Age_Group <chr> "Under30", "Under30", "Under30", "Under30", "Under30", "Under~
#> $ Age       <int> 25, 26, 27, 27, 29, 29, 35, 37, 49, 63, 65, 77, 22, 24, 24, ~
#> $ InOut     <chr> "Out", "Out", "Out", "Out", "Out", "Out", "Out", "Out", "Out", "Out~
```

- What are the observational units in this study?

- What type of variable is `Age`?

- What type of variable is `Age_Group`?

To further analyze the categorical variable, `From`, we can create either a frequency table:

```
#>   From  n
#> 1   CA 12
#> 2   CO  8
#> 3   MT 61
#> 4   WA 19
```

Or a relative frequency table:

```
#>   From  n freq
#> 1   CA 12 0.12
#> 2   CO  8 0.08
#> 3   MT 61 0.61
#> 4   WA 19 0.19
```

- How many home sales have buyers from WA?
- What proportion of sampled home sales have buyers from WA?
- What notation is used for the proportion of home sale buyers that that are from WA?

To look at the relationship between the variable, `Age_Group` and the variable, `From` create the following two-way table using the R output below. Note, we are using `From` as the explanatory variable to predict whether a home sale has a buyer that is over or under the age of 30.

```
moving %>%
  group_by(Age_Group) %>% count(From) %>% print(n=8)
```

```
#> # A tibble: 8 x 3
#> # Groups:   Age_Group [2]
#>   Age_Group From     n
#>   <chr>      <chr> <int>
#> 1 Over30    CA         6
#> 2 Over30    CO         2
#> 3 Over30    MT        47
#> 4 Over30    WA        10
#> 5 Under30    CA         6
#> 6 Under30    CO         6
#> 7 Under30    MT        14
#> 8 Under30    WA         9
```

	State				
Age Group	CA	CO	MT	WA	Total
Over30	6	2	47	10	65
Under30	6	6	14	9	35
Total	12	8	61	19	100

- Using the table above, how many of the sampled home sales have buyers who were under 30 years old and from Montana?

If we want to know what proportion of each age group is from each state, we would calculate the proportion of home sales with buyers from each state within each age group. In other words, divide the number of home sales from each state with buyers that are over 30 by the total for row 1, the total number of home sales with buyers over 30.

- What proportion of sampled home sales with buyers under 30-years-old were from California?
- What notation should be used for this value?

Additionally, we could find the proportion of home sales with buyers in each state for each age group. Here we would calculate the proportion of home sales with buyers in each age group within each state. Divide the number of home sales with buyers in each age group from CA by the total for column 1, the total number of home sales with buyers from CA.

	State				
Age Group	CA	CO	MT	WA	Total
Over30	6	2	47	10	65
Under30	6	6	14	9	35
Total	12	8	61	19	100


- Using the table, calculate the proportion of home sales in Gallatin County with in-state buyers who are over 30 years old? Use appropriate notation with informative subscripts.
- Using the table, calculate the proportion of home sales in Gallatin County with California buyers who are over 30 years old? Use appropriate notation with informative subscripts.
- Calculate the difference in proportion of home sales in Gallatin County over 30 years old from other parts of Montana and from California. Use MT - CA as the order of subtraction. Give appropriate notation.
- Interpret the difference in proportion in context of the study.

Displaying categorical variables - Video 4.2

- Types of plots for a single categorical variable
- Types of plots for two categorical variables

The following code in R will create a frequency bar plot of the variable, From.


```
moving %>%
  ggplot(aes(x = From))+ #Enter the variable to plot
  geom_bar(stat = "count") +
  labs(title = "Frequency Bar Plot of State of Origin for
    Gallatin County Home Sales",
    #Title your plot (type of plot, observational units, variable)
    y = "Frequency", #y-axis label
    x = "State of Origin") #x-axis label
```



- What can we see from this plot?

Additionally, we can create a relative frequency bar plot.


```
moving %>%
  ggplot(aes(x = From)) + #Enter the variable to plot
  geom_bar(aes(y = after_stat(prop), group = 1)) +
  labs(title = "Relative Frequency Bar Plot of State of Origin
    for Gallatin County Home Sales",
    #Title your plot
    y = "Relative Frequency", #y-axis label
    x = "State of Origin") #x-axis label
```



- Note: the x-axis is the _____ between the frequency bar plot and the relative frequency bar plot. However, the _____ differs. The scale for the frequency bar plot goes from _____ and the scale for the relative frequency bar plot is from _____.


In a segmented bar plot, the bar for each category will sum to 1. In this first plot, we are plotting the row proportions calculated conditional on the age group.

```
moving %>%
  ggplot(aes(x = Age_Group, fill = From))+ #Enter the variables to plot
  geom_bar(stat = "count", position = "fill") +
  labs(title = "Segmented bar plot of Age Group of Buyers by State of
    Origin for Gallatin County Home Sales",
    #Title your plot
    y = "Relative Frequency", #y-axis label
    x = "Age Group") + #x-axis label
  scale_fill_grey()
```




In this second plot, we are plotting the column proportions calculated conditional on the state of origin for the buyer.

```
moving %>%
  ggplot(aes(x = From , fill = Age_Group))+ #Enter variables to plot
  geom_bar(stat = "count", position = "fill") +
  labs(title = "Segmented bar plot of State of Origin of Buyers by Age
    Group for Gallatin County Home Sales",
    #Title your plot
    y = "Relative Frequency", #y-axis label
    x = "State of Origin") + #x-axis label
  scale_fill_grey()
```



Mosaic plot:

```
moving$Age_Group <- factor(moving$Age_Group, levels = c("Under30", "Over30"))
moving %>% # Data set piped into...
  ggplot() + # This specifies the variables
  geom_mosaic(aes(x=product(From), fill = Age_Group)) +
  # Tell it to make a mosaic plot
  labs(title = "Mosaic plot of State of Origin Segmented by
Age Group for Gallatin County Home Sales",
  # Title your plot
  x = "State of Origin", # Label the x axis
  y = "") + # Remove y axis label
  scale_fill_grey(guide = guide_legend(reverse = TRUE)) # Make figure color
```



- Why is the bar for MT the widest on the mosaic plot?


Simpson's paradox - Video 4.4

- When an apparent _____ between explanatory and response variables reverses when accounting for _____ variable.

Example: The “Berkeley Dataset” contains all 12,763 applicants to UC-Berkeley’s graduate programs in Fall 1973. This dataset was published by UC Berkeley researchers in an analysis to understand the possible gender bias in admissions and has now become a classic example of Simpson’s Paradox.

```
discrim <- read.csv ("data/berkeley.csv")
```


```
discrim %>%  
  ggplot(aes(x = Gender, fill = Admission)) +  
  geom_bar(stat = "count", position = "fill") +  
  labs(title = "Segmented bar plot of Sex of Berkley Applicants by  
    Admission Status",  
    y = "Relative Frequency",  
    x = "Sex") +  
  scale_fill_grey()
```




The data showed that 44% of male applicants were accepted and 35% of female applicants were accepted. Does it appear that the female students are discriminated against?

We can break down the data by major. A major code (either A, B, C, D, E, F, or Other) was used.


Here we look at the relationship between admission status and sex for Program A and for Program B.



Showing Program C and Program D.



And finally, Program E and F.



We can see in several programs the acceptance rate is higher for females than for males.


Summarizing quantitative data - Videos 5.2to5.4 and 5.5to5.6

Types of plots

We will revisit the moving to Montana data set and plot the age of the buyers.


Dotplot:

```
moving %>%  
  ggplot(aes(x = Age)) + #Enter variable to plot  
  geom_dotplot() +  
  labs(title = "Dotplot of Age of Buyers from Gallatin  
    County Home Sales", #Title your plot  
    x = "Age", #x-axis label  
    y = "Proportion") #y-axis label
```



Histogram:

```
moving %>%  
  ggplot(aes(x = Age))+  
  geom_histogram(binwidth = 7) +  
  labs(title = "Histogram of Age of Buyers from Gallatin  
    County Home Sales",  
        #Title your plot  
        x = "Age",  
        y = "Count")
```



Quantitative data can be numerically summarized by finding:

Two measures of center:

- Mean: _____ of all the _____ in the data set.
 - Sum the values in the data set and divide the sum by the sample size
- Notation used for the population mean:
 - Single quantitative variable:
 - One categorical and one quantitative variable:
 - Subscripts represent the _____ variable groups
- Notation used for the sample mean:
 - Single quantitative variable:
 - One categorical and one quantitative variable:

- Median: Value at the _____ percentile
 - _____ % of values are at and _____ and at and _____ the value of the _____.
 - Middle value in a list of ordered values

Two measures of spread:

- Standard deviation: Average _____ each data point is from the _____ of the data set.
 - Notation used for the population standard deviation
 - Notation used for the sample standard deviation

- Interquartile range: middle 50% of data values

Formula:

Quartile 3 (Q3) - value at the 75th percentile

- _____ % of values are at and _____ the value of Q3


Quartile 1 (Q1) - value at the 25th percentile

- _____ % of values are at and _____ the value of Q1

Boxplot (3rd type of plot for quantitative variables)

- Five number summary: minimum, Q1, median, Q3, maximum

```
moving %>%  
  ggplot(aes(x = Age))+ #Enter variable to plot  
  geom_boxplot() +  
  labs(title = "Boxplot of Age of Buyers from Gallatin  
    County Home Sales", #Title your plot  
        x = "Age", #x-axis label  
        y = "") #y-axis label
```



```
favstats(moving$Age)
```


```
#>   min  Q1 median    Q3 max  mean    sd   n missing  
#>   20  28   36 49.25  77 39.77 14.35471 100      0
```

Interpret the value of Q_3 for the age of buyers.

Interpret the value of s for the age of buyers.

Four characteristics of plots for quantitative variables

- Shape: overall pattern of the data



- What is the shape of the distribution of age of buyers for Gallatin County home sales?

- Center:

Mean or Median

- Report the measure of center for the boxplot of age of buyers for Gallatin County home sales.

- Spread (or variability):

Standard deviation or IQR

- Report the IQR for the distribution of age of buyers from Gallatin County home sales.

- Outliers?


values $< Q_1 - 1.5 \times IQR$

values $> Q_3 + 1.5 \times IQR$

- Use these formulas to show that there are no outliers in the distribution of age of buyers from Gallatin County home sales.

Let's look at side-by-side boxplot of the variable age by state of origin moved from.


```
moving %>% # Data set piped into...
  ggplot(aes(y = Age, x = From)) + # Identify variables
  geom_boxplot() + # Tell it to make a box plot
  labs(title = "Side by side box plot of Age by State of Origin
of Buyers from Gallatin County Home Sales", # Title
       x = "State of Origin", # x-axis label
       y = "Age") # y-axis label
```



- Which state of origin had the oldest median age of buyers from sampled home sales?
- Which state of origin had the most variability in age of buyers from sampled home sales?
- Which state of origin had the most symmetric distribution of ages of buyers from sampled home sales?
- Which state of origin had outliers for the age of buyers from sampled home sales?

Robust statistics - Video 5.7

Let's review the summary statistics and histogram of age of buyers from sampled home sales.



```
#>   min  Q1 median    Q3  max  mean    sd   n missing  
#>   20  28   36 49.25  77 39.77 14.35471 100      0
```

Notice that the _____ has been pulled in the direction of the _____.

- The _____ is a robust measure of center.
- The _____ is a robust measure of spread.
- Robust means not _____ by outliers.

When the distribution is symmetric use the _____ as the measure of center and the _____ as the measure of spread.

When the distribution is skewed with outliers use the _____ as the measure of center and the _____ as the measure of spread.

3.2.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. What summary statistic is calculated for a two categorical variables?
2. What four characteristics should be used to compare side-by-side boxplots?
3. Which measure of center is robust to outliers?

3.3 Guided Activity 4: Summarizing Categorical Data

3.3.1 Learning outcomes

- Identify and create appropriate summary statistics and plots given a data set or research question involving categorical variables.

3.3.2 Terminology review

In today's activity, we will review summary measures and plots for categorical variables. Some terms covered in this activity are:

- Categorical variable summary measure: proportion
- Plots of categorical variables: bar plots, segmented bar plots, mosaic plots

To review these concepts, see Chapter 4 in the textbook.

3.3.3 The Integrated Postsecondary Education Data System (IPEDS)

These data were collected on a subset of institutions that met the following selection criteria (Education Statistics 2018):

- Degree granting
- United States only
- Title IV participating
- Not for profit
- 2-year or 4-year or above
- Has full-time first-time undergraduates

Some of the variables collected and their descriptions are below. Note that several variables have missing values for some institutions (denoted by “NA”).

Variable	Description
UnitID	Unique institution identifier
Name	Institution name
State	State abbreviation
Sector	whether public or private
LandGrant	Is this a land-grant institution (Yes/No)
Size	Institution size category based on total student enrolled for credit, Fall 2018: Under 1,000, 1,000 - 4,999, 5,000 - 9,999, 10,000 - 19,999, 20,000 and above
Cost_OutofState	Cost of attendance for full-time out-of-state undergraduate students
Cost_InState	Cost of attendance for full-time in-state undergraduate students
Retention	Retention rate is the percent of the undergraduate students that re-enroll in the next year
Graduation_Rate	6-year graduation rate for undergraduate students
SATMath_75	75th percentile Math SAT score
ACT_75	75th percentile ACT score

Identifying Variables in a data set

Look through the provided chart showing the description of variables measured. The UnitID and Name are identifiers for each observational unit, *US degree granting institutions in 2018*.

1. Identify in the chart which variables collected on the US institutions are categorical (C) and which variables are quantitative (Q).

R code

For almost all activities and labs it will be necessary to upload the provided R script file from D2L for that day. Your instructor will highlight a few steps in uploading files to and using RStudio. We encourage you to complete the R script file posted on D2L prior to the next class.

The following are the steps to upload the necessary R script file for this activity:

- Download the IPEDS Activity R script file from D2L.
- Click “Upload” in the “Files” tab in the bottom right window of RStudio. In the pop-up window, click “Choose File”, and navigate to the folder where the IPEDS Activity R script file is saved (most likely in your downloads folder). Click “Open”; then click “Ok”.
- You should see the uploaded file appear in the list of files in the bottom right window. Click on the file name to open the file in the Editor window (upper left window).

Notice that the first four lines of code contain a prompt called, **library**. Packages needed to run functions in R are stored in directories called libraries. When using the MSU RStudio server, all the packages needed for the class are already installed. We simply must tell R which packages we need for each R script file. We use the prompt **library** to load each **package** (or library) needed for each activity. Note, these **library** lines MUST be run each time you open a R script file in order for the functions in R to work.

- Highlight and run lines 1–4 to load the packages needed for this activity. Notice the use of the # symbol in the R script file. This symbol is not part of the R code. It is used by these authors to add comments to the R code and explain what each call is telling the program to do. R will ignore everything after a # symbol when executing the code. Refer to the instructions following the # symbol to understand what you need to enter in the code.

```
library(tidyverse)
library(ggplot2)
library(mosaic)
library(ggmosaic)
```

Throughout activities, we will often include the R code you would use in order to produce output or plots. These “code chunks” appear in gray. In the code chunk below, we demonstrate how to read the data set into R using the `read.csv()` function. The line of code shown below (line 7 in the R script file) reads in the data set and names the data set `IPEDS`.

- Highlight and run line 7 in the R script file to load the data from the Stat 216 webpage.

```
# This will read in the data set
IPEDS <- read.csv("https://www.math.montana.edu/courses/s216/data/IPEDS_2018.csv")
```

- Click on the data set name (`IPEDS`) in the Environment tab (upper right window). This will open the data set in a 2nd tab in the Editor window (upper left window). R is case sensitive, which means that you must always type the name of a variable EXACTLY as it is written in the data set including upper and lower case letters and without misspellings!

2. Write down the number of observational units and the number of variables collected in the data set.

Since we are only interested in assessing 4-year institutions we will filter out the 2-year institutions from the data set.

- Highlight and run lines 9–12 to filter out the 2-year institutions.

```
IPEDS <- IPEDS %>%
  filter(Sector != "Public 2-year") #Filters the data set to remove Public 2-year
IPEDS <- IPEDS %>%
  filter(Sector != "Private 2-year") #Filters the data set to remove Private 2-year
```

3. Write down the number of 4-year institutions in the data set.

Summarizing categorical variables

An important part of understanding data is to create visual pictures of what the data represent. In this activity, we will create graphical representations of categorical data.

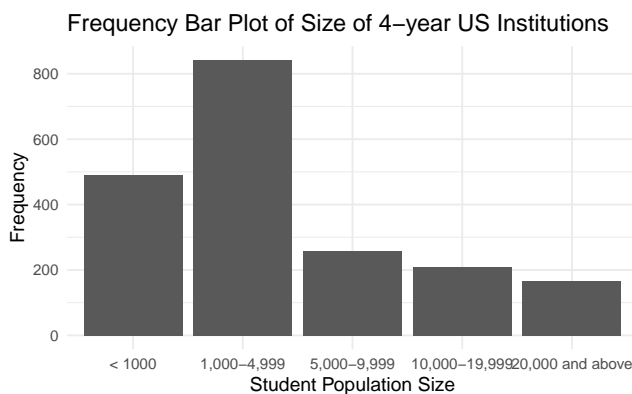
There are four categorical variables in this data set. **Sector**, **Size**, **LandGrant**, and **Tenure**.

Displaying a single categorical variable

If we wanted to know how many 4-year institutions in our data set were in each category of size, we could create a frequency bar plot of the variable **Size**.

- In the R code below (and the provided R script file), we will enter the variable name, **Size** (*note the capital S*), for **variable** into the **ggplot** code. This is in line 18 in the R script file.
- Highlight and run lines 16–22 to create the plot. Note: this is a **frequency** bar plot plotting counts (the number of institutions within each level of **Size** is displayed on the *y*-axis).

```
IPEDS$Size <- factor(IPEDS$Size, levels = c("< 1000", "1,000-4,999",
                                           "5,000-9,999", "10,000-19,999", "20,000 and above"))
IPEDS %>% # Data set piped into...
ggplot(aes(x = Size)) + # This specifies the variable
  geom_bar(stat = "count") + # Tell it to make a bar plot
  labs(title = "Frequency Bar Plot of Size of 4-year US Institutions",
       # Give your plot a title
       x = "Student Population Size", # Label the x axis
       y = "Frequency") # Label the y axis
```

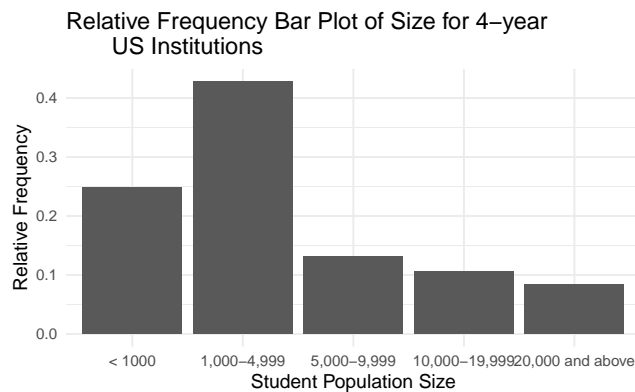


4. Using the bar chart created, estimate how many institutions have more than 10,000 students.

We could also choose to display the data as a proportion in a **relative frequency** bar plot. To find the relative frequency, the count in each level of **Size** is divided by the sample size. This calculation is the sample proportion for each level of size. Notice that in this code we told R to create a bar plot with proportions.

- In the R script file, enter the variable name **Size** for **variable** in line 27.
- Highlight and run lines 26 - 31 to create the relative frequency bar plot.

```
IPEDS %>% # Data set piped into...
ggplot(aes(x = Size)) + # This specifies the variable
  geom_bar(aes(y = after_stat(prop), group = 1)) + # Tell it to make a bar plot with proportions
  labs(title = "Relative Frequency Bar Plot of Size for 4-year
    US Institutions",
    # Give your plot a title
    x = "Student Population Size", # Label the x axis
    y = "Relative Frequency") # Label the y axis
```



5. Which features in the relative frequency bar plot are the same as the frequency bar plot? Which are different?

In R, to get the counts for each level of a variable use the function **count**.

- In the R script file, enter the variable name **Size** for **variable** in line 35.
- Highlight and run line 35.

```
IPEDS %>% count(Size)
```

```
#>      Size      n
#> 1    < 1000  490
#> 2 1,000-4,999 842
#> 3 5,000-9,999 259
#> 4 10,000-19,999 210
#> 5 20,000 and above 166
```

6. Calculate the proportion of 4-year Institutions that have a student population size of between 1,000-4,999? Use appropriate notation.

Displaying two categorical variables

Is there an association between the student population size of an institution and whether it is public or private institution? In this instance we will treat the variable **Sector** as the explanatory variable and **Size** as the response variable.

- Fill in the name of the explanatory variable, **Sector** for explanatory and name of the response variable, **Size** in line 39 in the R script file.
- Highlight and run line 39 to get the counts for each combination of levels of variables.

```
IPEDS %>% group_by(Sector) %>% count(Size) %>% print(n=10)
```

```
#> # A tibble: 10 x 3
#> # Groups:   Sector [2]
#>   Sector                Size                n
#>   <chr>                <fct>                <int>
#> 1 Private 4-year or above < 1000                457
#> 2 Private 4-year or above 1,000-4,999            639
#> 3 Private 4-year or above 5,000-9,999             95
#> 4 Private 4-year or above 10,000-19,999            54
#> 5 Private 4-year or above 20,000 and above         21
#> 6 Public 4-year or above < 1000                 33
#> 7 Public 4-year or above 1,000-4,999            203
#> 8 Public 4-year or above 5,000-9,999            164
#> 9 Public 4-year or above 10,000-19,999           156
#> 10 Public 4-year or above 20,000 and above       145
```

Below is a contingency table of the data.

	Sector		
Size	Public	Private	Total
<1000	33	457	490
1000-4999	203	639	842
5000-9999	164	95	259
10000-19999	156	54	210
20000 and above	145	21	166
Total	701	1266	1967

In the following questions, use the table to calculate the described proportions. Notation is important for each calculation. Since this is sample data, it is appropriate to use statistic notation for the proportion, \hat{p} . When calculating a proportion dependent on a single level of a variable, subscripts are needed when reporting the notation.

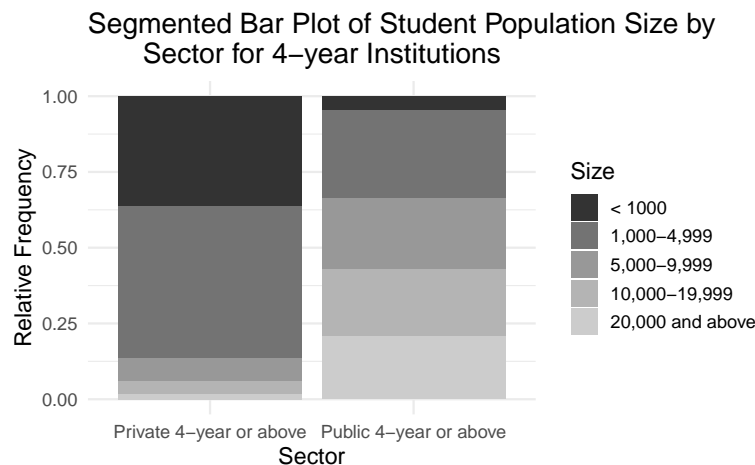
7. Calculate the proportion of institutions that are public. Use appropriate notation.
8. Calculate the proportion of public institutions that have 1000 - 4999 students. Use appropriate notation.

9. Calculate the proportion of private institutions that have 1000 - 4999 students. Use appropriate notation.
10. Calculate the difference in proportion of institutions that have 1000 - 4999 students for those that are public minus those that are private. Give the appropriate notation. Use public minus private as the order of subtraction.
11. Interpret the difference in proportion calculated in question 9 in context of the study.

Two types of plots can be created to display two categorical variables. To examine the association between the student population size of an institution and whether it is public or private institution, we will first create a segmented bar plot of **Sector** segmented by **Size**.

- To create the segmented bar plot enter the variable name, **Sector** for **explanatory** and the variable name, **Size** for **response** in line 44 in the R script file.
- Highlight and run lines 43–50.

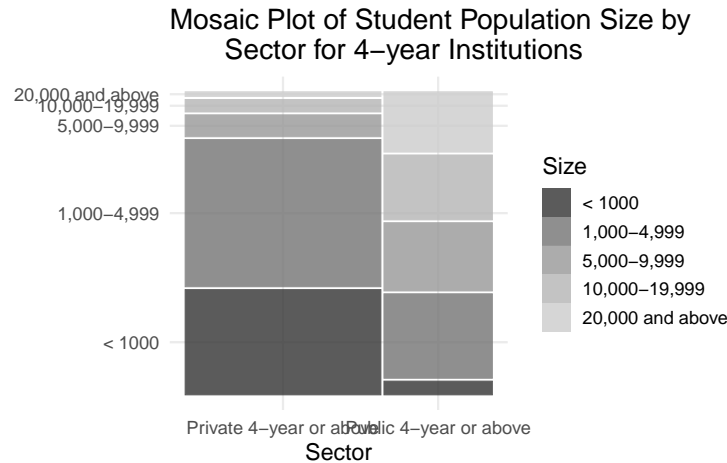
```
IPEDS %>% # Data set piped into...
ggplot(aes(x = Sector, fill = Size)) + # This specifies the variables
geom_bar(stat = "count", position = "fill") + # Tell it to make a stacked bar plot
labs(title = "Segmented Bar Plot of Student Population Size by
Sector for 4-year Institutions", # Make sure to title your plot
x = "Sector", # x axis label
y = "Relative Frequency") + # y axis label
scale_fill_grey() # Make figure color
```



12. From the segmented bar plot, which sector has the highest proportion of student population size of 1000 - 4999?

We could also create a mosaic plot of the data, shown below.

```
IPEDS %>% #Data set piped into...
  ggplot() + #This specifies the variables
  geom_mosaic(aes(x=product(Sector), fill = Size)) + #Tell it to make a mosaic plot
  labs(title = "Mosaic Plot of Student Population Size by
    Sector for 4-year Institutions", #Make sure to title your plot
    x = "Sector", #Label the x axis
    y = "") + #Remove y axis label
  scale_fill_grey() # Make figure bw
```



13. What is similar and what is different between the segmented bar chart and the mosaic bar chart?

3.3.4 Take-home messages

1. Bar charts can be used to graphically display a single categorical variable either as counts or proportions. Segmented bar charts and mosaic plots are used to display two categorical variables.
2. Segmented bar charts always have a scale from 0 – 100%. The bars represent the outcomes of the explanatory variable. Each bar is segmented by the response variable. If the heights of each segment are the same for each bar there is no association between variables.
3. Mosaic plots are similar to segmented bar charts but the widths of the bars also show the number of observations within each outcome. This allows assessment of the relative sizes of the levels of one variable in assessing changes in relative distributions of the levels of the other variable and the proportions of the totals in each combination of levels.

3.3.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

3.4 Guided Activity 5: Summarizing Quantitative Variables

3.4.1 Learning outcomes

- Identify and create appropriate summary statistics and plots given a data set or research question for quantitative data.
- Interpret the following summary statistics in context: median, lower quartile, upper quartile, standard deviation, interquartile range.

3.4.2 Terminology review

In today's activity, we will review summary measures and plots for quantitative variables. Some terms covered in this activity are:

- Two measures of center: mean, median
- Two measures of spread (variability): standard deviation, interquartile range (IQR)
- Plots of quantitative variables: dotplots, boxplots, histograms
- Given a plot or set of plots, describe and compare the distribution(s) of quantitative variables (center, variability, shape, outliers).

To review these concepts, see Chapter 5 in the textbook.

3.4.3 IPEDS continued

We will revisit the IPEDS data set now analyzing quantitative variables. These data were collected on a subset of institutions that met the following selection criteria (Education Statistics 2018):

- Degree granting
- United States only
- Title IV participating
- Not for profit
- 2-year or 4-year or above
- Has full-time first-time undergraduates

Some of the variables collected and their descriptions are below. Note that several variables have missing values for some institutions (denoted by "NA").

Variable	Description
UnitID	Unique institution identifier
Name	Institution name
State	State abbreviation
Sector	whether public or private
LandGrant	Is this a land-grant institution (Yes/No)
Size	Institution size category based on total student enrolled for credit, Fall 2018: Under 1,000, 1,000 - 4,999, 5,000 - 9,999, 10,000 - 19,999, 20,000 and above
Cost_OutofState	Cost of attendance for full-time out-of-state undergraduate students
Cost_InState	Cost of attendance for full-time in-state undergraduate students
Retention	Retention rate is the percent of the undergraduate students that re-enroll in the next year
Graduation_Rate	6-year graduation rate for undergraduate students

Variable	Description
SATMath_75	75th percentile Math SAT score
ACT_75	75th percentile ACT score

Summarizing quantitative variables

The `favstats()` function from the `mosaic` package gives the summary statistics for a quantitative variable. The R output below provides the summary statistics for the variable `Graduation_Rate`. The summary statistics provided are the two measures of center (mean and median) and two measures of spread (standard deviation and the quartile values to calculate the IQR) for IMDb score.

- Highlight and run lines 1 – 12 in the provided R script file to load the data set. Check that the summary statistics match the output given in the coursepack.

```
IPEDS <- read.csv("https://www.math.montana.edu/courses/s216/data/IPEDS_2018.csv")
IPEDS <- IPEDS %>%
  filter(Sector != "Public 2-year") #Filters the data set to remove Public 2-year
IPEDS <- IPEDS %>%
  filter(Sector != "Private 2-year") #Filters the data set to remove Private 2-year
IPEDS %>%
  summarize(favstats(Graduation_Rate))
```

```
#>   min Q1 median Q3 max      mean      sd    n missing
#> 1    0  38     53 67 100 52.48749 20.63192 1918        49
```

1. Calculate the interquartile range ($IQR = Q3 - Q1$) of Graduation Rates.
2. Report the value of the standard deviation and interpret this value in context of the problem.

Displaying a single quantitative variable

There are three type of plots used to plot a single quantitative variable: a dotplot, a histogram or a boxplot. A dotplot of graduation rate would plot a dot for the graduation rate for each 4-year US higher education institution.

We will create both a histogram and a boxplot of the variable `Graduation_Rate`.


- Enter the name of the variable in both line 19 and line 26 for `variable` in the R script file.
- Replace the word title for each plot (lines 21 and 28) between the quotations with a descriptive title. **A title should include: type of plot, variable or variables plotted, and observational units.**
- Highlight and run lines 18 – 30 to create a histogram and boxplot.

Notice that the **bin width** for the histogram is 10. For example the first bin consists of the number of institutions in the data set with a graduation rate of 0 to 10%. It is important to note that an graduation rate on the boundary of a bin will fall into the bin above it; for example, 20 would be counted in the bin 20–30.

```

IPEDS %>% # Data set piped into...
ggplot(aes(x = Graduation_Rate)) + # Name variable to plot
  geom_histogram(binwidth = 10) + # Create histogram with specified binwidth
  labs(title = "", # Title for plot
        x = "Graduation Rate", # Label for x axis
        y = "Frequency") # Label for y axis

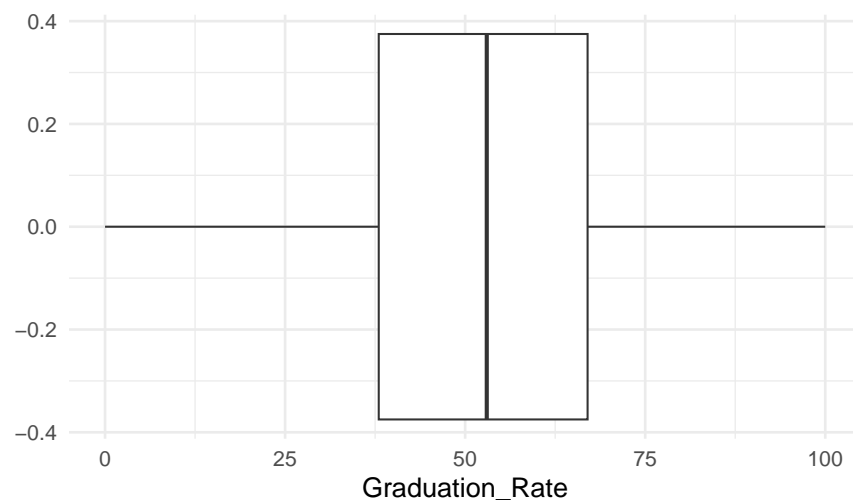
```



```

IPEDS %>% # Data set piped into...
ggplot(aes(x = Graduation_Rate)) + # Name variable to plot
  geom_boxplot() + # Create histogram with specified binwidth
  labs(title = "", # Title for plot
        x = "Graduation_Rate", # Label for x axis
        y = "") # Remove y axis label

```



Summarizing one categorical and one quantitative variable

Is there a difference in retention rates for public and private 4-year institutions? In this part of the activity we will compare retention rates for public and private 4-year institutions. Note that this variable (public or private) is labelled **Sector** in the data set.

3. Based on the research question, which variable will we treat as the explanatory variable? Response variable?

The `favstats()` function from the `mosaic` package gives the summary statistics for a quantitative variable. The summary statistics give the two measures of center and two measures of spread for retention rates between sector for 4-year institutions.

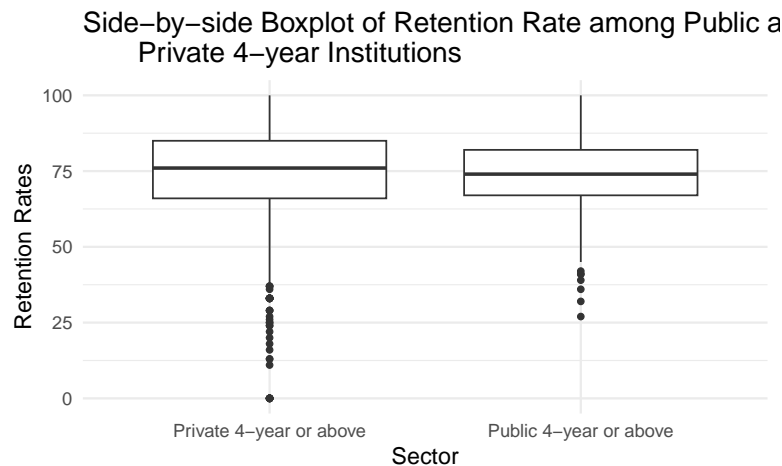
- To assess the research question described before question 3, enter the name of the explanatory variable and the name of the response variable in line 35 of the R script file. Remember that the variable name must be typed in EXACTLY as it is written in the data set.
- Highlight and run lines 34–35

```
IPEDS %>%  
  reframe(favstats(Retention~Sector)) #Gives the summary statistics
```

```
#>           Sector min Q1 median Q3 max      mean      sd    n missing  
#> 1 Private 4-year or above    0 66    76 85 100 73.67244 16.72878 1212      54  
#> 2 Public 4-year or above   27 67    74 82 100 74.21322 11.32172  605      96
```

- To create a side-by-side boxplot, enter the name of the explanatory variable and the name of the response variable in line 41 of the R script file.
- Highlight and run lines 39 – 45.

```
IPEDS %>% # Data set piped into...  
  ggplot(aes(y = Retention, x = Sector))+ # Identify variables  
  geom_boxplot()+ # Create box plot  
  labs(title = "Side-by-side Boxplot of Retention Rate among Public and  
    Private 4-year Institutions",  
    # A title should include: type of plot, variable or variables plotted, and observational units  
    x = "Sector", # x-axis label  
    y = "Retention Rates") # y-axis label
```



4. Compare the two boxplots.

Which sector of university has the highest center?

Largest spread?

What is the shape of each distribution?

Does either distribution have potential outliers?

5. Report the difference in mean retention rates for private and public universities. Use private minus public as the order of subtraction. Use the appropriate notation.
6. Interpret the difference in mean retention rates found in question 5.
7. Does there appear to be an association between retention rates and type of university? Explain your answer using the boxplots and summary statistic.

3.4.4 Take-home messages

1. Histograms, box plots, and dot plots can all be used to graphically display a single quantitative variable.
2. The box plot is created using the five number summary: minimum value, quartile 1, median, quartile 3, and maximum value. Values in the data set that are less than $Q_1 - 1.5 \times \text{IQR}$ and greater than $Q_3 + 1.5 \times \text{IQR}$ are considered outliers and are graphically represented by a dot outside of the whiskers on the box plot.
3. Data should be summarized numerically and displayed graphically to give us information about the study.
4. When comparing distributions of quantitative variables we look at the shape, center, spread, and for outliers. There are two measures of center: mean and the median and two measures of spread: standard deviation and the interquartile range, $\text{IQR} = Q_3 - Q_1$.

3.4.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

3.5 Module 3 Lab: EDA Review

3.5.1 Learning outcomes

- Identify and create appropriate summary statistics and plots given a data set or research question for different types of data.
- Interpret the following summary statistics in context: median, lower quartile, upper quartile, standard deviation, interquartile range.

3.5.2 General information labs

Throughout the semester, labs will be completed at the end of specific modules. Questions are selected from each lab to be turned in on Gradescope. The questions to be submitted on Gradescope are bolded in the lab. As you work through the lab have the Gradescope lab assignment open so that you can answer those questions as you go.

3.5.3 Movies released in 2016

A data set was collected on movies released in 2016 (“IMDb Movies Extensive Dataset” 2016). Here is a list of some of the variables collected on the observational units, movies released in 2016.

Variable	Description
budget_mil	Amount of money (in US \$ millions) budgeted for the production of the movie
revenue_mil	Amount of money (in US \$ millions) the movie made after release
duration	Length of the movie (in minutes)
content_rating	Rating of the movie (G, PG, PG-13, R, Not Rated)
imdb_score	IMDb user rating score from 1 to 10
genres	Categories the movie falls into (e.g., Action, Drama, etc.)
facebook_likes	Number of likes a movie receives on Facebook

Robust Statistics


Let's examine how the presence of outliers affect the values of center and spread. For this part of the activity we will look at the variable revenue in the movies data set.

```
movies <- read.csv("https://www.math.montana.edu/courses/s216/data/Movies2016_F24.csv")
```

```
movies %>% # Data set piped into...  
  summarise(favstats(revenue_mil))
```

```
#>   min      Q1  median      Q3     max    mean      sd  n missing  
#> 1    0 1.467318 33.02703 73.13457 407.1973 61.87334 89.57824 92      0
```

```
movies %>% # Data set piped into...  
  ggplot(aes(x = revenue_mil)) + # Name variable to plot  
  geom_boxplot() + # Create histogram with specified binwidth  
  labs(title = "Histogram of Revenue of Movies in 2016", # Title for plot  
        x = "Revenue (in Millions)", # Label for x axis  
        y = "Frequency") # Label for y axis
```



1. Report the two measures of center for these data.
2. Report the two measures of spread for these data.

To show the effect of outliers on the measures of center and spread, the largest values in the data set were reduced by 100 \$MM to decrease the range of the data set. This variable is called `revenue_reduced`


```
movies %>% # Data set piped into...  
  summarise(favstats(revenue_reduced))
```

```
#>   min      Q1  median      Q3     max    mean      sd  n missing  
#> 1    0 1.467318 33.02703 73.13457 307.1973 55.3516 68.79572 92      0
```

```

movies %>% # Data set piped into...
ggplot(aes(x = revenue_reduced)) + # Name variable to plot
geom_boxplot() + # Create histogram with specified binwidth
labs(title = "Histogram of Decreased Revenue of Movies in 2016", # Title for plot
x = "Revenue (in Millions)", # Label for x axis
y = "Frequency") # Label for y axis

```



3. Report the two measures of center for this new data set.
4. Report the two measures of spread for this new data set.
5. Which measure of center is robust to outliers? Explain your answer.
6. Which measure of spread is robust to outliers? Explain your answer.

Summary statistics for a single categorical and single quantitative variable

Is there an association between content rating and budget for movies released in 2016? To use the `favstats()` function in the `mosaic` package with two variables, we will enter the variables as a formula, `response~explanatory`. This function will give the summary statistics for budget for each content rating.

- Enter the variable name `content_rating` for explanatory and `budget` for response in line 33
- Highlight and run lines 31–33 in the provided R script file.

```
movies %>% # Data set piped into...  
  filter(content_rating != "Not Rated") %>% # Remove Not Rated movies  
  reframe(favstats(response~explanatory)) # Find the summary measures for each content rating
```

7. Which content rating has the largest IQR?
8. Report the mean budget amount for the PG rating. Use appropriate notation.
9. Report the mean budget amount for the R rating. Use appropriate notation.
10. Calculate the difference in mean budget amount for movies released in 2016 with a PG rating minus those with a R rating. Use appropriate notation with informative subscripts.
11. Interpret the difference in means calculated in question 10 in context of the problem.
12. Report and interpret the standard deviation of the budget amount for movies released in 2016 with a PG rating.
13. Calculate the $IQR = Q_3 - Q_1$ of the budget amount for movies released in 2016 with a PG rating.

Displaying a single categorical and single quantitative variable

The boxplot of movie budgets (in millions) by content rating is plotted using the code below.

- Enter the variable `budget_mil` for `response` and the variable `content_rating` for explanatory at line 30.
- Highlight and run code lines 37–44. This plot compares the budget for different levels of content rating.

```
movies %>% # Data set piped into...
  filter(content_rating != "Not Rated") %>% # Remove Not Rated movies
  ggplot(aes(y = response, x = explanatory))+ # Identify variables
  geom_boxplot()+ # Tell it to make a box plot
  labs(title = "Side-by-side Box Plot of Budget by Content Rating for Movies Released in 2016",
        # Title
        x = "Content Rating", # x-axis label
        y = "Budget (in Millions)") # y-axis label
```

14. Answer the following questions about the box plots created.

- Which content rating has the highest center?
- Which content rating has the largest spread?
- Which content rating has the most skewed distribution?
- Fifty percent of movies released in 2016 with a PG-13 content rating fall below what value? What is the name of this value?
- What is the value for the first quartile (Q1) for the PG-13 rating? Interpret this value in context.

Nightlight use and myopia

In the previous part of the activity, the code was provided to import the data set needed directly from the Stat 216 website. Follow these steps to upload and import the data set for the rest of the lab.

- Download the provided data set `ChildrenLightSight.csv` from D2L
- Upload the data set `ChildrenLightSight` to the RStudio server using the same steps to upload the R script file.
- Enter the name of the data set for `datasetname` in the R script file in line 48. Don't forget to include `.csv` (`ChildrenLightSight.csv`).
- Highlight and run line 48 in the provided R script file to load the data set.

```
# Read in data set
myopia <- read.csv("datasetname")
```

In a study reported in Nature (Quinn et al. 1999), a survey of 479 children found that those who had slept with a nightlight or in a fully lit room before the age of two had a higher incidence of nearsightedness (myopia) later in childhood.

In this study, there are two variables studied: **Light**: level of light in room at night (no light, nightlight, full light) and **Sight**: level of myopia developed later in childhood (high myopia, myopia, no myopia).

Displaying two categorical variables

Is there an association between the level of light in a room and the development of myopia?

- Fill in the name of the explanatory variable, **Light** for explanatory and name of the response variable, **Sight** in line 52 in the R script file.
- Highlight and run line 52 to get the counts for each combination of levels of variables.

```
myopia %>% group_by(explanatory) %>% count(response)
```

15. Fill in the following table with the values from the R output.

	Light Level			
Myopia Level	Full Light	Nightlight	No Light	Total
High Myopia				
Myopia				
No Myopia				
Total				

In the following questions, use the table to calculate the described proportions. Notation is important for each calculation. Since this is sample data, it is appropriate to use statistic notation for the proportion, \hat{p} . When calculating a proportion dependent on a single level of a variable, subscripts are needed when reporting the notation.

16. Calculate the proportion of children with no myopia. Use appropriate notation.
17. Calculate the proportion of children with no myopia among those that slept with full light. Use appropriate notation.
18. Calculate the proportion of children with no myopia among those that slept with no light. Use appropriate notation.
19. Calculate the difference in proportion of children with no myopia for those that slept with full light minus those who slept with no light. Give the appropriate notation. Use full light minus no light as the order of subtraction.

To examine the differences in level of myopia for the level of light, we will create a segmented bar plot of **Light** segmented by **Sight**.

- To create the segmented bar plot enter the variable name, **Light** for **explanatory** and the variable name, **Sight** for **response** in the R script file in line 57.
- Highlight and run lines 56–63.

```
myopia %>% # Data set piped into...
ggplot(aes(x = explanatory, fill = response)) + # This specifies the variables
geom_bar(stat = "count", position = "fill") + # Tell it to make a stacked bar plot
labs(title = "Segmented Bar Plot of Night Light Use by Level of Myopia for Children",
      # Make sure to title your plot
      x = "Level of Light", # Label the x axis
      y = "") + # Remove y axis label
scale_fill_viridis_d() # Make figure color
```

20. From the segmented bar plot, which level of light has the highest proportion of No Myopia?

3.5.4 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

Exploring Multivariable Data

4.1 Reading Guide: Two Quantitative Variables and Multivariable Concepts

Textbook Chapter 6: Correlation and regression

Section 6.1 (Fitting a line, residuals, and correlation)

Reminders from Section 5.1

Scatterplot: displays two quantitative variables; one dot = two measurements (x , y) on one observational unit.

Four characteristics of a scatterplot:

- *Form*: pattern of the dots plotted. Is the trend generally linear (you can fit a straight line to the data) or non-linear?
- *Strength*: how closely do the points follow a trend? Very closely (strong)? No pattern (weak)?
- *Direction*: as the x values increase, do the y -values tend to increase (positive) or decrease (negative)?
- Unusual observations or *outliers*: points that do not fit the overall pattern of the data.

Vocabulary

Predictor:

Residual:

Formula:

Residual plot:

Correlation:

Notes

General equation of a linear model for a *population*: $y = \beta_0 + \beta_1 x + \epsilon$, where

x represents

y represents

β_0 represents

β_1 represents

ϵ represents

General equation of a linear regression model from *sample* data: $\hat{y} = b_0 + b_1x$, where

x represents

\hat{y} represents

b_0 represents

b_1 represents

Fill in the following table with the appropriate notation for each summary measure.

Summary Measure	Parameter	Statistic
Correlation		
Slope		
y -intercept		

Fill in the blanks below to define some of the properties of correlation:

The value of correlation must be between _____. (Includes the endpoints of the interval)

The sign of correlation gives the _____ of the linear relationship.

The magnitude of correlation gives the _____ of the linear relationship.

True or false: A scatterplot that shows random scatter would be considered non-linear.

True or false: If the correlation between two quantitative variables is equal to zero, then the two variables are not associated.

True or false: To calculate a predicted y -value from a given x -value, just look at the scatterplot and estimate the y -value.

True or false: A positive residual indicates the data point is above the regression line.

Example: Brushtail possums

1. What are the observational units?

2. Look at the scatterplot in Figure 6.5.
 - a) What is the explanatory variable? The response variable? What type is each?
 - b) What is the form of the scatterplot?
 - c) What is the direction of the scatterplot?
 - d) What is the strength of the scatterplot?
 - e) Are there any outliers on the scatterplot?
3. Write the equation of the regression line, in context (do not use x and y , use variable names instead).
4. Calculate the predicted head length for a possum with a 76.0 cm total length.
5. One of the possums in the data set has a total length of 76.0 cm and a head length of 85.1 mm. Calculate the residual for this possum. Does this possum lie above or below the regression line?

Section 6.2: Least squares regression

You may skip the special topic sections (6.2.7)

Vocabulary

Least squares criterion:

Least squares line:

`lm()` R function: `name_of_model <- lm(response ~ explanatory, data = data_set_name)`

slope:

y -intercept:

Extrapolation:

Coefficient of determination:

s_y^2 (or SST) represents

s_{RES}^2 (or SSE) represents

Notes

Two methods for determining the best line:

- 1.
- 2.

Notation for the coefficient of determination:

Formulas for calculating the coefficient of determination:

True or false: A correlation between two quantitative variables implies a causal relationship exists between the variables.

True or false: The slope of the line tells us how much to expect the y variable to increase or decrease when the x variable increases by 1 unit.

True or false: The coefficient of determination is just the square of the correlation.

Example: Elmhurst College

1. What are the observational units?
2. Look at the scatterplot in Figure 6.13.
 - a) What is the explanatory variable? The response variable?
 - b) What is the form of the scatterplot?
 - c) What is the direction of the scatterplot?
 - d) What is the strength of the scatterplot?
 - e) Are there any outliers on the scatterplot?
3. Write the equation of the regression line, in context (do not use x and y , use variable names instead).
4. Interpret the slope of the line, in the context of the problem. Remember that both family income and gift aid from the university are measured in \$1000s.

5. Interpret the y -intercept of the line, in the context of the problem. Remember that both family income and gift aid from the university are measured in \$1000s.
6. Is your interpretation in question 5 an example of extrapolation? Justify your answer.
7. Give and interpret, in context, the value of the coefficient of determination.

Section 6.3: Outliers in linear regression

Vocabulary

Outlier:

Leverage:

Influential point:

Notes

Investigate, but do not remove, outliers. Unless you find there was an actual error in the data collection, ignoring outliers can make models poor predictors!

True or false: All high leverage outliers are influential.

True or false: An outlier is considered high leverage if it is extreme in its x -value.

Section 6.4: Chapter 6 review

Look at the table of vocabulary terms in the final section of each chapter. If there are any you do not know, be sure to review the appropriate section of your text.

Notes

Statistics summarize:

Parameters summarize:

Determine whether each of the following statements about the correlation coefficient are true or false:

1. The correlation coefficient must be a positive number.
2. Stronger linear relationships are indicated by correlation coefficients far from 0.
3. The correlation coefficient is a robust statistic.
4. When two variables are highly correlated, that indicates a causal relationship exists between the variables.

5. The sign of the correlation coefficient will be the same as the sign of the regression line slope, though the values are typically different.

Fill in the blanks to correctly interpret:

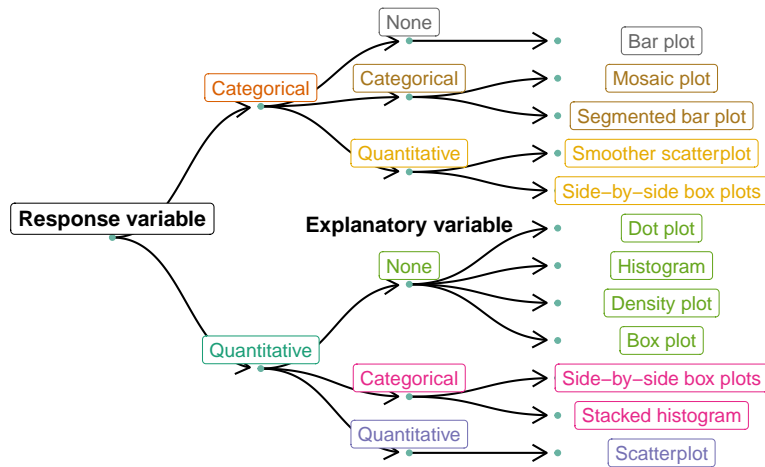
- Slope:

For every _____, we expect _____ to increase (if slope is _____) or decrease (if slope is _____) by the absolute value of the _____.

- y -intercept:

If _____, we predict the _____ to equal _____.

Decision tree for determining an appropriate plot given a number of variables and their types from Chapter review:



Section 7.1: Gapminder world

Vocabulary

Interaction:

Aesthetic:

Notes

Use color and a legend to add a third variable to a scatterplot. E.g., Color the dots to represent different levels of a categorical variable or use shading of the dots to represent different values of a quantitative variable.

If the response and one predictor are quantitative and the other predictor is categorical, we fit a regression line for each level of the categorical predictor.

- Parallel slopes would indicate that the two predictors _____ in explaining the response.
- Non-parallel slopes would indicate that the two predictors _____ in explaining the response.

True or false: Scatterplots can only display two variables at a time.

Section 7.2: Simpson's Paradox, revisited

Reminder from Section 4.4

Simpson's Paradox: when the relationship between the explanatory and response variable is reversed when looking at the relationship within different levels of a confounding variable.

Notes

True or false: Simpson's Paradox can only occur when the explanatory, response, and confounding variables are all categorical.

Example: SAT scores

1. What are the observational units?
2. Look at the scatterplot in Figure 7.5.
 - a) What is the explanatory variable? The response variable?
 - b) What is the form of the scatterplot?
 - c) What is the direction of the scatterplot?
 - d) What is the strength of the scatterplot?
 - e) Are there any outliers on the scatterplot? \vspace{0.1in}
3. What would need to be done to the study design in order to eliminate the confounding variable: percent of eligible students taking the SAT?
4. What features of the scatterplots in Figure 7.6 demonstrate that the percent of eligible students taking the SAT is a confounding variable?
5. How does Figure 7.7 demonstrate Simpson's Paradox?

4.2 Video Notes: Regression and Correlation

Read Chapters 6, 7, and 8 in the course textbook. Use the following videos to complete the lecture notes for Module 4.

4.2.1 Course Videos

- 6.1
- 6.2
- 6.3
- Ch 7

Summary measures and plots for two quantitative variables - Videos 6.1 - 6.3

Example: Data were collected from 1236 births between 1960 and 1967 in the San Francisco East Bay area to better understand what variables contributed to child birthweight, as children with low birthweight often suffer from an array of complications later in life (“Child Health and Development Studies,” n.d.). There were some missing values in the study and with those observations removed we have a total of 1223 births.

```
babies<-read.csv("data/babies.csv") %>%
  drop_na(bwt) %>%
  drop_na(gestation)
glimpse(babies)
#> Rows: 1,223
#> Columns: 8
#> $ case      <int> 1, 2, 3, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, ~
#> $ bwt        <int> 120, 113, 128, 108, 136, 138, 132, 120, 143, 140, 144, 141, ~
#> $ gestation  <int> 284, 282, 279, 282, 286, 244, 245, 289, 299, 351, 282, 279, ~
#> $ parity     <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
#> $ age        <int> 27, 33, 28, 23, 25, 33, 23, 25, 30, 27, 32, 23, 36, 30, 38, ~
#> $ height     <int> 62, 64, 64, 67, 62, 62, 65, 62, 66, 68, 64, 63, 61, 63, 63, ~
#> $ weight     <int> 100, 135, 115, 125, 93, 178, 140, 125, 136, 120, 124, 128, 9~
#> $ smoke      <int> 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, ~
```

Here you see a glimpse of the data. The 1223 rows correspond to the sample size. The case variable is labeling each pregnancy 1 through 1223. Then 7 variables are recorded. birthweight (bwt), length of gestation in days, parity is called an indicator variable telling us if the pregnancy was a first pregnancy (labeled as 0) or not (labeled as 1) were recorded about the child and pregnancy. The age, height, and weight were recorded for the mother giving birth, as was smoke, another indicator variable where 0 means the mother did not smoke during pregnancy, and 1 indicates that she did smoke while pregnant.

Type of plot


A _____ is used to display the relationship between two _____ variables.

Four characteristics of the scatterplot:

- Form:
- Direction:
- Strength:
- Outliers:
 - Influential points: outliers that change the regression line; far from the line of regression
 - High leverage points: outliers that are extreme in the x- axis; far from the mean of the x-axis

The following shows a scatterplot of length of gestation as a predictor of birthweight.

```
babies %>% # Data set pipes into...
ggplot(aes(x = gestation, y = bwt))+ # Specify variables
  geom_point(alpha=0.5) + # Add scatterplot of points
  labs(x = "number of days of gestation", # Label x-axis
       y = "birthweight (oz)", # Label y-axis
       title = "Scatterplot of Gestation vs. Birthweight for Births
               between 1960 and 1967 in San Francisco") +
  # Be sure to title your plots with the type of plot, observational units, variable(s)
  geom_smooth(method = "lm", se = FALSE) + # Add regression line
  theme_bw()
```



Describe the scatterplot using the four characteristics of a scatterplot.

The summary measures for two quantitative variables are:

- _____
- _____
- _____

Notation:

- Population slope:
- Population correlation:
- Sample slope:
- Sample correlation:

Correlation

Correlation is always between the values of _____ and _____.

- Measures the _____ and _____ of the linear relationship between two quantitative variables.
- The stronger the relationship between the variables the closer the value of _____ is to _____ or _____.
- The sign gives the _____.

The following code creates a correlation matrix between different quantitative variables in the data set.

```
babies %>%  
  select(c("gestation", "age", "height", "weight", "bwt")) %>%  
  cor(use="pairwise.complete.obs") %>%  
  round(3)
```

```
#>      gestation    age height weight    bwt  
#> gestation      1.000 -0.056  0.064  0.022 0.408  
#> age           -0.056  1.000 -0.005  0.147 0.029  
#> height         0.064 -0.005  1.000  0.436 0.201  
#> weight         0.022  0.147  0.436  1.000 0.154  
#> bwt            0.408  0.029  0.201  0.154 1.000
```

The value of correlation between gestation and birthweight is _____. This shows a _____, _____ relationship between gestation and birthweight.

Slope

- Least-squares regression line: $\hat{y} = b_0 + b_1 \times x$ (put y and x in the context of the problem) or $\widehat{response} = b_0 + b_1 \times \text{explanatory}$
- \hat{y} or $\widehat{response}$ is
- b_0 is
- b_1 is
- x or explanatory is

- The estimates for the linear model output will give the value of the _____ and the _____.
- Interpretation of slope: an increase in the _____ variable of 1 unit is associated with an increase/decrease in the _____ variable by the value of slope, on average.
- Interpretation of the y-intercept: for a value of 0 for the _____ variable, the predicted value for the _____ variable would be the value of y-intercept.
- We can predict values of the _____ variable by plugging in a given _____ variable value using the least squares equation line.
- A prediction of a response variable value for an explanatory value outside the range of x values is called _____.
- To find how far the predicted value deviates from the actual value we find the _____.

- To find the least squares regression line the line with the _____ SSE is found.

SSE = sum of squared errors

- To find SSE, the residual for each data point is found, squared and all the squared residuals are summed together

The linear model output for this study is given below:

```
# Fit linear model: y ~ x
babiesLM <- lm(bwt ~ gestation, data=babies)
round(summary(babiesLM)$coefficients,3) # Display coefficient summary
```

```
#>           Estimate Std. Error t value Pr(>|t|)
#> (Intercept)  -10.064      8.322  -1.209   0.227
#> gestation      0.464      0.030  15.609   0.000
```

Write the least squares equation of the line.

Interpret the slope in context of the problem.

Interpret the y-intercept in context of the problem.

Predict the birthweight for a birth with a baby born at 310 days gestation.

Calculate the residual for a birth of a baby with a birthweight of 151 ounces and born at 310 days gestation.


Is this value (310, 151) above or below the line of regression? Did the line of regression overestimate or underestimate the birthweight?

Coefficient of Determination

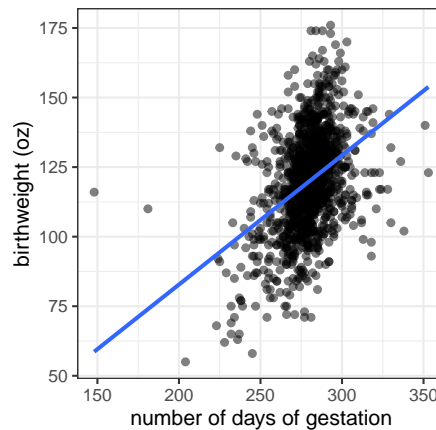
The coefficient of determination can be found by squaring the value of correlation, using the variances for each variable or using the SSE (sum of squares error) and SST (sum of squares total)

- $r^2 = (r)^2 = \frac{SST - SSE}{SST} = \frac{s_y^2 - s_{residual}^2}{s_y^2}$
- The coefficient of determination measures the _____ of total variation in the _____ variable that is explained by the changes in the _____ variable.

A Scatterplot of Gestation vs. Birthweight for Births between 1960 and 1967 in SF with Horizontal Line



B Scatterplot of Gestation vs. Birthweight for Births between 1960 and 1967 in SF with Regression Line



The value for SST was calculated as 406753.48. The value for SSE was calculated as 339092.13.

Calculate the coefficient of determination between gestation and birthweight.

Interpret the coefficient of determination between gestation and birthweight.

Multivariable plots - Video Chapter7


Aesthetics: visual property of the objects in your plot

- Position on the axes: groups for _____ variables, or a number line if the variable is _____
- Color or shape - to represent _____ variables
- Size - to represent _____ variables

Adding the quantitative variable maternal age to the scatterplot between gestation and birthweight.


```
babies %>% # Data set pipes into...
ggplot(aes(x = gestation, y = bwt))+ # Specify variables
  geom_point(alpha=0.5, shape=1, aes(size=age)) + # Add scatterplot of points
  labs(x = "number of days of gestation", # Label x-axis
       y = "birthweight (oz)", # Label y-axis
       title = "Scatterplot of Gestation vs. Birthweight by Age
               for Births between 1960 and 1967 in San Francisco") +
  # Be sure to title your plots
  geom_smooth(method = "lm", se = FALSE) # Add regression line
```

Scatterplot of Gestation vs. Birthweight by Age
for Births between 1960 and 1967 in San Francisco



Let's add the categorical variable, whether a mother smoked, to the scatterplot between gestation and birthweight.

```
babies <- babies %>%  
  mutate(smoke = factor(smoke)) %>%  
  na.omit()  
  
babies %>% # Data set pipes into...  
  ggplot(aes(x = gestation, y = bwt, color = smoke)) + #Specify variables  
  geom_point(aes(shape = smoke), size = 2) + #Add scatterplot of points  
  labs(x = "number of days of gestation", #Label x-axis  
       y = "birthweight (oz)", #Label y-axis  
       title = "Scatterplot of Gestation vs. Birthweight by  
               Smoking Status for Births between 1960 and 1967  
               in San Francisco") +  
  #Be sure to title your plots  
  geom_smooth(method = "lm", se = FALSE) + #Add regression line  
  scale_color_grey()
```



Does the relationship between length of gestation and birthweight appear to depend upon maternal smoking status?

Is the variable smoking status a potential confounding variable?

Adding a categorical predictor:

- Look at the regression line for each level of the _____
- If the slopes are _____, the two predictor variables do not _____ to help explain the response
- If the slopes _____, there is an interaction between the categorical predictor and the relationship between the two quantitative variables.

4.2.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. What are the three summary measures for two quantitative variables?
2. What are the four characteristics used to describe a scatterplot?
3. When we add a categorical predictor variable to a scatterplot of two quantitative variables, what summary measure will we compare across the categories to assess the change in the relationship between the two quantitative variables.

4.3 Guided Activity 6: IPEDS (continued)

4.3.1 Learning outcomes

- Identify and create appropriate summary statistics and plots given a data set with two quantitative variables.
- Find the estimated line of regression using summary statistics and R linear model (`lm()`) output.
- Interpret the slope coefficient in context of the problem.
- Calculate and interpret r^2 , the coefficient of determination, in context of the problem.
- Find the correlation coefficient from R output or from r^2 and the sign of the slope.

4.3.2 Terminology review

In today's activity, we will review summary measures and plots for two quantitative variables. Some terms covered in this activity are:

- Least-squares line of regression
- Slope and y -intercept
- Residuals
- Correlation (r)
- Coefficient of determination (r -squared)

To review these concepts, see Chapter 6 in the textbook.

4.3.3 The Integrated Postsecondary Education Data System (IPEDS)

We will continue to assess the IPEDS data set collected on a subset of institutions that met the following selection criteria (Education Statistics 2018):

- Degree granting
- United States only
- Title IV participating
- Not for profit
- 2-year or 4-year or above
- Has full-time first-time undergraduates

Some of the variables collected and their descriptions are below. Note that several variables have missing values for some institutions (denoted by "NA").

Variable	Description
UnitID	Unique institution identifier
Name	Institution name
State	State abbreviation
Sector	whether public or private
LandGrant	Is this a land-grant institution (Yes/No)
Size	Institution size category based on total student enrolled for credit, Fall 2018: Under 1,000, 1,000 - 4,999, 5,000 - 9,999, 10,000 - 19,999, 20,000 and above
Cost_OutofState	Cost of attendance for full-time out-of-state undergraduate students

Variable	Description
Cost_InState	Cost of attendance for full-time in-state undergraduate students
Retention	Retention rate is the percent of the undergraduate students that re-enroll in the next year
Graduation_Rate	6-year graduation rate for undergraduate students
SATMath_75	75th percentile Math SAT score
ACT_75	75th percentile ACT score


The code below reads in the needed data set, IPEDS_2018.csv, and filters out the 2-year institutions. Highlight lines 1 – 11 to load the data set and filter out the 2-year institutions.

```
IPEDS <- read.csv("https://www.math.montana.edu/courses/s216/data/IPEDS_2018.csv")
IPEDS <- IPEDS %>%
  filter(Sector != "Public 2-year") #Filters the data set to remove Public 2-year
IPEDS <- IPEDS %>%
  filter(Sector != "Private 2-year") #Filters the data set to remove Private 2-year
IPEDS <- na.omit(IPEDS)
```

The create a scatterplot of the 75th percentile Math SAT score by retention rate for 4-year US Higher Education Institutions...

- Enter the variable SATMath_75 for explanatory and Retention for response in line 16.
- Highlight and run lines 15–21.

```
IPEDS %>% # Data set pipes into...
  ggplot(aes(x = SATMath_75, y = Retention)) + # Specify variables
  geom_point(alpha=0.5) + # Add scatterplot of points
  labs(x = "75th Percentile SAT Math Score", # Label x-axis
       y = "Retention Rate (%)", # Label y-axis
       title = "Scatterplot of SAT Math Score vs. Retention Rate for
               4-year US Higher Education Institutions") +
  # Be sure to title your plots with the type of plot, observational units, variable(s)
  geom_smooth(method = "lm", se = FALSE) + # Add regression line
  theme_bw()
```



1. Describe the relationship between 75th percentile SAT Math score and retention rate.

Slope of the Least Squares Linear Regression Line

There are three summary measures calculated from two quantitative variables: slope, correlation, and the coefficient of determination. We will first assess the slope of the least squares regression line between 75th percentile SAT Math score and retention rate.

- Enter **Retention** for response and **SATMath_75** for explanatory in line 25
- Highlight and run lines 25 – 26 to fit the linear model.

```
# Fit linear model: y ~ x
IPEDSLM <- lm(Retention~SATMath_75, data=IPEDS)
round(summary(IPEDSLM)$coefficients,3) # Display coefficient summary
```

```
#>               Estimate Std. Error t value Pr(>|t|)
#> (Intercept)    0.059      1.898    0.031   0.975
#> SATMath_75     0.125      0.003   40.485   0.000
```

2. Write out the least squares regression line using the summary statistics from the R output in context of the problem.

Slope Interpretation: An increase of one point in SAT Math 75th percentile score is associated with an increase in retention rate, on average, of 0.125 percentage points for 4-year higher education institutions.

3. Predict the retention rate for a 4-year US higher education institution with a 75th percentile SAT Math score of 440.
4. Calculate the residual for a 4-year US higher education institution with a 75th percentile SAT Math score of 440 and a retention rate of 24%.

Correlation

Correlation measures the strength and the direction of the linear relationship between two quantitative variables. The closer the value of correlation to +1 or -1, the stronger the linear relationship. Values close to zero indicate a very weak linear relationship between the two variables.

The following output creates a correlation matrix between several pairs of quantitative variables.

```
IPEDS %>% # Data set pipes into
  select(c("Retention", "Cost_InState",
           "Graduation_Rate", "Salary",
           "SATMath_75", "ACT_75")) %>%
```

```
cor(use="pairwise.complete.obs") %>%
round(3)
```

```
#>           Retention Cost_InState Graduation_Rate Salary SATMath_75 ACT_75
#> Retention           1.000         0.388           0.832  0.698         0.767  0.768
#> Cost_InState        0.388         1.000           0.563  0.365         0.502  0.514
#> Graduation_Rate     0.832         0.563           1.000  0.683         0.817  0.833
#> Salary              0.698         0.365           0.683  1.000         0.747  0.706
#> SATMath_75          0.767         0.502           0.817  0.747         1.000  0.920
#> ACT_75              0.768         0.514           0.833  0.706         0.920  1.000
```

5. What is the value of correlation between SATMath_75 and Retention?

Coefficient of determination (squared correlation)

Another summary measure used to explain the linear relationship between two quantitative variables is the coefficient of determination (r^2). The coefficient of determination, r^2 , can also be used to describe the strength of the linear relationship between two quantitative variables. The value of r^2 (a value between 0 and 1) represents the **proportion of variation in the response that is explained by the least squares line with the explanatory variable**. There are two ways to calculate the coefficient of determination:

Square the correlation coefficient: $r^2 = (r)^2$


Use the variances of the response and the residuals: $r^2 = \frac{s_y^2 - s_{RES}^2}{s_y^2} = \frac{SST - SSE}{SST}$

6. Use the correlation, r , found in question 5, to calculate the coefficient of determination between SATMath_75 and Retention, r^2 .

7. The variance of the response variable, Retention in \$MM, is $s_{Retention}^2 = 138.386 \%^2$ and the variability in the residuals is $s_{RES}^2 = 56.934 \%^2$. Use these values to calculate the coefficient of determination.


In the next part of the activity we will explore what the coefficient of determination measures.

In the first scatterplot, we see the data plotted with a horizontal line. Note that the regression line in this plot has a slope of zero; this assumes there is no relationship between SATMath_75 and Retention. The value of the y-intercept, 76.387, is the mean of the response variable when there is no relationship between the two variables. To find the sum of squares total (SST) we find the residual ($residual = y - \hat{y}$) for each response value from the horizontal line (from the value of 76.387). Each residual is squared and the sum of the squared values is calculated. The SST gives the **total variability in the response variable, Retention**.



The calculated value for the SST is 93318.731.

This next scatterplot, shows the plotted data with the best fit regression line. This is the line of best fit between budget and revenue and has the smallest sum of squares error (SSE). The SSE is calculated by finding the residual from each response value to the regression line. Each residual is squared and the sum of the squared values is calculated.



The calculated value for the SSE is 65133.022.


8. Calculate the value for r^2 using the values for SST and SSE provided below each of the previous graphs.

9. Write a sentence interpreting the coefficient of determination in context of the problem.

Multivariable plots

When adding another categorical predictor, we can add that variable as shape or color to the plot. In the following code we have added the variable **Sector**, whether the 4-year institution is public or private.

```
IPEDS %>% # Data set pipes into...
  ggplot(aes(x = SATMath_75, y = Retention, shape = Sector, color=Sector))+ # Specify variables
  geom_point(alpha=0.5) + # Add scatterplot of points
  labs(x = "75th Percentile SAT Math Score", # Label x-axis
       y = "Retention Rate (%)", # Label y-axis
       title = "Scatterplot of SAT Math Score vs. Retention Rate for
       4-year US Higher Education Institutions") +
  # Be sure to title your plots with the type of plot, observational units, variable(s)
  geom_smooth(method = "lm", se = FALSE) + # Add regression line
  scale_color_grey()
```



10. Does the relationship between 75th percentile SAT math score and retention rate of 4-year institutions change depending on the level of sector?

4.3.4 Take-home messages

1. The sign of correlation and the sign of the slope will always be the same. The closer the value of correlation is to -1 or $+1$, the stronger the linear relationship between the explanatory and the response variable.
2. The coefficient of determination multiplied by 100 ($r^2 \times 100$) measures the percent of variation in the response variable that is explained by the relationship with the explanatory variable. The closer the value of the coefficient of determination is to 100%, the stronger the relationship.
3. We can use the line of regression to predict values of the response variable for values of the explanatory variable. Do not use values of the explanatory variable that are outside of the range of values in the data set to predict values of the response variable (reflect on why this is true.). This is called **extrapolation**.

4.3.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

4.4 Activity 3: Moneyball — Linear Regression

4.4.1 Learning outcomes

- Identify and create appropriate summary statistics and plots given a data set with two quantitative variables.
- Use scatterplots to assess the relationship between two quantitative variables.
- Find the estimated line of regression using summary statistics and R linear model (`lm()`) output.
- Interpret the slope coefficient in context of the problem.

4.4.2 Terminology review

In today's activity, we will review summary measures and plots for two quantitative variables. Some terms covered in this activity are:

- Scatterplot
- Least-squares line of regression
- Slope and y -intercept
- Residuals

To review these concepts, see Chapter 6 & 7 in the textbook.

4.4.3 Movies released in 2016

The goal of a Major League baseball team is to make the playoffs. In 2002, the manager of the Oakland A's, Billy Bean, with the help of Paul DePodesta began to use statistics to determine which players to choose for their season. Based on past data, DePodesta determined that to make it to the playoffs, the A's would need to win at least 95 games in the regular season. In order to win more games, they would need to score more runs than they allowed. The Oakland A's won 20 consecutive games and a total of 103 games for the season. The success of this use of sports analytics was portrayed by the 2011 movie, Moneyball. In this study, we will see if there is evidence of a positive linear relationship between the difference in the number of runs scored minus the number of runs allowed (RD) and the number of wins for Major League baseball teams in the years before 2002. Some of the variables collected in the data set baseball consist of the following:

Variable	Description
RA	Runs allowed
RS	Runs scored
OBP	On-base percentage
SLG	Slugging percentage
BA	Batting average
OOBP	Opponent's on-base percentage
OSLG	Opponent's slugging percentage
W	Number of wins in the season
RD	Difference of runs scored minus runs allowed

```
moneyball <- read.csv("data/baseball.csv") # Reads in data set
moneyball$RD <- moneyball$RS - moneyball$RA
moneyball <-
  moneyball %>% # Pipe data set into
  subset(Year < 2002) # Select only years before 2002
```

Vocabulary review

- Use the provided R script file to create a scatterplot to examine the relationship between the difference in number of runs scored minus number of runs allowed and the number of wins by filling in the variable names (RD and W) for explanatory and response in line 14. Note, we are using the difference in runs scores minus runs allowed to predict the number of season wins.
- Highlight and run lines 1–20.

```
moneyball %>% # Data set pipes into...
  ggplot(aes(x = explanatory, y = response))+ # Specify variables
  geom_point() + # Add scatterplot of points
  labs(x = "Difference in number of runs", # Label x-axis
       y = "Number of Season wins", # Label y-axis
       title = "Scatterplot of Run Difference vs. Number of Season Wins for MLB Teams") +
  # Be sure to tile your plots
  geom_smooth(method = "lm", se = FALSE) # Add regression line
```

1. Assess the four features of the scatterplot that describe this relationship.
 - Form (linear, non-linear)
 - Direction (positive, negative)
 - Strength
 - Unusual observations or outliers
2. Based on the plot, does there appear to be an association between run difference and number of season wins? Explain your answer.

Slope

The linear model function in R (`lm()`) gives us the summary for the least squares regression line. The estimate for (**Intercept**) is the y -intercept for the line of least squares, and the estimate for **budget_mil** (the x -variable name) is the value of b_1 , the slope.

- Run lines 24–25 in the R script file to reproduce the linear model output found in the coursepack.

```
# Fit linear model: y ~ x
moneyballLM <- lm(W~RD, data=moneyball)
round(summary(moneyballLM)$coefficients, 3) # Display coefficient summary
```

```
#>               Estimate Std. Error t value Pr(>|t|)
#> (Intercept)    80.881      0.131 616.675      0
#> RD              0.106      0.001  81.554      0
```

3. Write out the least squares regression line using the summary statistics provided above in context of the problem.

4. Interpret the value of slope in context of the problem.
5. Using the least squares line from question 3, predict the number of season wins for a MLB team that has a run difference of -66 runs.
6. Predict the number of season wins for a MLB team that has a run difference of 400 runs.
7. The prediction in question 6 is an example of what?

Residuals

The model we are using assumes the relationship between the two variables follows a straight line. The residuals are the errors, or the variability in the response that hasn't been modeled by the regression line.

$$\Rightarrow \text{Residual} = \text{actual y value} - \text{predicted y value}$$

$$e = y - \hat{y}$$

8. The MLB team *Florida Marlins* had a run difference of -66 runs and 79 wins for the season. Find the residual for this MLB team.
9. Did the line of regression overestimate or underestimate the number of wins for the season for this team?

Correlation

The following output shows a correlation matrix between several pairs of quantitative variables.

- Highlight and run lines 29–33 to produce the same table as below.

```
moneyball %>% # Data set pipes into
  select(c("RD", "BA",
           "SLG", "W")) %>%
  cor(use="pairwise.complete.obs") %>%
  round(3)
```

```
#>      RD    BA   SLG    W
#> RD  1.000 0.442 0.428 0.939
#> BA  0.442 1.000 0.814 0.416
#> SLG 0.428 0.814 1.000 0.406
#> W   0.939 0.416 0.406 1.000
```

10. Report the value of correlation between the run difference and the number of season wins.
11. Calculate the coefficient of determination between the run difference and the number of season wins.
12. Interpret the value of coefficient of determination in context of the study.

4.4.4 Take-home messages

1. Two quantitative variables are graphically displayed in a scatterplot. The explanatory variable is on the x -axis and the response variable is on the y -axis. When describing the relationship between two quantitative variables we look at the form (linear or non-linear), direction (positive or negative), strength, and for the presence of outliers.
2. There are three summary statistics used to summarize the relationship between two quantitative variables: correlation (r), slope of the regression line (b_1), and the coefficient of determination (r^2).

4.4.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

4.5 Module 4 Lab: Penguins

4.5.1 Learning outcomes

- Identify and create appropriate summary statistics and plots given a data set with two quantitative variables.
- Use scatterplots to assess the relationship between two quantitative variables.
- Find the estimated line of regression using summary statistics and R linear model (`lm()`) output.
- Interpret the slope coefficient in context of the problem.
- Calculate and interpret r^2 , the coefficient of determination, in context of the problem.
- Find the correlation coefficient from R output or from r^2 and the sign of the slope.

4.5.2 Penguins

The Palmer Station Long Term Ecological Research Program sampled three penguin species on islands in the Palmer Archipelago in Antarctica. Researchers took various body measurements on the penguins, including bill depth and body mass. The researchers were interested in the relationship between bill depth and body mass and wondered if bill depth could be used to accurately predict the body mass of these three penguin species.

- Upload the `Antarctica_Penguins` csv file
- Upload and open the provided R script file for the module 4 lab.
- Enter the name of the data set (`Antarctica_Penguins`) for `datasetname` in the R script file in line 5.
*Note that there are some missing values in the dataset. Line 6 in the R script file removes the NAs.

First we will create a scatterplot of the bill depth and body mass. Notice that we are using bill depth (mm) to predict body mass (g). This makes bill depth the explanatory variable.

- **Make sure to give your plot a descriptive title between the quotations in line 16. Remember that the title should include the type of plot, the observational units, and the variable(s) plotted.**
- Highlight and run lines 1–17 in the R script file.
- **Upload a copy of your scatterplot to Gradescope.**

```
penguins <- read.csv("datasetname") %>% #Creates the object penguins
  drop_na(body_mass_g) %>% drop_na(bill_depth_mm) #Removes data points without values
penguins %>%
  ggplot(aes(x = bill_depth_mm, y = body_mass_g))+ # Specify variables
  geom_point(alpha=0.5) + # Add scatterplot of points
  labs(x = "bill depth (mm)", # Label x-axis
       y = "body mass (g)", # Label y-axis
       title = "Title") + # Be sure to title your plot
  geom_smooth(method = "lm", se = FALSE) # Add regression line
```

1. Assess the four features of the scatterplot that describe this relationship.

- Form (linear, non-linear)
- Direction (positive, negative)
- Strength
- Unusual observations or outliers

To create the correlation matrix...

- Highlight and run lines 21–25 in the R script file.

```
penguins %>% # Data set pipes into
select(c("bill_length_mm", "bill_depth_mm",
         "flipper_length_mm", "body_mass_g")) %>%
cor(use="pairwise.complete.obs") %>%
round(3)
```

2. Using the R output, report the value of correlation between bill depth and body mass.

3. Using the value of correlation found in question 2, calculate the value of the coefficient of determination.

4. **Interpret the coefficient of determination in context of the problem.**

To get the linear model output...

- Enter the variable name `body_mass_g` for response and the variable name `bill_depth_mm` for explanatory in line 30 in the R script file.
- Highlight and run lines 30–31.

```
# Fit linear model: y ~ x
penguinsLM <- lm(response~explanatory, data=penguins)
round(summary(penguinsLM)$coefficients, 3) # Display coefficient summary
```

5. Write out the least squares regression line using the summary statistics from the R output in context of the problem.

6. **Interpret the value of slope in context of the problem.**

7. Using the least squares regression line from question 5, predict the body mass for a penguin with a bill depth of 19.6 mm.
8. One penguin had a bill depth of 19.6 mm and a body mass of 4675 g. Find the residual for this penguin.
9. Did the line of regression overestimate or underestimate the body mass for this penguin?

Multivariable plot

Does species change the relationship between bill depth and body mass?

- Highlight and run lines 35–44 to get the multivariable plot.

```
penguins %>%
  ggplot(aes(x = bill_depth_mm, y = body_mass_g, color=species))+ # Specify variables
  geom_point(aes(shape = species), size = 2, alpha=0.5) + # Add scatterplot of points
  labs(x = "bill depth (mm)", # Label x-axis
       y = "body mass (g)", # Label y-axis
       color = "species",
       shape = "species",
       title = "Scatterplot of Bill Depth and Body Mass by Penguin Species") +
  # Enter the title for the plot between the quotations
  geom_smooth(method = "lm", se = FALSE) + # Add regression line
  scale_color_viridis_d(end=0.8)
```

10. What three variables are plotted on this plot?
11. Do species and bill depth appear to interact when predicting body mass of Antarctic penguins? Explain your answer.
12. Explain the association between species and each of the other two variables.
13. Notice that the slope of the line between bill depth and body mass for each species is positive while the slope for the line not accounting for species is negative. What phenomena is this an example of?

4.5.3 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

Inference for a Single Categorical Variable: Simulation-based Methods

5.1 Reading Guide: Categorical Inference

Textbook Chapter 9: Hypothesis testing with randomization

Reminders from previous sections

n = sample size

\hat{p} = sample proportion

π = population proportion

Vocabulary

Statistical inference:

Hypothesis test:

Also called a ‘significance test’.

Simulation-based method:

Theory-based method:

This will be discussed in detail next week.

Observed statistic:

Null statistics:

Point estimate:

Null hypothesis (H_0):

Alternative hypothesis (H_A):

Test statistic:

P-value:

Statistically significant:

Significance level (α):

Sampling variability (found in the chapter 9 review):

Notes

Consider the US judicial system:

What is the null hypothesis?

What is the alternative hypothesis?

The jury is presented with evidence.

- If the evidence is strong (beyond a reasonable doubt), the jury will find the defendant:
- If the evidence is not strong (not beyond a reasonable doubt), the jury will find the defendant:

To create a simulation, which hypothesis (null or alternative) do we assume is true?

Lower p-values indicate (stronger or weaker?) evidence (for or against?) the null hypothesis.

General steps of a hypothesis test:

Example from section 9.1: Martian alphabet

1. What is the sample statistic presented in this example? What notation would be used to represent this value?
2. What are the two possible explanations for how these data could have occurred?
3. Of the two explanations, which is the null and which is the alternative hypothesis?
4. How could coins be used to create a simulation of what should happen if everyone in the class was just guessing?

5. How can we use the simulation to determine which of the two possibilities is more believable?

Example from section 9.2: Sex discrimination

1. Identify the explanatory and response variables in this study.
2. Is this a randomized experiment or an observation study? Justify your answer.
3. What is the sample statistic presented in this example?
4. What are the null and alternative hypotheses?
5. Based on the simulation pictured in Figure 9.6, which hypothesis seems more believable? Justify your answer, being sure to reference (and explain the meaning of) the p-value.

Chapter 10: Confidence intervals with bootstrapping

Reminders from previous sections

Parameter: a value summarizing a variable(s) for a population.

Statistic: a value summarizing a variable(s) for a sample. Also called a **point estimate**.

Vocabulary

Confidence interval:

Bootstrapping:

Bootstrapped resample:

Bootstrapped statistic:

Bootstrap percentile confidence interval:

Notes

What is the purpose of bootstrapping:

How is bootstrapping used?

If we want to find a 90% confidence interval, what percentiles of the bootstrap distribution would we need?

Example: Medical consultant

1. What is the sample statistic presented in this example? What notation would be used to represent this value?
2. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
3. How could we use cards to simulate **one** bootstrapped resample? How many blue cards — to represent what? How many red cards — to represent what? How many times would we draw a card and replace it back in the deck? What would you record once you completed the draw-with-replacement process?
4. Provide and interpret the 95% confidence interval provided in the textbook.

Section 14.1: Simulation-based test for $H_0 : \pi = \pi_0$

Reminders from previous sections

n = sample size

\hat{p} = sample proportion

π = population proportion

General steps of a hypothesis test:

1. Frame the research question in terms of hypotheses.
2. Collect and summarize data using a test statistic.
3. Assume the null hypothesis is true, and simulate or mathematically model a null distribution for the test statistic.
4. Compare the observed test (standardized) statistic to the null distribution to calculate a p-value.
5. Make a conclusion based on the p-value and write the conclusion in context.

Parameter: a value summarizing a variable(s) for a population.

Statistic: a value summarizing a variable(s) for a sample.

Sampling distribution: plot of statistics from 1000s of samples of the same size taken from the same population.

Hypothesis test: a process to determine how strong the evidence of an effect is.

Also called a ‘significance test’.

Simulation-based method: Simulate lots of samples of size n under assumption of the null hypothesis, then find the proportion of the simulations that are at least as extreme as the observed sample statistic.

Null hypothesis (H_0): the skeptical perspective; no difference; no change; no effect; random chance; what the researcher hopes to prove is **wrong**.

Alternative hypothesis (H_A): the new perspective; a difference/increase/decrease; an effect; not random chance; what the researcher hopes to prove is **correct**.

P-value: probability of seeing the observed sample data, or something more extreme, assuming the null hypothesis is true.

\Rightarrow Lower the p-value the stronger the evidence AGAINST the null hypothesis and FOR the alternative hypothesis.

Significance level (α): a threshold used to determine if a p-value provides enough evidence to reject the null hypothesis or not.

Common levels of α include 0.01, 0.05, and 0.10.

Statistically significant: results are considered statistically significant if the p-value is below the significance level.

Vocabulary

Null value:

Null distribution:

Notes

Which hypothesis must we assume is true in order to simulate a null distribution?

How is a p-value calculated in a simulation-based hypothesis test?

How is a null distribution created in a simulation-based hypothesis test for a single proportion?

True or false: The sign in the alternative hypothesis is based off of the researcher’s question.

Example: Medical consultant

As a reminder from section 10.1, $\hat{p} = 0.048$.

1. Write the null and alternative hypotheses in words.
2. Write the null and alternative hypotheses in notation.
3. To simulate the null distribution, we would not be able to use coins. Why not?
4. How could we use cards to simulate 1 sample which assumes the null hypothesis is true? How many blue cards — to represent what? How many red cards — to represent what? How many times would we draw a card and replace it back in the deck? What would you record once you completed the draw-with-replacement process?
5. How can we calculate a p-value from the simulated null distribution for this example in figure 14.1?
6. What was the p-value of the test?
7. What conclusion should the researcher make?
8. Are the results in this example statistically significant? Justify your answer.
9. To what population can we generalize the results of this study? Justify your answer.

Section 14.2: Bootstrap confidence interval for π

Reminders from previous sections

Confidence interval: a process to determine how large an effect is; a range of plausible values for the parameter. Also called ‘estimation’.

Bootstrapping: the process of drawing with replacement n times from the original sample.

Bootstrapped resample: a random sample of size n from the original sample, selected with replacement.

Bootstrapped statistic: the statistic recorded from the bootstrapped resample.

Bootstrap percentile confidence interval: An $X\%$ confidence interval made by finding the bounds of the middle $X\%$ of bootstrap statistics.

Notes

Explain how to use a confidence interval to test a null hypothesis.

Example: Medical consultant

1. What is the 95% confidence interval for the parameter of interest, based on the bootstrap distribution shown in figure 10.6?
2. Does the 95% confidence interval indicate that the consults has a lower rate of complications than the national average (10%)? Justify your answer.

5.2 Video Notes: Inference for One Categorical Variable using Simulation-based Methods

Read Chapters 9 and 10 and Sections 14.1 and 14.2 in the course textbook. Use the following videos to complete the lecture notes for Module 5.

5.2.1 Course Videos

- Chapter9
- 14.1
- Chapter10
- 14.2

Hypothesis Testing - Video Chapter9

Purpose of a hypothesis test:

- Use data collected on a sample to give information about the population.
- Determines _____ of _____ of an effect

General steps of a hypothesis test

1. Write a research question and hypotheses.
2. Collect data and calculate a summary statistic.
3. Model a sampling distribution which assumes the null hypothesis is true.
4. Calculate a p-value.
5. Draw conclusions based on a p-value.

Hypothesis Testing/Justice System

- Two possible outcomes if the observed statistic is unusual:
 - Strong evidence against _____ -> _____
 - Not enough evidence against _____ -> _____
- Always written about the _____ (population)

Null hypothesis

- Skeptical perspective, no difference, no effect, random chance
- What the researcher hopes is _____.

Notation:

Alternative hypothesis

- New perspective, a chance, a difference, an effect
- What the researcher hopes is _____.

Notation:

Simulation vs. Theory-based Methods

Simulation-based method

Creation of the null distribution


- Simulate many samples assuming _____
- Find the proportion of _____ at least as extreme as the observed sample _____
- The null distribution estimates the sample to sample variability expected in the population

Theory-based method

- Use a mathematical model to determine a distribution under the null hypothesis
- Compare the observed sample statistic to the model to calculate a probability
- *Theory-based methods will be discussed in the next module*

P-value

- What does the p-value measure?
 - Probability of observing the sample _____ or more _____ assuming the _____ hypothesis is _____.
- How much evidence does the p-value provide against the null hypothesis?



- The _____ the p-value, the _____ the evidence against the null hypothesis.

- Write a conclusion based on the p-value.
 - Answers the _____ question.
 - Amount of _____ in support of the _____ hypothesis.
- Decision: can we reject or fail to reject the null hypothesis?
 - Significance level: cut-off of “small” vs “large” p-value
 - p-value $\leq \alpha$
 - Strong enough evidence against the null hypothesis
 - Decision:
 - Results are _____ significant.
 - p-value $> \alpha$
 - Not enough evidence against the null hypothesis
 - Decision:
 - Results are not _____ significant.

One proportion test

- Reminder: review summary measures and plots discussed in the Week 3 material and Chapter 4 of the textbook.
- The summary measure for a single categorical variable is a _____.

Notation:

- Population proportion:
- Sample proportion:

Parameter of Interest:

- Include:
 - Reference of the population (true, long-run, population, all)
 - Summary measure
 - Context
 - * Observational units/cases
 - * Response variable (and explanatory variable if present)
 - If the response variable is categorical, define a ‘success’ in context

Hypothesis testing

Conditions:

- Independence:

Null hypothesis assumes “no effect”, “no difference”, “nothing interesting happening”, etc.

Always of form: “parameter” = null value

H_0 :

H_A :

- Research question determines the direction of the alternative hypothesis.

Video 14.1 Example: A 2007 study published in the Behavioral Ecology and Sociobiology Journal was titled “Why do blue-eyed men prefer blue-eyed women?” (Laeng 2007) In this study, conducted in Norway, 114 volunteer heterosexual blue-eyed males rated the attractiveness of 120 pictures of females. The researchers recorded which eye-color (blue, green, or brown) was rated the highest, on average. In the sample, 51 of the volunteers rated the blue-eyed women the most attractive. Do blue-eyed heterosexual men tend to find blue-eyed women the most attractive?

Parameter of interest:

Write the null and alternative hypotheses for the blue-eyed study:

In notation:

H_0 :

H_A :


Statistic:

Is the independence condition met to analyze these data using a simulation-based approach?

Simulation-based method

- Simulate many samples assuming $H_0 : \pi = \pi_0$
 - Create a spinner with that represents the null value
 - Spin the spinner n times
 - Calculate and plot the simulated sample proportion from each simulation
 - Repeat 1000 times (simulations) to create the null distribution
 - Find the proportion of simulations at least as extreme as \hat{p}

```
set.seed(216)
one_proportion_test(probability_success = 0.333, # Null hypothesis value
  sample_size = 114, # Enter sample size
  number_repetitions = 1000, # Enter number of simulations
  as_extreme_as = 0.447, # Observed statistic
  direction = "greater", # Specify direction of alternative hypothesis
  summary_measure = "proportion") # Reporting proportion or number of successes?
```



Explain why the null distribution is centered at the value of approximately 0.333:

Interpretation of the p-value:

- Statement about probability or proportion of samples
- Statistic (summary measure and value)
- Direction of the alternative
- Null hypothesis (in context)

Conclusion:

- Amount of evidence
- Parameter of interest
- Direction of the alternative hypothesis

Generalization:

- Can the results of the study be generalized to the target population?

Confidence interval - Video Chapter10

statistic \pm margin of error

Vocabulary:

- Point estimate:
- Margin of error:

Purpose of a confidence interval

- To give an _____ for the parameter of interest
- Determines how _____ an effect is

Sampling distribution

- Ideally, we would take many samples of the same _____ from the same population to create a sampling distribution
- But only have 1 sample, so we will _____ with _____ from the one sample.
- Need to estimate the sampling distribution to see the _____ in the sample

Simulation-based methods

Bootstrap distribution:

- Write the response variable values on cards
- Sample with replacement n times (bootstrapping)
- Calculate and plot the simulated difference in sample means from each simulation

- Repeat 1000 times (simulations) to create the bootstrap distribution
- Find the cut-offs for the middle X% (confidence level) in a bootstrap distribution.

What is bootstrapping?


- Assume the “population” is many, many copies of the original sample.
- Randomly sample with replacement from the original sample n times.

Video 14.2

Let’s revisit the blue-eyed male study to estimate the *proportion of ALL heterosexual blue-eyed males who tend to find blue-eyed women the most attractive* by creating a 90% confidence interval.

Bootstrap distribution:

```
set.seed(216)
one_proportion_bootstrap_CI(sample_size = 114, # Sample size
                             number_successes = 51, # Observed number of successes
                             number_repetitions = 1000, # Number of bootstrap samples to use
                             confidence_level = 0.90) # Confidence level as a decimal
```




Confidence interval interpretation:

- How confident you are (e.g., 90%, 95%, 98%, 99%)
- Parameter of interest
- Calculated interval
- Order of subtraction when comparing two groups

How does changing the confidence level impact the width of the confidence interval?


95% Confidence Interval:

```
set.seed(216)
one_proportion_bootstrap_CI(sample_size = 114, # Sample size
                             number_successes = 51, # Observed number of successes
                             number_repetitions = 1000, # Number of bootstrap samples to use
                             confidence_level = 0.95) # Confidence level as a decimal
```



99% Confidence Interval:

```
set.seed(216)
one_proportion_bootstrap_CI(sample_size = 114, # Sample size
                             number_successes = 51, # Observed number of successes
                             number_repetitions = 1000, # Number of bootstrap samples to use
                             confidence_level = 0.99) # Confidence level as a decimal
```



5.2.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. What was the null value for the study discussed?
2. Write the alternative hypothesis for this study in notation? How was the direction of the alternative hypothesis determined?
3. Do the results of the confidence interval *match* the results based on the p-value?

5.3 Guided Activity 7: Sleep Study — Simulation-based Inference Methods

5.3.1 Learning outcomes

- Identify the two possible explanations (one assuming the null hypothesis and one assuming the alternative hypothesis) for a relationship seen in sample data.
- Given a research question involving a single categorical variable, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Describe and perform a simulation-based hypothesis test for a single proportion.

5.3.2 Terminology review

In today's guided activity, we will use simulation based methods for inference of a single categorical variable. Some terms covered in this activity are:

- Parameter of interest
- Null hypothesis
- Alternative hypothesis
- Simulation
- Bootstrapping
- Confidence interval

To review these concepts, see Chapters 9, 10, & 14 in your textbook.

5.3.3 Steps of the statistical investigation process

We will work through a five-step process to complete a hypothesis test for a single proportion, first introduced in the activity in week 1.

- **Ask a research question** that can be addressed by collecting data. What are the researchers trying to show?
- **Design a study and collect data.** This step involves selecting the people or objects to be studied and how to gather relevant data on them.
- **Summarize and visualize the data.** Calculate summary statistics and create graphical plots that best represent the research question.
- **Use statistical analysis methods to draw inferences from the data.** Choose a statistical inference method appropriate for the data and identify the p-value and/or confidence interval after checking assumptions. In this study, we will focus on using randomization to generate a simulated p-value.
- **Communicate the results and answer the research question.** Using the p-value and confidence interval from the analysis, determine whether the data provide statistical evidence against the null hypothesis. Write a conclusion that addresses the research question.

5.3.4 Sleep Study

The Gallup Poll reports that 59% of Americans get seven or more hours of sleep at night. Recent research (Wahlstrom 2014) on the sleep-wake cycle of teenagers has indicated that biological changes associated with puberty may lead adolescents to require more sleep and feel sleepier at later times. One external factor that can impact their sleep is the early start time of high school. A study was conducted, involving a random sample of 446 students from public schools across three different states, to investigate whether these students sleep for at least 7 hours daily. The results revealed that 296 students sleep for at least 7 hours daily, while 150 sleep for less than 7 hours daily. Is there evidence that the true proportion of public high school students who sleep for at least 7 hours daily is greater than what was reported by the Gallup Poll?

Ask a research question

Is there evidence that the true proportion of public high school students who sleep for at least 7 hours daily is greater than what was reported by the Gallup Poll?

Design a study and collect data

Before using statistical inference methods, we must check that the cases are independent. The sample observations are independent if the outcome of one observation does not influence the outcome of another. One way this condition is met is if data come from a simple random sample of the target population.

1. Are the cases independent? Justify your answer.

Summarize and visualize the data

The following code reads in the data set and gives the number of public high school students in each level of the variable, whether the student gets at least 7 hours of sleep daily.

```
# Read in data set
sleep <- read.csv("data/LosingSleep.csv")
sleep %>% count(Outcome) # Count number in each category
```

```
#>   Outcome    n
#> 1      No 150
#> 2     Yes 296
```

$$\hat{p} = \frac{\text{number of successes}}{\text{total number of observational units}}$$

2. Using the R output and the formula given, calculate the summary statistic (sample proportion) to represent the research question. Use appropriate notation.
3. How far is the summary statistic from the value of 0.59? Is this a meaningful difference?

We cannot use the statistic alone to assess that this is a meaningful difference. Instead we will perform a hypothesis test to determine the strength of evidence of an effect.

Use statistical analysis methods to draw inferences from the data

Hypothesis testing Ideally, we would take many samples of size 446 (sample size) from the population of public high school students, calculating the proportion that sleep at least 7 hours daily. This would create a **sampling distribution** of statistics from the population. From this sampling distribution we would find the sample to sample variability, the population standard deviation. Rather than gathering all these samples from the population we will create a Null Distribution to estimate the population standard deviation. This way we can discover if the sample statistic is likely to occur under the assumption the null hypothesis is true.

When performing a hypothesis test, we must first identify the null hypothesis. The null hypothesis is written about the parameter of interest, or the value that summarizes the variable in the population.

For this study, the parameter of interest is the **true or population proportion of public high school students who sleep at least 7 hours daily**.

4. What assumption are we making about the true proportion of public high school students who sleep at least 7 hours daily in the null hypothesis?

The **alternative hypothesis** is the claim to be tested and the direction of the claim (less than, greater than, or not equal to) is based on the research question.

5. Write out the alternative hypothesis in words.

Notation:

Remember that when utilizing a hypothesis test, we are evaluating two competing possibilities if the observed statistic is unusual. For this study the **two possibilities** are either...


- The true proportion of public high school students who sleep at least 7 hours daily is 0.59 and our results just occurred by random chance; or,
- The true proportion of public high school students who sleep at least 7 hours daily is greater than 0.59 and our results reflect this.

Notice that these two competing possibilities represent the null and alternative hypotheses.


We will now simulate a one sample of a **null distribution** of sample proportions. The null distribution is created under the assumption the null hypothesis is true. In this case, we assume the true proportion of public high school students who sleep at least 7 hours daily is 0.59, so we will create 1000 (or more) different simulations of 446 under this assumption.

We will use the `one_proportion_test()` function in R (in the `catstats` package) to simulate the null distribution of sample proportions and compute a p-value.

```
set.seed(216)
one_proportion_test(probability_success = 0.59, # Null hypothesis value
  sample_size = 446, # Enter sample size
  number_repetitions = 1000, # Enter number of simulations
  as_extreme_as = 0.664, # Observed statistic
  direction = "greater", # Specify direction of alternative hypothesis
  summary_measure = "proportion") # Reporting proportion or number of successes?
```




6. Write down the standard deviation of the null distribution. This value estimates the sample to sample variability in the population of public high school students.
7. Explain why the null distribution is centered around 0.59.
8. Interpret the p-value in the context of the problem.



9. How much evidence does the p-value provide against the null hypothesis?

Estimation To create an interval estimate of plausible values for the parameter of interest, we will use the `one_proportion_bootstrap_CI()` function in R (in the `catstats` package) to simulate the bootstrap distribution of sample proportions and calculate a confidence interval.

```
set.seed(216)
one_proportion_bootstrap_CI(sample_size = 446, # Sample size
                             number_successes = 296, # Observed number of successes
                             number_repetitions = 1000, # Number of bootstrap samples to use
                             confidence_level = 0.95) # Confidence level as a decimal
```



10. Explain why the bootstrap distribution is centered around the value of 0.664.

11. Interpret the confidence interval in context of the problem.

12. Is the null value contained within the 95% confidence interval? Explain how this indicates that there is evidence that the true proportion of public high school students who sleep for at least 7 hours daily is greater than what was reported by the Gallup Poll.

13. What is the purpose of completing a hypothesis test?

14. What is the purpose of finding a confidence interval?

5.3.5 Take-home messages

1. A hypothesis test determines the strength of evidence of an effect. The smaller the p-value the more evidence the statistic provides against the null hypothesis.
2. A confidence interval gives an interval estimate for the parameter of interest - a range of plausible values.
3. The direction of the alternative hypothesis depends on the research question.

5.3.6 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

5.4 Activity 4: Helper-Hinderer Part 1 — Simulation-based Hypothesis Test

5.4.1 Learning outcomes

- Identify the two possible explanations (one assuming the null hypothesis and one assuming the alternative hypothesis) for a relationship seen in sample data.
- Given a research question involving a single categorical variable, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Describe and perform a simulation-based hypothesis test for a single proportion.

5.4.2 Terminology review

In today's activity, we will work through a simulation-based hypothesis testing for a single categorical variable. Some terms covered in this activity are:

- Parameter of interest
- Null hypothesis
- Alternative hypothesis
- Simulation

To review these concepts, see Chapters 9 & 14 in your textbook.

5.4.3 Steps of the statistical investigation process

We will work through a five-step process to complete a hypothesis test for a single proportion, first introduced in the activity in week 1.

- **Ask a research question** that can be addressed by collecting data. What are the researchers trying to show?
- **Design a study and collect data.** This step involves selecting the people or objects to be studied and how to gather relevant data on them.
- **Summarize and visualize the data.** Calculate summary statistics and create graphical plots that best represent the research question.
- **Use statistical analysis methods to draw inferences from the data.** Choose a statistical inference method appropriate for the data and identify the p-value and/or confidence interval after checking assumptions. In this study, we will focus on using randomization to generate a simulated p-value.
- **Communicate the results and answer the research question.** Using the p-value and confidence interval from the analysis, determine whether the data provide statistical evidence against the null hypothesis. Write a conclusion that addresses the research question.

5.4.4 Helper-Hinderer

A study by Hamblin, Wynn, and Bloom reported in Nature (Hamblin, Wynn, and Bloom 2007) was intended to check young kids' feelings about helpful and non-helpful behavior. Non-verbal infants ages 6 to 10 months were shown short videos with different shapes either helping or hindering the climber. As a class we will watch this short video to see how the experiment was run: <https://youtu.be/anCaGBsBOxM>. Researchers were hoping to assess: Are infants more likely to choose the helper toy over the hinderer toy? In the study, of the 16 infants age 6 to 10 months, 14 chose the *helper* toy and 2 chose the *hinderer* toy.

In this study, the **observational units are the infants ages 6 to 10 months**. The **variable measured on each observational unit (infant) is whether they chose the helper or the hinderer toy**. This is a categorical variable so we will be assessing the proportion of infants ages 6 to 10 months that choose the helper toy. Choosing the helper toy in this study will be considered a success.

Ask a research question

1. Identify the research question for this study. What are the researchers hoping to show?

Design a study and collect data

Before using statistical inference methods, we must check that the cases are independent. The sample observations are independent if the outcome of one observation does not influence the outcome of another. One way this condition is met is if data come from a simple random sample of the target population.

2. Are the cases independent? Justify your answer.

Summarize and visualize the data

The following code reads in the data set and gives the number of infants in each level of the variable, whether the infant chose the helper or the hinderer. Remember to visually display this data we can use either a frequency bar plot or a relative frequency bar plot.

```
# Read in data set
infants <- read.csv("https://math.montana.edu/courses/s216/data/infantchoice.csv")
infants %>% count(choice) # Count number in each choice category

#>      choice    n
#> 1    helper  14
#> 2 hinderer   2
```

$$\hat{p} = \frac{\text{number of successes}}{\text{total number of observational units}}$$

3. Using the R output and the formula given, calculate the summary statistic (sample proportion) to represent the research question. Recall that **choosing the helper toy** is a considered a success. Use appropriate notation.
4. Sketch a **relative frequency** bar plot of these data.

We cannot assess whether infants are more likely to choose the helper toy based on the statistic and plot alone. The next step is to analyze the data by using a hypothesis test to discover if there is evidence against the null hypothesis.

Use statistical analysis methods to draw inferences from the data

When performing a hypothesis test, we must first identify the null hypothesis. The null hypothesis is written about the parameter of interest, or the value that summarizes the variable in the population.

For this study, the parameter of interest is the **true or population proportion of infants ages 6–10 months who will choose the helper toy**.

If the children are just randomly choosing the toy, we would expect half (0.5) of the infants to choose the helper toy. This is the null value for our study.

5. Using the parameter of interest given above, write out the null hypothesis in words. That is, what do we assume to be true about the parameter of interest when we perform our simulation?

The notation used for a population proportion (or probability, or true proportion) is π . Since this summarizes a population, it is a parameter. When writing the **null hypothesis** in notation, we set the parameter equal to the null value, $H_0 : \pi = \pi_0$.

6. Write the null hypothesis in notation using the null value of 0.5 in place of π_0 in the equation given on the previous page.

The **alternative hypothesis** is the claim to be tested and the direction of the claim (less than, greater than, or not equal to) is based on the research question.

7. Based on the research question from question 1, are we testing that the parameter is greater than 0.5, less than 0.5 or different than 0.5?
8. Write out the alternative hypothesis in notation.

Remember that when utilizing a hypothesis test, we are evaluating two competing possibilities. For this study the **two possibilities** are either...

- The true proportion of infants who choose the helper is 0.5 and our results just occurred by random chance; or,
- The true proportion of infants who choose the helper is greater than 0.5 and our results reflect this.

Notice that these two competing possibilities represent the null and alternative hypotheses.

We will now simulate a one sample of a **null distribution** of sample proportions. The null distribution is created under the assumption the null hypothesis is true. In this case, we assume the true proportion of infants who choose the helper is 0.5, so we will create 1000 (or more) different simulations of 16 infants under this assumption.

Let's think about how to use a coin to create one simulation of 16 infants under the assumption the null hypothesis is true. Let heads equal infant chose the helper toy and tails equal infant chose the hinderer toy.

9. How many times would you flip a coin to simulate the sample of infants?
10. Flip a coin 16 times recording the number of times the coin lands on heads. This represents one simulated sample of 16 infants randomly choosing the toy. Calculate the proportion of coin flips that resulted in heads.
11. Is the value from question 10 closer to 0.5, the null value, or closer to the sample proportion, 0.875?

Report the number of coin flips you got in the Google sheet on D2L.

12. Sketch the graph created by your instructor of the proportion of heads out of 16 coin flips.
13. Circle the observed statistic (value from question 3) on the distribution shown above. Where does this statistic fall in this distribution: Is it near the center of the distribution (near 0.5) or in one of the tails of the distribution?

14. Is the observed statistic likely to happen or unlikely to happen if the true proportion of infants who choose the helper is 0.5? Explain your answer using the plot.

In the next class, we will continue to assess the strength of evidence against the null hypothesis by using a computer to simulate 1000 samples when we assume the null hypothesis is true.

5.4.5 Take-home messages

1. In a hypothesis test we have two competing hypotheses, the null hypothesis and the alternative hypothesis. The null hypothesis represents either a skeptical perspective or a perspective of no difference or no effect. The alternative hypothesis represents a new perspective such as the possibility that there has been a change or that there is a treatment effect in an experiment.
2. In a simulation-based test, we create a distribution of possible simulated statistics for our sample if the null hypothesis is true. Then we see if the calculated observed statistic from the data is likely or unlikely to occur when compared to the null distribution.
3. To create one simulated sample on the null distribution for a sample proportion, spin a spinner with probability equal to π_0 (the null value), n times or draw with replacement n times from a deck of cards created to reflect π_0 as the probability of success. Calculate and plot the proportion of successes from the simulated sample.

5.4.6 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

5.5 Activity 5: Helper-Hinderer (continued)

5.5.1 Learning outcomes

- Describe and perform a simulation-based hypothesis test for a single proportion.
- Interpret and evaluate a p-value for a simulation-based hypothesis test for a single proportion.
- Explore what a p-value represents

5.5.2 Steps of the statistical investigation process

In today's activity we will continue with steps 4 and 5 in the statistical investigation process. We will continue to assess the Helper-Hinderer study from last class.

- **Ask a research question** that can be addressed by collecting data. What are the researchers trying to show?
- **Design a study and collect data.** This step involves selecting the people or objects to be studied and how to gather relevant data on them.
- **Summarize and visualize the data.** Calculate summary statistics and create graphical plots that best represent the research question.
- **Use statistical analysis methods to draw inferences from the data.** Choose a statistical inference method appropriate for the data and identify the p-value and/or confidence interval after checking assumptions. In this study, we will focus on using randomization to generate a simulated p-value.
- **Communicate the results and answer the research question.** Using the p-value and confidence interval from the analysis, determine whether the data provide statistical evidence against the null hypothesis. Write a conclusion that addresses the research question.

5.5.3 Helper-Hinderer

A study by Hamblin, Wynn, and Bloom reported in Nature (Hamblin, Wynn, and Bloom 2007) was intended to check young kids' feelings about helpful and non-helpful behavior. Non-verbal infants ages 6 to 10 months were shown short videos with different shapes either helping or hindering the climber. As a class we will watch this short video to see how the experiment was run: <https://youtu.be/anCaGBsBOxM>. Researchers were hoping to assess: Are infants more likely to choose the helper toy over the hinderer toy? In the study, of the 16 infants age 6 to 10 months, 14 chose the *helper* toy and 2 chose the *hinderer* toy.

1. Report the sample proportion (summary statistic) calculated in the previous activity.
2. Write the alternative hypothesis in words in context of the problem. Remember the direction we are testing is dependent on the research question.

Today, we will use the computer to simulate a null distribution of 1000 different samples of 16 infants, plotting the proportion who chose the helper in each sample, based on the assumption that the true proportion of infants who choose the helper is 0.5 (or that the null hypothesis is true).

To use the computer simulation, we will need to enter the

- assumed “probability of success” (π_0),
 - “sample size” (the number of observational units or cases in the sample),
 - “number of repetitions” (the number of samples to be generated),
 - “as extreme as” (the observed statistic), and
 - the “direction” (matches the direction of the alternative hypothesis).
3. What values should be entered for each of the following into the one proportion test to create 1000 simulations?
- Probability of success:
 - Sample size:
 - Number of repetitions:
 - As extreme as:
 - Direction ("greater", "less", or "two-sided"):

We will use the `one_proportion_test()` function in R (in the `catstats` package) to simulate the null distribution of sample proportions and compute a p-value. Using the provided R script file, fill in the values/words for each `xx` with your answers from question 3 in the one proportion test to create a null distribution with 1000 simulations. Then highlight and run lines 1–16.


```
one_proportion_test(probability_success = xx, # Null hypothesis value
  sample_size = xx, # Enter sample size
  number_repetitions = 1000, # Enter number of simulations
  as_extreme_as = xx, # Observed statistic
  direction = "xx", # Specify direction of alternative hypothesis
  summary_measure = "proportion") # Reporting proportion or number of successes?
```

4. Sketch the null distribution created from the R code here.
5. Around what value is the null distribution centered? Why does that make sense?

6. Circle the observed statistic (value from question 1) on the distribution you drew in question 4. Where does this statistic fall in the null distribution: Is it near the center of the distribution (near 0.5) or in one of the tails of the distribution?
7. Is the observed statistic likely to happen or unlikely to happen if the true proportion of infants who choose the helper is 0.5? Explain your answer using the plot.
8. Using the simulation, what is the proportion of simulated samples that generated a sample proportion at the observed statistic or greater, if the true proportion of infants who choose the helper is 0.5? *Hint: Look under the simulation.*

The value in question 8 is the **p-value**. The smaller the p-value, the more evidence we have against the null hypothesis.

9. Using the following guidelines for the strength of evidence, how much evidence do the data provide against the null hypothesis? (Circle one of the five descriptions.)



Interpret the p-value

The p-value measures the probability that we observe a sample proportion as extreme as what was seen in the data or more extreme (matching the direction of the H_a) IF the null hypothesis is true.

10. What did we assume to create the null distribution?
11. What value did we compare to the null distribution to find the p-value?
12. In what direction (greater or less than) did we count from the statistic to find the number of simulations?

13. Fill in the blanks below to interpret the p-value.

We would observe a sample proportion of (value of the sample proportion)_____

or (greater, less, more extreme) _____

with a probability of (value of p-value) _____

If we assume (H_0 in context) _____.

Communicate the results and answer the research question

When we write a conclusion we answer the research question by stating how much evidence there is for the alternative hypothesis.

14. Write a conclusion in context of the study. How much evidence does the data provide in support of the alternative hypothesis?

5.5.4 Take-home messages

1. The null distribution is created based on the assumption the null hypothesis is true. We compare the sample statistic to the distribution to find the likelihood of observing this statistic.
2. The p-value measures the probability of observing the sample statistic or more extreme (in direction of the alternative hypothesis) if the null hypothesis is true.

5.5.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

5.6 Activity 6: Helper-Hinderer — Simulation-based Confidence Interval

5.6.1 Learning outcomes

- Use bootstrapping to find a confidence interval for a single proportion.
- Interpret a confidence interval for a single proportion.

5.6.2 Terminology review

In today's activity, we will introduce simulation-based confidence intervals for a single proportion. Some terms covered in this activity are:

- Parameter of interest
- Bootstrapping
- Confidence interval

To review these concepts, see Chapters 10 & 14 in your textbook.

5.6.3 Helper-Hinderer

In the last class, we found very strong evidence that the true proportion of infants who will choose the helper character is greater than 0.5. But what *is* the true proportion of infants who will choose the helper character? We will use this same study to estimate this parameter of interest by creating a confidence interval.

As a reminder: A study by Hamblin, Wynn, and Bloom reported in *Nature* (Hamblin, Wynn, and Bloom 2007) was intended to check young kids' feelings about helpful and non-helpful behavior. Non-verbal infants ages 6 to 10 months were shown short videos with different shapes either helping or hindering the climber. Researchers were hoping to assess: Are infants more likely to preferentially choose the helper toy over the hinderer toy? In the study, of the 16 infants age 6 to 10 months, 14 chose the *helper* toy and 2 chose the *hinderer* toy.

A **point estimate** (our observed statistic) provides a single plausible value for a parameter. However, a point estimate is rarely perfect; usually there is some error in the estimate. In addition to supplying a point estimate of a parameter, a next logical step would be to provide a plausible *range* of values for the parameter. This plausible range of values for the population parameter is called an **interval estimate** or **confidence interval**.

Activity intro

1. What is the value of the point estimate?
2. If we took another random sample of 16 infants, would we get the exact same point estimate? Explain why or why not.

In today's activity, we will use bootstrapping to find a 95% confidence interval for π , the parameter of interest.

3. In your own words, explain the bootstrapping process.

Use statistical analysis methods to draw inferences from the data

4. Write out the parameter of interest in words, in context of the study. *Hint: this is the same as in Activity 4 and 5.*

To use the computer simulation to create a bootstrap distribution, we will need to enter the

- “sample size” (the number of observational units or cases in the sample),
 - “number of successes” (the number of cases that choose the helper character),
 - “number of repetitions” (the number of samples to be generated), and
 - the “confidence level” (which level of confidence are we using to create the confidence interval).
5. What values should be entered for each of the following into the simulation to create the bootstrap distribution of sample proportions to find a 95% confidence interval?
 - Sample size:
 - Number of successes:
 - Number of repetitions:
 - Confidence level (as a decimal):

We will use the `one_proportion_bootstrap_CI()` function in R (in the `catstats` package) to simulate the bootstrap distribution of sample proportions and calculate a confidence interval. Using the provided R script file, fill in the values/words for each `xx` with your answers from question 5 in the one proportion bootstrap confidence interval (CI) code to create a bootstrap distribution with 1000 simulations. Then highlight and run lines 1–9.

```
one_proportion_bootstrap_CI(sample_size = xx, # Sample size
                             number_successes = xx, # Observed number of successes
                             number_repetitions = 1000, # Number of bootstrap samples to use
                             confidence_level = 0.95) # Confidence level as a decimal
```

6. Sketch the bootstrap distribution created below.
-
-
-
-
-
-
-
-
-
-
7. What is the value at the center of this bootstrap distribution? Why does this make sense?

8. Explain why the two vertical lines are at the 2.5th percentile and the 97.5th percentile.
9. Report the 95% bootstrapped confidence interval for π . Use interval notation: (lower value, upper value).
10. Interpret the 95% confidence interval in context.

Communicate the results and answer the research question

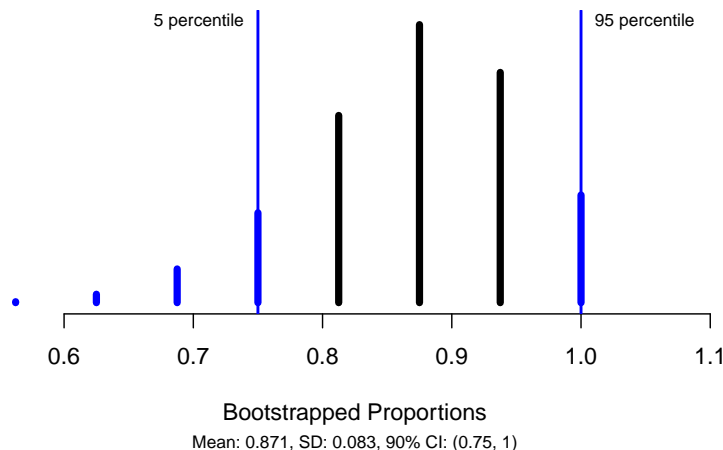
11. Is the value 0.5 (the null value) in the 95% confidence interval?

Explain how this indicates that the p-value provides strong evidence against the null.

Effect of confidence level

12. Suppose instead of finding a 95% confidence interval, we found a 90% confidence interval. Would you expect the 90% confidence interval to be narrower or wider? Explain your answer.
13. The following R code produced the bootstrap distribution with 1000 simulations that follows. Circle the value that changed in the code.

```
one_proportion_bootstrap_CI(sample_size = 16, # Sample size
                             number_successes = 14, # Observed number of successes
                             number_repetitions = 1000, # Number of bootstrap samples to use
                             confidence_level = 0.90) # Confidence level as a decimal
```



14. Report both the 95% confidence interval (question 9) and the 90% confidence interval (question 13). Is the 90% confidence interval narrower or wider than the 95% confidence interval?

15. Explain why the upper value of the confidence interval is truncated at 1.

16. Fill in the blanks below to write a paragraph summarizing the results of the study as if writing a press release.

Researchers were interested if infants observe social cues and would be more likely to choose the helper toy over the hinderer toy. In a sample of (sample size) _____ infants, (number of successes) _____ chose the helper toy. A simulation null distribution with 1000 simulations was created in RStudio. The p-value was found by calculating the proportion of simulations in the null distribution at the sample statistic of 0.875 and greater. This resulted in a p-value of (value of p-value) _____. We would observe a sample proportion of (value of the sample proportion) _____ or (greater, less, more extreme) _____ with a probability of (value of p-value) _____.

IF we assume (H_0 in context) _____. Based on this p-value, there is (very strong/little to no) _____ evidence that the (sample/true) _____ proportion of infants age 6 to 10 months who will choose the helper toy is (greater than, less than, not equal to) _____ 0.5. In addition, a 95% confidence interval was found for the parameter of interest. We are 95% confident that the (true/sample) _____ proportion of infants age 6 to 10 months who will choose the helper toy is between (lower value) _____ and (upper value) _____. The results of this study can be generalized to (all infants age 6 to 10 months/infants similar to those in this study) _____ as the researchers (did/did not) _____ select a random sample.

5.6.4 Take-home messages

1. The goal in a hypothesis test is to assess the strength of evidence for an effect, while the goal in creating a confidence interval is to determine how large the effect is. A **confidence interval** is a range of *plausible* values for the parameter of interest.
2. A confidence interval is built around the point estimate or observed calculated statistic from the sample. This means that the sample statistic is always the center of the confidence interval. A confidence interval includes a measure of sample to sample variability represented by the **margin of error**.
3. In simulation-based methods (bootstrapping), a simulated distribution of possible sample statistics is created showing the possible sample-to-sample variability. Then we find the middle X percent of the distribution around the sample statistic using the percentile method to give the range of values for the confidence interval. This shows us that we are $X\%$ confident that the parameter is within this range, where X represents the level of confidence.
4. When the null value is within the confidence interval, it is a plausible value for the parameter of interest; thus, we would find a larger p-value for a hypothesis test of that null value. Conversely, if the null value is NOT within the confidence interval, we would find a small p-value for the hypothesis test and strong evidence against this null hypothesis.
5. To create one simulated sample on the bootstrap distribution for a sample proportion, label n cards with the original responses. Draw with replacement n times. Calculate and plot the resampled proportion of successes.

5.6.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

Inference for a Single Categorical Variable: Theory-based Methods + Errors and Power

6.1 Reading Guide: Categorical Inference

Textbook Chapter 11: Inference with mathematical models

Reminders from previous sections

n_1 = sample size of group 1

n_2 = sample size of group 2

\bar{x} = sample mean

s = sample standard deviation

μ = population mean

σ = population standard deviation

General steps of a hypothesis test:

1. Frame the research question in terms of hypotheses.
2. Collect and summarize data using a test statistic.
3. Assume the null hypothesis is true, and simulate or mathematically model a null distribution for the test statistic.
4. Compare the observed test statistic to the null distribution to calculate a p-value.
5. Make a conclusion based on the p-value and write the conclusion in context.

Parameter: a value summarizing a variable(s) for a population.

Statistic: a value summarizing a variable(s) for a sample.

Hypothesis test: a process to determine how strong the evidence of an effect is. Also called a ‘significance test’.

Simulation-based method: Simulate lots of samples of size n under assumption of the null hypothesis, then find the proportion of the simulations that are at least as extreme as the observed sample statistic.

Theory-based method: Develop a mathematical model for the sampling distribution of the statistic under the null hypothesis and use the model to calculate the probability of the observed sample statistic (or one more extreme) occurring.

Null hypothesis (H_0): the skeptical perspective; no difference; no change; no effect; random chance; what the researcher hopes to prove is **wrong**.

Alternative hypothesis (H_A): the new perspective; a difference/increase/decrease; an effect; not random chance; what the researcher hopes to prove is **correct**.

Null value: the value of the parameter when we assume the null hypothesis is true (labeled as $parameter_0$).

P-value: probability of seeing the observed sample data, or something more extreme, assuming the null hypothesis is true.

\Rightarrow Lower the p-value the stronger the evidence AGAINST the null hypothesis and FOR the alternative hypothesis.

Significance level (α): a threshold used to determine if a p-value provides enough evidence to reject the null hypothesis or not.

Common levels of α include 0.01, 0.05, and 0.10.

Statistically significant: results are considered statistically significant if the p-value is below the significance level.

Confidence interval: a process to determine how large an effect is; a range of plausible values for the parameter. Also called 'estimation'.

Vocabulary

Central Limit Theorem:

Sampling distribution:

Normal distribution (Also known as: normal curve, normal model, Gaussian distribution):

Notation:

Standard normal distribution:

Notation:

Z-score:

Xth percentile:

68-95-99.7 rule:

Standard error of a statistic:

Standard deviation of a statistic:

Margin of error:

Notes

The two general conditions for the sampling distribution for a sample proportion (or difference in sample proportions) to be approximately normally distributed are:

1)

2)

Interpretation of a Z-score:

True or False: The more unusual observation will be the observation with the largest Z-score.

Approximately what percent of a normal distribution is in the interval

(mean – standard deviation, mean + standard deviation):

(mean – 2×(standard deviation), mean + 2×(standard deviation)):

(mean – 3×(standard deviation), mean + 3×(standard deviation)):

Given a mean and standard deviation, what function in R would help us find the percent of the normal distribution above (or below) a specific value?

Given a mean and standard deviation, what function in R would help us find the value at a given percentile?

How is the standard deviation of a statistic ($SD(statistic)$) different from the standard error of a statistic ($SE(statistic)$)?

How is the standard deviation of a statistic ($SD(statistic)$) different from the standard deviation of a sample (s)?

Formulas

$Z =$

$SD(\hat{p}) =$

General form of a theory-based confidence interval =

General form for margin of error =

R coding

Calculating normal probabilities When using the `pnorm()` R function, you will need to enter values for the arguments `mean`, `sd`, and `q` to match the question.

```
pnorm(mean = mu, sd = sigma, q = x, lower.tail = TRUE)
```

This function will return the proportion of the $N(\mu, \sigma)$ distribution which is *below* the value x .

Example: `pnorm(mean = 5, sd = 2, q = 3, lower.tail = TRUE)` will give us the proportion of a $N(5, 2)$ distribution which is below 3, which equals 0.159:

```
pnorm(mean = 5, sd = 2, q = 3, lower.tail = TRUE)
#> [1] 0.1586553
```

Changing to `lower.tail = FALSE` will give the proportion of the distribution which is *above* the value x .


```
pnorm(mean = 5, sd = 2, q = 3, lower.tail = FALSE)
#> [1] 0.8413447
```

Displaying normal probabilities When using the `normTail()` R function, you will need to enter values for the arguments `m`, `s`, and `L` (or `U`) to match the question.

```
normTail(m = mu, s = sigma, L = x)
```

This function (in the `openintro` package) will plot a $N(\mu, \sigma)$ distribution and shade the area that is below the value x .

Example: `normTail(m = 5, s = 2, L = 3)` creates the plot pictured below.



Changing `L` to `U` will shade the area *above* x .

Example: `normTail(m = 5, s = 2, U = 3)` plots a $N(5, 2)$ distribution with the area above 3 shaded.

Calculating normal percentiles When using the `qnorm()` R function, you will need to enter values for the arguments `mean`, `sd`, and `p` to match the question.

```
qnorm(mean = mu, sd = sigma, p = x, lower.tail = TRUE)
```

This function will return the value on the $N(\mu, \sigma)$ distribution which has x area of the distribution *below* it.

Example: `qnorm(mean = 5, sd = 2, p = 0.159, lower.tail = TRUE)` will give us the value on a $N(5,2)$ distribution which has 0.159 (15.9%) of the distribution below it, which equals 3 (from the R output above).

Changing to `lower.tail = FALSE` will give the value which has x area of the distribution *above* it.

We would recommend you work through each of the examples in Section 5.2.4 using R.

Section 14.3: Theory-based inferential methods for π

Vocabulary

Reminders from previous sections

n = sample size

\hat{p} = sample proportion

π = population proportion

General steps of a hypothesis test:

1. Frame the research question in terms of hypotheses.
2. Collect and summarize data using a test statistic.
3. Assume the null hypothesis is true, and simulate or mathematically model a null distribution for the test statistic.
4. Compare the observed test statistic to the null distribution to calculate a p-value.
5. Make a conclusion based on the p-value and write the conclusion in context.

Parameter: a value summarizing a variable(s) for a population.

Statistic: a value summarizing a variable(s) for a sample.

Sampling distribution: plot of statistics from 1000s of samples of the same size taken from the same population.

Standard deviation of a statistic: the variability of statistics from 1000s of samples; how far, on average, each statistic is from the true value of the parameter.

Standard error of a statistic: estimated standard deviation of a statistic.

Hypothesis test: a process to determine how strong the evidence of an effect is.

Also called a ‘significance test’.

Theory-based method: Develop a mathematical model for the sampling distribution of the statistic under the null hypothesis and use the model to calculate the probability of the observed sample statistic (or one more extreme) occurring.

Null hypothesis (H_0): the skeptical perspective; no difference; no change; no effect; random chance; what the researcher hopes to prove is **wrong**.

Alternative hypothesis (H_A): the new perspective; a difference/increase/decrease; an effect; not random chance; what the researcher hopes to prove is **correct**.

P-value: probability of seeing the observed sample data, or something more extreme, assuming the null hypothesis is true.

\Rightarrow Lower the p-value the stronger the evidence AGAINST the null hypothesis and FOR the alternative hypothesis.

Decision: a determination of whether to 'reject' or 'fail to reject' a null hypothesis based on a p-value and a pre-set level of significance.

Significance level (α): a threshold used to determine if a p-value provides enough evidence to reject the null hypothesis or not.

Common levels of α include 0.01, 0.05, and 0.10.

Statistically significant: results are considered statistically significant if the p-value is below the significance level.

Central Limit Theorem: For large sample sizes, the sampling distribution of a sample proportion (or mean) will be approximately normal (bell-shaped and symmetric).

Confidence interval: a process to determine how large an effect is; a range of plausible values for the parameter; also called 'estimation'.

Margin of error: the value that is added to and subtracted from the sample statistic to create a confidence interval; half the width of a confidence interval.

Vocabulary

Null standard error:

Standardized statistic:

Confidence level:

Notes

Conditions for the Central Limit Theorem to apply (for the sampling distribution of \hat{p} to be approximately normal)

Independence:

Checked by:

Success-failure condition:

Checked by:

How can we determine the value of z^* to use as the multiplier in a confidence interval?

In R, use `qnorm(mean = __, sd = __, p = __)`.

Select one answer in each set of parentheses: The higher the confidence level, the (larger/smaller) the multiplier, meaning the confidence interval will be (wider/narrower).

If the success-failure condition for the Central Limit Theorem is not met, what is the appropriate method of analysis? Select one: A. Theory-based approach B. Simulation based approach.

Formulas

$$SD(\hat{p}) =$$

Null standard error of the sample proportion:

$$SE_0(\hat{p}) =$$

Standardized statistic (in this case, standardized sample proportion):

$$Z =$$

Standard error of the sample proportion when we do not assume the null hypothesis is true:

$$SE(\hat{p}) =$$

Theory-based confidence interval for a sample proportion:

Margin of error of a confidence interval for a sample proportion:

Example: Payday loans

1. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
2. Write the null and alternative hypotheses in words.

3. Write the null and alternative hypotheses in notation.
4. Are the conditions met to use theoretical methods to analyze these data? Show your calculations to justify your answer.
5. Calculate the null standard error of the sample proportion.
6. What is the sample statistic presented in this example? What notation would be used to represent this value?
7. Calculate the standardized sample proportion (standardized statistic).
8. How can we calculate a p-value from the normal distribution for this example?
9. What was the p-value of the test?
10. What conclusion should the researcher make?
11. Are the results in this example statistically significant? Justify your answer.
12. Calculate the standard error of the sample proportion when we do not assume the null hypothesis is true.
13. Calculate the margin of error for a 95% confidence interval for π using 1.96 as the multiplier.
14. Calculate a 95% confidence interval for π using your margin of error calculated above.
15. Interpret the 95% confidence interval provided in the textbook.
16. Does the 95% confidence interval support the same conclusion as the p-value from the hypothesis test? Justify your answer.

Chapter 12: Errors, power, and practical importance

Reminders from previous sections

Significance level (α): a threshold used to determine if a p-value provides enough evidence to reject the null hypothesis or not.

Common levels of α include 0.01, 0.05, and 0.10.

Statistically significant: results are considered statistically significant if the p-value is below the significance level.

Vocabulary

Decision:

- If the p-value is small (less than or equal to the significance level), the decision will be to _____ the null hypothesis.
- If the p-value is large (greater than the significance level), the decision will be to _____ the null hypothesis.

Type 1 Error:

Type 2 Error:

Confirmation bias:

One-sided hypothesis tests:

Two-sided hypothesis tests:

Power:

Practical importance:

Notes

Fill in the following table with whether the decision was correct or not, and if not, what type of error was made.

	Test conclusion (based on data)	
Truth (unknown)	Reject null hyp.	Fail to reject null hyp.
H_0 is true		
H_A is true (H_0 is false)		

How are the significance level and Type 1 Error rate related?

How are the significance level and Type 2 Error rate related?

Explain the differences between a one-sided and two-sided hypothesis test.

How will the research questions differ?

How will the notation in the alternative hypothesis differ?

How does the p-value calculation differ?

How does the p-value in a two-sided test compare to the p-value in a one-sided test?

Should the default in research be a one-sided or two-sided hypothesis test? Explain why.

After collecting data, a researcher decides to change from a two-sided test to a one-sided test. Why is this a bad idea?

1. It _____ (increases/decreases) the chance of a Type 1 Error.
2. This can result in _____.

How are power and Type 1 Error rate related?

How are power and Type 2 Error rate related?

How can we increase the power of a test?

1. _____ (Increase/Decrease) the significance level
2. _____ (Increase/Decrease) the sample size

3. Change from a ____ (one/two)-sided to a ____ (one/two)-sided test
4. Have a ____ (larger/smaller) standard deviation of the statistic
5. Have the alternative parameter value ____ (closer/farther) from the null value

Results are likely to be statistically significant (but may not be practically important) if the sample size is ____ (large/small).

Results are unlikely to be statistically significant (but may be practically important) if the sample size is ____ (large/small).

Examples:

1. In the Martian Alphabet study section 9.1 of the textbook,
 - a. What was the p-value of the test?
 - b. At the 5% significance level, what decision would you make?
 - c. What type of error might have occurred in these data?
 - d. Interpret that error in the context of the problem.
2. In the Medical Consultant study in section 10.1 of the textbook,
 - a. What was the p-value of the test?
 - b. At the 5% significance level, what decision would you make?
 - c. What type of error might have occurred in these data?
 - d. Interpret that error in the context of the problem.
3. In the Payday Loans study section 14.3 of the textbook,
 - a. What was the p-value of the test?
 - b. At the 5% significance level, what decision would you make?

- c. What type of error might have occurred in these data?
- d. Interpret that error in the context of the problem.

6.2 Lecture Notes: Inference for One Categorical Variable using Theory-based Methods

Read Chapters 11, 12, and 13 and Sections 14.3 and 14.4 in the course textbook. Use the following videos to complete the lecture notes for Module 6.

6.2.1 Course Videos

- Chapter11
- 14.3TheoryTests
- 14.3TheoryIntervals
- Chapter12

Theory-based methods

Central limit theorem - Video Chapter11

The Central Limit Theorem tells us that the _____ distribution of a sample proportion (and sample mean and sample differences) will be approximately _____ if the sample size is _____.

The _____ of the distribution of sample proportions (sampling distribution) from thousands of samples will be bell-shaped/symmetric (Normal), if the sample size is large enough and the observations are _____.


- $\hat{p} \sim N(\pi, \sqrt{\frac{\pi \times (1-\pi)}{n}})$

Conditions of the CLT:

- Independence (*also must be met to use simulation methods*): the response for one observational unit will not influence another observational unit
- Large enough sample size:

Normal distribution:

- Bell-shaped and _____
- Standard normal distribution: $N(0, 1)$



Standardized statistic: Z - score

- $$Z = \frac{\text{statistic} - \text{null value}}{\text{standard error of the statistic}}$$


- Measures the _____ of standard _____ the statistic is from the null value

Example(s): Heights of Caucasian American adult males are roughly Normally distributed with a mean of 1.72 m and a standard deviation of 0.28 m. Find and interpret the z-score for a man who is 5'4" (1.626 m) tall. Round your answer to three decimal places.

Heights of Caucasian American adult females are roughly Normally distributed with a mean of 1.59 meters and a standard deviation of 0.22 meters. Which is more unusual: a 5'4" (1.626 m) tall male or a 5'9" (1.753 m) tall female?


In a Normal curve, the area under the curve is equal to 1, representing a probability. Therefore the shaded area represents the probability of a man being under 1.626 meters tall.

```
library(openintro)
normTail(m = 1.72, s = 0.28, L = 1.626)
pnorm(mean = 1.72, sd = 0.28, q = 1.626)
#> [1] 0.3685432
```




We can also reverse that order. Given a percentage, we can find the associated percentile, or quantile. Here we display calculating the value that cuts off the lower 0.75 proportion of male adult Caucasian heights using the `qnorm()` function.

```
qnorm(mean = 1.72, sd = 0.28, p = 0.75)
#> [1] 1.908857
normTail(m = 1.72, s = 0.28, L = 1.909)
```



68-95-99.7 Rule

- 68% of Normal distribution within 1 SD of the mean (mean $-$ SD, mean $+$ SD)
- 95% within (mean $-$ 2SD, mean $+$ 2SD)
- 99.7% within (mean $-$ 3SD, mean $+$ 3SD)



General steps of a hypothesis test

1. Write a research question and hypotheses.
2. Collect data and calculate a summary statistic.
3. Model a sampling distribution which assumes the null hypothesis is true.
4. Calculate a p-value.
5. Draw conclusions based on a p-value.

Example in Video 14.3 Theory Tests

Example: The American Red Cross reports that 10% of US residents eligible to donate blood actually do donate. A poll conducted on a representative of 200 Montana residents eligible to donate blood found that 33 had donated blood sometime in their life. Do Montana residents donate at a different rate than US population?

Hypotheses:

In notation:

H_0 :

H_A :

Parameter of interest:

Conditions for inference using theory-based methods:

- Independence:
 - The outcome of one observation does not influence the outcome of another.
 - Taking a random sample is one way to satisfy this condition.
- Large enough sample size:

Are the conditions met to analyze the blood donations data using theory-based methods?

To use theory-based methods to perform a hypothesis test:

- 1st: Calculate the standardized statistic
- 2nd: Find the area under the standard normal distribution at least as extreme as the standardized statistic

Equation for the standard error of the sample proportion assuming the null hypothesis is true:


- This value measures how far each possible sample statistic is from the null value, on average.

Equation for the standardized sample proportion:

- This value measures how many standard deviations the sample proportion is above/below the null value.

Calculate the standardized sample proportion of Montana residents that have donated blood sometime in their life.

- First calculate the standard error of the sample proportion assuming the null hypothesis is true
- Then calculate the Z score.



Interpret the standardized statistic

To find the p-value, find the area under the standard normal distribution at the standardized statistic and more extreme.

```
pnorm(3.064, lower.tail = FALSE)*2
```

```
#> [1] 0.002183989
```

Interpretation of the p-value:

- Statement about probability or proportion of samples
- Statistic (summary measure and value)
- Direction of the alternative
- Null hypothesis (in context)

Conclusion:

- Amount of evidence
- Parameter of interest
- Direction of the alternative hypothesis

Decision at a significance level of 0.05 ($\alpha = 0.05$):

Generalization:


- Can the results of the study be generalized to the target population?

Confidence interval - 14.3TheoryIntervals

- Interval of _____ values for the parameter of interest
- $CI = \text{statistic} \pm \text{margin of error}$

Theory-based method for a single categorical variable

- $CI = \hat{p} \pm (z^* \times SE(\hat{p}))$
- Multiplier (z^*) is the value at a certain _____ under the standard normal distribution



For a 95% confidence interval:

```
qnorm(0.975, lower.tail=TRUE)
```

```
#> [1] 1.959964
```

- When creating a confidence interval, we no longer assume the _____ hypothesis is true.
Use _____ to calculate the sample to sample variability, rather than π_0 .

Equation for the standard error of the sample proportion *NOT* assuming the null is true:

Example: Estimate the true proportion of Montana residents that have donated blood at least once in their life.
Find a 95% confidence interval:

Confidence interval interpretation:

- How confident you are (e.g., 90%, 95%, 98%, 99%)
- Parameter of interest
- Calculated interval
- Order of subtraction when comparing two groups

Interpreting confidence level

What does it mean to be 95% confident in a created confidence interval?

- Our goal is to only take one sample from the population to create a confidence interval.
- Based on the 68-95-99.7 rule, we know that approximately _____% of sample _____ will fall within _____ from the parameter.
- If we create 95% confidence intervals, _____% of samples will create a 95% _____ interval that will contain the _____ of interest.
- 95% of samples accurately _____ the parameter of interest
 - When we create one confidence interval, we are 95% _____ that we have a “good” sample that created a confidence interval that contains the _____ of interest.

Interpret the confidence **level** for the blood donation study.

Errors, power, and practical importance - Video Chapter12

Type 1 Error: _____ the null hypothesis, when the null is _____.

- Only can have a Type 1 Error when we make the _____ to _____ the null hypothesis.
- The probability of a Type 1 Error is α , the _____ level

Type 2 Error: _____ to reject the null hypothesis, when the null is _____.

- Only can have a Type 2 Error when we make the _____ to _____ to reject the null hypothesis.

Power: probability of _____ the null hypothesis, when the null is _____.

Increasing power:

- Increase _____
- Increase _____
- Use a _____ alternative vs. a _____ alternative
- Increase the _____ size, the _____ between the believed true value and the null value

Confirmation bias: looking for _____ that supports our ideas

- Always should write H_A based on the _____ prior to _____ collection!

Recall from the blood donation study, that we concluded there was very strong evidence that the true proportion of Montana residents who are eligible to donate blood differs from 0.10.

Since, we made the decision to _____ the null hypothesis, we have the possibility of a _____ error.

- What is the probability of this error?
- Write the error in context of the problem.

For each of the following changes to the blood donation study, determine whether the power of the test would increase or decrease.

- If we decreased the sample size from 200 to 100, power would _____.
- If we decreased the significance level from 0.05 to 0.01, power would _____.

- If we changed the research question to only asking if the probability a Montana resident eligible to donate blood actually does so is greater than 0.10, power would _____.
- This is an example of _____.

Practical importance

- A result can be _____ significant but not _____ important.
- Statistically significant: $p\text{-value} < \alpha$
 - Depends on the _____, the _____, and the selected _____ level.
- Practically important: the _____ seen in the data is meaningful and has _____ applications.
 - Depends on the _____ and subjective opinion.

Example: An Austrian study of heights of 507,125 military recruits reported that men born in spring were statistically significantly taller than men born in the fall ($p\text{-value} < 0.0001$). A confidence interval for the true difference in mean height between men born in spring and men born in fall was (0.598, 0.602) cm.

Is there statistical significance?

Is there practical importance?

6.2.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. What conditions must be met to use the Normal Distribution to approximate the sampling distribution of sampling proportions?
2. Should the conclusion include a population word like *true* or *long-run*? Explain your answer.
3. What is the difference between statistical significance and practical importance?

6.3 Guided Activity 8: Cashless Society

6.3.1 Learning outcomes

- Describe and perform a theory-based hypothesis test for a single proportion.
- Check the appropriate conditions to use a theory-based hypothesis test.
- Calculate and interpret the standardized sample proportion.
- Interpret and evaluate a p-value for a theory-based hypothesis test for a single proportion.
- Use the normal distribution to find the p-value.
- Calculate a theory-based confidence interval for a single proportion.
- Check the appropriate conditions to find a theory-based confidence interval.
- Interpret a confidence interval for a single proportion.
- Use the normal distribution to find the multiplier needed for a confidence interval

6.3.2 Terminology review

In this activity, we will introduce theory-based hypothesis tests and confidence intervals for a single categorical variable. Some terms covered in this activity are:

- Parameter of interest
- Standardized statistic
- Normal distribution
- p-value
- Parameter of interest
- Multiplier
- Normal distribution

To review these concepts, see Chapter 11 & 14 in your textbook.

Activities from module 5 covered simulation-based methods for hypothesis tests involving a single categorical variable. This module covers theory-based methods for testing a single categorical variable.

6.3.3 Cashless Society

According to published results in 2020, 34% of U.S. adults reported using only cashless purchases. In a 2022 Pew Research Center study (“More Americans Are Joining the ‘Cashless’ Economy” 2022), a random sample of 6034 U.S. adults was surveyed about their cashless purchase habits. In this 2022 study, 2474 U.S. adults reported using only cashless purchases in a typical week. Is there evidence that the rate of using only cashless purchases among U.S. adults differs from 0.34?

```
# Read in data set
cashless <- read.csv("data/cashless.csv")
cashless %>% count(Transaction_Habits) # Count number in each Stance category
```

```
#> Transaction_Habits    n
#> 1           Cashless 2474
#> 2       Some cash used 3560
```

In this study, the parameter of interest is the **population proportion of U.S. adults who report using only cashless purchases**.

1. What assumption are we making about the value for this parameter?
2. Write out the null hypothesis in words.
3. Write out the alternative hypothesis in notation.
4. Calculate the summary statistic (sample proportion) for this study. Use proper notation.

Theory-based methods

The sampling distribution of a single proportion — how that proportion varies from sample to sample — can be mathematically modeled using the normal distribution if certain conditions are met.

Conditions for the sampling distribution of \hat{p} to follow an approximate normal distribution:

- **Independence:** The sample's observations are independent, e.g., are from a simple random sample. (*Remember:* This also must be true to use simulation methods!)
- **Success-failure condition:** We *expect* to see at least 10 successes and 10 failures in the sample, $n\hat{p} \geq 10$ and $n(1 - \hat{p}) \geq 10$.

5. Verify that the independence condition is satisfied.
6. Is the success-failure condition met to model the data with the normal distribution? Explain your answer in context of the problem.




Figure 6.1: Simulated Null Distribution with Normal Distribution Overlay

To calculate the standardized statistic we use the general formula:

$$Z = \frac{\text{point estimate} - \text{null value}}{SE_0(\text{point estimate})}.$$

For a single categorical variable the standardized sample proportion is calculated using:

$$Z = \frac{\hat{p} - \pi_0}{SE_0(\hat{p})},$$

where the standard error is calculated using the null value:

$$SE_0(\hat{p}) = \sqrt{\frac{\pi_0 \times (1 - \pi_0)}{n}}$$

The standard error of the sample proportion measures the variability of possible sample proportions from the actual proportion. In other words, how far each possible sample proportion is from the actual proportion on average.

7. Calculate the null standard error of the sample proportion.

8. Using the null standard error of the sample proportion, calculate the standardized sample proportion (Z).

The standardized statistic is used as a ruler to measure how far the sample statistic is from the null value. Essentially, we are converting the sample proportion into a measure of standard errors to compare to the standard normal distribution.

The standardized statistic measures the *number of standard errors the sample statistic is from the null value*.




Figure 6.2: Standard Normal Curve

The value of 11.480 measures how far the sample proportion (0.410) is from the null value of 0.34. In other words, the sample proportion of U.S. adults that reported using only cashless purchases of 0.41 is 11.48 standard errors above the null value of 0.34.

We will use the `pnorm()` function in R to find the p-value. The value for Z was entered into the code below to get the p-value. Notice that we used `lower.tail = FALSE` to find the p-value. R will calculate the p-value *greater* than the value of the standardized statistic.

Notes:

- Use `lower.tail = TRUE` when doing a left-sided test.
- Use `lower.tail = FALSE` when doing a right-sided test.
- To find a two-sided p-value, use a left-sided test for negative Z or a right-sided test for positive Z, then multiply the value found by 2 to get the p-value.

```
pnorm(11.480, # Enter value of standardized statistic
      m=0, s=1, # Using the standard normal mean = 0, sd = 1
      lower.tail=FALSE) # Gives a p-value greater than the standardized statistic
```

```
#> [1] 8.314026e-31
```

9. Write a conclusion based on the p-value.

Theory-based confidence interval

To calculate a theory-based 95% confidence interval for π , we will first find the **standard error** of \hat{p} by plugging in the value of \hat{p} for π in $SD(\hat{p})$:

$$SE(\hat{p}) = \sqrt{\frac{\hat{p} \times (1 - \hat{p})}{n}}.$$

Note that we do not include a “0” subscript, since we are not assuming a null hypothesis.

10. Calculate the standard error of the sample proportion to find a 95% confidence interval.

The margin of error (ME) is the value of the z^* multiplier times the standard error of the statistic.

$$ME = z^* \times SE(\hat{p})$$

The z^* multiplier is the percentile of a standard normal distribution that corresponds to our confidence level. If our confidence level is 95%, we find the Z values that encompass the middle 95% of the standard normal distribution. If 95% of the standard normal distribution should be in the middle, that leaves 5% in the tails, or 2.5% in each tail.

The `qnorm()` function in R will tell us the z^* value for the desired percentile (in this case, $95\% + 2.5\% = 97.5\%$ percentile).

```
qnorm(0.975, lower.tail = TRUE) # Multiplier for 95% confidence interval
```

```
#> [1] 1.959964
```




Figure 6.3: Standard Normal Curve

11. Calculate the margin of error for the 95% confidence interval.

To find the confidence interval, we will add and subtract the **margin of error** to the point estimate:

$$\text{point estimate} \pm \text{margin of error}$$

$$\hat{p} \pm z^* \times SE(\hat{p})$$

12. Calculate the 95% confidence interval for the parameter of interest.

13. Interpret the 95% confidence interval in the context of the problem.

6.3.4 Take-home messages

1. Both simulation and theory-based methods can be used to find a p-value for a hypothesis test. In order to use theory-based methods we need to check that both the independence and the success-failure conditions are met.
2. The standardized statistic measures how many standard errors the statistic is from the null value. The larger the standardized statistic the more evidence there is against the null hypothesis.

6.3.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

6.4 Activity 7: Handedness of Male Boxers

6.4.1 Learning outcomes

- Describe and perform a theory-based hypothesis test for a single proportion.
- Check the appropriate conditions to use a theory-based hypothesis test.
- Calculate and interpret the standardized sample proportion.
- Interpret and evaluate a p-value for a theory-based hypothesis test for a single proportion.
- Use the normal distribution to find the p-value.

6.4.2 Terminology review

In this activity, we will introduce theory-based hypothesis tests for a single categorical variable. Some terms covered in this activity are:

- Parameter of interest
- Standardized statistic
- Normal distribution
- p-value

To review these concepts, see Chapter 11 & 14 in your textbook.

Activities from module 5 covered simulation-based methods for hypothesis tests involving a single categorical variable. This activity covers theory-based methods for testing a single categorical variable.

6.4.3 Handedness of male boxers

Left-handedness is a trait that is found in about 10% of the general population. Past studies have shown that left-handed men are over-represented among professional boxers (Richardson and Gilman 2019). The fighting claim states that left-handed men have an advantage in competition. In this random sample of 500 male professional boxers, we want to see if there is an over-prevalence of left-handed fighters. In the sample of 500 male boxers, 81 were left-handed.

```
# Read in data set
boxers <- read.csv("https://math.montana.edu/courses/s216/data/Male_boxers_sample.csv")
boxers %>% count(Stance) # Count number in each Stance category
```

```
#>      Stance    n
#> 1 left-handed  81
#> 2 right-handed 419
```

Review of summary statistics

1. Write out the parameter of interest in words, in context of the study.
2. Write out the null hypothesis in words.

3. Write out the alternative hypothesis in notation.
4. Give the value of the summary statistic (sample proportion) for this study. Use proper notation.

Theory-based methods

The sampling distribution of a single proportion — how that proportion varies from sample to sample — can be mathematically modeled using the normal distribution if certain conditions are met.

Conditions for the sampling distribution of \hat{p} to follow an approximate normal distribution:

- **Independence:** The sample's observations are independent, e.g., are from a simple random sample. (*Remember:* This also must be true to use simulation methods!)
- **Success-failure condition:** We *expect* to see at least 10 successes and 10 failures in the sample, $n\hat{p} \geq 10$ and $n(1 - \hat{p}) \geq 10$.

5. Verify that the independence condition is satisfied.
6. Is the success-failure condition met to model the data with the normal distribution? Explain your answer in context of the problem.

To calculate the standardized statistic we use the general formula

$$Z = \frac{\text{point estimate} - \text{null value}}{SE_0(\text{point estimate})}.$$

For a single categorical variable the standardized sample proportion is calculated using

$$Z = \frac{\hat{p} - \pi_0}{SE_0(\hat{p})},$$

where the standard error is calculated using the null value:

$$SE_0(\hat{p}) = \sqrt{\frac{\pi_0 \times (1 - \pi_0)}{n}}$$

For this study, the null standard error of the sample proportion is calculated using the null value, 0.1.

$$SE_0(\hat{p}) = \sqrt{\frac{0.1 \times (1 - 0.1)}{500}} = 0.013$$

Each sample proportion of male boxers that are left-handed is 0.013 from the true proportion of male boxers that are left-handed, on average.

7. Label the standard normal distribution shown below with the null value as the center value (below the value of zero). Label the tick marks to the right of the null value by adding 1 standard error to the null value to represent 1 standard error, 2 standard errors, and 3 standard errors from the null. Repeat this process to the left of the null value by subtracting 1 standard error for each tick mark.




Figure 6.4: Standard Normal Curve

8. Using the null standard error of the sample proportion, calculate the standardized sample proportion (Z). Mark this value on the standard normal distribution above.

The standardized statistic is used as a ruler to measure how far the sample statistic is from the null value. Essentially, we are converting the sample proportion into a measure of standard errors to compare to the standard normal distribution.

The standardized statistic measures the *number of standard errors the sample statistic is from the null value*.

9. Interpret the standardized sample proportion from question 8 in context of the problem.

We will use the `pnorm()` function in R to find the p-value. The value for Z was entered into the code below to get the p-value. Check that this answer matches what you calculated in question 7. Notice that we used `lower.tail = FALSE` to find the p-value. R will calculate the p-value *greater* than the value of the standardized statistic.

Notes:

- Use `lower.tail = TRUE` when doing a left-sided test.
- Use `lower.tail = FALSE` when doing a right-sided test.
- To find a two-sided p-value, use a left-sided test for negative Z or a right-sided test for positive Z , then multiply the value found by 2 to get the p-value.

```
pnorm(4.769, # Enter value of standardized statistic
      m=0, s=1, # Using the standard normal mean = 0, sd = 1
      lower.tail=FALSE) # Gives a p-value greater than the standardized statistic
```

```
#> [1] 9.257133e-07
```

10. Report the p-value obtained from the R output.
11. Write a conclusion based on the value of the p-value.

Theory-based confidence interval

To calculate a theory-based 95% confidence interval for π , we will first find the **standard error** of \hat{p} by plugging in the value of \hat{p} for π in $SD(\hat{p})$:

$$SE(\hat{p}) = \sqrt{\frac{\hat{p} \times (1 - \hat{p})}{n}}.$$

Note that we do not include a “0” subscript, since we are not assuming a null hypothesis.

12. Calculate the standard error of the sample proportion to find a 95% confidence interval.

We will calculate the margin of error and confidence interval in questions 4 and 5 of this activity. **The margin of error (ME)** is the value of the z^* multiplier times the standard error of the statistic.

$$ME = z^* \times SE(\hat{p})$$

The z^* multiplier is the percentile of a standard normal distribution that corresponds to our confidence level. If our confidence level is 95%, we find the Z values that encompass the middle 95% of the standard normal distribution. If 95% of the standard normal distribution should be in the middle, that leaves 5% in the tails, or 2.5% in each tail.

The `qnorm()` function in R will tell us the z^* value for the desired percentile (in this case, $95\% + 2.5\% = 97.5\%$ percentile).

- Enter the value of 0.975 for `xx` in the provided R script file.
- Highlight and run line 15. This will give the value of the multiplier for a 95% confidence interval.

```
qnorm(xx, lower.tail = TRUE) # Multiplier for 95% confidence interval
```

13. Report the value of the multiplier needed to calculate the 95% confidence interval for the true proportion of male boxers that are left-handed.

14. Fill in the normal distribution shown below to show how R found the z^* multiplier.




Figure 6.5: Standard Normal Curve

15. Calculate the margin of error for the 95% confidence interval.

To find the confidence interval, we will add and subtract the **margin of error** to the point estimate:

point estimate \pm margin of error

$$\hat{p} \pm z^* \times SE(\hat{p})$$

16. Calculate the 95% confidence interval for the parameter of interest.
17. Interpret the 95% confidence interval in the context of the problem.

6.4.4 Take-home messages

1. Both simulation and theory-based methods can be used to find a p-value for a hypothesis test. In order to use theory-based methods we need to check that both the independence and the success-failure conditions are met.
2. The standardized statistic measures how many standard errors the statistic is from the null value. The larger the standardized statistic the more evidence there is against the null hypothesis.

6.4.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

6.5 Activity 8: What does confidence mean?

6.5.1 Learning outcomes

- Explore what confidence means
- Interpret the confidence level
- Explore impact of sample size, direction of the alternative hypothesis, and value of the sample statistic on the p-value.

6.5.2 Terminology review

In this activity, we will explore what being 95% confidence means. Some terms covered in this activity are:

- Parameter of interest
- Two-sided vs. one-sided tests
- Confidence level

6.5.3 Handedness of male boxers continued

We will use the male boxer study to look at what confidence means.

Left-handedness is a trait that is found in about 10% of the general population. Past studies have shown that left-handed men are over-represented among professional boxers (Richardson and Gilman 2019). The fighting claim states that left-handed men have an advantage in competition. In this random sample of 500 male professional boxers, we want to see if there is an over-prevalence of left-handed fighters. In the sample of 500 male boxers, 81 were left-handed.

```
# Read in data set
boxers <- read.csv("https://math.montana.edu/courses/s216/data/Male_boxers_sample.csv")
boxers %>% count(Stance) # Count number in each Stance category
```

```
#>      Stance    n
#> 1 left-handed  81
#> 2 right-handed 419
```

What does *confidence* mean?

In the interpretation of a 95% confidence interval, we say that we are 95% confident that the parameter is within the confidence interval. Why are we able to make that claim? What does it mean to say “we are 95% confident”?

1. In the last activity we found a 95% confidence interval for the parameter of interest. As a class, determine a plausible value for the true proportion of male boxers that are left-handed. *Note: we are making assumptions about the population here. This is not based on our calculated data, but we will use this applet to better understand what happens when we take many, many samples from this believed population.*
2. Go to this website, <http://www.rossmanchance.com/ISIApplets.html> and choose ‘Simulating Confidence Intervals’. In the input on the left-hand side of the screen enter the value from question 1 for π (the true value), 500 for n , and 100 for ‘Number of intervals’. Click ‘sample’.
- In the graph on the bottom right, click on a green dot. Write down the confidence interval for this sample given on the graph on the left. Does this confidence interval contain the true value chosen in question 1?

- Now click on a red dot. Write down the confidence interval for this sample. Does this confidence interval contain the true value chosen in question 1?
 - How many intervals out of 100 contain π , the true value chosen in question 1? *Hint:* This is given to the left of the graph of green and red intervals.
3. Click on ‘sample’ nine more times. Write down the ‘Running Total’ for the proportion of intervals that contain π .
 4. Change the confidence level to 90%. What happened to the width of the intervals?
 5. Write down the **Running Total** for the proportion of intervals that contain π using a 90% confidence level.
 6. Interpret the level of confidence. *Hint:* What proportion of samples would we expect to give a confidence interval that contains the parameter of interest?

Impacts on the P-value

Suppose that we want to show that the true proportion of male boxers **differs** from that in the general population.

7. Write out the alternative hypothesis in notation for this new research question.
8. How would this impact the p-value?
9. How much evidence would this p-value provide against the null hypothesis?

Suppose instead of 500 male boxers the researchers only took a sample of 300 male boxers and found the same proportion ($\hat{p} = 0.182$) of male boxers that are left-handed. Since we are still assuming the same null value, 0.1, the standard error would be calculated as below:

$$SE_0(\hat{p}) = \sqrt{\frac{0.1 \times (1 - 0.1)}{300}} = 0.017$$

10. Calculate the standardized statistic for this new sample.

Use Rstudio to find the p-value for this new sample. Enter the value of the standardized statistic found in question 10 for xx in line 4. Highlight and run lines 4–6.

```
pnorm(xx, # Enter value of standardized statistic
      m=0, s=1, # Using the standard normal mean = 0, sd = 1
      lower.tail=FALSE) # Gives a p-value greater than the standardized statistic
```

11. How does the decrease in sample size affect the p-value?

Suppose another sample of 500 male boxers was taken and 68 were found to be left-handed. Since we are still assuming the same null value, 0.1, the standard error would be calculated as before:

$$SE_0(\hat{p}) = \sqrt{\frac{0.1 \times (1 - 0.1)}{500}} = 0.013$$

12. Calculate the standardized statistic for this new sample.

Use Rstudio to find the p-value for this new sample.

```
pnorm(xx, # Enter value of standardized statistic
      m=0, s=1 # Using the standard normal mean = 0, sd = 1
      lower.tail=FALSE) # Gives a p-value greater than the standardized statistic
```

13. How does a statistic closer to the null value affect the p-value?
14. Summarize how each of the following affected the p-value:
- a) Switching to a two-sided test.
 - b) Using a smaller sample size.
 - c) Using a sample statistic closer to the null value.

6.5.4 Take-home messages

1. If repeat samples of the same size are selected from the population, approximately 95% of samples will create a 95% confidence interval that contains the parameter of interest.
2. The p-value for a two-sided test is approximately two times the value for a one-sided test. A two-sided test provides less evidence against the null hypothesis.
3. The larger the sample size, the smaller the sample to sample variability. This will result in a larger standardized statistic and more evidence against the null hypothesis.
4. The farther the statistic is from the null value, the larger the standardized statistic. This will result in a smaller p-value and more evidence against the null hypothesis.

6.5.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

6.6 Activity 9: Errors and Power

6.6.1 Learning outcomes

- Explain Type I and Type 2 Errors in the context of a study.
- Explain the power of a test in the context of a study.
- Understand how changes in sample size, significance level, and the difference between the null value and the parameter value impact the power of a test.
- Understand how significance level impacts the probability of a Type 1 Error.
- Understand the relationship between the probability of a Type 2 Error and power.
- Be able to distinguish between practical importance and statistical significance.

6.6.2 Terminology review

In this activity, we will examine the possible errors that can be made based on the decision in a hypothesis test as well as factors influencing the power of the test. Some terms covered in this activity are:

- Significance level
- Type 1 Error
- Type 2 Error
- Power

To review these concepts, see Chapter 12 in the textbook.

6.6.3 ACL recovery

It is widely reported that the median recovery time for athletes who undergo surgery to repair a torn anterior cruciate ligament (ACL) is 8 months, indicating that 50% of athletes return to their sport within 8 months after an ACL surgery. Suppose a local physical therapy company hopes to advertise that their rehabilitation program can increase this percentage.

1. Write the parameter of interest (π) in words, in the context of this problem.
2. Use proper notation to write the null and alternative hypothesis the company would need to test in order to check their advertisement claim.

After determining hypotheses and prior to collecting data, researchers should set a **significance level** for a hypothesis test. The significance level, represented by α and most commonly 0.01, 0.05, or 0.10, is a cut-off for determining whether a p-value is small or not. The *smaller* the p-value, the *stronger* the evidence against the null hypothesis, so a p-value that is smaller than or equal to the significance level is strong enough evidence to *reject the null hypothesis*. Similarly, the *larger* the p-value, the *weaker* the evidence against the null hypothesis, so a p-value that is larger than the significance level does not provide enough evidence against the null hypothesis and the researcher would *fail to reject the null hypothesis*. Rejecting the null hypothesis or failing to reject the null hypothesis are the two **decisions** that can be made based on the data collected.

As you have already learned in this course, sample size of a study is extremely important. Often times, researchers will conduct what is called a power analysis to determine the appropriate sample size based on the goals of

their research, including a desired **power** of their test. Power is the probability of correctly rejecting the null hypothesis, or the probability of the data providing strong evidence against the null hypothesis *when the null hypothesis is false*.

The remainder of this lab will be spent investigating how different factors influence the power of a test, after which you will complete a power analysis for this physical therapy company.

- Navigate to <https://istats.shinyapps.io/power/>. Please note that this applet uses p_0 to represent the null value rather than π_0 .
- Use the scale under “Null Hypothesis value p_0 ” to change the value to your null value from question 2.
- Change the “Alternative Hypothesis” to the direction you wrote in question 2.
- Leave all boxes un-checked. Do not change the scales under “True value of p_0 ”, “Sample size n ”, or “Type I Error α ”

The red distribution you see is the scaled-Normal distribution representing the null distribution for this hypothesis test, if the sample size was 50 and the significance level was 0.05. This means the red distribution is showing the probability of each possible sample proportion of athletes who returned to their sport within 8 months (\hat{p}) if we assume the null hypothesis is true.

3. Based off this distribution and your alternative hypothesis, give one possible sample proportion which you think would lead to rejecting the null hypothesis. Explain how you decided on your value.
4. Check the box for “Show Critical Value(s) and Rejection Region(s)”. You will now see a vertical line on the plot indicating the *minimum* sample proportion which would lead to reject the null hypothesis. What is this value?
5. Notice that there are some sample proportions under the red line (when the null hypothesis is true) which would lead us to reject the null hypothesis. Give the range of sample proportions which would lead to rejecting the null hypothesis when the null hypothesis is true? What is the statistical name for this mistake?

Check the “Type I Error” box under **Display**. This should verify (or correct) your answer to question 5! The area shaded in red represents the probability of making a **Type 1 Error** in our hypothesis test. Recall that a Type 1 Error is when we reject the null hypothesis even though the null hypothesis is true. To reject the null hypothesis, the p-value, which was found assuming the null hypothesis is true, must be less than or equal to the significance level. Therefore the significance level is the maximum probability of rejecting the null hypothesis when the null hypothesis is true, so the significance level IS the probability of making a Type 1 Error in a hypothesis test!

6. Based on the current applet settings, What percent of the null distribution is shaded red (what is the probability of making a Type 1 Error)?

Let’s say this physical therapist company believes their program can get 70% of athletes back to their sport within 8 months of an ACL surgery. In the applet, set the scale under “True value of p ” to 0.7.

7. Where is the blue distribution centered?

The blue distribution that appears represents what the company believes, that 0.7 (not 0.5) is the true proportion of its clients who return to their sport within 8 months of ACL surgery. This blue distribution represents the idea that the **null hypothesis is false**.

8. Consider the definition of power provided earlier in this lab. Do you believe the power of the test will be an area within the blue distribution or red distribution? How do you know? What about the probability of making a Type 2 Error?

- Check the “Type II Error” and “Power” boxes under **Display**. This should verify (or correct) your answers to question 8! The area shaded in blue represents the probability of making a **Type 2 Error** in our hypothesis test (failing to reject the null hypothesis even though the null hypothesis is false). The area shaded in green represents the power of the test. Notice that the Type 1 and Type 2 Error rates and the power of the test are provided above the distribution.
9. Complete the following equation: Power + Type 2 Error Rate = . Explain why that equation makes sense. *Hint: Consider what power and Type 2 Error are conditional on.*

Now let’s investigate how changes in different factors influence the power of a test.

10. Using the same sample size and significance level, change the “True value of p ” to see the effect on Power.

True value of p	0.60	0.65	0.70	0.75	0.80
Power					

11. What is changing about the simulated distributions pictured as you change the “True value of p ”?
12. How does increasing the distance between the null and believed true probability of success affect the power of the test?
13. Using the same significance level, set the “True value of p ” to 0.7 and change the sample size to see the effect on Power.

Sample Size	20	40	50	60	80
Power					

14. What is changing about the simulated distributions pictured as you change the sample size?

15. How does increasing the sample size affect the power of the test?
16. Using the same “True value of p ”, set the sample size to 50 and change the “Type I Error α ” to see the effect on Power.

Type I Error α	0.01	0.03	0.05	0.10	0.15
Power					

17. What is changing about the simulated distributions pictured as you change the significance level?
18. How does increasing the significance level affect the power of the test?
19. Complete the power analysis for this physical therapy company. The company believes 70% of their patients will return to their sport within 8 months of ACL surgery. They want to limit the probability of a type 1 error to 10% and the probability of a type 2 error to 15%. What is the minimum number of athletes the company will need to collect data from in order to meet these goals? Use the applet to answer this question, then download your image created and upload the file to Gradescope.
20. Based on the goals outlined in question 19, which mistake below is the company more concerned about? In other words, which error were the researchers trying to minimize. Explain your answer.
- Not being able to advertise their ACL recovery program is better than average when their program really is better.
 - Advertising their ACL recovery program is better even though it is not.

6.6.4 Take-home messages

1. There is a possibility of Type I Error when we make the decision to reject the null hypothesis. Type I Error - reject the null hypothesis when the null hypothesis is true.
2. There is a possibility of Type II Error when we make the decision to fail to reject the null hypothesis. Type II Error - fail to reject the null hypothesis when the null hypothesis is false.
3. Increasing the sample size will increase the power of the test.

6.6.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today’s activity and material covered.

6.7 Module 6 Lab: Mixed Breed Dogs in the U.S.

6.7.1 Learning outcomes

- Determine whether simulation or theory-based methods of inference can be used.
- Analyze and interpret a study involving a single categorical variable.

6.7.2 Mixed Breed Dogs in the U.S.

The American Veterinary Medical Association estimated in 2010 that approximately 49% of dog owners in the U.S. own dogs that are classified as “mixed breed.” As part of a larger 2022 international study (Banton 2022) about overall dog health, survey participants were asked, among other things, to report whether their dog was purebred or a mixed breed. Seven hundred and fifty (750) dog owners from the U.S. were recruited to complete an online survey via an email indicating they had been randomly selected by Qualtrics (an “experience management” company that specializes in surveys). Three hundred sixty-four (364) out of 675 respondents from the U.S. reported they owned a mixed breed dog. Is there evidence that, in the last decade, the proportion of dog owners in the U.S. that own a mixed breed dog has changed from the value reported in 2010?

Activity intro

1. What is the value of the point estimate?

Use statistical analysis methods to draw inferences from the data

2. Write out the parameter of interest in words, in context of the study.

3. Write out the null and alternative hypotheses in notation.

H_0 :

H_A :

4. Can theory-based methods be used to analyze these data? Explain your answer.

To use the computer simulation, we will need to enter the

- assumed “probability of success” (π_0),
 - “sample size” (the number of observational units or cases in the sample),
 - “number of repetitions” (the number of samples to be generated),
 - “as extreme as” (the observed statistic), and
 - the “direction” (matches the direction of the alternative hypothesis).
5. What values should be entered for each of the following into the one proportion test to create 1000 simulations?
- Probability of success:
 - Sample size:
 - Number of repetitions:
 - As extreme as:
 - Direction ("greater", "less", or "two-sided"):

We will use the `one_proportion_test()` function in R (in the `catstats` package) to simulate the null distribution of sample proportions and compute a p-value. Using the provided R script file, fill in the values/words for each `xx` with your answers from question 5 in the one proportion test to create a null distribution with 1000 simulations. Then highlight and run lines 1–12.

```
one_proportion_test(probability_success = xx, # Null hypothesis value
  sample_size = xx, # Enter sample size
  number_repetitions = 1000, # Enter number of simulations
  as_extreme_as = xx, # Observed statistic
  direction = "xx", # Specify direction of alternative hypothesis
  summary_measure = "proportion") # Reporting proportion or number of successes?
```

6. Report the p-value from the study.

The z^* multiplier is the percentile of a standard normal distribution that corresponds to our confidence level. If our confidence level is 95%, we find the Z values that encompass the middle 95% of the standard normal distribution. If 95% of the standard normal distribution should be in the middle, that leaves 5% in the tails, or 2.5% in each tail.

The `qnorm()` function in R will tell us the z^* value for the desired percentile (in this case, $90\% + 5\% = 95\%$ percentile).

- Enter the value of 0.95 for `xx` in the provided R script file.
- Highlight and run line 17. This will give the value of the multiplier for a 90% confidence interval.

```
qnorm(xx, lower.tail = TRUE) # Multiplier for 90% confidence interval
```

7. Calculate the margin of error for a 90% confidence interval.

8. Calculate a 90% confidence interval.
9. What type of error may have occurred?
10. Describe this error in the context of the study?

Summarize the results of the study

11. Write a paragraph summarizing the results of the study. Be sure to describe:
 - Summary statistic and interpretation
 - Summary measure (in context)
 - Value of the statistic
 - Order of subtraction when comparing two groups
 - P-value and interpretation
 - Statement about probability or proportion of samples
 - Statistic (summary measure and value)
 - Direction of the alternative
 - Null hypothesis (in context)
 - Confidence interval and interpretation
 - How confident you are (e.g., 90%, 95%, 98%, 99%)
 - Parameter of interest
 - Calculated interval
 - Order of subtraction when comparing two groups
 - Conclusion (written to answer the research question)
 - Amount of evidence
 - Parameter of interest
 - Direction of the alternative hypothesis
 - Scope of inference
 - To what group of observational units do the results apply (target population or observational units similar to the sample)?
 - What type of inference is appropriate (causal or non-causal)?

Upload a copy of your group's paragraph to Gradescope.

Paragraph (continued):

Inference for Two Categorical Variables: Simulation-based Methods

7.1 Reading Guide: Hypothesis Testing for a Difference in Proportions

Textbook Chapter 15: Inference for comparing two proportions

Section 15.1: Randomization test for $H_0 : \pi_1 - \pi_2 = 0$) and Section 15.2: Bootstrap confidence interval for $\pi_1 - \pi_2$

Reminders from previous sections

n = sample size

\hat{p} = sample proportion

π = population proportion

General steps of a hypothesis test:

1. Frame the research question in terms of hypotheses.
2. Collect and summarize data using a test statistic.
3. Assume the null hypothesis is true, and simulate or mathematically model a null distribution for the test statistic.
4. Compare the observed test (standardized) statistic to the null distribution to calculate a p-value.
5. Make a conclusion based on the p-value and write the conclusion in context.

Parameter: a value summarizing a variable(s) for a population.

Statistic: a value summarizing a variable(s) for a sample.

Sampling distribution: plot of statistics from 1000s of samples of the same size taken from the same population.

Standard deviation of a statistic: the variability of statistics from 1000s of samples; how far, on average, each statistic is from the true value of the parameter.

Standard error of a statistic: estimated standard deviation of a statistic.

Hypothesis test: a process to determine how strong the evidence of an effect is.

Also called a ‘significance test’.

Simulation-based method: Simulate lots of samples of size n under assumption of the null hypothesis, then find the proportion of the simulations that are at least as extreme as the observed sample statistic.

Null hypothesis (H_0): the skeptical perspective; no difference; no change; no effect; random chance; what the researcher hopes to prove is **wrong**.

Alternative hypothesis (H_A): the new perspective; a difference/increase/decrease; an effect; not random chance; what the researcher hopes to prove is **correct**.

Null value: the value of the parameter when we assume the null hypothesis is true (labeled as $parameter_0$).

Null distribution: the simulated or modeled distribution of statistics (sampling distribution) we would expect to occur if the null hypothesis is true.

P-value: probability of seeing the observed sample data, or something more extreme, assuming the null hypothesis is true.

\Rightarrow Lower the p-value the stronger the evidence AGAINST the null hypothesis and FOR the alternative hypothesis.

Decision: a determination of whether to ‘reject’ or ‘fail to reject’ a null hypothesis based on a p-value and a pre-set level of significance.

Significance level (α): a threshold used to determine if a p-value provides enough evidence to reject the null hypothesis or not.

Common levels of α include 0.01, 0.05, and 0.10.

Statistically significant: results are considered statistically significant if the p-value is below the significance level.

Confidence interval: a process to determine how large an effect is; a range of plausible values for the parameter. Also called ‘estimation’.

Margin of error: the value that is added to and subtracted from the sample statistic to create a confidence interval; half the width of a confidence interval.

Bootstrapping: the process of drawing with replacement n times from the original sample.

Bootstrapped resample: a random sample of size n from the original sample, selected with replacement.

Bootstrapped statistic: the statistic recorded from the bootstrapped resample.

Confidence level: how confident we are that the confidence interval will capture the parameter.

Vocabulary

Randomization test:

Notes

In a randomization test involving two categorical variables,

how many cards will you need and how will the cards be labeled?

Why, in the randomization test, are the cards all shuffled together and randomly dealt into two new groups?

After shuffling, how many cards are dealt into each pile?

To create a single bootstrap resample for two categorical variables,
how many cards will you need and how will the cards be labeled?

What is done with the cards once they are labeled?

Interpretations of confidence level must include:

How do you determine if the results of a hypothesis test agree with a confidence interval?

How are the confidence level and the significance level related (for a two-sided test)?

Notation

Sample size of group 1:

Sample size of group 2:

Sample proportion of group 1:

Sample proportion of group 2:

Population proportion of group 1:

Population proportion of group 2:

Example: Gender discrimination

1. What is the research question?
2. What are the observational units?
3. What type of study design was used? Justify your answer.

4. What is the appropriate scope of inference for these data?
5. What is the sample statistic presented in this example? What notation would be used to represent this value?
6. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
7. Write the null and the alternative hypotheses in words.
8. Write the null and the alternative hypotheses in notation.
9. How could we use cards to simulate **one** sample *which assumes the null hypothesis is true*? How many blue cards — to represent what? How many red cards — to represent what? What would we do with the cards? What would you record once you have a simulated sample?
10. How can we calculate a p-value from the simulated null distribution for this example?
11. What was the p-value of the test?
12. At the 5% significance level, what decision would you make?
13. What conclusion should the researcher make?
14. Are the results in this example statistically significant? Justify your answer.

Example: Opportunity cost

1. What is the research question?
2. What are the observational units?
3. What type of study design was used? Justify your answer.
4. What is the appropriate scope of inference for these data?

5. What is the sample statistic presented in this example? What notation would be used to represent this value?
6. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
7. Write the null and the alternative hypotheses in words.
8. Write the null and the alternative hypotheses in notation.
9. How could we use cards to simulate **one** sample *which assumes the null hypothesis is true*? How many blue cards — to represent what? How many red cards — to represent what? What would we do with the cards? What would you record once you have a simulated sample?
10. How can we calculate a p-value from the simulated null distribution for this example?
11. What was the p-value of the test?
12. Interpret the p-value in the context of the problem.
13. At the 5% significance level, what decision would you make?
14. What conclusion should the researcher make?
15. Are the results in this example statistically significant? Justify your answer.

Example: CPR and blood thinners

1. What is the research question?
2. What are the observational units?
3. What type of study design was used? Justify your answer.
4. What is the appropriate scope of inference for these data?

5. What is the sample difference in proportions presented in this example? What notation would be used to represent this value?
6. What is the parameter (using a difference in proportions) representing in the context of this problem? What notation would be used to represent this parameter?
7. Write the null and the alternative hypotheses in words.
8. Write the null and the alternative hypotheses in notation.
9. How could we use cards to simulate **one** sample *which assumes the null hypothesis is true*? How many blue cards — to represent what? How many red cards — to represent what? What would we do with the cards? What would you record once you have a simulated sample?
10. How can we calculate a p-value from the simulated null distribution for this example?
11. What was the p-value of the test?
12. Interpret the p-value in the context of the problem.
13. At the 5% significance level, what decision would you make?
14. What conclusion should the researcher make?
15. Are the results in this example statistically significant? Justify your answer.
16. How could we use cards to simulate **one** bootstrap resample? How many blue cards — to represent what? How many red cards — to represent what? What would we do with the cards? What would you record once you have a simulated sample?
17. How can we calculate a 90% confidence interval from the bootstrap distribution for this example?
18. What was the 90% confidence interval?

19. Interpret the confidence *interval* $((-0.03, 0.28))$ in the context of the problem.
20. Interpret the confidence *level* (90%) in the context of the problem.
21. Does the conclusion of the hypothesis test match the confidence interval?

7.2 Video Notes: Inference for Two Categorical Variables using Simulation-based Methods

Read Sections 15.1 and 15.2 in the course textbook. Use the following videos to complete the lecture notes for Module 7.

7.2.1 Course Videos

- 15.1
- 15.2

Two categorical variables - Video 15.1

- In this module, we will study inference for a _____ explanatory variable and a _____ response.
- The summary measure for two categorical variables is the _____ in _____.

Example: In a double-blind experiment (Weiss 1988) on 48 cocaine addicts hoping to overcome their addiction, half were randomly assigned to a drug called desipramine and the other half a placebo. The addicts were followed for 6 weeks to see whether they were still clean. Is desipramine more effective at helping cocaine addicts overcome their addiction than the placebo?

Observational units:

Explanatory variable:

Response variable:

Notation:

- Population proportion for group 1:
- Population proportion for group 2:
- Sample proportion for group 1:
- Sample proportion for group 2:
- Sample difference in proportions:
- Sample size for group 1:
- Sample size for group 2:

Hypothesis Testing

Conditions:

- Independence: the response for one observational unit will not influence another observational unit

Null hypothesis assumes “no effect”, “no difference”, “nothing interesting happening”, etc.

Always of form: “parameter” = null value

H_0 :

H_A :

- Research question determines the direction of the alternative hypothesis.

Write the null and alternative hypotheses for the cocaine study:

In notation:

H_0 :

H_A :

Summary statistics and plot


```
cocaine %>% group_by(drug) %>% count(outcome)
```

```
#> # A tibble: 4 x 3
#> # Groups:   drug [2]
#>   drug      outcome      n
#>   <chr>      <chr>   <int>
#> 1 desipramine clean     14
#> 2 desipramine relapsed  10
#> 3 placebo     clean      4
#> 4 placebo     relapsed   20
```

Summary statistic:

Interpretation:

```
cocaine%>%
  ggplot(aes(x = drug, fill = outcome))+
  geom_bar(stat = "count", position = "fill") +
  labs(title = "Bar plot of Type of Drug, Segmented by
    Outcome for Cocaine Addicts",
    y = "Relative Frequency",
    x = "Drug or Placebo") +
  scale_fill_grey()
```




Is the independence condition met for simulation inference?

Simulation-based method

- Simulate many samples assuming $H_0 : \pi_1 = \pi_2$
 - Write the response variable values on cards
 - Mix the explanatory variable groups together
 - Shuffle cards into two explanatory variable groups to represent the sample size in each group (n_1 and n_2)
 - Calculate and plot the simulated difference in sample proportions from each simulation
 - Repeat 1000 times (simulations) to create the null distribution
 - Find the proportion of simulations at least as extreme as $\hat{p}_1 - \hat{p}_2$

```
set.seed(216)
two_proportion_test(formula = outcome~drug, # response ~ explanatory
  data = cocaine, # Name of data set
  first_in_subtraction = "desipramine", # Order of subtraction: enter the name of Group 1
  number_repetitions = 1000, # Always use a minimum of 1000 repetitions
  response_value_numerator = "clean", # Define which outcome is a success
  as_extreme_as = 0.417, # Calculated observed statistic (difference in sample proportions)
  direction="greater") # Alternative hypothesis direction ("greater", "less", "two-sided")
```

Explain why the null distribution is centered at the value of zero:

Interpretation of the p-value:

- Statement about probability or proportion of samples
- Statistic (summary measure and value)
- Direction of the alternative
- Null hypothesis (in context)

Conclusion with scope of inference:

- Amount of evidence
- Parameter of interest
- Direction of the alternative hypothesis
- Generalization
- Causation

Confidence interval - Video 15.2


To estimate the difference in true proportion we will create a confidence interval.

Simulation-based method

- Write the response variable values on cards
- Keep explanatory variable groups separate
- Sample with replacement n_1 times in explanatory variable group 1 and n_2 times in explanatory variable group 2
- Calculate and plot the simulated difference in sample proportions from each simulation
- Repeat 1000 times (simulations) to create the bootstrap distribution
- Find the cut-offs for the middle X% (confidence level) in a bootstrap distribution.

Returning to the cocaine example, we will estimate the difference in true proportion of cocaine addicts that stay clean for those on the desipramine and those on the placebo.

```
set.seed(216)
two_proportion_bootstrap_CI(formula = outcome ~ drug,
  data=cocaine, # Name of data set
  first_in_subtraction = "desipramine", # Order of subtraction: enter the name of Group 1
  response_value_numerator = "clean", # Define which outcome is a success
  number_repetitions = 1000, # Always use a minimum of 1000 repetitions
  confidence_level = 0.99) # Enter the level of confidence as a decimal
```



Confidence interval interpretation:

- How confident you are (e.g., 90%, 95%, 98%, 99%)
- Parameter of interest
- Calculated interval
- Order of subtraction when comparing two groups

7.2.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. Explain why the null distribution is centered at the value of zero.
2. Does the confidence interval agree with the p-value?

7.3 Guided Activity 9: The Good Samaritan

7.3.1 Learning outcomes

- Given a research question involving two categorical variables, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Investigate the process of creating a null distribution for two categorical variables

7.3.2 Terminology review

In today's activity, we will use simulation-based methods to analyze two categorical variables. Some terms covered in this activity are:

- Conditional proportion
- Null hypothesis
- Alternative hypothesis

To review these concepts, see Chapter 15 in your textbook.

7.3.3 The Good Samaritan

Researchers at the Princeton University wanted to investigate influences on behavior (Darley and Batson 1973). The researchers randomly selected 67 students from the Princeton Theological Seminary to participate in a study. Only 47 students chose to participate in the study, and the data below includes 40 of those students (7 students were removed from the study for various reasons). As all participants were theology majors planning a career as a preacher, the expectation was that all would have a similar disposition when it comes to helping behavior. Each student was then shown a 5-minute presentation on the Good Samaritan, a parable in the Bible which emphasizes the importance of helping others. After the presentation, the students were told they needed to give a talk on the Good Samaritan parable at a building across campus. Half the students were told they were late for the presentation; the other half told they could take their time getting across campus (the condition was randomly assigned). On the way between buildings, an actor pretending to be a homeless person in distress asked the student for help. The researchers recorded whether the student helped the actor or not. The results of the study are shown in the table below. Do these data provide evidence that those in a hurry will be less likely to help people in need in this situation? Use the order of subtraction hurry – no hurry.

	Hurry Condition	No Hurry Condition	Total
Helped Actor	2	11	13
Did Not Help Actor	18	9	27
Total	20	20	40

These counts can be found in R by using the `count()` function:

```
# Read data set in
good <- read.csv("https://math.montana.edu/courses/s216/data/goodsam.csv")
good %>% group_by(Condition) %>% count(Behavior)
```

```
#> # A tibble: 4 x 3
#> # Groups:   Condition [2]
#>   Condition Behavior    n
#>   <chr>      <chr>  <int>
#> 1 Hurry      Help      2
#> 2 Hurry      No help    18
#> 3 No hurry   Help      11
#> 4 No hurry   No help     9
```

Ask a research question

The research question as stated above is: Do these data provide evidence that those in a hurry will be less likely to help people in need in this situation? In order to set up our hypotheses, we need to express this research question in terms of parameters.

Remember, we define the parameter for a single categorical variable as the true proportion of observational units that are labeled as a “success” in the response variable.

For this study we are identifying two parameters and looking at the difference between these two parameters.

- π_{hurry} = long-run proportion of Princeton Theological Seminary students assigned to hurry that helped the actor
- $\pi_{\text{no hurry}}$ = long-run proportion of Princeton Theological Seminary students assigned not to hurry that helped the actor
- $\pi_{\text{hurry}} - \pi_{\text{no hurry}}$ = the difference in long-run proportion of Princeton Theological Seminary Students that helped the actor between those who were assigned to hurry and those who were not assigned to hurry


When comparing two groups, we assume the two parameters are equal in the null hypothesis—there is no association between the variables.

1. Write the null hypothesis out in words.
2. Based on the research question, fill in the appropriate sign for the alternative hypothesis ($<$, $>$, or \neq):

$$H_A : \pi_{\text{hurry}} - \pi_{\text{no hurry}} \quad \underline{\hspace{1cm}} \quad 0$$

Summarize and visualize the data

```
good %>%
  ggplot(aes(x = Condition, fill = Behavior)) + #Enter the variables to plot
  geom_bar(stat = "count", position = "fill") +
  labs(title = "Segmented bar plot of Condition of Seminary \n Students by Behavior", #Title your plot
        y = "Relative Frequency", #y-axis label
        x = "Condition") + #x-axis label
  scale_fill_grey()
```



- Using the provided segmented bar plot, is there an association between whether a Seminary student helps the actor and condition assigned?
- Using the two-way table given in the introduction, calculate the conditional proportion of students in the hurry condition who helped the actor.
- Using the two-way table given in the introduction, calculate the conditional proportion of students in the no hurry condition who helped the actor.
- Calculate the summary statistic (difference in sample proportion) for this study. Use Hurry - No hurry as the order of subtraction. Use appropriate notation.

The proportion of Princeton Theological Seminary students that helped the actor is 0.45 less for those assigned to hurry compared to those assigned not to hurry.


Hypothesis Test

We will now simulate a **null distribution** of sample differences in proportions. The null distribution is created under the assumption the null hypothesis is true.

Using the cards provided by your instructor, simulate one sample under the assumption the null hypothesis is true.

- Start with 40 cards (13 labeled helped, 27 labeled did not help)
- Mix the cards together
- Shuffle the cards into two piles (20 in hurry, 20 in no hurry)
- Calculate the proportion of simulated students that helped in each group.
- Report the difference in proportion of simulated students that helped (hurry - no hurry)


The segmented bar plot below shows the relationship between the variables for one simulation assuming the null hypothesis is true.



To create the null distribution of differences in sample proportions, we will use the `two_proportion_test()` function in R (in the `catstats` package). We will need to enter the response variable name and the explanatory variable name for the formula, the data set name (identified above as `good`), the outcome for the explanatory variable that is first in subtraction, number of repetitions, the outcome for the response variable that is a success (what the numerator counts when calculating a sample proportion), and the direction of the alternative hypothesis.

The response variable name is `Behavior` and the explanatory variable name is `Condition`.

```
two_proportion_test(formula = Behavior~Condition, # response ~ explanatory
  data = good, # Name of data set
  first_in_subtraction = "Hurry", # Order of subtraction: enter the name of Group 1
  number_repetitions = 1000, # Always use a minimum of 1000 repetitions
  response_value_numerator = "Help", # Define which outcome is a success
  as_extreme_as = -0.45, # Calculated observed statistic (difference in sample proportions)
  direction="less") # Alternative hypothesis direction ("greater", "less", "two-sided")
```




7. Explain why the null distribution is centered around the value of zero?

8. Interpret the p-value in context of the study.

Confidence interval

We will use the `two_proportion_bootstrap_CI()` function in R (in the `catstats` package) to simulate the bootstrap distribution of differences in sample proportions and calculate a 90% confidence interval. We will need to enter the response variable name and the explanatory variable name for the formula, the data set name (identified above as `good`), the outcome for the explanatory variable that is first in subtraction, number of repetitions, the outcome for the response variable that is a success (what the numerator counts when calculating a sample proportion), and the confidence level as a decimal.

```
two_proportion_bootstrap_CI(formula = Behavior ~ Condition,
  data=good, # Name of data set
  first_in_subtraction = "Hurry", # Order of subtraction: enter the name of Group 1
  response_value_numerator = "Help", # Define which outcome is a success
  number_repetitions = 1000, # Always use a minimum of 1000 repetitions
  confidence_level = 0.90) # Enter the level of confidence as a decimal
```



9. Interpret the confidence interval in context of the problem.

10. Write a conclusion to the test.

7.3.4 Take-home messages

1. When comparing two groups, we are looking at the difference between two parameters. In the null hypothesis, we assume the two parameters are equal, or that there is no difference between the two proportions.
2. To create one simulated sample on the null distribution for a difference in sample proportions, label $n_1 + n_2$ cards with the response variable outcomes from the original data. Mix cards together and shuffle into two new groups of sizes n_1 and n_2 , representing the explanatory variable groups. Calculate and plot the difference in proportion of successes.

7.3.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

7.4 Activity 10: Poisonous Mushrooms

7.4.1 Learning outcomes

- Given a research question involving two categorical variables, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Describe and perform a simulation-based hypothesis test for a difference in proportions.
- Interpret and evaluate a p-value for a simulation-based hypothesis test for a difference in proportions.
- Interpret and evaluate a confidence interval for a simulation-based confidence interval for a difference in proportions.

7.4.2 Poisonous Mushrooms

Wild mushrooms, such as chanterelles or morels, are delicious, but eating wild mushrooms carries the risk of accidental poisoning. Even a single bite of the wrong mushroom can be enough to cause fatal poisoning. An amateur mushroom hunter is interested in finding an easy rule to differentiate poisonous and edible mushrooms. They think that the mushroom's gills (the part which holds and releases spores) might be related to a mushroom's edibility. They used a data set of 8124 mushrooms and their descriptions. For each mushroom, the data set includes whether it is edible (e) or poisonous (p) and the size of the gills (broad (b) or narrow (n)). Is there evidence gill size is associated with whether a mushroom is poisonous? PLEASE NOTE: According to The Audubon Society Field Guide to North American Mushrooms, there is no simple rule for determining the edibility of a mushroom; no rule like "leaflets three, let it be" for Poisonous Oak and Ivy.

- Upload and open the R script file for the Module 7 Activity. Upload the csv file, `mushrooms_edibility`.
- Enter the name of the data set for `datasetname` in the R script file in line 8.
- Highlight and run lines 1–9 to get the counts for each combination of categories.

```
mushrooms <- read.csv("datasetname") # Read data set in
mushrooms %>% group_by(gill_size) %>% count(edibility) #finds the counts in each group
```

1. What is the explanatory variable? How are the two levels of the explanatory variable written in the data set?
2. What is the response variable? How are the two levels of the response variable written in the data set?
3. Write the parameter of interest in words, in context of the study.
4. Write the null hypothesis for this study in notation.

5. Using the research question, write the alternative hypothesis in words.

6. Fill in the following two-way table using the R output.

	Gill Size		
Edibility	Broad (b)	Narrow (n)	Total
Poisonous (p)			
Edible (e)			
Total			

7. Calculate the difference in proportion of mushrooms that are poisonous for broad gill mushrooms and narrow gill mushrooms. Use broad - narrow for the order of subtraction. Use appropriate notation.

- Fill in the missing values/names in the R script file for the `two-proportion_test` function to create the null distribution and find the p-value for the test.

```
two_proportion_test(formula = response~explanatory, # response ~ explanatory
  data= mushrooms, # Name of data set
  first_in_subtraction = "xx", # Order of subtraction: enter the name of Group 1
  number_repetitions = 1000, # Always use a minimum of 1000 repetitions
  response_value_numerator = "xx", # Define which outcome is a success
  as_extreme_as = xx, # Calculated observed statistic (difference in sample proportions)
  direction="xx") # Alternative hypothesis direction ("greater","less","two-sided")
```

8. Report the p-value for the study.

9. Do you expect that a 90% confidence interval would contain the null value of zero? Explain your answer.

- Fill in the missing values/names in the R script file in the `two_proportion_bootstrap_CI` function to create a simulation 90% confidence interval.

```
two_proportion_bootstrap_CI(formula = response~explanatory,
  data=mushrooms, # Name of data set
  first_in_subtraction = "xx", # Order of subtraction: enter the name of Group 1
  response_value_numerator = "xx", # Define which outcome is a success
  number_repetitions = 1000, # Always use a minimum of 1000 repetitions
  confidence_level = xx) # Enter the level of confidence as a decimal
```

10. Report the 90% confidence interval.
11. Write a paragraph summarizing the results of the study as if writing a press release. Be sure to describe:
 - Summary statistic and interpretation
 - Summary measure (in context)
 - Value of the statistic
 - Order of subtraction when comparing two groups
 - P-value and interpretation
 - Statement about probability or proportion of samples
 - Statistic (summary measure and value)
 - Direction of the alternative
 - Null hypothesis (in context)
 - Confidence interval and interpretation
 - How confident you are (e.g., 90%, 95%, 98%, 99%)
 - Parameter of interest
 - Calculated interval
 - Order of subtraction when comparing two groups
 - Conclusion (written to answer the research question)
 - Amount of evidence
 - Parameter of interest
 - Direction of the alternative hypothesis
 - Scope of inference
 - To what group of observational units do the results apply (target population or observational units similar to the sample)?
 - What type of inference is appropriate (causal or non-causal)?

Paragraph:

Inference for Two Categorical Variables: Theory-based Methods

8.1 Reading Guide: Theory-based Inference for a Difference in Proportions

Textbook Chapter 15: Inference for comparing two proportions

Section 15.3: Theory-based inferential methods for $\pi_1 - \pi_2$

Reminders from previous sections

n = sample size

\hat{p} = sample proportion

π = population proportion

General steps of a hypothesis test:

1. Frame the research question in terms of hypotheses.
2. Collect and summarize data using a test statistic.
3. Assume the null hypothesis is true, and simulate or mathematically model a null distribution for the test statistic.
4. Compare the observed test (standardized) statistic to the null distribution to calculate a p-value.
5. Make a conclusion based on the p-value and write the conclusion in context.

Parameter: a value summarizing a variable(s) for a population.

Statistic: a value summarizing a variable(s) for a sample.

Sampling distribution: plot of statistics from 1000s of samples of the same size taken from the same population.

Standard deviation of a statistic: the variability of statistics from 1000s of samples; how far, on average, each statistic is from the true value of the parameter.

Standard error of a statistic: estimated standard deviation of a statistic.

Hypothesis test: a process to determine how strong the evidence of an effect is.

Also called a ‘significance test’.

Theory-based method: Develop a mathematical model for the sampling distribution of the statistic under the null hypothesis and use the model to calculate the probability of the observed sample statistic (or one more extreme) occurring.

Null hypothesis (H_0): the skeptical perspective; no difference; no change; no effect; random chance; what the researcher hopes to prove is **wrong**.

Alternative hypothesis (H_A): the new perspective; a difference/increase/decrease; an effect; not random chance; what the researcher hopes to prove is **correct**.

Null value: the value of the parameter when we assume the null hypothesis is true (labeled as *parameter*₀).

Null distribution: the simulated or modeled distribution of statistics (sampling distribution) we would expect to occur if the null hypothesis is true.

P-value: probability of seeing the observed sample data, or something more extreme, assuming the null hypothesis is true.

⇒ Lower the p-value the stronger the evidence AGAINST the null hypothesis and FOR the alternative hypothesis.

Decision: a determination of whether to ‘reject’ or ‘fail to reject’ a null hypothesis based on a p-value and a pre-set level of significance.

Significance level (α): a threshold used to determine if a p-value provides enough evidence to reject the null hypothesis or not.

Common levels of α include 0.01, 0.05, and 0.10.

Statistically significant: results are considered statistically significant if the p-value is below the significance level.

Confidence interval: a process to determine how large an effect is; a range of plausible values for the parameter. Also called ‘estimation’.

Margin of error: the value that is added to and subtracted from the sample statistic to create a confidence interval; half the width of a confidence interval.

Confidence level: how confident we are that the confidence interval will capture the parameter.

Notes

Conditions for the Central Limit Theorem to apply for a difference in proportions

Independence:

Checked by:

Success-failure condition:

Checked by:

Formulas

$$SD(\hat{p}_1 - \hat{p}_2) =$$

Null standard error of the difference in sample proportions: $SE_0(\hat{p}_1 - \hat{p}_2) =$

Standardized statistic (or standardized difference in sample proportions): $Z =$

Standard error of the difference in sample proportions when we do not assume the null hypothesis is true:
 $SE(\hat{p}_1 - \hat{p}_2) =$

Theory-based confidence interval for a difference in proportions:

Margin of error of a confidence interval for a difference in proportions:

Notation

Overall (pooled) proportion of successes:

Example: CPR and blood thinners

1. What are the observational units?
2. What type of study design was used? Justify your answer.
3. What is the appropriate scope of inference for these data?
4. What is the sample difference in proportions presented in this example? What notation would be used to represent this value?
5. What is the parameter (using a difference in proportions) representing in the context of this problem? What notation would be used to represent this parameter?
6. Is it valid to use theory-based methods to analyze these data?
7. Calculate the standard error of the difference in sample proportions without assuming a null hypothesis.
8. Calculate the 90% confidence interval using $z^* = 1.65$ as the multiplier.

Note: A confidence interval interpretation and confidence level interpretation for this example can be found in the Reading Guide solutions for Sections 15.1 and 15.2.

Example: Mammograms

1. What are the observational units?
2. What type of study design was used? Justify your answer.
3. What is the appropriate scope of inference for these data?
4. What is the sample difference in proportions presented in this example? What notation would be used to represent this value?
5. What is the parameter (using a difference in proportions) representing in the context of this problem? What notation would be used to represent this parameter?
6. Write the null and the alternative hypotheses in words.
7. Write the null and the alternative hypotheses in notation.
8. Is it valid to use theory-based methods to analyze these data?
9. Calculate the pooled or overall proportion of successes. What notation would be used to represent this value?
10. Calculate the null standard error of the difference in sample proportions.
11. Calculate the standardized statistic.
12. Interpret the standardized statistic in the context of the problem.
13. Explain how the p-value for this test was calculated.
14. Interpret the p-value in the context of the study.

15. At the 10% significance level, what decision should be made?
16. Write a conclusion for the research question.

8.2 Video Notes: Theoretical Inference for Two Categorical Variables

Read Sections 15.3 and 15.4 and Chapter 16 in the course textbook. Use the following videos to complete the lecture notes for Module 8.

8.2.1 Course Videos

- 15.3TheoryTests
- 15.3TheoryIntervals
- RelativeRisk

Hypothesis testing using theory-based methods - Video 15.4TheoryTests

Example: In Modules 3 and 4, we investigated data on higher education institutions in the United States, collected by the Integrated Postsecondary Education Data System (IPEDS) for the National Center for Education Statistics (NCES) (Education Statistics 2018). A random sample of 2900+ higher education institutions in the United States was collected in 2018. Two variables measured on this data set is whether the institution is a land grant university and whether the institution offers tenure. Does the proportion of universities that offer tenure differ between land grant and non-land-grant institutions?

What is the explanatory variable?

What is the response variable?

Write the parameter of interest:

Hypotheses:

In notation:

H_0 :

H_A :

```
IPED <- read.csv("https://math.montana.edu/courses/s216/data/IPEDS_2018.csv")

IPEDS <- IPED %>%
  drop_na(Tenure)


IPEDS %>% # Data set piped into...
  ggplot(aes(x = LandGrant, fill = Tenure)) + # This specifies the variables
  geom_bar(stat = "count", position = "fill") + # Tell it to make a stacked bar plot
  labs(title = "Segmented Bar Plot of Tenure Availability")
```

```
by Type of Institution for Higher Ed Institutions",
# Make sure to title your plot
x = "Land Grant", # Label the x axis
y = "") + # Remove y axis label
scale_fill_grey()
```

```
IPEDS %>% group_by(LandGrant) %>% count(Tenure)
```

```
#> # A tibble: 4 x 3
#> # Groups:   LandGrant [2]
#>   LandGrant Tenure     n
#>   <chr>      <chr> <int>
#> 1 No        No      976
#> 2 No        Yes     1829
#> 3 Yes       No       31
#> 4 Yes       Yes       72
```

Segmented Bar Plot of Tenure Availability
by Type of Institution for Higher Ed Institutions



Report the summary statistic:

Conditions for inference using theory-based methods for two categorical variables:

- Independence: the response for one observational unit will not influence another observational unit
- Large enough sample size:

Are the conditions met to analyze the university data using theory-based methods?

Steps to use theory-based methods:

- Calculate the standardized statistic
- Find the area under the standard normal distribution at least as extreme as the standardized statistic

Equation for the standard error of the difference in sample proportions assuming the null hypothesis is true:


- This value measures how far each possible sample difference in proportions is from the null value, on average.

Equation for the standardized difference in sample proportions:

- This value measures how many standard errors the sample difference in proportions is above/below the null value.

Calculate the standardized difference in sample proportion of higher education institutions that offer tenure between land grant universities and non-land grant universities.

- First calculate the standard error of the difference in proportion assuming the null hypothesis is true
- Then calculate the Z score



Interpret the standardized statistic

To find the p-value, find the area under the standard normal distribution at the standardized statistic and more extreme.

```
pnorm(0.985, lower.tail = FALSE)*2
```

```
#> [1] 0.3246241
```

Interpretation of the p-value:

- Statement about probability or proportion of samples
- Statistic (summary measure and value)
- Direction of the alternative
- Null hypothesis (in context)

Conclusion with scope of inference:

- Amount of evidence
- Parameter of interest
- Direction of the alternative hypothesis
- Generalization
- Causation

Confidence interval - Video 15.3 Theory Intervals

- Estimate the _____ in true _____
- $CI = \text{statistic} \pm \text{margin of error}$

Theory-based method for a two categorical variables

- $CI = \hat{p}_1 - \hat{p}_2 \pm (z^* \times SE(\hat{p}_1 - \hat{p}_2))$
- When creating a confidence interval, we no longer assume the _____ hypothesis is true.
Use the sample _____ to calculate the sample to sample variability, rather than \hat{p}_{pooled} .

Equation for the standard error of the difference in sample proportions *NOT* assuming the null is true:

Example: Estimate the difference in true proportions of higher education institutions that offer tenure between land grant universities and non-land grant universities.

Find a 90% confidence interval:

- 1st find the z^* multiplier

```
qnorm(0.95, lower.tail=TRUE)
```

```
#> [1] 1.644854
```

- Next, calculate the standard error for the difference in proportions **NOT** assuming the null hypothesis is true

- Calculate the margin of error

- Calculate the endpoints of the 90% confidence interval

Confidence interval interpretation:

- How confident you are (e.g., 90%, 95%, 98%, 99%)
- Parameter of interest
- Calculated interval
- Order of subtraction when comparing two groups

Relative Risk - Video RelativeRisk

- Relative risk is the ratio of the risks in two different categories of an explanatory variable.

Relative Risk:

Example: In a study reported in the New England Journal of Medicine (Du Toit 2015), one-hundred fifty (150) children who had shown sensitivity to peanuts were randomized to receive a flour containing a peanut protein or a placebo flour for 2.5 years. At age 5 years, children were tested with a standard skin prick to see if they had an allergic reaction to peanut protein (yes or no). 71% of those in the peanut flour group no longer demonstrated a peanut allergy compared to 2% of those in the placebo group.

- Calculate the relative risk of desensitization comparing the peanut flour group to the placebo group.

- Interpretation:

- The proportion of successes in group 1 is the RR _____ the proportion of successes in group 2.

Increase in risk:

- Interpretation:

- The proportion of successes in group 1 is the $(RR - 1)$ _____ higher/lower than the proportion of successes in group 2.

Percent increase in risk:

- Interpretation:

- The proportion of successes in group 1 is the $(RR - 1) \times 100$ _____ higher/lower than the proportion of successes in group 2.

- Interpret the value of relative risk from the peanut study in context of the problem.

- Find the increase (or decrease) in risk of desensitization and interpret this value in context of the problem.

- Find the percent increase (or decrease) in risk of desensitization and interpret this value in context of the problem.

Within the peanut flour group, the percent desensitized within each age group (at start of study) is as follows:

1-year-olds: 71%; 2-year-olds: 35%; 3-year-olds: 19%

- Calculate the relative risk of desensitization comparing the 3 year olds to the 2 year olds within the peanut flour group.

- Interpret the percent increase (or decrease) in risk of desensitization comparing the 3 year olds to the 2 year olds within the peanut flour group.

Relative risk in the news

People 50 and older who have had a mild case of covid-19 are 15% more likely to develop shingles (herpes zoster) within six months than are those who have not been infected by the coronavirus, according to research published in the journal Open Forum Infectious Diseases (Bhavsar 2022).

- What was the calculated relative risk of developing shingles when comparing those who has mild COVID-19 to those who had not had COVID-19, among the 50 and older population?

Testing Relative Risk

In Unit 2, we tested for a difference in proportion. We could also test for relative risk.

Null Hypothesis:

$H_0 :$

Alternative Hypothesis:

$H_A :$

8.2.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. What conditions must be met to use the Normal Distribution to approximate the sampling distribution for the difference in sample proportions?
2. How is the value of relative risk calculated?
3. Explain why a theory-based confidence interval for the Good Samaritan study from last module would NOT be similar to the bootstrap interval created.

8.3 Guided Activity 10: Winter Sports Helmet Use and Head Injuries — Theory-based Methods

8.3.1 Learning outcomes

- Assess the conditions to use the normal distribution model for a difference in proportions.
- Create and interpret a theory-based confidence interval for a difference in proportions.

8.3.2 Terminology review

In today's activity, we will use theory-based methods to estimate the difference in two proportions. Some terms covered in this activity are:

- Standard normal distribution
- Independence and success-failure conditions

To review these concepts, see Chapter 15 in your textbook.

8.3.3 Winter sports helmet use and head injury

In this activity we will focus on theory-based methods to calculate a confidence interval. The sampling distribution of a difference in proportions can be mathematically modeled using the normal distribution if certain conditions are met.

Conditions for the sampling distribution of $\hat{p}_1 - \hat{p}_2$ to follow an approximate normal distribution:

- **Independence:** The data are independent within and between the two groups. (*Remember:* This also must be true to use simulation methods!)
- **Success-failure condition:** This condition is met if we have at least 10 successes and 10 failures in each sample. Equivalently, we check that all cells in the table have at least 10 observations.

A study was reported in “Helmet Use and Risk of Head Injuries in Alpine Skiers and Snowboarders” by Sullheim et. al., (Sulheim et al. 2017), on the use of helmets and head injuries for skiers and snowboarders involved in accidents. The summary results from a random sample of 3562 skiers and snowboarders involved in accidents is shown in the two-way table below.

	Helmet Use	No Helmet Use	Total
Head Injury	96	480	576
No Head Injury	656	2330	2986
Total	752	2810	3562

1. Calculate the difference in sample proportion of skiers and snowboarders involved in accidents with a head injury for those who wear helmets and those who do not. Use appropriate notation with informative subscripts.

Confidence interval

To find a confidence interval for the difference in proportions we will add and subtract the margin of error from the point estimate to find the two endpoints.

$$\hat{p}_1 - \hat{p}_2 \pm z^* \times SE(\hat{p}_1 - \hat{p}_2), \text{ where}$$
$$SE(\hat{p}_1 - \hat{p}_2) = \sqrt{\frac{\hat{p}_1 \times (1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2 \times (1 - \hat{p}_2)}{n_2}}$$


In this formula, we use the sample proportions for each group to calculate the standard error for the difference in proportions since we are not assuming that the true difference is zero.

2. Calculate the standard error for a difference in proportions to find a 90% confidence interval.

Recall that the z^* multiplier is the percentile of a standard normal distribution that corresponds to our confidence level. If our confidence level is 90%, we find the Z values that encompass the middle 90% of the standard normal distribution. If 90% of the standard normal distribution should be in the middle, that leaves 10% in the tails, or 5% in each tail. The `qnorm()` function in R will tell us the z^* value for the desired percentile (in this case, 90% + 5% = 95% percentile).

```
qnorm(0.95, lower.tail = TRUE) # Multiplier for 90% confidence interval
```

```
#> [1] 1.644854
```



Remember that the margin of error is the value added and subtracted to the sample difference in proportions to find the endpoints for the confidence interval.

$$ME = z^* \times SE(\hat{p}_1 - \hat{p}_2)$$

3. Using the multiplier of $z^* = 1.645$ and the calculated standard error, calculate the **margin of error** for a 90% confidence interval.
4. Calculate the 90% confidence interval for the parameter of interest.

5. Interpret the confidence interval found in question 4 in context of the problem.

Hypothesis test

6. Write the null and alternative hypotheses in notation.

H_0 :

H_A :

Use statistical analysis methods to draw inferences from the data

To test the null hypothesis, we could use simulation-based methods as we did in the activities in Module 7. In this activity, we will focus on theory-based methods. Like with a single proportion, the sampling distribution of a difference in sample proportions can be mathematically modeled using the normal distribution if certain conditions are met.

To calculate the standardized statistic we use:

$$Z = \frac{(\hat{p}_1 - \hat{p}_2) - \text{null value}}{SE_0(\hat{p}_1 - \hat{p}_2)},$$

where the null standard error is calculated using the pooled proportion of successes:

$$SE_0(\hat{p}_1 - \hat{p}_2) = \sqrt{\hat{p}_{pool} \times (1 - \hat{p}_{pool}) \times \left(\frac{1}{n_1} + \frac{1}{n_2} \right)}.$$


For this study we would first calculate the pooled proportion of successes.

$$\hat{p}_{pool} = \frac{\text{number of "successes"}}{\text{number of cases}} = \frac{576}{3562} = 0.162$$

We use the value for the pooled proportion of successes to calculate the $SE_0(\hat{p}_1 - \hat{p}_2)$.

$$SE_0(\hat{p}_1 - \hat{p}_2) = \sqrt{0.162 \times (1 - 0.162) \times \left(\frac{1}{752} + \frac{1}{2810} \right)} = 0.015$$

7. Use the value of the null standard error to calculate the standardized statistic (standardized difference in proportion).



We will use the `pnorm()` function in R to find the p-value.

```
pnorm(-2.878, # Enter value of standardized statistic
      m=0, s=1, # Using the standard normal mean = 0, sd = 1
      lower.tail=TRUE) # Gives a p-value less than the standardized statistic
```

```
#> [1] 0.002001026
```

8. Write a conclusion to the test.

8.3.4 Effect of sample size

How would an increase in sample size impact the width of the confidence interval. Suppose in another sample of skiers and snowboards involved in accidents we saw these results:

	Helmet Use	No Helmet Use	Total
Head Injury	135	674	809
No Head Injury	921	3270	4191
Total	1056	3944	5000

Note that the sample proportions for each group are the same as the smaller sample size.

$$\hat{p}_h = \frac{135}{1056} = 0.128, \quad \hat{p}_n = \frac{674}{3944} = 0.171$$

$$SE(\hat{p}_1 - \hat{p}_2) = \sqrt{\frac{0.128 \times (1 - 0.128)}{1056} + \frac{0.171 \times (1 - 0.171)}{3944}} = 0.012$$

Margin of Error for 90% confidence interval:

$$ME = 1.645 \times 0.012 = 0.020$$

90% Confidence Interval:

$$(0.128 - 0.171) \pm 0.02$$
$$(-0.063, -0.023)$$

9. How did an increase in sample size impact the width of the confidence interval?

How would an increase in sample size impact the p-value of the test? First calculate the pooled proportion of successes.

$$\hat{p}_{pool} = \frac{\text{number of "successes"}}{\text{number of cases}} = \frac{809}{5000} = 0.162$$

We use the value for the pooled proportion of successes to calculate the $SE_0(\hat{p}_1 - \hat{p}_2)$.

$$SE_0(\hat{p}_1 - \hat{p}_2) = \sqrt{0.162 \times (1 - 0.162) \times \left(\frac{1}{1056} + \frac{1}{3944} \right)} = 0.013$$

Standardized Statistic Calculation:

$$Z = \frac{0.128 - 0.171 - 0}{0.013} = -3.308$$

Use Rstudio to find the p-value for this new sample.

```
pnorm(-3.308, # Enter value of standardized statistic
      m=0, s=1, # Using the standard normal mean = 0, sd = 1
      lower.tail=TRUE) # Gives a p-value greater than the standardized statistic
```

```
#> [1] 0.000469824
```

10. How does the increase in sample size affect the p-value?

8.3.5 Take-home messages

1. Simulation-based methods and theory-based methods should give similar results for a study *if the validity conditions are met*. For both methods, observational units need to be independent. To use theory-based methods, additionally, the success-failure condition must be met. Check the validity conditions for each type of test to determine if theory-based methods can be used.
2. When calculating the standard error for the difference in sample proportions when doing a hypothesis test, we use the pooled proportion of successes, the best estimate for calculating the variability *under the assumption the null hypothesis is true*. For a confidence interval, we are not assuming a null hypothesis, so we use the values of the two conditional proportions to calculate the standard error. Make note of the difference in these two formulas.
3. Increasing sample size will result in less sample-to-sample variability in statistics, which will result in a smaller standard error, and thus a narrower confidence interval.

8.3.6 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

8.4 Activity 11: Diabetes

8.4.1 Learning outcomes

- Given a research question involving two categorical variables, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Assess the conditions to use the normal distribution model for a difference in proportions.
- Describe and perform a simulation-based hypothesis test for a difference in proportions.
- Calculate the Z test statistic for a difference in proportions.
- Find, interpret, and evaluate the p-value for a hypothesis test for a difference in proportions.
- Create and interpret a theory-based confidence interval for a difference in proportions.

8.4.2 Glycemic control in diabetic adolescents

Researchers compared the efficacy of two treatment regimens to achieve durable glycemic control in children and adolescents with recent-onset type 2 diabetes (Group 2012). A convenience sample of patients 10 to 17 years of age with recent-onset type 2 diabetes were randomly assigned to either a medication (rosiglitazone) or a lifestyle-intervention program focusing on weight loss through eating and activity. Researchers measured whether the patient still needs insulin (failure) or had glycemic control (success). Of the 233 children who received the Rosiglitazone treatment, 143 had glycemic control, while of the 234 who went through the lifestyle-intervention program, 125 had glycemic control. Is there evidence that there is difference in proportion of patients that achieve durable glycemic control between the two treatments? Use Rosiglitazone – Lifestyle as the order of subtraction.

- Upload and open the R script file for this activity. Upload the csv file, **diabetes**.
- Enter the name of the data set **datasetname** in the R script file in line 7.
- Highlight and run lines 1–8 to get the counts for each combination of categories.

```
glycemic <- read.csv("datasetname")
glycemic %>% group_by(treatment) %>% count(outcome)
```

1. Complete the following two-way table using the R output.

Outcome	Treatment		Total
	rosiglitazone	lifestyle	
glycemic control (success)			
insulin required (failure)			
Total			

2. Is the independence condition met for this study? Explain your answer.

3. Write the null hypothesis in words.

4. Using the research question, write the alternative hypothesis in notation.
5. Calculate the summary statistic (difference in proportions). Use appropriate notation.
6. Calculate the null standard error for the difference in sample proportions.
7. How many standard errors is the sample difference in proportions from the null value?

Use the `pnorm()` function in R to find the p-value.

- Enter the value of Z calculated in question 7 for xx.

```
2*pnorm(xx, # Enter value of standardized statistic
        m=0, s=1, # Using the standard normal mean = 0, sd = 1
        lower.tail=FALSE) # Gives a p-value less than the standardized statistic
```

8. Report the p-value. How much evidence does the p-value provide against the null hypothesis?
9. Will a 95% simulation confidence interval contain the null value of zero? Explain your answer.
10. Calculate the standard error for a difference in proportions to create a 95% confidence interval.
11. Use the multiplier of $z^* = 1.96$ and the standard error found in question 10 to calculate a 95% confidence interval for the parameter of interest.
12. Write a conclusion to the research question.

Relative Risk

Another summary statistic that can be calculated for two categorical variables is the relative risk. The relative risk is calculated as the ratio of the conditional proportions:

$$\text{relative risk} = \frac{\hat{p}_1}{\hat{p}_2}.$$

13. Calculate the relative risk of glycemic control for patients given Rosiglitazone compared to patients given the placebo.
14. Interpret the value of relative risk in context of the problem.
15. Calculate the increase or decrease in risk of glycemic control for patients given Rosiglitazone compared to patients given the placebo.
16. Interpret the increase or decrease in risk of glycemic control in context of the problem.
17. Calculate the percent increase or percent decrease in glycemic control.
18. Interpret the value of relative risk as a percent increase or percent decrease in glycemic control.

8.4.3 Take-home messages

1. Relative risk calculates the ratio of the proportion of successes in group 1 compared to the proportion of successes in group 2.
2. Relative risk evaluates the percent increase or percent decrease in the response variable attributed to the explanatory variable. To find the percent increase or percent decrease we calculate the following percent change $= (RR - 1) \times 100\%$. If relative risk is less than 1 there is a percent decrease. If relative risk is greater than 1 there is a percent increase.

8.4.4 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

8.5 Module 8 Lab: Is yawning contagious?

8.5.1 Learning outcomes

- Given a research question involving two categorical variables, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Assess the conditions to use the normal distribution model for a difference in proportions.
- Describe and perform a simulation-based hypothesis test for a difference in proportions.
- Calculate the Z test statistic for a difference in proportions.
- Find, interpret, and evaluate the p-value for a hypothesis test for a difference in proportions.
- Create and interpret a confidence interval for a difference in proportions.

8.5.2 Is yawning contagious?

Conventional wisdom says yawns are contagious; when you see someone else yawn, you're prone to feel sleepy and let out a yawn yourself. How many times have you caught yourself in this situation, or noticed it in someone else?

The folks at MythBusters, a popular television program on the Discovery Channel, investigated this issue by using a two-way mirror and a hidden camera. Fifty participants who volunteered to be on the show sat in a booth, accompanied only by an experimental attendee. For some of the participants, the attendee yawned (planting a yawn “seed”), while for other participants the attendee did not yawn. The researchers decided in advance, with a random mechanism, which participants would receive the yawn seed and which would not. As time passed, the researchers watched to see which participants yawned.

Is there evidence that planting a “yawn seed” increases the likelihood of a subject to yawn?

- Upload and open the R script file for the Two Categorical Variable Lab.
- Upload the csv file, `YawnData.csv`. Enter the name of the data set for `datasetname` in the R script file in line 7.
- Highlight and run lines 1–8 to get the counts for each combination of categories.

```
yawn <- read.csv("datasetname")
yawn %>% group_by(YawnSeed) %>% count(Outcome)
```

1. Is this an experiment or an observational study?
2. Complete the following two-way table using the R output.

Outcome	Treatment		Total
	Yawn Seed	No Yawn Seed	
Yawned			
Did Not Yawn			
Total			

3. Write the parameter of interest for the research question.

4. Write the null and alternative hypotheses in notation.

H_0 :

H_A :

5. Calculate the summary statistic (difference in proportions). Use appropriate notation.

- Fill in the missing values/names in the R script file in the two-proportion_test function to create the null distribution and find the simulation p-value for the test.

```
two_proportion_test(formula = response~explanatory, # response ~ explanatory
  data= yawn, # Name of data set
  first_in_subtraction = "xx", # Order of subtraction: enter the name of Group 1
  number_repetitions = 1000, # Always use a minimum of 1000 repetitions
  response_value_numerator = "xx", # Define which outcome is a success
  as_extreme_as = xx, # Calculated observed statistic (difference in sample proportions)
  direction="xx") # Alternative hypothesis direction ("greater","less","two-sided")
```

6. Report the p-value. How much evidence does the p-value provide against the null hypothesis?
7. Calculate the number of standard errors the sample difference in proportion is from the null value of zero.

- Enter the value of Z in the pnorm function in the R script file to find the theory-based p-value.

```
pnorm(xx, lower.tail = FALSE)
```

8. Explain why the theory-based p-value does not match the p-value found using simulation methods.

Confidence Interval

9. Will a 95% simulation confidence interval contain the null value of zero? Explain your answer.

- Fill in the missing values/names in the R script file in the two-proportion_bootstrap_CI function to create the null distribution and find the 90% simulation confidence interval.

```
two_proportion_bootstrap_CI(formula = response~explanatory, # response~explanatory,
                             data = yawn, # Name of data set
                             first_in_subtraction = "xx", # Order of subtraction; enter the name of Group
                             number_repetitions = 1000, #Always use a minimum of 1000 repetitions
                             response_value_numerator = "xx", # Define which outcome is a success
                             confidence_level = xx) #Enter the level of confidence as a decimal)
```

10. Report the 90% confidence interval.
11. What type of error may have occurred?
12. Write a paragraph summarizing the results of the study. Be sure to describe:
 - Summary statistic and interpretation
 - Summary measure (in context)
 - Value of the statistic
 - Order of subtraction when comparing two groups
 - P-value and interpretation
 - Statement about probability or proportion of samples
 - Statistic (summary measure and value)
 - Direction of the alternative
 - Null hypothesis (in context)
 - Confidence interval and interpretation
 - How confident you are (e.g., 90%, 95%, 98%, 99%)
 - Parameter of interest
 - Calculated interval
 - Order of subtraction when comparing two groups
 - Conclusion (written to answer the research question)
 - Amount of evidence
 - Parameter of interest
 - Direction of the alternative hypothesis
 - Scope of inference
 - To what group of observational units do the results apply (target population or observational units similar to the sample)?
 - What type of inference is appropriate (causal or non-causal)?

Upload a copy of your group's paragraph to Gradescope.

Paragraph (continued):

Inference for Inference for a Single Mean or Paired Mean Difference

9.1 Reading Guide: Inference for a Single Mean or Paired Mean Difference

Textbook Chapter 17: Inference for a single mean

Reminders from previous sections

n = sample size

\bar{x} = sample mean

s = sample standard deviation

μ = population mean

σ = population standard deviation

General steps of a hypothesis test:

1. Frame the research question in terms of hypotheses.
2. Collect and summarize data using a test statistic.
3. Assume the null hypothesis is true, and simulate or mathematically model a null distribution for the test statistic.
4. Compare the observed test statistic to the null distribution to calculate a p-value.
5. Make a conclusion based on the p-value and write the conclusion in context.

Parameter: a value summarizing a variable(s) for a population.

Statistic: a value summarizing a variable(s) for a sample.

Sampling distribution: plot of statistics from 1000s of samples of the same size taken from the same population.

Standard deviation of a statistic: the variability of statistics from 1000s of samples; how far, on average, each statistic is from the true value of the parameter.

Standard error of a statistic: estimated standard deviation of a statistic.

Hypothesis test: a process to determine how strong the evidence of an effect is. Also called a 'significance test'.

Simulation-based method: Simulate lots of samples of size n under assumption of the null hypothesis, then find the proportion of the simulations that are at least as extreme as the observed sample statistic.

Theory-based method: Develop a mathematical model for the sampling distribution of the statistic under the null hypothesis and use the model to calculate the probability of the observed sample statistic (or one more extreme) occurring.

Null hypothesis (H_0): the skeptical perspective; no difference; no change; no effect; random chance; what the researcher hopes to prove is **wrong**.

Alternative hypothesis (H_A): the new perspective; a difference/increase/decrease; an effect; not random chance; what the researcher hopes to prove is **correct**.

Null value: the value of the parameter when we assume the null hypothesis is true (labeled as $parameter_0$).

Null distribution: the simulated or modeled distribution of statistics (sampling distribution) we would expect to occur if the null hypothesis is true.

P-value: probability of seeing the observed sample data, or something more extreme, assuming the null hypothesis is true.

\Rightarrow Lower the p-value the stronger the evidence AGAINST the null hypothesis and FOR the alternative hypothesis.

Decision: a determination of whether to reject or fail to reject a null hypothesis based on a p-value and a pre-set level of significance.

- If p-value $\leq \alpha$, then reject H_0 .
- If p-value $> \alpha$, then fail to reject H_0 .

Significance level (α): a threshold used to determine if a p-value provides enough evidence to reject the null hypothesis or not.

Common levels of α include 0.01, 0.05, and 0.10.

Statistically significant: results are considered statistically significant if the p-value is below the significance level.

Confidence interval: a process to determine how large an effect is; a range of plausible values for the parameter. Also called 'estimation'.

Margin of error: the value that is added to and subtracted from the sample statistic to create a confidence interval; half the width of a confidence interval.

Bootstrapping: the process of drawing with replacement n times from the original sample.

Bootstrapped resample: a random sample of size n from the original sample, selected with replacement.

Bootstrapped statistic: the statistic recorded from the bootstrapped resample.

Confidence level: how confident we are that the confidence interval will capture the parameter.

Bootstrap $X\%$ confidence interval: $((\frac{1-X}{2})^{th} \text{ percentile}, (X + (\frac{1-X}{2})^{th} \text{ percentile}))$ of a bootstrap distribution.

Central Limit Theorem: For large sample sizes, the sampling distribution of a sample mean (or proportion) will be approximately normal (bell-shaped and symmetric).

Vocabulary

Shifted bootstrap test:

t -distribution:

- The variability in the t -distribution depends on the sample size (used to calculate degrees of freedom — df for short).
- The larger df, the closer the t distribution is to the standard normal distribution.

Degrees of freedom (df):

T-score:

Notes

To create a shifted bootstrap distribution test,

How many cards will you need and how will the cards be labeled?

Why are the data values shifted prior to being written on the cards?

What do you do with the cards after labeling them?

After resampling, what value will be plotted on the bootstrap distribution?

True or false: Bootstrapping can only be used if the sample size is small.

Why do we use a t -distribution rather than the normal distribution when analyzing quantitative data?

How do we calculate degrees of freedom for the t -distribution?

Conditions to use the CLT for means:

Independence:

Checked by:

Normality:

Checked by:

Formulas

$$SE(\bar{x}) =$$

$$T =$$

Confidence interval for a single mean:

Notation

μ_0 represents

Example from section 17.1: Edinburgh rentals

1. What are the observational units?
2. What are the sample statistics presented in this example? What notation would be used to represent each value?
3. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
4. How could we use cards to simulate **one** bootstrap resample *which does not assume the null hypothesis is true*? How many cards? What is written on the cards? What would we do with the cards? What would you record once you have a simulated sample?
5. After 1000 resamples are generated, where is the resulting bootstrap distribution centered? Why does that make sense?
6. Based on Figure 17.3, give the confidence interval for the true mean for each of the following confidence levels.

90% confidence interval =

95% confidence interval =

99% confidence interval =

7. Interpret your 99% confidence interval in the context of the problem.

Example from section 17.2: Sleep times of MSU students

1. What is the research question?
2. What are the observational units?
3. Can the results of this study be generalized to a larger population? Why or why not?
4. What are the sample statistics presented in this example? What notation would be used to represent each value?
5. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
6. Write the null and the alternative hypotheses in words.
7. Write the null and the alternative hypotheses in notation.
8. How could we use cards to simulate **one** shifted bootstrap resample *which assumes the null hypothesis is true*? How many cards? What is written on the cards (be sure to include the amount and direction of the shift)? What would we do with the cards? What would you record once you have a simulated sample?
9. What was the p-value of the test?
10. Interpret the p-value in the context of the problem.
11. At the 5% significance level, what decision would you make? What type of error might that be?
12. What conclusion should the researcher make?
13. Are the results in this example statistically significant? Justify your answer.

Example from section 17.3: Mercury content of dolphin muscle

1. What is the research question?
2. What are the observational units?
3. Can the results of this study be generalized to a larger population? Why or why not?
4. What are the sample statistics presented in this example? What notation would be used to represent each value?
5. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
6. Are the independence and normality conditions satisfied?
7. Calculate the standard error of the sample mean.
8. What distribution should be referenced to find the multiplier for a 95% confidence interval?
9. Using $t^* = 2.10$, calculate a 95% confidence interval for μ .
10. Interpret the interval calculated in the context of the problem.

Example from section 17.3: Cherry Blossom Race

1. What is the research question?
2. What are the observational units?
3. Can the results of this study be generalized to a larger population? Why or why not?
4. What are the sample statistics presented in this example? What notation would be used to represent each value?

5. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
6. Are the independence and normality conditions satisfied?
7. Write the null and the alternative hypotheses in words.
8. Write the null and the alternative hypotheses in notation.
9. Calculate the standard error of the sample mean.
10. Calculate the T-score (the standardized statistic for the sample mean).
11. What distribution should the T-score be compared to in order to calculate a p-value?
12. What was the p-value of the test?
13. Interpret the p-value in the context of the problem.
14. At the 5% significance level, what decision would you make? What type of error might that be?
15. What conclusion should the researcher make?
16. Are the results in this example statistically significant? Justify your answer.

Textbook Chapter 18: Inference for paired mean difference

Vocabulary

Paired data:

Paired with repeated measures:

Paired with matching:

Notes

For each of the following scenarios, determine if the two sets of observations are paired or independent.

1. To test whether the IQ is related to genetics, researchers measured the IQ of two biological parents and the IQ of their first-born child. The average parent IQ was compared to the IQ of the first born child.
2. Hoping to see how exercise is related to heart rates, researchers asked a group of 30 volunteers to do either bicycle kicks or jumping jacks for 30 seconds. Each volunteer's heart rate was measured at the end of 30 seconds, then the volunteer sat for a 5 minute rest period. At the end of the rest period, the volunteer performed the other activity and their heart rate was measured again. Which activity was done first was randomly assigned.
3. Researchers hoping to look into the effectiveness of blended learning gathered two random samples of 50 8th graders (one at Belgrade Middle School which had 5 full-day instruction at the time of the study, the other from Chief Joseph Middle School which utilized a 2-day on, 3-day off blended learning structure). All 8th graders were given the same lessons and same homework, then asked to take the same end-of-unit test.

Conditions to use the CLT for paired mean difference:

Independence:

Checked by:

Normality:

Checked by:

Formulas

$$SE(\overline{x_d}) =$$

$$T =$$

Confidence interval for a paired mean difference:

Notation

$$\overline{x_d} =$$

$$s_d =$$

$$\mu_d =$$

$$\sigma_d =$$

Example from section 18.1: Tires

1. What are the observational units?
2. Why should we treat these data as paired rather than two independent samples?
3. What are the sample statistics presented in this example? What notation would be used to represent each value?
4. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
5. Write the null and alternative hypotheses in appropriate notation.
6. How could we use cards to simulate **one** shifted bootstrap resample *which assumes the null hypothesis is true*? How many cards? What is written on the cards? What would we do with the cards? What would you record once you have a simulated sample?
7. After 1000 resamples are generated, where is the resulting null distribution centered? Why does that make sense?
8. What was the p-value of the test? Interpret this p-value in the context of the problem.
9. Write a conclusion in the context of the problem.

Example from sections 18.2 and 18.3: UCLA textbook prices

1. What is the research question?
2. What are the observational units?
3. Why should we treat these data as paired rather than two independent samples?

4. What are the sample statistics presented in this example? What notation would be used to represent each value?
5. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
6. How could we use cards to simulate **one** bootstrap resample *which does not assume the null hypothesis is true*? How many cards? What is written on the cards? What would we do with the cards? What would you record once you have a simulated sample?
7. After 1000 resamples are generated, where is the resulting bootstrap distribution centered? Why does that make sense?
8. Give the 95% confidence interval for μ_d .
9. Interpret your 95% confidence interval in the context of the problem.
10. Are the independence and normality conditions satisfied?
11. Write the null and the alternative hypotheses in words.
12. Calculate the standard error of the sample mean difference.
13. Calculate the T-score (the standardized statistic for the sample mean difference).
14. What distribution should the T-score be compared to in order to calculate a p-value?
15. What was the p-value of the test?
16. At the 5% significance level, what decision would you make? What type of error might that be?
17. What conclusion should the researcher make?

18. Are the results in this example statistically significant? Justify your answer.
19. Using $t^* = 2.00$, calculate a 95% confidence interval for μ_d .
20. Interpret the interval calculated in the context of the problem.

9.2 Video Notes: Inference for Paired Data

Read Chapters 17 and 18 in the course textbook. Use the following videos to complete the lecture notes for Module 9.

9.2.1 Course Videos

- PairedData
- 18.1and18.2
- 18.3

Single categorical, single quantitative variables Video Paired_Data

- In this module, we will study inference for a _____ explanatory variable and a _____ response variable where the two groups are _____.

Paired vs. Independent Samples

Two groups are paired if an observational unit in one group is connected to an observational unit in another group

Data are paired if the samples are _____

Examples:

- Change in test score from pre and post test
- Weight of college students before and after 1st year
- Change in blood pressure

Independent Samples		Paired Data		
Sample 1	Sample 2	Sample 1	Sample 2	Difference
x_{1a}	x_{2a}	x_{1a}	x_{2a}	$x_{1a}-x_{2a}$
x_{1b}	x_{2b}	x_{1b}	x_{2b}	$x_{1b}-x_{2b}$
x_{1c}	x_{2c}	x_{1c}	x_{2c}	$x_{1c}-x_{2c}$
.
.
.
x_{1g}	x_{2g}	x_{1g}	x_{2g}	$x_{1g}-x_{2g}$
\bar{x}_1	\bar{x}_2	Mean of the Differences		\bar{x}_d
Difference in Means	$\bar{x}_1-\bar{x}_2$			

Figure 9.1: Illustration of Independent vs. Paired Samples

Example 1: Three hundred registered voters were selected at random to participate in a study on attitudes about how well the president is performing. They were each asked to answer a short multiple-choice questionnaire and then they watched a 20-minute video that presented information about the job description of the president. After watching the video, the same 300 selected voters were asked to answer a follow-up multiple-choice questionnaire.

- Is this an example of a paired samples or independent samples study?

Example 2: Thirty dogs were selected at random from those residing at the humane society last month. The 30 dogs were split at random into two groups. The first group of 15 dogs was trained to perform a certain task using a reward method. The second group of 15 dogs was trained to perform the same task using a reward-punishment method.

- Is this an example of a paired samples or independent samples study?


Example 3: Fifty skiers volunteered to study how different waxes impacted their downhill race times. The participants were split into groups of two based on similar race times from the previous race. One of the two then had their skis treated with Wax A while the other was treated with Wax B. The downhill ski race times were then measured for each of the 25 volunteers who used Wax A as well as for each of the 25 volunteers who used Wax B.

- Is this an example of a paired samples or independent samples study?

Example: Is there a difference in heights between husbands and wives? The heights were measured on the husband and wife in a random sample of 199 married couples from Great Britain ("Great Britain Married Couples: Great Britain Office of Population Census and Surveys," n.d.).

For a paired experiment, we look at the difference between responses for each unit (pair), rather than just the average difference between treatment groups

```
hw <-read.csv("data/husbands_wives_ht.csv")
paired_observed_plot(hw)
```



```
hw_diff %>%
  summarise(fav_stats(ht_diff))
```

```
#>   min    Q1 median   Q3 max    mean      sd    n missing
#> 1  -96  83.5   131  179  303 130.5427 74.13608 199      0
```

- The summary measure for paired data is the _____.

- Mean difference: the average _____ in the _____ variable outcomes for observational units between _____ variable groups

Notation for the Paired differences

- Population mean of the differences:
- Population standard deviation of the differences:
- Sample mean of the differences:
- Sample standard deviation of the differences:

Conditions for inference for paired data:

- Independence:

Is the independence condition met for the height study?

Hypothesis testing

Null hypothesis assumes “no effect”, “no difference”, “nothing interesting happening”, etc.

- Treat the differences like a single mean
- Always of form: “parameter” = null value

H_0 :

H_A :

- Research question determines the direction of the alternative hypothesis.

Write the null and alternative for the height study:

In notation:

H_0 :

H_A :

Simulation-based method

- Simulate many samples assuming $H_0 : \mu_d = 0$
 - Shift the data by the difference between μ_0 and \bar{x}_d
 - Sample with replacement n times from the shifted data
 - Plot the simulated shifted sample mean from each simulation
 - Repeat 1000 times (simulations) to create the null distribution

- Find the proportion of simulations at least as extreme as \bar{x}_d

Reminder of summary statistics:

```
hw_diff %>%
  summarise(fav_stats(ht_diff))
```


```
#>   min    Q1 median   Q3 max    mean      sd    n missing
#> 1  -96  83.5    131 179 303 130.5427 74.13608 199      0
```

Find the difference:

$$\mu_0 - \bar{x}_d =$$

Simulated null distribution:

```
set.seed(216)
paired_test(data = hw_diff$ht_diff,    # Vector of differences
            # or data set with column for each group
            shift = -130.543,          # Shift needed for bootstrap hypothesis test
            as_extreme_as = 130.543,   # Observed statistic
            direction = "two-sided",    # Direction of alternative
            number_repetitions = 1000,  # Number of simulated samples for null distribution
            which_first = 1)           # Not needed when using calculated differences
```



Interpret the p-value:

- Statement about probability or proportion of samples
- Statistic (summary measure and value)
- Direction of the alternative
- Null hypothesis (in context)

Conclusion:

- Amount of evidence
- Parameter of interest
- Direction of the alternative hypothesis


Confidence interval

Simulation-based method

- Label cards with the values (differences) from the data set
- Sample with replacement (bootstrap) from the original sample n times
- Plot the simulated sample mean on the bootstrap distribution
- Repeat at least 1000 times (simulations)
- Find the cut-offs for the middle X% (confidence level) in a bootstrap distribution.
 - i.e., 95% CI = (2.5th percentile, 97.5th percentile)

Simulated bootstrap distribution:

```
set.seed(216)
paired_bootstrap_CI(data = hw_diff$ht_diff, # Enter vector of differences
  number_repetitions = 1000, # Number of bootstrap samples for CI
  confidence_level = 0.95, # Confidence level in decimal form
  which_first = 1) # Not needed when entering vector of differences
```



Interpret the 99% confidence interval:

- How confident you are (e.g., 90%, 95%, 98%, 99%)
- Parameter of interest
- Calculated interval
- Order of subtraction when comparing two groups

Theory-based method - Video 18.3

t-distribution

In the theoretical approach, we use the CLT to tell us that the distribution of sample means will be approximately normal, centered at the assumed true mean under H_0 and with standard deviation $\frac{\sigma}{\sqrt{n}}$.

$$\bar{x} \sim N(\mu_0, \frac{\sigma_d}{\sqrt{n}})$$

- Estimate the population standard deviation, σ_d , with the _____ standard deviation, _____.
- For a single quantitative variable we use the _____ - distribution with _____ degrees of freedom to approximate the sampling distribution.

Conditions for inference using theory-based methods for paired data (categorical explanatory and quantitative response):

- Independence: (same as for simulation); the difference in outcome for one observational unit will not influence another observation.
- Large enough sample size:
 - Normality: The data should be approximately normal or the sample size should be large.

$n < 30$:

$30 \leq n < 100$:

$n \geq 100$:

Theory-based Hypothesis Test:

- Calculate the standardized statistic
- Find the area under the t-distribution with $n - 1$ df at least as extreme as the standardized statistic

Equation for the standard error for the sample mean difference:

Equation for the standardized sample mean difference:


Reminder of summary statistics for height data:

```
hw_diff %>%  
  summarise(fav_stats(ht_diff))  
  
#>   min    Q1 median   Q3 max    mean      sd    n missing  
#> 1 -96 83.5    131 179 303 130.5427 74.13608 199      0
```

Calculate the standardized sample mean difference in height:

- 1st calculate the standard error of the sample mean difference
- Then calculate the T score

What theoretical distribution should we use to find the p-value using the value of the standardized statistic?




To find the p-value:

```
pt(24.84, df = 198, lower.tail=FALSE)*2  
  
#> [1] 9.477617e-63
```

Theory-based Confidence Interval:

statistic \pm margin of error

The t^* multiplier is the value at the given percentile of the t-distribution with $n - 1$ degrees of freedom. For the height data, we will use a t-distribution with _____ df.



To find the t^* multiplier for a 99% confidence interval:

```
qt(0.975, df=198, lower.tail = TRUE)
```

```
#> [1] 1.972017
```

Calculate the margin of error:

Calculate the theory-based confidence interval.

9.2.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. What theoretical distribution is used to approximate paired quantitative data?
2. What is the difference between a paired and independent study design?

9.3 Guided Activity 11: Color Interference

9.3.1 Learning outcomes

- Given a research question involving paired differences, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Describe and perform a theory-based hypothesis test for a paired mean difference.
- Interpret and evaluate a p-value for a theory-based hypothesis test for a paired mean difference.
- Use theory-based methods to find a confidence interval for a paired mean difference.
- Interpret a confidence interval for a paired mean difference.
- Use a confidence interval to determine the conclusion of a hypothesis test.

9.3.2 Terminology review

In today's activity, we will analyze paired quantitative data using theory-based methods. Some terms covered in this activity are:

- Paired data
- Mean difference
- Independent observational units
- Normality
- t -distribution
- Degrees of freedom
- T-score

To review these concepts, see Chapter 18 in the textbook.

9.3.3 Paired vs. Independent Samples

For each of the following scenarios, determine whether the samples are paired or independent.

1. Researchers interested in studying the effect of a medical treatment on insulin rate measured insulin rates of 30 patients before and after the medical treatment.
2. A university is planning to bring emotional support animals to campus during finals week and wants to determine which type of animals are more effective at calming students. Anxiety levels will be measured before and after each student interacts with either a dog or a cat. The university will then compare change in anxiety levels between the 'dog' people and the 'cat' people.
3. An industry leader is investigating a possible wage gap between male and non-male employees. Twenty companies within the industry are randomly selected and the average salary for all males and non-males in mid-management positions is recorded for each company.

9.3.4 Color Interference

The abstract of the article “Studies of interference in serial verbal reactions” in the *Journal of Experimental Psychology* (Stroop 1935) reads:

In this study pairs of conflicting stimuli, both being inherent aspects of the same symbols, were presented simultaneously (a name of one color printed in the ink of another color—a word stimulus and a color stimulus). The difference in time for reading the words printed in colors and the same words printed in black is the measure of interference of color stimuli upon reading words. ... The interference of conflicting color stimuli upon the time for reading 100 words (each word naming a color unlike the ink-color of its print) caused an increase of 2.3 seconds or 5.6% over the normal time for reading the same words printed in black.

The article reports on the results of a study in which seventy college undergraduates were given forms with 100 names of colors written in black ink, and the same 100 names of colors written in another color (i.e., the word purple written in green ink). The total time (in seconds) for reading the 100 words printed in black, and the total time (in seconds) for reading the 100 words printed in different colors were recorded for each subject. The order in which the forms (black or color) were given was randomized to the subjects. Does printing the name of colors in a different color increase the time it takes to read the words? Use color – black as the order of subtraction.

Identify the scenario

1. Should these observations be considered paired or independent? Explain your answer.

Ask a research question

2. Write out the null hypothesis in words, in the context of this study.
3. Write out the alternative hypothesis in proper notation for this study.

In general, the sampling distribution for a sample mean, \bar{x} , based on a sample of size n from a population with a true mean μ and true standard deviation σ can be modeled using a Normal distribution when certain conditions are met.

Conditions for the sampling distribution of \bar{x} to follow an approximate Normal distribution:

- **Independence:** The sample’s observations are independent. For paired data, that means each pairwise difference should be independent.
- **Normality:** The data should be approximately normal or the sample size should be large.
 - $n < 30$: If the sample size n is less than 30 and the distribution of the data is approximately normal with no clear outliers in the data, then we typically assume the data come from a nearly normal distribution to satisfy the condition.

- $30 \leq n < 100$: If the sample size n is between 30 and 100 and there are no particularly extreme outliers in the data, then we typically assume the sampling distribution of \bar{x} is nearly normal, even if the underlying distribution of individual observations is not.
- $n \geq 100$: If the sample size n is at least 100 (regardless of the presence of skew or outliers), we typically assume the sampling distribution of \bar{x} is nearly normal, even if the underlying distribution of individual observations is not.




Figure 9.2: Comparison of the standard Normal vs t -distribution with various degrees of freedom

Like we saw in Chapter 5, we will not know the values of the parameters and must use the sample data to estimate them. Unlike with proportions, in which we only needed to estimate the population proportion, π , quantitative sample data must be used to estimate both a population mean μ and a population standard deviation σ . This additional uncertainty will require us to use a theoretical distribution that is just a bit wider than the Normal distribution. Enter the **t -distribution**!

As you can see from Figure 9.2, the t -distributions (dashed and dotted lines) are centered at 0 just like a standard Normal distribution (solid line), but are slightly wider. The variability of a t -distribution depends on its degrees of freedom, which is calculated from the sample size of a study. (For a single sample of n observations or paired differences, the degrees of freedom is equal to $n - 1$.) Recall from previous classes that larger sample sizes tend to result in narrower sampling distributions. We see that here as well. The larger the sample size, the larger the degrees of freedom, the narrower the t -distribution. (In fact, a t -distribution with infinite degrees of freedom actually IS the standard Normal distribution!)

Summarize and visualize the data




Since the original data from the study are not available, we simulated data to match the means and standard deviations reported in the article. We will use these simulated data in the analysis below.

The following code plots each subject's time to read the colored words (above) and time to read the black words (below) connected by a grey line, a histogram of the differences in time to read words between the two conditions, and a boxplot displaying the pairwise differences in time (color – black).

```
color <- read.csv("https://math.montana.edu/courses/s216/data/interference.csv")
paired_observed_plot(color)

color_diff <- color %>%
```

```
mutate(differences = DiffCol-Black)
color_diff %>%
  ggplot(aes(x = differences))+
  geom_boxplot()+
  labs(title="Boxplot of the Difference in Time to Read Words
           Between Color and Black for College Undergraduates",
        x = "Differences in time to read words (Color - Black)")
```



The following code gives the summary statistics for the pairwise differences.

```
color_diff %>%
  summarise(favstats(differences))
```

```
#>      min      Q1 median      Q3     max mean      sd  n missing
#> 1 -16.42 -2.925   2.15  7.0325  17.27   2.3  7.810196 70      0
```

Check theoretical conditions

4. How do you know the independence condition is met for these data?
5. Is the normality condition met to use the theory-based methods for analysis? Explain your answer.

Use statistical inferential methods to draw inferences from the data

To find the standardized statistic for the paired differences we will use the following formula:

$$T = \frac{\bar{x}_d - \mu_0}{SE(\bar{x}_d)},$$

where the standard error of the sample mean difference is:

$$SE(\bar{x}_d) = \frac{s_d}{\sqrt{n}}.$$

6. Calculate the standard error of the sample mean difference.
7. How many standard errors is the observed mean difference from the null mean difference?

To find the p-value we enter the value for the standardized statistic ($T = 2.464$) into the `pt` function in R. If you did not get his answer for question 7, double check your work. For a single sample or paired data, degrees of freedom are found by subtracting 1 from the sample size. You should therefore use `df = n - 1 = 70 - 1 = 69` and `lower.tail = FALSE` to find the p-value.

```
pt(2.464, df=69, lower.tail=FALSE)
```

```
#> [1] 0.008117801
```

8. What does this p-value mean, in the context of the study? Hint: it is the probability of what...assuming what?

Next we will calculate a theory-based confidence interval. To calculate a theory-based confidence interval for the paired mean difference, use the following formula:

$$\bar{x}_d \pm t^* \times SE(\bar{x}_d).$$

We will need to find the t^* multiplier using the function `qt()`. The code below will return the 95th percentile of the t distribution with $df = n_d - 1 = 70 - 1 = 69$.

```
qt(0.95, df = 69, lower.tail=TRUE)
```

```
#> [1] 1.667239
```




Figure 9.3: t-distribution with 69 degrees of freedom

9. Calculate the margin of error for the true paired mean difference using theory-based methods.
10. Calculate the confidence interval for the true paired mean difference using theory-based methods.
11. Interpret the confidence interval in context of the study.
12. Write a conclusion to the test in context of the study.

9.3.5 Take-home messages

1. In order to use theory-based methods for dependent groups (paired data), the independent observational units and normality conditions must be met.
2. A T-score is compared to a t -distribution with $n - 1$ df in order to calculate a one-sided p-value. To find a two-sided p-value using theory-based methods we need to multiply the one-sided p-value by 2.
3. A t^* multiplier is found by obtaining the bounds of the middle X% (X being the desired confidence level) of a t -distribution with $n - 1$ df.

9.3.6 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered

9.4 Guided Activity 12: Snakes

9.4.1 Learning outcomes

- Given a research question involving paired differences, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Describe and perform a simulation-based hypothesis test for a paired mean difference.
- Interpret and evaluate a p-value for a simulation-based hypothesis test for a paired mean difference.
- Use bootstrapping to find a confidence interval for a paired mean difference.
- Interpret a confidence interval for a paired mean difference.
- Use a confidence interval to determine the conclusion of a hypothesis test.

9.4.2 Terminology review

In today's activity, we will analyze paired quantitative data using simulation-based methods. Some terms covered in this activity are:


- Mean difference
- Paired data
- Independent groups
- Shifted bootstrap (null) distribution

To review these concepts, see Section 18 in the textbook.

9.4.3 Snake mazes

“Sidewinding” refers to a method of locomotion used by snakes to travel over slippery or loose terrain, such as sand. This manner of movement is a form of lateral undulation, in which an animal uses wave-like movement patterns to propel themselves forward. The desert horned viper is one of a handful of snake species that primarily uses sidewinding. As part of a recent study (Subach 2022), researchers exploring foraging behaviors of these vipers in the Sahara and western Negev deserts captured 27 unique vipers with the goal of examining how environment complexity affected movement. Each snake was placed in a circular “maze” of 1045 poles stuck into a sand dune (meant to simulate a dense-vegetation area). The snake was timed to see how long it took to get out of the maze. Researchers also measured how far the snake had “traveled” within the maze by examining the tracks left in the sand. Once the snake had completed the maze, it was allowed to travel in an open area for the same length of time as it spent in the maze; distance was again measured by the tracks in the sand. Is there evidence the distance traveled (in meters) by snakes is less in areas of dense vegetation than in open spaces, on average? Use circular maze - open area as the order of subtraction.

```
snakes <- read.csv("data/vipers.csv")
paired_observed_plot(snakes)
```




To find the difference in distance traveled in the circular maze vs in open spaces for each snake we will create the variable differences.

```
snakes_diff <- snakes %>%
  mutate(differences = DistanceMaze - DistanceOpen)
snakes_diff %>%
  summarise(favstats(differences))
```

```
#>   min      Q1 median      Q3    max      mean      sd  n missing
#> 1 -62 -17.15   -9.2  -6.1  -0.3 -14.45444 14.23788 27      0
```

```
snakes_diff %>%
  ggplot(aes(x = differences)) +
  geom_boxplot() +
  labs(title="Boxplot of the Difference in Distance for Snakes to
  Complete the Open Area vs the Maze
  (Open - Maze)")
```



1. Explain why simulation methods should be used to analyze these data.
2. Explain why this is a paired study design.

Ask a research question

3. Write the null hypothesis in words.
4. Write the alternative hypothesis in notation.


Use statistical inferential methods to draw inferences from the data

Hypothesis test To simulate the null distribution of paired sample mean differences we will use a bootstrapping method. Recall that the null distribution must be created under the assumption that the null hypothesis is true. Therefore, before bootstrapping, we will need to *shift* each data point by the difference $\mu_0 - \bar{x}_d$. This will ensure that the mean of the shifted data is μ_0 (rather than the mean of the original data, \bar{x}_d), and that the simulated null distribution will be centered at the null value.

5. Calculate the difference $\mu_0 - \bar{x}_d$. Will we need to shift the data up or down?

We will use the `paired_test()` function in R (in the `catstats` package) to simulate the shifted bootstrap (null) distribution of sample mean differences and compute a p-value.


```
paired_test(data = snakes_diff$differences,  # Vector of differences
             # or data set with column for each group
             shift = 14.454,  # Shift needed for bootstrap hypothesis test
             as_extreme_as = -14.454,  # Observed statistic
             direction = "less",  # Direction of alternative
             number_repetitions = 1000,  # Number of simulated samples for null distribution
             which_first = 1)  # Not needed when using calculated differences
```



6. Interpret the p-value in the context of the problem.

Confidence interval We will use the `paired_bootstrap_CI()` function in R (in the `catstats` package) to simulate the bootstrap distribution of sample mean differences and calculate a 99% confidence interval.

```
paired_bootstrap_CI(data = snakes_diff$differences, # Enter vector of differences
  number_repetitions = 1000, # Number of bootstrap samples for CI
  confidence_level = 0.99, # Confidence level in decimal form
  which_first = 1) # Not needed when entering vector of differences
```



7. Interpret the 99% confidence interval in the context of the problem.

8. Do the results of your confidence interval and hypothesis test agree? What does each tell you about the null hypothesis?
9. Write a conclusion to the test.

9.4.4 Take-home messages

1. The differences in a paired data set are treated like a single quantitative variable when performing a statistical analysis. Paired data (or paired samples) occur when pairs of measurements are collected. We are only interested in the population (and sample) of differences, and not in the original data.
2. When using bootstrapping to create a null distribution centered at the null value for both paired data and a single quantitative variable, we first need to shift the data by the difference $\mu_0 - \bar{x}_d$, and then sample with replacement from the shifted data.
3. When analyzing paired data, the summary statistic is the ‘mean difference’ NOT the ‘difference in means’¹. This terminology will be *very* important in interpretations.
4. To create one simulated sample on the null distribution for a sample mean or mean difference, shift the original data by adding $(\mu_0 - \bar{x})$ or $(0 - \bar{x}_d)$. Sample with replacement from the shifted data n times. Calculate and plot the sample mean or the sample mean difference.
5. To create one simulated sample on the bootstrap distribution for a sample mean or mean difference, label n cards with the original response values. Randomly draw with replacement n times. Calculate and plot the resampled mean or the resampled mean difference.

9.4.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today’s activity and material covered.

¹Technically, if we calculate the differences and then take the mean (mean difference), and we calculate the two means and then take the difference (difference in means), the value will be the same. However, the *sampling variability* of the two statistics will differ, as we will see in Week 12.

9.5 Activity 12: Swearing

9.5.1 Learning outcomes

- Identify whether a study is a paired design or independent groups
- Given a research question involving paired data, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Describe and perform a simulation-based hypothesis test for a mean difference.
- Interpret and evaluate a p-value for a hypothesis test for a mean difference.
- Use bootstrapping methods to find a confidence interval for a mean difference.
- Interpret a confidence interval for a mean difference.

9.5.2 Swearing

Profanity (language considered obscene or taboo) and society’s attitude about its acceptableness is a highly debated topic, but does swearing serve a physiological purpose or function? Previous research has shown that swearing produces increased heart rates and higher levels of skin conductivity. It is theorized that since swearing provokes intense emotional responses, it acts as a distracter, allowing a person to withstand higher levels of pain. To explore the relationship between swearing and increased pain tolerance, researchers from Keele University (Staffordshire, UK) recruited 83 native English-speaking participants (Stephens and Robertson 2020). Each volunteer performed two trials holding a hand in an ice-water bath, once while repeating the “f-word” every three seconds, and once while repeating a neutral word (“table”). The order of the word to repeat was randomly assigned. Researchers recorded the length of time, in seconds, from the moment the participant indicated they were in pain until they removed their hand from the ice water for each trial. They hope to find evidence that pain tolerance is greater (longer times) when a person swears compared to when they say a neutral word, on average. Use Swear – Neutral as the order of subtraction.

1. What is the explanatory variable for this study? What is the response?
 2. What does μ_d represent in the context of this study?
 3. Write out the null hypothesis in proper notation for this study.
 4. What sign ($<$, $>$, or \neq) would you use in the alternative hypothesis for this study? Explain your choice.
-
- Upload and open the R script file for this activity.
 - Upload the csv file, `pain_tolerance`.
 - Enter the name of the data set for `datasetname` in the R script file in line 8.
 - Highlight and run lines 1–9 to load the data and create a paired plot of the data.

```
swearing <- read.csv("datasetname")
paired_observed_plot(swearing)
```

5. Based on the plots, does there appear to be some evidence in favor of the alternative hypothesis? How do you know?

- Enter the outcome for group 1 (Swear) for `group_1` and the outcome for group 2 (Neutral) for `group_2` in line 16.
- Highlight and run lines 14–25 to get the summary statistics and boxplot of the differences.

```
swearing_diff <- swearing %>%
  mutate(differences = group_1 - group_2)
swearing_diff %>%
  summarise(favstats(differences))

swearing_diff %>%
  ggplot(aes(x = differences)) +
  geom_boxplot() +
  labs(title="Boxplot of the Difference in Time Participants Held Their Hand
           in Ice Water while Swearing or while Saying a Neutral Word (Swearing - Neutral)")
```

6. What is the value of \bar{x}_d ? What is the sample size?
7. How far, on average, is each difference in time the participant holds their hand in ice water from the mean of the differences in time? What is the appropriate notation for this value?

Use statistical inferential methods to draw inferences from the data

8. Using the provided graphs and summary statistics, determine if both theory-based methods and simulation methods could be used to analyze the data. Explain your reasoning.

Hypothesis test

Remember that the null distribution is created based on the assumption the null hypothesis is true. In this study, the null hypothesis states that swearing does not affect pain tolerance, or that the length of time a subject kept their hand in the water would be the same whether the patient was swearing or not.

We will use the `paired_test()` function in R (in the `catstats` package) to simulate the null distribution of sample mean differences and compute a p-value.

9. When using the `paired_test()` function, we need to enter the name of the data set, either the order of subtraction (if the data set has both measurements) or the name of the differences (if the data set contains them). We will also need to provide R with the observed mean difference, the direction of the alternative hypothesis, and the shift required in order to force the null hypothesis to be true. The name of the data set as shown above is `swearing_diff` and the column of differences is called `differences`. What values should be entered for each of the following to create 1000 simulated samples?
- shift:
 - As extreme as:
 - Direction ("`greater`", "`less`", or "`two-sided`"):
 - Number of repetitions:
10. Simulate a null distribution and compute the p-value. Using the R script file for this lab, enter your answers for question 9 in place of the `xx`'s to produce the null distribution with 1000 simulations. Highlight and run lines 30–36.

```
paired_test(data = swearing$differences,    # Vector of differences
             # or data set with column for each group
             shift = xx,    # Shift needed for bootstrap hypothesis test
             as_extreme_as = xx, # Observed statistic
             direction = "xx", # Direction of alternative
             number_repetitions = xx, # Number of simulated samples for null distribution
             which_first = 1) # Not needed when using calculated differences
```

Sketch the null distribution created using the `paired_test` code.

Communicate the results and answer the research question

11. Report the p-value. Based off of this p-value and a 1% significance level, what decision would you make about the null hypothesis? What potential error might you be making based on that decision?
12. Do you expect the 98% confidence interval to contain the null value of zero? Explain.

Confidence interval

We will use the `paired_bootstrap_CI()` function in R (in the `catstats` package) to simulate the bootstrap distribution of sample mean differences and calculate a confidence interval.

13. Using bootstrapping and the provided R script file, find a 98% confidence interval. Fill in the missing values/numbers in the `paired_bootstrap_CI()` function to create the 98% confidence interval. Highlight and run lines 41–44.

```
paired_bootstrap_CI(data = swearing_diff$differences, # Enter vector of differences
                    number_repetitions = 1000, # Number of bootstrap samples for CI
                    confidence_level = xx, # Confidence level in decimal form
                    which_first = 1) # Not needed when entering vector of differences
```

Report the 98% confidence interval in interval notation.

14. Interpret the *confidence level* of the interval found in question 13.

15. Write a paragraph summarizing the results of the study. Be sure to describe:

- Summary statistic and interpretation
 - Summary measure (in context)
 - Value of the statistic
 - Order of subtraction when comparing two groups
- P-value and interpretation
 - Statement about probability or proportion of samples
 - Statistic (summary measure and value)
 - Direction of the alternative
 - Null hypothesis (in context)
- Confidence interval and interpretation
 - How confident you are (e.g., 90%, 95%, 98%, 99%)
 - Parameter of interest
 - Calculated interval
 - Order of subtraction when comparing two groups
- Conclusion (written to answer the research question)
 - Amount of evidence
 - Parameter of interest
 - Direction of the alternative hypothesis

- Scope of inference
 - To what group of observational units do the results apply (target population or observational units similar to the sample)?
 - What type of inference is appropriate (causal or non-causal)?

Inference for a Quantitative Response with Independent Samples

10.1 Reading Guide: Inference for a Difference in Two Means

Textbook Chapter 19 Inference for comparing two independent means

Reminders from previous sections

n_1 = sample size of group 1

n_2 = sample size of group 2

\bar{x} = sample mean

s = sample standard deviation

μ = population mean

σ = population standard deviation

General steps of a hypothesis test:

1. Frame the research question in terms of hypotheses.
2. Collect and summarize data using a test statistic.
3. Assume the null hypothesis is true, and simulate or mathematically model a null distribution for the test statistic.
4. Compare the observed test statistic to the null distribution to calculate a p-value.
5. Make a conclusion based on the p-value and write the conclusion in context.

Parameter: a value summarizing a variable(s) for a population.

Statistic: a value summarizing a variable(s) for a sample.

Sampling distribution: plot of statistics from 1000s of samples of the same size taken from the same population.

Standard deviation of a statistic: the variability of statistics from 1000s of samples; how far, on average, each statistic is from the true value of the parameter.

Standard error of a statistic: estimated standard deviation of a statistic.

Hypothesis test: a process to determine how strong the evidence of an effect is. Also called a ‘significance test’.

Simulation-based method: Simulate lots of samples of size n under assumption of the null hypothesis, then find the proportion of the simulations that are at least as extreme as the observed sample statistic.

Theory-based method: Develop a mathematical model for the sampling distribution of the statistic under the null hypothesis and use the model to calculate the probability of the observed sample statistic (or one more extreme) occurring.

Null hypothesis (H_0): the skeptical perspective; no difference; no change; no effect; random chance; what the researcher hopes to prove is **wrong**.

Alternative hypothesis (H_A): the new perspective; a difference/increase/decrease; an effect; not random chance; what the researcher hopes to prove is **correct**.

Null value: the value of the parameter when we assume the null hypothesis is true (labeled as $parameter_0$).

Null distribution: the simulated or modeled distribution of statistics (sampling distribution) we would expect to occur if the null hypothesis is true.

P-value: probability of seeing the observed sample data, or something more extreme, assuming the null hypothesis is true.

\Rightarrow Lower the p-value the stronger the evidence AGAINST the null hypothesis and FOR the alternative hypothesis.

Decision: a determination of whether to reject or fail to reject a null hypothesis based on a p-value and a pre-set level of significance.

- If p-value $\leq \alpha$, then reject H_0 .
- If p-value $> \alpha$, then fail to reject H_0 .

Significance level (α): a threshold used to determine if a p-value provides enough evidence to reject the null hypothesis or not.

Common levels of α include 0.01, 0.05, and 0.10.

Statistically significant: results are considered statistically significant if the p-value is below the significance level.

Confidence interval: a process to determine how large an effect is; a range of plausible values for the parameter. Also called 'estimation'.

Margin of error: the value that is added to and subtracted from the sample statistic to create a confidence interval; half the width of a confidence interval.

Bootstrapping: the process of drawing with replacement n times from the original sample.

Bootstrapped resample: a random sample of size n from the original sample, selected with replacement.

Bootstrapped statistic: the statistic recorded from the bootstrapped resample.

Confidence level: how confident we are that the confidence interval will capture the parameter.

Bootstrap $X\%$ confidence interval: $((\frac{1-X}{2})^{th} \text{ percentile}, (X + (\frac{1-X}{2})^{th} \text{ percentile}))$ of a bootstrap distribution.

Central Limit Theorem: For large sample sizes, the sampling distribution of a sample mean (or proportion) will be approximately normal (bell-shaped and symmetric).

t-distribution: A bell-shaped symmetric distribution, centered at 0, wider than the standard normal distribution.

- The variability in a *t*-distribution depends on the sample size (used to calculate degrees of freedom — df for short).
- The *t*-distribution gets closer to the standard normal distribution as df increases.

Degrees of freedom (df): describes the variability of the *t*-distribution.

T-score: the name for a standardized statistic which is compared to a *t*-distribution.

Notes

To create a **simulated null distribution** of differences in independent sample means,

How many cards will you need and how will the cards be labeled?

What do you do with the cards after labeling them?

After shuffling, what value will be plotted on the simulated null distribution?

To create a **bootstrap distribution** of differences in independent sample means,

How many cards will you need and how will the cards be labeled?

What do you do with the cards after labeling them?

After shuffling, what value will be plotted on the bootstrap distribution?

Conditions to use the CLT for a difference in independent sample means:

Independence:

Checked by:

Normality:

Checked by:

In a two-sample *t*-test, how are the degrees of freedom determined?

True or false: A large p-value indicates that the null hypothesis is true.

Formulas

$$SE(\bar{x}_1 - \bar{x}_2) =$$

$$T =$$

Confidence interval for a difference in independent sample means:

Notation

μ_1 represents

μ_2 represents

σ_1 represents

σ_2 represents

\bar{x}_1 represents

\bar{x}_2 represents

s_1 represents

s_2 represents

Example from section 19.1: Test scores

1. What are the observational units?
2. What are the sample statistics presented in this example? What notation would be used to represent each value?
3. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
4. What is the research question?
5. Write the null and alternative hypothesis in appropriate notation.

6. How could we use cards to simulate **one** sample *which assumes the null hypothesis is true*? How many cards? What is written on the cards? What would we do with the cards? What would you record once you have a simulated sample?
7. After 1000 shuffles are generated, where is the resulting simulated distribution centered? Why does that make sense?
8. How was the p-value for this test found? The proportion of simulated null samples at _____ or _____.
9. Interpret the p-value in the context of the problem.
10. From these data, can we conclude the exams are equally difficult?
11. What type of error may have occurred at the 5% significance level? Interpret that error in context.

Example from section 19.2: ESC and heart attacks

1. What is the research question?
2. What are the observational units?
3. What variables are recorded? Give the type (categorical or quantitative) and role (explanatory or response) of each.
4. What are the sample statistics presented in this example? What notation would be used to represent each value?
5. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
6. How could we use cards to simulate **one** bootstrap resample *which does not assume the null hypothesis is true*? How many cards? What is written on the cards? What would we do with the cards? What would you record once you have a simulated sample?

7. After 1000 resamples are generated, where is the resulting bootstrap distribution centered? Why does that make sense?
8. Does the 90% confidence interval provide evidence of a difference across the two treatments?

Example from section 19.3: North Carolina births

1. What is the research question?
2. What are the observational units?
3. What variables will be analyzed? Give the type and role of each.
4. Can the results of this study be generalized to a larger population?
5. Are causal conclusions appropriate for these data?
6. Write the null and the alternative hypotheses in words.
7. Write the null and the alternative hypotheses in notation.
8. What are the sample statistics presented in this example? What notation would be used to represent each value?
9. Are the independence and normality conditions satisfied?
10. Calculate the standard error of the difference in sample means.
11. Calculate the T-score (the standardized statistic for the sample mean).
12. What distribution should the T-score be compared to in order to calculate a p-value?
13. What was the p-value of the test?

14. What conclusion should the researcher make?
15. Calculate a 95% confidence interval for the parameter of interest using $\text{qt}(0.975, \text{df} = 49) = 1.677$ as the t^* value.
16. Interpret your interval in the context of the problem.

10.2 Video Notes: Inference for Independent Samples

Read Chapters 19 and 20 in the course textbook. Use the following videos to complete the lecture notes for Module 10.

10.2.1 Course Videos

- 19.1
- 19.2
- 19.3TheoryTests
- 19.4TheoryInterval

Single categorical, single quantitative variable with independent samples

- In this module, we will study inference for a _____ explanatory variable and a _____ response variable where the two groups are _____.
- Independent groups: When the measurements in one sample are not related to the measurements in the other sample.
- Two random samples taken separately from two populations and the same response variable is recorded. Compare the average number of sick days off from work for people who had a flu shot and people who didn't.
- Participants are randomly assigned to one of two treatment conditions, and the same response variable is recorded.

Rather than analyzing the differences as a single mean we will calculate summary statistics on each sample.

Example: Fifty-one (51) college students volunteered to look at impacts on memorization, specifically if putting letters into recognizable patterns (like FBI, CIA, EDA, CDC, etc.) would increase the number letters memorized. (Miller 1956) The college students were randomly assigned to either a recognizable or non-recognizable letter group. After a period of study time, the number of letters memorized was collected on each study. Is there evidence that putting letters into recognizable letter groups improve memory?

- The summary measure for two independent groups is the _____ in _____.

Notation for Independent Groups

- Population mean for group 1:
- Population mean for group 2:
- Sample mean for group 1:
- Sample mean for group 2:

- Sample difference in means:
- Population standard deviation for group 1:
- Population standard deviation for group 2:
- Sample standard deviation for group 1:
- Sample standard deviation for group 2:
- Sample size for group 1:
- Sample size for group 2:

Why should we treat this as two independent groups rather than paired data?

Hypothesis Testing

Conditions:

- Independence: the response for one observational unit will not influence the outcome for another observational unit

Null hypothesis assumes “no effect”, “no difference”, “nothing interesting happening”, etc.

Always of form: “parameter” = null value

H_0 :

H_A :

- Research question determines the alternative hypothesis.

Write the null and alternative hypotheses for the letters study:

In notation:

H_0 :

H_A :


```
letters<-read.csv("data/letters.csv")
letters %>%
  reframe(favstats(Memorized~Grouped))
```

```
#>           Grouped min Q1 median   Q3 max   mean      sd n missing
#> 1 NotRecognizable  1  6    12 14.75  24 11.15385 6.576883 26      0
#> 2   Recognizable  1  6    15 21.00  30 14.32000 8.518216 25      0
```

Summary statistic:

Interpret the summary statistic in context of the problem:


```
letters%>%  
  ggplot(aes(y = Memorized, x = Grouped)) + #Enter the name of the explanatory and response variable  
  geom_boxplot()+  
  labs(title = "Boxplot of Number of Letters memorized by Type  
    of Grouping for College Students", #Title your plot  
    y = "Number of letters memorized", #y-axis label  
    x = "Letter Grouping") #x-axis label
```



Simulation-based method

- Simulate many samples assuming $H_0 : \mu_1 = \mu_2$
 - Write the response variable values on cards
 - Mix the explanatory variable groups together
 - Shuffle cards into two explanatory variable groups to represent the sample size in each group (n_1 and n_2)
 - Calculate and plot the simulated difference in sample means from each simulation
 - Repeat 1000 times (simulations) to create the null distribution
 - Find the proportion of simulations at least as extreme as $\bar{x}_1 - \bar{x}_2$

```
set.seed(216)  
two_mean_test(Memorized~Grouped, #Enter the names of the variables  
  data = letters, # Enter the name of the dataset  
  first_in_subtraction = "Recognizable", # First outcome in order of subtraction  
  number_repetitions = 1000, # Number of simulations  
  as_extreme_as = 3.166, # Observed statistic  
  direction = "greater") # Direction of alternative: "greater", "less", or "two-sided"
```



Explain why the null distribution is centered at the value of zero:

Interpretation of the p-value:

- Statement about probability or proportion of samples
- Statistic (summary measure and value)
- Direction of the alternative
- Null hypothesis (in context)

Conclusion:

- Amount of evidence
- Parameter of interest
- Direction of the alternative hypothesis

Confidence interval


To estimate the difference in true mean we will create a confidence interval.

Simulation-based method - Video 19.2

- Write the response variable values on cards
- Keep explanatory variable groups separate
- Sample with replacement n_1 times in explanatory variable group 1 and n_2 times in explanatory variable group 2
- Calculate and plot the simulated difference in sample means from each simulation
- Repeat 1000 times (simulations) to create the bootstrap distribution
- Find the cut-offs for the middle X% (confidence level) in a bootstrap distribution.

For the letters example, we will estimate the difference in true mean number of letters recognized for students given recognizable letter groupings and students given non-recognizable letter groupings.

```
set.seed(216)
two_mean_bootstrap_CI(Memorized ~ Grouped, #Enter the name of the variables
  data = letters, # Enter the name of the data set
  first_in_subtraction = "Recognizable", # First value in order of subtraction
  number_repetitions = 1000, # Number of simulations
  confidence_level = 0.95)
```



Confidence interval interpretation:

- How confident you are (e.g., 90%, 95%, 98%, 99%)
- Parameter of interest
- Calculated interval
- Order of subtraction when comparing two groups

Theory-based method - Video 19.3 Theory Tests

Example: Every year, orange and black monarch butterflies migrate from their summer breeding grounds in the US and Canada to mountain forests in central Mexico, where they hibernate for the winter. Due to abnormal weather patterns and drought affecting monarch habitats and feeding grounds, the population of monarch butterflies is estimated to have decreased by 53% from the 2018-2019 wintering season to the 2019-2020 wintering season (WWF, 2020). While conservationists often resort to captive-rearing with the goal of raising biologically indistinct individuals for release into the wild, tagging studies have shown that captive-reared monarchs have lower migratory success compared to wild monarchs. For this study, the researchers raised 67 monarchs (descended from wild monarchs) from eggs to maturity and then compared them to a group of 40 wild-caught monarchs. The researchers want to explore whether the maximum grip strength (how many Newtons a butterfly exerts at the moment of release when gently tugged from a mesh-covered perch) differs between captive-reared and wild-caught monarchs. Use Captive – Wild for order of subtraction.

Write the null and alternative hypotheses in notation.

H_0 :


H_A :

```
butterfly <- read.csv("data/butterfly1.csv")

butterflies <- butterfly %>% na.omit() %>%
  rename(Monarch_Group = "Monarch.Group",
         MaxGrip = "Max.Grip.Strength..N.") %>%
  mutate(Monarch_Group = factor(Monarch_Group),
         Sex = factor(Sex)) %>%
  mutate(Monarch_Group = fct_collapse(Monarch_Group, "Captive" = c("Incubator - Fall conditions", "Reari

butterflies %>%
  reframe(favstats(MaxGrip~Monarch_Group))
```

```
#>   Monarch_Group   min    Q1 median    Q3   max    mean    sd n missing
#> 1      Captive 0.081 0.162  0.217 0.2845 0.596 0.2363731 0.09412948 67      0
#> 2        Wild 0.108 0.271  0.352 0.4330 0.650 0.3607500 0.14066796 40      0
```



Conditions:

- Independence: the response for one observational unit will not influence the outcome for another observational unit
- Large enough sample size

Like with paired data the t-distribution can be used to model the difference in means.

- For independent samples we use the _____ - distribution with _____ degrees of freedom to approximate the sampling distribution.

Theory-based test:

- Calculate the standardized statistic
- Find the area under the t-distribution with the smallest $n - 1$ df $[\min(n_1 - 1, n_2 - 1)]$ at least as extreme as the standardized statistic

Equation for the standard error of the difference in sample mean:

Equation for the standardized difference in sample mean:

Are the conditions met to analyze the butterfly data using theory based-methods?

Calculate the standardized difference in mean max grip strength.

- First calculate the $SE(\bar{x}_1 - \bar{x}_2)$

- Then calculate the T-score

What theoretical distribution should we use to find the p-value?

To find the theory-based p-value:

```
pt(-5, df=39, lower.tail=FALSE)*2
```

```
#> [1] 1.999987
```

Conclusion:

- Amount of evidence
- Parameter of interest
- Direction of the alternative hypothesis

Confidence Interval - Video 19.3 Theory Intervals

- Calculate the interval centered at the sample statistic
statistic \pm margin of error

Using the butterfly data, calculate the 99% confidence interval.

```
butterflies %>%  
  reframe(favstats(MaxGrip~Monarch_Group))
```

```
#>   Monarch_Group  min    Q1 median    Q3   max      mean      sd n missing  
#> 1      Captive 0.081 0.162  0.217 0.2845 0.596 0.2363731 0.09412948 67      0  
#> 2       Wild 0.108 0.271  0.352 0.4330 0.650 0.3607500 0.14066796 40      0
```

- Need the t^* multiplier for a 99% confidence interval from a t-distribution with _____ df.

```
qt(0.995, df=39, lower.tail = TRUE)
```

```
#> [1] 2.707913
```

- We will use the same value for the $SE(\bar{x}_1 - \bar{x}_2)$ as calculated for the standardized statistic.

Calculate the margin of error for a 99% confidence interval for the parameter of interest.

Calculate a 99% confidence interval for the parameter of interest.

10.2.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. Why is the recognizable letter study analyzed as two independent groups rather than paired data?
2. Write out the equation for the standard error for a difference in sample means.

10.3 Guided Activity 13: Does behavior impact performance?

10.3.1 Learning outcomes

- Given a research question involving one categorical explanatory variable and one quantitative response variable, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Describe and perform a simulation-based hypothesis test for a difference in means.
- Interpret and evaluate a p-value for a simulation-based hypothesis test for a difference in means.
- Use bootstrapping to find a confidence interval for a difference in means.
- Interpret a confidence interval for a difference in means.
- Use a confidence interval to determine the conclusion of a hypothesis test.

10.3.2 Terminology review

In today's activity, we will use simulation-based methods to analyze the association between one categorical explanatory variable and one quantitative response variable, where the groups formed by the categorical variable are independent. Some terms covered in this activity are:

- Independent groups
- Difference in means


To review these concepts, see Chapter 19 in the textbook.

10.3.3 Behavior and Performance

A study in the Academy of Management Journal (Porath 2017) investigated how rude behaviors influence a victim's task performance. Randomly selected college students enrolled in a management course were randomly assigned to one of two experimental conditions: rudeness condition (45 students) and control group (53 students). Each student was asked to write down as many uses for a brick as possible in five minutes; this value (total number of uses) was used as a performance measure for each student, where higher values indicate better performance. During this time another individual showed up late for class. For those students in the rudeness condition, the facilitator displayed rudeness by berating the students in general for being irresponsible and unprofessional (due to the late-arriving person). No comments were made about the late-arriving person for students in the control group. Is there evidence that the average performance score for students in the rudeness condition is lower than for students in the control group? Use the order of subtraction of rudeness – control.

```
# Read in data set
rude <- read.csv("https://math.montana.edu/courses/s216/data/rude.csv")
```

```
# Side-by-side box plots
rude %>%
ggplot(aes(x = condition, y = number_of_uses)) +
  geom_boxplot() +
  labs(title = "Number of Uses for a Brick based on Behavior Condition
for College Students in a Management Course",
x = "Behavior")
```



```
# Summary statistics
rude %>%
  reframe(favstats(number_of_uses ~ condition))
```

```
#>   condition min Q1 median Q3 max   mean   sd  n missing
#> 1  control   0  6   12  17  30 11.811321 7.382559 53      0
#> 2  rudeness   0  6    9  11  18  8.511111 3.992164 45      0
```

Quantitative variables review

1. Compare the distributions of the number of bricks between the two treatment conditions.
2. Is this an experiment or an observational study? Justify your answer.
3. Explain why this is two independent samples and not paired data.

Ask a research question

In this study we are assessing the difference in true mean number of uses for a brick given by college students enrolled in a management course assigned to a rudeness condition and for those assigned to a control group.

4. What assumption are we making about the difference in true mean?

5. Write the alternative hypothesis in notation.

Summarize and visualize the data

6. Calculate the summary statistic of interest (difference in means). What is the appropriate notation for this statistic?

Interpret this calculated value.

Use statistical inferential methods to draw inferences from the data


Hypothesis test Remember that the null distribution is created based on the assumption the null hypothesis is true. In this study, the null hypothesis states that there is no association between the two variables. This means that the values observed in the data set would have been the same regardless of the behavior condition.

To demonstrate this simulation, we could create cards to simulate a sample.

- Write the number of uses for a brick given by each student on one card.
- Mix together and shuffle into two piles, one with 45 cards to represent the rudeness condition and one with 53 cards to represent the control group.
- Calculate the difference in mean number of uses for a brick (rudeness - control)

We will use the `two_mean_test()` function in R (in the `catstats` package) to simulate the null distribution of differences in sample means and compute a p-value.


```
set.seed(216)
two_mean_test(number_of_uses~condition, #Enter the names of the variables
               data = rude,             # Enter the name of the dataset
               first_in_subtraction = "rudeness", # First outcome in order of subtraction
               number_repetitions = 1000, # Number of simulations
               as_extreme_as = -3.3,      # Observed statistic
               direction = "less")       # Direction of alternative: "greater", "less", or "two-sided"
```



7. Report the p-value. Based off of this p-value, write a conclusion to the hypothesis test.

Confidence interval We will use the `two_mean_bootstrap_CI()` function in R (in the `catstats` package) to simulate the bootstrap distribution of differences in sample means and calculate a 95% confidence interval.

```
set.seed(216)
two_mean_bootstrap_CI(number_of_uses ~ condition, #Enter the name of the variables
  data = rude, # Enter the name of the data set
  first_in_subtraction = "rudeness", # First value in order of subtraction
  number_repetitions = 1000, # Number of simulations
  confidence_level = 0.95)
```



8. Interpret the 95% confidence interval.

10.3.4 Take-home messages

1. This activity differs from the activities in Module 9 because the responses are independent, not paired. These data are analyzed as a difference in means, not a mean difference.
2. To create one simulated sample on the null distribution for a difference in sample means, label cards with the response variable values from the original data. Mix cards together and shuffle into two new groups of sizes n_1 and n_2 . Calculate and plot the difference in means.
3. To create one simulated sample on the bootstrap distribution for a difference in sample means, label $n_1 + n_2$ cards with the original response values. Keep groups separate and randomly draw with replacement n_1 times from group 1 and n_2 times from group 2. Calculate and plot the resampled difference in means.

10.3.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered

10.4 Guided Activity 14: Moon Phases and Virtual Reality

10.4.1 Learning outcomes

- Given a research question involving one categorical explanatory variable and one quantitative response variable, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Describe and perform a theory-based hypothesis test for a difference in means.
- Interpret and evaluate a p-value for a theory-based hypothesis test for a difference in means.
- Use theory-based methods to find a confidence interval for a difference in means.
- Interpret a confidence interval for a difference in means.
- Use a confidence interval to determine the conclusion of a hypothesis test.

10.4.2 Terminology review

In today's activity, we will use theory-based methods to analyze the association between one categorical explanatory variable and one quantitative response variable, where the groups formed by the categorical variable are independent. Some terms covered in this activity are:

- Difference in means
- Independence within and between groups
- Normality

To review these concepts, see Chapter 19 in the textbook.

10.4.3 Moon Phases and Virtual Reality

In a study comparing immersive virtual reality (VR) to traditional hands-on methods, researchers recruited 115 undergraduate students to assess the effectiveness of these approaches in teaching complex scientific concepts like Moon phases (Madden 2020). Participants were randomly assigned to experience either a VR simulation replicating the Sun-Earth-Moon system or a hands-on activity where they physically manipulated models to observe Moon phases. The students were given a 14 multiple choice question quiz about Moon phases and the Moon's motion relative to the Earth to evaluate their understanding of Moon phases and the Moon's motion. Each question had only one correct answer, and the participant's score was the sum of the number of correct answers, with all questions weighted equally (with a maximum score of 14). Is there evidence of a difference, on average, in student learning comparing those using VR methods to those using the traditional method? Use order of subtraction VR – Hands-on.


1. Write out the parameter of interest in words in context of the study.
2. Write out the null hypothesis in notation for this study. Be sure to clearly identify the subscripts.
3. Write out the alternative hypothesis in words for this study.

The sampling distribution for $\bar{x}_1 - \bar{x}_2$ can be modeled using a normal distribution when certain conditions are met.

Conditions for the sampling distribution of $\bar{x}_1 - \bar{x}_2$ to follow an approximate normal distribution:

- **Independence:** The sample's observations are independent
- **Normality:** Each sample should be approximately normal or have a large sample size. For *each* sample:
 - $n < 30$: If the sample size n is less than 30 and there are no clear outliers in the data, then we typically assume the data come from a nearly normal distribution to satisfy the condition.
 - $30 \leq n < 100$: If the sample size n is between 30 and 100 and there are no particularly extreme outliers, then we typically assume the sampling distribution of \bar{x} is nearly normal, even if the underlying distribution of individual observations is not.
 - $n \geq 100$: If the sample size n is at least 100 (regardless of the presence of skew or outliers), we typically assume the sampling distribution of \bar{x} is nearly normal, even if the underlying distribution of individual observations is not.

```
moon <- read.csv("data/Moon_VR.csv")
moon %>% # Data set piped into...
  ggplot(aes(y = TestScore, x = Method))+ # Identify variables
  geom_boxplot()+ # Tell it to make a box plot
  labs(title = "Boxplots of Test Scores for Undergraduate Students Comparing VR
    Teaching Methods and Traditional Teaching Methods", # Title
    x = "Methods", # x-axis label
    y = "Test Score (points)" # y-axis label
```



```
moon %>%
  reframe(favstats(TestScore~Method))
```

#>	Method	min	Q1	median	Q3	max	mean	sd	n	missing
#> 1	Hands-on	2	6.00	8.0	9.5	12	7.694915	2.647408	59	0
#> 2	VR	3	5.75	8.5	10.0	12	7.982143	2.370202	56	0

4. Can theory-based methods be used to analyze these data?

5. Calculate the summary statistic (difference in means) for this study. Use appropriate notation with clearly defined subscripts.

Use statistical inferential methods to draw inferences from the data

To find the standardized statistic for the difference in means we will calculate:

$$T = \frac{\bar{x}_1 - \bar{x}_2 - 0}{SE(\bar{x}_1 - \bar{x}_2)},$$

where the standard error of the difference in means is calculated using:

$$SE(\bar{x}_1 - \bar{x}_2) = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}.$$

6. Calculate the standard error for the difference in sample means.

7. Calculate the standardized statistic for the difference in sample means.

When we are comparing two quantitative variables to find the degrees of freedom to use for the t-distribution, we need to use the group with the smallest sample size and subtract 1. (**df** = minimum of $n_1 - 1$ or $n_2 - 1$).

```
2*pt(0.613, df=55, lower.tail=FALSE)
```

```
#> [1] 0.542403
```

To calculate a theory-based 95% confidence interval for a difference in means, use the formula:

$$(\bar{x}_1 - \bar{x}_2) \pm (t^* \times SE(\bar{x}_1 - \bar{x}_2))$$

We will need to find the t^* multiplier using the function `qt()`. For a 95% confidence level, we are finding the t^* value at the 97.5th percentile with (**df** = minimum of $n_1 - 1$ or $n_2 - 1$).

```
qt(0.975, df = 55, lower.tail=TRUE)
```

```
#> [1] 2.004045
```

8. Calculate the 95% confidence interval using theory-based methods.

9. Write a conclusion to the test.

10.4.4 Take-home messages

1. In order to use theory-based methods for independent groups, the normality condition must be met for each sample.
2. A T-score is compared to a t -distribution with the minimum $n - 1$ df in order to calculate a one-sided p-value. To find a two-sided p-value using theory-based methods we need to multiply the one-sided p-value by 2.
3. A t^* multiplier is found by obtaining the bounds of the middle X% (X being the desired confidence level) of a t -distribution with the minimum $n - 1$ df.

10.4.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered

10.5 Activity 13: Dinosaurs

10.5.1 Learning outcomes

- Identify whether a study is a paired design or independent groups
- Given a research question involving one categorical explanatory variable and one quantitative response variable, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Describe and perform a hypothesis test for a difference in means.
- Interpret and evaluate a p-value for a hypothesis test for a difference in means.
- Find a confidence interval for a difference in means.
- Interpret a confidence interval for a difference in means.

10.5.2 Dinosaurs

The backbone of heavy, two-legged, carnivorous dinosaurs, such as the *T. rex*, is subject to stress. Intriguingly, these dinosaurs have protrusions (rugosity) at the top and sides of their spinal vertebrae, potentially for extra support. These protrusions do not seem to be present in smaller carnivorous dinosaurs. MSU paleontologists hypothesize that the presence of the protrusions is associated with the size of the two-legged carnivorous dinosaurs, potentially allowing them to grow big (Wilson 2016). To test this hypothesis, the researchers collected multiple scientific papers describing the fossil bones of 57 two-legged carnivorous dinosaur species. Then, they checked for the presence or absence of these rugose protrusions from photographs published in the papers and collected measurements of the length in centimeters of the femur (or thigh) bone. Femur length is a proxy for dinosaur size. Is there evidence that the presence of the protrusions result in larger dinosaurs? Use present – absent as the order of subtraction.

1. Write out the null hypothesis in proper notation for this study.
2. What sign ($<$, $>$, or \neq) would you use in the alternative hypothesis for this study? Explain your choice.

- Upload and open the R script file for Week 12 lab.
- Upload the csv file, `dinosaur`.
- Enter the name of the data set for `datasetname` in the R script file in line 10 and the name of the explanatory and response variable in line 12.
- Highlight and run lines 1–16 to load the data and create a plot of the data.

```
dinos <- read.csv("datasetname")
dinos %>%
  ggplot(aes(y = response, x = explanatory)) +
  geom_boxplot() +
  labs(title = "Side-by-side Box Plots of Femur Length by Rugosity
    for Carnivorous Dinosaur",
    x = "Rugose structures on the backbone",
    y = "Femur length (cm)")
```

3. Based on the plots, does there appear to be some evidence in favor of the alternative hypothesis? How do you know?

- Enter the name of the explanatory variable for `explanatory` and the response variable for `response` in line 23.
- Run lines 22–23 to find the summary statistics.

```
dinos %>%
  reframe(favstats(response~explanatory))
```

4. Calculate the summary statistic for the research question. Use proper notation.
5. How far, on average, is each difference in mean femur length from the difference in true mean femur length? What is the appropriate notation for this value?

Use statistical inferential methods to draw inferences from the data

6. Using the provided graphs and summary statistics, determine if both theory-based methods and simulation methods could be used to analyze the data. Explain your reasoning.

Hypothesis test

Remember that the null distribution is created based on the assumption the null hypothesis is true. In this study, the null hypothesis states that there is no difference in true mean femur length for dinosaurs with protrusions and dinosaurs without protrusions.

We will use the `two_mean_test()` function in R (in the `catstats` package) to simulate the null distribution of differences in sample means and compute a p-value.

7. Simulate a null distribution and compute the p-value.
 - Using the R script file for this lab, enter the correct values in place of the `xx`'s to produce the null distribution with 1000 simulations.
 - Highlight and run lines 28–33.

```
two_mean_test(response~explanatory, #Enter the names of the variables
  data = dinos, # Enter the name of the dataset
  first_in_subtraction = "xx", # First outcome in order of subtraction
  number_repetitions = 1000, # Number of simulations
  as_extreme_as = xx, # Observed statistic
  direction = "xx") # Direction of alternative: "greater", "less", or "two-sided"
```

Sketch the null distribution created using the `two_mean_test` code.

Communicate the results and answer the research question

8. Report the p-value. Based off of this p-value and a 1% significance level, what decision would you make about the null hypothesis? What potential error might you be making based on that decision?
9. Do you expect the 98% confidence interval to contain the null value of zero? Explain.

Confidence interval

We will use the `two_mean_bootstrap_CI()` function in R (in the `catstats` package) to simulate the bootstrap distribution of sample differences in means and calculate a confidence interval.

10. Using bootstrapping and the provided R script file, find a 98% confidence interval.
 - Fill in the missing values/numbers in the `two_mean_bootstrap_CI()` function to create the 98% confidence interval.
 - Highlight and run lines 38–42.

```
two_mean_bootstrap_CI(response ~ explanatory, #Enter the name of the variables
                      data = dinos, # Enter the name of the data set
                      first_in_subtraction = "xx", # First value in order of subtraction
                      number_repetitions = 1000, # Number of simulations
                      confidence_level = xx)
```

Report the 98% confidence interval in interval notation.

11. Write a paragraph summarizing the results of this study as if you were describing the results to your roommate. Be sure to describe:
 - Summary statistic and interpretation
 - Summary measure (in context)
 - Value of the statistic
 - Order of subtraction when comparing two groups
 - P-value and interpretation
 - Statement about probability or proportion of samples
 - Statistic (summary measure and value)
 - Direction of the alternative
 - Null hypothesis (in context)
 - Confidence interval and interpretation
 - How confident you are (e.g., 90%, 95%, 98%, 99%)
 - Parameter of interest
 - Calculated interval

- Order of subtraction when comparing two groups
- Conclusion (written to answer the research question)
 - Amount of evidence
 - Parameter of interest
 - Direction of the alternative hypothesis
- Scope of inference
 - To what group of observational units do the results apply (target population or observational units similar to the sample)?
 - What type of inference is appropriate (causal or non-causal)?

10.6 Module 10 Lab: Trustworthiness

10.6.1 Learning outcomes

- Given a research question involving one categorical explanatory variable and one quantitative response variable, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Describe and perform a theory-based hypothesis test for a difference in means.
- Interpret and evaluate a p-value for a theory-based hypothesis test for a difference in means.
- Use theory-based methods to find a confidence interval for a difference in means.
- Interpret a confidence interval for a difference in means.
- Use a confidence interval to determine the conclusion of a hypothesis test.

10.6.2 Trustworthiness

Researchers in India wanted to find out how trustworthy famous YouTubers are (Kalra 2022). They went through a process in which they collected data on many videos from famous YouTubers to determine a trustworthiness score. Scientists randomly selected videos from famous YouTubers (>1000 subscribers) to include in the study. There were many different factors that went into calculating the trustworthiness score. Researchers also recorded if YouTubers were a subject matter expert (SME) or not a subject matter expert (non-SME). An example of an SME would be if one of your statistics professors made a YouTube video of how to do hypothesis testing. An example of someone who isn't an SME would be if one of your friends who has never taken a civil engineering class in their life decided to make a YouTube video about how to build a bridge. There were 621 Youtubers who are SMEs in the sample and 1026 who aren't SMEs. Is there evidence of a difference in trustworthiness score between subject matter experts (SME) YouTubers and non-SME YouTubers? Use SME – Non -SME as the order of subtraction

1. **Write out the parameter of interest in words in context of the study.**
2. Write out the null hypothesis in notation for this study. Be sure to clearly identify the subscripts.
3. Write out the alternative hypothesis in words for this study.

The sampling distribution for $\bar{x}_1 - \bar{x}_2$ can be modeled using a normal distribution when certain conditions are met.

Conditions for the sampling distribution of $\bar{x}_1 - \bar{x}_2$ to follow an approximate normal distribution:

- **Independence:** The sample's observations are independent
- **Normality:** Each sample should be approximately normal or have a large sample size. For *each* sample:

- $n < 30$: If the sample size n is less than 30 and there are no clear outliers in the data, then we typically assume the data come from a nearly normal distribution to satisfy the condition.
- $30 \leq n < 100$: If the sample size n is between 30 and 100 and there are no particularly extreme outliers, then we typically assume the sampling distribution of \bar{x} is nearly normal, even if the underlying distribution of individual observations is not.
- $n \geq 100$: If the sample size n is at least 100 (regardless of the presence of skew or outliers), we typically assume the sampling distribution of \bar{x} is nearly normal, even if the underlying distribution of individual observations is not.

- Upload and open the R script file for Module 10 lab. Upload the csv file, `Trustworthiness.csv`.
- Enter the name of the data set for `datasetname` in the R script file in line 10.
- Write a title for the boxplots in line 14.
- Highlight and run lines 1–16 to load the data and create plots of the data.

```
trust <- read.csv("datasetname")
trust %>% # Data set piped into...
  ggplot(aes(y = Trustworthiness_Video, x = Creator_SME))+ # Identify variables
  geom_boxplot()+ # Tell it to make a box plot
  labs(title = "Don't forget to include a title", # Title: should include the type of plot,
        # observational units, variables
        x = "Whether the Creator is SME", # x-axis label
        y = "Trustworthiness Score") # y-axis label
```

4. Is the independence condition met? Explain your answer.

5. Check that the normality condition is met to use theory-based methods to analyze these data.

- Enter the name of the explanatory variable for `explanatory` and the name of the response variable for `response` in line 22.
- Highlight and run lines 21–22 to get the summary statistics for the data.

```
trust %>%
  reframe(favstats(response~explanatory))
```

6. Calculate the summary measure (difference in means) for this study. Use appropriate notation with clearly defined subscripts.

Use statistical inferential methods to draw inferences from the data

To find the standardized statistic for the difference in means we will calculate:

$$T = \frac{\bar{x}_1 - \bar{x}_2 - 0}{SE(\bar{x}_1 - \bar{x}_2)},$$

where the standard error of the difference in means is calculated using:

$$SE(\bar{x}_1 - \bar{x}_2) = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}.$$

7. Calculate the standard error for the difference in sample means.

8. Calculate the standardized statistic for the difference in sample means.

9. When we are comparing two quantitative variables to find the degrees of freedom to use for the t-distribution, we need to use the group with the smallest sample size and subtract 1. (**df** = minimum of $n_1 - 1$ or $n_2 - 1$). Calculate the **df** for this study.

10. Using the provided R script file, enter the T-score (for **xx**) and the **df** calculated in question 9 for **yy** into the **pt()** function to find the p-value. Highlight and run line 27. Report the p-value calculated.

```
2*pt(xx, df=yy, lower.tail=FALSE)
```

11. Explain why we multiplied by 2 in the code above.

12. Do you expect the 95% confidence interval to contain the null value of zero? Explain your answer.

To calculate a theory-based 95% confidence interval for a difference in means, use the formula:

$$(\bar{x}_1 - \bar{x}_2) \pm (t^* \times SE(\bar{x}_1 - \bar{x}_2))$$

We will need to find the t^* multiplier using the function **qt()**. For a 95% confidence level, we are finding the t^* value at the 97.5th percentile with (**df** = minimum of $n_1 - 1$ or $n_2 - 1$).

- Enter the appropriate percentile value (as a decimal) for **xx** and degrees of freedom for **yy** into the **qt()** function at line 32 to find the appropriate t^* multiplier

```
qt(xx, df = yy, lower.tail=FALSE)
```

13. Report the t^* multiplier for the 95% confidence interval.

14. Calculate the 95% confidence interval using theory-based methods.
15. Do the results of the CI agree with the p-value? Explain your answer.
16. What type of error may be possible?
17. Write a paragraph summarizing the results of the study as if you are reporting the results to your supervisor.
Upload a copy of your paragraph to Gradescope for your group. Be sure to describe:
 - Summary statistic and interpretation
 - P-value and interpretation
 - Statement about probability or proportion of samples
 - Statistic (summary measure and value)
 - Direction of the alternative
 - Null hypothesis (in context)
 - Confidence interval and interpretation
 - How confident you are (e.g., 90%, 95%, 98%, 99%)
 - Parameter of interest
 - Calculated interval
 - Order of subtraction when comparing two groups
 - Conclusion (written to answer the research question)
 - Amount of evidence
 - Parameter of interest
 - Direction of the alternative hypothesis
 - Scope of inference

Paragraph continued:

Inference for Two Quantitative Variables

11.1 Reading Guide: Inference for Slope and Correlation

Chapter 21 (Inference for regression and model conditions)

Reminders from previous sections

β_0 : population y -intercept

β_1 : population slope

ρ : population correlation

b_0 : sample y -intercept

b_1 : sample slope

r : sample correlation

Scatterplot: displays two quantitative variables; one dot = two measurements (x, y) on one observational unit.

Four characteristics of a scatterplot:

- *Form*: pattern of the dots plotted. Is the trend generally linear (you can fit a straight line to the data) or non-linear?
- *Strength*: how closely do the points follow a trend? Very closely (strong)? No pattern (weak)?
- *Direction*: as the x values increase, do the y -values tend to increase (positive) or decrease (negative)?
- Unusual observations or *outliers*: points that do not fit the overall pattern of the data.

Least squares regression line: $\hat{y} = b_0 + b_1 \times x$, where b_0 is the sample y -intercept (the estimate for the **(Intercept)** row in the R regression output), and b_1 is the sample slope (the estimate for the **x-variable_name** row in the R).

Sample slope interpretation: a 1 unit increase in the x variable is associated with a $|b_1|$ unit *predicted* increase/decrease in the y -variable.

General steps of a hypothesis test:

1. Frame the research question in terms of hypotheses.
2. Collect and summarize data using a test statistic.
3. Assume the null hypothesis is true, and simulate or mathematically model a null distribution for the test statistic.
4. Compare the observed test statistic to the null distribution to calculate a p-value.

5. Make a conclusion based on the p-value and write the conclusion in context.

Parameter: a value summarizing a variable(s) for a population.

Statistic: a value summarizing a variable(s) for a sample.

Sampling distribution: plot of statistics from 1000s of samples of the same size taken from the same population.

Standard deviation of a statistic: the variability of statistics from 1000s of samples; how far, on average, each statistic is from the true value of the parameter.

Standard error of a statistic: estimated standard deviation of a statistic.

Hypothesis test: a process to determine how strong the evidence of an effect is. Also called a ‘significance test’.

Simulation-based method: Simulate lots of samples of size n under assumption of the null hypothesis, then find the proportion of the simulations that are at least as extreme as the observed sample statistic.

Theory-based method: Develop a mathematical model for the sampling distribution of the statistic under the null hypothesis and use the model to calculate the probability of the observed sample statistic (or one more extreme) occurring.

Null hypothesis (H_0): the skeptical perspective; no difference; no change; no effect; random chance; what the researcher hopes to prove is **wrong**.

Alternative hypothesis (H_A): the new perspective; a difference/increase/decrease; an effect; not random chance; what the researcher hopes to prove is **correct**.

Null value: the value of the parameter when we assume the null hypothesis is true (labeled as $parameter_0$).

Null distribution: the simulated or modeled distribution of statistics (sampling distribution) we would expect to occur if the null hypothesis is true.

P-value: probability of seeing the observed sample data, or something more extreme, assuming the null hypothesis is true.

\Rightarrow Lower the p-value the stronger the evidence AGAINST the null hypothesis and FOR the alternative hypothesis.

Decision: a determination of whether to reject or fail to reject a null hypothesis based on a p-value and a pre-set level of significance.

- If p-value $\leq \alpha$, then reject H_0 .
- If p-value $> \alpha$, then fail to reject H_0 .

Significance level (α): a threshold used to determine if a p-value provides enough evidence to reject the null hypothesis or not.

Common levels of α include 0.01, 0.05, and 0.10.

Statistically significant: results are considered statistically significant if the p-value is below the significance level.

Confidence interval: a process to determine how large an effect is; a range of plausible values for the parameter. Also called ‘estimation’.

Margin of error: the value that is added to and subtracted from the sample statistic to create a confidence interval; half the width of a confidence interval.

Bootstrapping: the process of drawing with replacement n times from the original sample.

Bootstrapped resample: a random sample of size n from the original sample, selected with replacement.

Bootstrapped statistic: the statistic recorded from the bootstrapped resample.

Confidence level: how confident we are that the confidence interval will capture the parameter.

Bootstrap $X\%$ confidence interval: $((\frac{1-X}{2})^{th} \text{ percentile}, (X + (\frac{1-X}{2})^{th} \text{ percentile}))$ of a bootstrap distribution

t -distribution: A bell-shaped symmetric distribution, centered at 0, wider than the standard normal distribution.

- The variability in a t -distribution depends on the sample size (used to calculate degrees of freedom — df for short).
- The t -distribution gets closer to the standard normal distribution as df increases.

Degrees of freedom (df): describes the variability of the t -distribution.

T-score: the name for a standardized statistic which is compared to a t -distribution.

Notes

To create a **simulated null distribution** of sample slopes or sample correlations,

How many cards will you need and how will the cards be labeled?

What do you do with the cards after labeling them?

After shuffling, what value will be plotted on the simulated null distribution?

To create a **bootstrap distribution** of sample slopes or sample correlations,

How many cards will you need and how will the cards be labeled?

What do you do with the cards after labeling them?

After shuffling, what value will be plotted on the bootstrap distribution?

Conditions to use the CLT for testing slope (or correlation):

Linearity:

Checked by:

Independent observations:

Checked by:

Nearly normal residuals:

Checked by:

Constant or equal variance:

Checked by:

In a theory-based test of slope or correlation, how are the degrees of freedom determined?

Explain why testing for slope is equivalent to testing for correlation.

Where in the R output can $SE(b_1)$ be found?

Formulas

$T =$

Confidence interval:

Example from sections 21.2 and 21.3: Crop yields

1. What are the observational units?
2. What is the parameter representing in the context of this problem? What notation would be used to represent this parameter?
3. What is the research question?
4. Write the null and alternative hypotheses in appropriate notation.
5. How could we use cards to simulate **one** sample which assumes *the null hypothesis is true*? How many cards? What is written on the cards? What would we do with the cards? What would you record once you have a simulated sample?

6. After 1000 shuffles are generated, where is the resulting simulated distribution centered? Why does that make sense?
7. What are the sample statistics presented in this example? What notation would be used to represent each value?
8. Write the least squares regression line for these data in appropriate notation.
9. How was the p-value for this test found? The proportion of simulated null samples at _____ or _____.
10. Interpret the p-value in the context of the problem.
11. What conclusion can be drawn from these data?
12. How could we use cards to simulate **one** bootstrap resample *which does not assume the null hypothesis is true*? How many cards? What is written on the cards? What would we do with the cards? What would you record once you have a simulated sample?
13. Interpret the 95% confidence interval provided.

Example from section 21.4: Midterm elections and unemployment

1. What is the research question?
2. What are the observational units?
3. What variables will be analyzed? Give the type and role of each.
4. Can the results of this study be generalized to a larger population?
5. Are causal conclusions appropriate for these data?
6. Write the null and the alternative hypotheses in words.

7. Write the null and the alternative hypotheses in notation.
8. What are the sample statistics presented in this example? What notation would be used to represent each value?
9. Write the least squares regression line for these data in appropriate notation.
10. From the R output provided in table 21.2, what is the standard error of the slope estimate?
11. Calculate the T-score (the standardized statistic for the slope).
12. What distribution should the T-score be compared to in order to calculate a p-value?
13. What was the p-value of the test?
14. What conclusion should the researcher make?
15. Calculate a 95% confidence interval for the parameter of interest using $\text{qt}(0.975, \text{df} = 27) = 2.052$ as the t^* value.
16. Interpret your interval in the context of the problem.

11.2 Video Notes: Inference for Two Quantitative Variables

Read Chapters 21 and 22 in the course textbook. Use the following videos to complete the lecture notes for Module 11.

11.2.1 Course Videos

- 21.1
- 21.3
- 21.4TheoryTests
- 21.4TheoryIntervals

Summary measures and plots for two quantitative variables.

Scatterplot:

- Form: linear or non-linear?
- Direction: positive or negative?
- Strength: how clear is the pattern between the two variables?
- Outliers: points that are far from the pattern or bulk of the data
 - Influential points: outliers that are extreme in the x- variable.

The summary measures for two quantitative variables are:

- _____, interpreted as the on average change in the response variable for a one unit increase in the explanatory variable.
- _____, which measures the strength and direction of the linear relationship between two quantitative variables.
- _____, interpreted as the percent of variability in the response variable that is explained by the relationship with the explanatory variable.
- Least-squares regression line: $\hat{y} = b_0 + b_1 \times x$ (put y and x in the context of the problem)

Notation:

- Population slope:
- Population correlation:
- Sample slope:
- Sample correlation:

Example: Oceanic temperature is important for sea life. The California Cooperative Oceanic Fisheries Investigations has measured several variables on the Pacific Ocean for more than 70 years hoping to better understand weather patterns and impacts on ocean life. (“Ocean Temperature and Salinity Study,” n.d.) For this example, we will look at the most recent 100 measurements of salt water salinity (measured in PSUs or practical salinity units) and the temperature of the ocean measured in degrees Celsius. Is there evidence that water temperature in the Pacific Ocean tends to decrease with higher levels of salinity?

Hypothesis Testing

Null hypothesis assumes “no effect”, “no difference”, “nothing interesting happening”, etc.

Always of form: “parameter” = null value

H_0 :

H_A :

- Research question determines the alternative hypothesis.


Write the null and alternative for the ocean study:

In notation:

H_0 :

H_A :

```
water %>% # Pipe data set into...
ggplot(aes(x = Salnty, y = T_degC))+ # Specify variables
  geom_point(alpha=0.5) + # Add scatterplot of points
  labs(x = "salinity (PSUs)", # Label x-axis
       y = "temperature (C)", # Label y-axis
       title = "Scatterplot of Pacific Ocean Salinity vs Temperature") +
  # Be sure to title your plots
  geom_smooth(method = "lm", se = FALSE) # Add regression line
```



Describe the four characteristics of the scatterplot:

Linear model output:

```
lm.water <- lm(T_degC~Salnty, data=water) # lm(response~explanatory)
round(summary(lm.water)$coefficients, 3)
```

```
#>               Estimate Std. Error t value Pr(>|t|)
#> (Intercept)  197.156      21.478    9.18     0
#> Salnty       -5.514       0.636   -8.67     0
```

Correlation:

```
cor(T_degC~Salnty, data=water)
```

```
#> [1] -0.6588365
```

Write the least squares equation of the line in context of the problem:

Interpret the value of slope in the context of the problem:

Report and describe the correlation value:

Calculate and interpret the coefficient of determination:


Simulation-based method

Conditions:

- Independence: the response for one observational unit will not influence another observational unit
- Linear relationship:
- Simulate many samples assuming $H_0 : \beta_1 = 0$ or $H_0 : \rho = 0$
 - Write the response variable values on cards
 - Hold the explanatory variable values constant
 - Shuffle a new response variable to an explanatory variable
 - Plot the shuffled data points to find the least squares line of regression
 - Calculate and plot the simulated slope or correlation from each simulation
 - Repeat 1000 times (simulations) to create the null distribution
 - Find the proportion of simulations at least as extreme as b_1 or r


To test slope:

```
set.seed(216)
regression_test(T_degC ~ Salnty, # response ~ explanatory
  data = water, # Name of data set
  direction = "less", # Sign in alternative ("greater", "less", "two-sided")
  summary_measure = "slope", # "slope" or "correlation"
  as_extreme_as = -5.514, # Observed slope or correlation
  number_repetitions = 1000) # Number of simulated samples for null distribution
```



To test correlation:

```
set.seed(216)
regression_test(T_degC~Salnty, # response ~ explanatory
  data = water, # Name of data set
  direction = "less", # Sign in alternative ("greater", "less", "two-sided")
  summary_measure = "correlation", # "slope" or "correlation"
  as_extreme_as = -0.659, # Observed slope or correlation
  number_repetitions = 1000) # Number of simulated samples for null distribution
```



Explain why the null distribution is centered at the value of zero:

Interpretation of the p-value:

- Statement about probability or proportion of samples
- Statistic (summary measure and value)
- Direction of the alternative
- Null hypothesis (in context)

Conclusion:

- Amount of evidence
- Parameter of interest
- Direction of the alternative hypothesis

Confidence interval - Video 21.3


To estimate the true slope (or true correlation) we will create a confidence interval.

Simulation-based method

- Write the explanatory and response value pairs on cards
- Sample pairs with replacement n times
- Plot the resampled data points to find the least squares line of regression
- Calculate and plot the simulated slope (or correlation) from each simulation
- Repeat 1000 times (simulations) to create the bootstrap distribution
- Find the cut-offs for the middle X% (confidence level) in a bootstrap distribution.

Returning to the ocean example, we will estimate the true slope between salinity and temperature of the Pacific Ocean.

```
set.seed(216)
regression_bootstrap_CI(T_degC~Salnty, # response ~ explanatory
  data = water, # Name of data set
  confidence_level = 0.95, # Confidence level as decimal
  summary_measure = "slope", # Slope or correlation
  number_repetitions = 1000) # Number of simulated samples for bootstrap distribution
```




Confidence interval interpretation:

- How confident you are (e.g., 90%, 95%, 98%, 99%)
- Parameter of interest
- Calculated interval
- Order of subtraction when comparing two groups

Now we will estimate the true correlation between salinity and temperature of the Pacific Ocean.

```
set.seed(216)
regression_bootstrap_CI(T_degC~Salnty, # response ~ explanatory
  data = water, # Name of data set
  confidence_level = 0.95, # Confidence level as decimal
  summary_measure = "correlation", # Slope or correlation
  number_repetitions = 1000) # Number of simulated samples for bootstrap distribution
```



Confidence interval interpretation:

- How confident you are (e.g., 90%, 95%, 98%, 99%)
- Parameter of interest
- Calculated interval
- Order of subtraction when comparing two groups

Theory-based method - Video 21.4to21.5TheoryTests

Conditions:

- Linearity (for both simulation-based and theory-based methods): the data should follow a linear trend.
 - Check this assumption by examining the _____ of the two variables, and _____. The pattern in the residual plot should display a horizontal line.


- Independence (for both simulation-based and theory-based methods)
 - One _____ for an observational unit has no impact on _____.
- Constant variability (for theory-based methods only): the variability of points around the least squares line remains roughly constant
 - Check this assumption by examining the _____. The variability in the residuals around zero should be approximately the same for all fitted values.
- Nearly normal residuals (for theory-based methods only): residuals must be nearly normal
 - Check this assumption by examining a _____, which should appear approximately normal

Example:


It is a generally accepted fact that the more carats a diamond has, the more expensive that diamond will be. The question is, how much more expensive? Data on thousands of diamonds were collected for this data set. We will only look at one type of cut (“Ideal”) and diamonds less than 1 carat. Does the association between carat size and price have a linear relationship for these types of diamonds? What can we state about the association between carat size and price?

Scatterplot:

```
Diamonds %>% # Pipe data set into...
  ggplot(aes(x = carat, y = price)) + # Specify variables
  geom_point(alpha=0.5) + # Add scatterplot of points
  labs(x = "carat", # Label x-axis
       y = "price ($)", # Label y-axis
       title = "Scatterplot of Diamonds Carats vs Price") +
  # Be sure to title your plots
  geom_smooth(method = "lm", se = FALSE) # Add regression line
```




Diagnostic plots:




Check the conditions for the ocean data:

Scatterplot:

```
water %>% # Pipe data set into...
ggplot(aes(x = Salnty, y = T_degC)) + # Specify variables
  geom_point(alpha=0.5) + # Add scatterplot of points
  labs(x = "salinity (PSUs)", # Label x-axis
       y = "temperature (C)", # Label y-axis
       title = "Scatterplot of Pacific Ocean Salinity vs Temperature") +
  # Be sure to title your plots
  geom_smooth(method = "lm", se = FALSE) # Add regression line
```



Diagnostic plots:



Like with paired data the t -distribution can be used to model slope and correlation.

- For two quantitative variables we use the _____-distribution with _____ degrees of freedom to approximate the sampling distribution.

Theory-based test:


- Calculate the standardized statistic
- Find the area under the t -distribution with $n - 2$ df at least as extreme as the standardized statistic

Equation for the standardized slope:

Calculate the standardized slope for the ocean data

```
lm.water <- lm(T_degC~Salnty, data=water) # lm(response~explanatory)
round(summary(lm.water)$coefficients,3)
```

```
#>           Estimate Std. Error t value Pr(>|t|)
#> (Intercept)  197.156     21.478   9.18    0
#> Salnty       -5.514      0.636  -8.67    0
```



Interpret the standardized statistic:

To find the theory-based p-value:

```
lm.water <- lm(T_degC~Salnty, data=water) # lm(response~explanatory)
round(summary(lm.water)$coefficients,3)
```

```
#>           Estimate Std. Error t value Pr(>|t|)
#> (Intercept)  197.156    21.478    9.18     0
#> Salnty       -5.514     0.636   -8.67     0
```

or

```
pt(-8.670, df = 98, lower.tail=TRUE)
#> [1] 4.623445e-14
```

Theory-based method

- Calculate the interval centered at the sample statistic
statistic \pm margin of error

```
lm.water <- lm(T_degC~Salnty, data=water) # lm(response~explanatory)
round(summary(lm.water)$coefficients, 3)
```

```
#>               Estimate Std. Error t value Pr(>|t|)
#> (Intercept)   197.156      21.478    9.18     0
#> Salnty        -5.514       0.636   -8.67     0
```

Using the ocean data, calculate a 95% confidence interval for the true slope.

- Need the t^* multiplier for a 95% confidence interval from a t-distribution with _____ df.

```
qt(0.975, df=98, lower.tail = TRUE)
```

```
#> [1] 1.984467
```

11.2.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. Explain why theory-based methods should not be used to analyze the salinity study?
2. What is the proper notation for the population slope? Population correlation?

11.3 Guided Activity 15: Prediction of Crocodilian Body Size

11.3.1 Learning outcomes

- Given a research question involving two quantitative variables, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Describe and perform a simulation-based hypothesis test for slope or correlation.
- Interpret and evaluate a p-value for a simulation-based hypothesis test for a slope or correlation.
- Use bootstrapping to find a confidence interval for the slope or correlation.
- Interpret a confidence interval for a slope or correlation.

11.3.2 Terminology review

In today's activity, we will use simulation-based methods for hypothesis tests and confidence intervals for a linear regression slope or correlation. Some terms covered in this activity are:

- Correlation
- Slope
- Regression line

To review these concepts, see Chapter 21 in the textbook.


11.3.3 Crocodilian Body Size

Much research surrounds using measurements of animals to estimate body-size of extinct animals. Many challenges exist in making accurate estimates for extinct crocodilians. The term crocodilians refers to all members of the family Crocodylidae ("true" crocodiles), family Alligatoridae (alligators and caimans) and family Gavialidae (gharial, Tomistoma). The researchers in this study (O'Brien 2019) state, "Among extinct crocodilians and their precursors (e.g., suchians), several methods have been developed to predict body size from suites of hard-tissue proxies. Nevertheless, many have limited applications due to the disparity of some major suchian groups and biases in the fossil record. Here, we test the utility of head width (HW) as a broadly applicable body-size estimator in living and fossil suchians." Data were collected on 76 male and female individuals of different species. Is there evidence that head width (measured in cm) is a good predictor of total body length (measured in cm) for crocodilians?

```
# Read in data set
croc <- read.csv("https://math.montana.edu/courses/s216/data/Crocodylian_headwidth.csv")
croc <- croc %>%
  na.omit()
```

To create a scatterplot to examine the relationship between head width and total body length we will use HW_cm as the explanatory variable and TL_cm as the response variable.

```
croc %>% # Pipe data set into...
ggplot(aes(x = HW_cm, y = TL_cm)) + # Specify variables
  geom_point(alpha=0.5) + # Add scatterplot of points
  labs(x = "head width (cm)", # Label x-axis
       y = "total length (cm)", # Label y-axis
       title = "Scatterplot of Crocodilian Head Width vs. Total Length") +
  # Be sure to title your plots
  geom_smooth(method = "lm", se = FALSE) # Add regression line
```

1. Describe the features of the plot above, addressing all four characteristics of a scatterplot.

If you indicated there are potential outliers, which points are they?

Hypotheses

When analyzing two quantitative variables we can either test regression slope or correlation. In both cases, we are testing that there is a linear relationship between variables.

2. Write the null hypothesis in words.
3. Write the null hypothesis to test slope in notation.
4. Write the null hypothesis to test correlation in notation.
5. Write the alternative hypothesis in words.

Summarize and visualize the data

The linear model output for the data is given below.

```
lm.croc <- lm(TL_cm~HW_cm, data=croc) #lm(response~explanatory)
round(summary(lm.croc)$coefficients, 5)
```

```
#>               Estimate Std. Error  t value Pr(>|t|)
#> (Intercept)  17.61250    11.36269   1.55003  0.12687
#> HW_cm        10.59983     0.51294  20.66494  0.00000
```

The value of correlation is given below.

```
cor(croc$HW_cm, croc$TL_cm)
```

```
#> [1] 0.9412234
```

6. Using the output from the evaluated R code above, write the equation of the regression line in the context of the problem using appropriate statistical notation.

7. Interpret the estimated slope in context of the problem.

Use statistical inferential methods to draw inferences from the data

In this activity, we will focus on using simulation-based methods for inference in regression.

Simulation-based hypothesis test


Let's start by thinking about how one simulation would be created on the null distribution using cards.

- First, we would write the values for the response variable, total length, on each card.
- Next, we would shuffle these y values while keeping the x values (explanatory variable) in the same order.
- Then, find the line of regression for the shuffled (x, y) pairs and calculate either the slope or correlation of the shuffled sample.

We will use the `regression_test()` function in R (in the `catstats` package) to simulate the null distribution of shuffled slopes (or shuffled correlations) and compute a p-value. We will need to enter the response variable name and the explanatory variable name for the formula, the data set name (identified above as `croc`), the summary measure for the test (either slope or correlation), number of repetitions, the sample statistic (value of slope or correlation), and the direction of the alternative hypothesis.


The response variable name is TL_cm and the explanatory variable name is HW_cm for these data.

```
set.seed(216)
regression_test(TL_cm~HW_cm, # response ~ explanatory
  data = croc, # Name of data set
  direction = "two-sided", # Sign in alternative ("greater", "less", "two-sided")
  summary_measure = "slope", # "slope" or "correlation"
  as_extreme_as = 10.600, # Observed slope or correlation
  number_repetitions = 1000) # Number of simulated samples for null distribution
```



8. Suppose we wanted to complete the simulation test using correlation as the summary measure, instead of slope. Which two inputs in the regression_test function would need to be changed to test for correlation? What inputs should you use instead?


```
set.seed(216)
regression_test(TL_cm~HW_cm, # response ~ explanatory
  data = croc, # Name of data set
  direction = "two-sided", # Sign in alternative ("greater", "less", "two-sided")
  summary_measure = "correlation", # "slope" or "correlation"
  as_extreme_as = 0.941, # Observed slope or correlation
  number_repetitions = 1000) # Number of simulated samples for null distribution
```



Simulation-based confidence interval

We will use the `regression_bootstrap_CI()` function in R (in the `catstats` package) to simulate the bootstrap distribution of sample slopes (or sample correlations) and calculate a confidence interval. The following code gives the 95% confidence interval for the true slope.

```
set.seed(216)
regression_bootstrap_CI(TL_cm~HW_cm, # response ~ explanatory
  data = croc, # Name of data set
  confidence_level = 0.95, # Confidence level as decimal
  summary_measure = "slope", # Slope or correlation
  number_repetitions = 1000) # Number of simulated samples for bootstrap distribution
```



9. Interpret the confidence interval in context of the problem. *Hint: use the interpretation of slope in your confidence interval interpretation.*

Communicate the results and answer the research question

10. Based on the p-value and confidence interval, write a conclusion in context of the problem.

11.3.4 Take-home messages

1. The p-value for a test for correlation should be approximately the same as the p-value for the test of slope. In the simulation test, we just change the statistic type from slope to correlation and use the appropriate sample statistic value.
2. To interpret a confidence interval for the slope, think about how to interpret the sample slope and use that information in the confidence interval interpretation for slope.
3. To create one simulated sample on the null distribution when testing for a relationship between two quantitative variables, hold the x values constant and shuffle the y values to new x values. Find the regression line for the shuffled data and plot the slope or the correlation for the shuffled data.
4. To create one simulated sample on the bootstrap distribution when assessing two quantitative variables, label n cards with the original (response, explanatory) values. Randomly draw with replacement n times. Find the regression line for the resampled data and plot the resampled slope or correlation.

11.3.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

11.4 Guided Activity 16: Golf Driving Distance

11.4.1 Learning outcomes

- Given a research question involving two quantitative variables, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Assess the conditions to use the normal distribution model for a slope.
- Find the T test statistic (T-score) for a slope based off of `lm()` output in R.
- Find, interpret, and evaluate the p-value for a theory-based hypothesis test for a slope.
- Create and interpret a theory-based confidence interval for a slope.
- Use a confidence interval to determine the conclusion of a hypothesis test.

11.4.2 Terminology review

In this week's in-class activity, we will use theory-based methods for hypothesis tests and confidence intervals for a linear regression slope. Some terms covered in this activity are:

- Slope
- Regression line

To review these concepts, see Chapter 21 in the textbook.

11.4.3 Golf driving distance


In golf the goal is to complete a hole with as few strokes as possible. A long driving distance to start a hole can help minimize the strokes necessary to complete the hole, as long as that drive stays on the fairway. Data were collected on 354 PGA and LPGA players in 2008 ("Average Driving Distance and Fairway Accuracy" 2008). For each player, the average driving distance (yards), fairway accuracy (percentage), and sex was measured. Use these data to assess, "Does a professional golfer give up accuracy when they hit the ball farther?"

```
# Read in data set
golf <- read.csv("https://math.montana.edu/courses/s216/data/golf.csv")
```

Plot review.

Following is a scatterplot showing the relationship between the driving distance and percent accuracy for professional golfers.

```
golf %>% # Pipe data set into...
ggplot(aes(x = Driving_Distance, y = Percent_Accuracy)) + # Specify variables
  geom_point(alpha=0.5) + # Add scatterplot of points
  labs(x = "Driving Distance (yards)", # Label x-axis
       y = "Percent Accuracy", # Label y-axis
       title = "Scatterplot of Driving Distance by Percent Accuracy
for Professional Golfers") +
  # Be sure to tile your plots
  geom_smooth(method = "lm", se = FALSE) # Add regression line
```



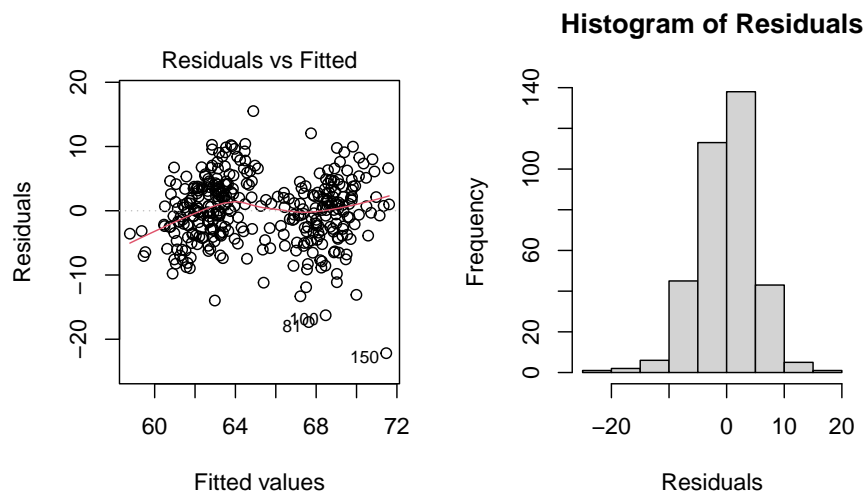
Conditions for the least squares line

When performing inference on a least squares line, the follow conditions are generally required:

- *Independent observations* (for both simulation-based and theory-based methods): individual data points must be independent.
 - Check this assumption by investigating the sampling method and determining if the observational units are related in any way.
- *Linearity* (for both simulation-based and theory-based methods): the data should follow a linear trend.
 - Check this assumption by examining the scatterplot of the two variables, and a scatterplot of the residuals (on the y -axis) versus the fitted values (on the x -axis). The pattern in the residual plot should display a horizontal line.
- *Constant variability* (for theory-based methods only): the variability of points around the least squares line remains roughly constant
 - Check this assumption by examining a scatterplot of the residuals (on the y -axis) versus the fitted values (on the x -axis). The variability in the residuals around zero should be approximately the same for all fitted values.
- *Nearly normal residuals* (for theory-based methods only: residuals must be nearly normal.
 - Check this assumption by examining a histogram of the residuals, which should appear approximately normal¹.

The scatterplot generated earlier and the residual plots shown below will be used to assess these conditions for approximating the data with the t -distribution.

¹A better plot for checking the normality assumption is called a *normal quantile-quantile plot* (or QQ-plot). However, this type of plot will be covered in a future course



1. Are the conditions met to use the t -distribution to approximate the sampling distribution of the standardized statistic? Justify your answer.

Ask a research question

2. Write out the null hypothesis in words to test the slope.
3. Using the research question, write the alternative hypothesis in notation to test the slope.

Summarize and visualize the data

The linear model output for this study is shown below.

```
lm.golf <- lm(Percent_Accuracy~Driving_Distance, data=golf) # lm(response~explanatory)
round(summary(lm.golf)$coefficients, 3)
```

```
#>               Estimate Std. Error t value Pr(>|t|)
#> (Intercept)      103.586      3.329   31.119      0
#> Driving_Distance   -0.142      0.012  -11.553      0
```

Use statistical inferential methods to draw inferences from the data

Hypothesis test To find the value of the standardized statistic to test the slope we will use,

$$T = \frac{\text{slope estimate} - \text{nullvalue}}{SE} = \frac{b_1 - 0}{SE(b_1)}.$$

We will use the linear model R output above to get the estimate for slope and the standard error of the slope.

4. Calculate the standardized statistic for slope. Identify where this calculated value is in the linear model R output.
5. The p-value in the linear model R output is the two-sided p-value for the test of significance for slope. Report the p-value to answer the research question.
6. Based on the p-value, how much evidence is there against the null hypothesis?

Confidence interval Recall that a confidence interval is calculated by adding and subtracting the margin of error to the point estimate.

$$\text{point estimate} \pm t^* \times SE(\text{estimate}).$$

When the point estimate is a regression slope, this formula becomes

$$b_1 \pm t^* \times SE(b_1).$$

The t^* multiplier comes from a t -distribution with $n - 2$ degrees of freedom. Recall for a 95% confidence interval, we use the 97.5% percentile (95% of the distribution is in the middle, leaving 2.5% in each tail). The sample size for this study is 354 so we will use the degrees of freedom 352 ($n - 2$).

```
qt(0.975, 352, lower.tail = TRUE) # 95% t* multiplier
```

```
#> [1] 1.966726
```

7. Calculate the 95% confidence interval for the true slope.

- Interpret the 95% confidence interval in context of the problem.

Communicate the results and answer the research question

- Write a conclusion to answer the research question in context of the problem.

Multivariable plots

Another variable that may affect the percent accuracy is the which league the golfer is part of. We will look at how this variable may change the relationship between driving distance and percent accuracy.

```
golf %>%  
  ggplot(aes(x = Driving_Distance, y = Percent_Accuracy, color=League))+ # Specify variables  
  geom_point(aes(shape = League), size = 2, alpha=0.5) + # Add scatterplot of points  
  labs(x = "Driving Distance (yards)", # Label x-axis  
       y = "Percent Accuracy", # Label y-axis  
       color = "League", shape = "League",  
       title = "Scatterplot of Golf Driving Distance and Percent  
               Accuracy by League for Professional Golfers") + # Be sure to title your plots  
  geom_smooth(method = "lm", se = FALSE) + # Add regression line  
  scale_color_viridis_d(end=0.8)
```

- Does the association between driving distance and percent accuracy change depending on which league the golfer is a part of? Explain your answer.
- Explain the association between league and each of the other two variables.

11.4.4 Take-home messages

1. To check the validity conditions for using theory-based methods we must use the residual diagnostic plots to check for normality of residuals and constant variability, and the scatterplot to check for linearity.
2. To interpret a confidence interval for the slope, think about how to interpret the sample slope and use that information in the confidence interval interpretation for slope.
3. Use the explanatory variable row in the linear model R output to obtain the slope estimate (**estimate** column) and standard error of the slope (**Std. Error** column) to calculate the standardized slope, or T-score. The calculated T-score should match the **t value** column in the explanatory variable row. The standardized slope tells the number of standard errors the observed slope is above or below 0.
4. The explanatory variable row in the linear model R output provides a **two-sided** p-value under the **Pr(>|t|)** column.
5. The standardized slope is compared to a t -distribution with $n - 2$ degrees of freedom in order to obtain a p-value. The t -distribution with $n - 2$ degrees of freedom is also used to find the appropriate multiplier for a given confidence level.

11.4.5 Additional notes

Use this space to summarize your thoughts and take additional notes on this week's activity and material covered.

11.5 Module 11 Lab: Big Mac Index

11.5.1 Learning outcomes

- Given a research question involving two quantitative variables, construct the null and alternative hypotheses in words and using appropriate statistical symbols.
- Assess the conditions to determine in theory or simulation-based methods should be used.
- Find, interpret, and evaluate the p-value for a hypothesis test for a slope or correlation.
- Create and interpret a confidence interval for a slope or correlation.

11.5.2 Big Mac Index

Can the relative cost of a Big Mac across different countries be used to predict the Gross Domestic Product (GDP) per person for that country? The log GDP per person and the adjusted dollar equivalent to purchase a Big Mac was found on a random sample of 55 countries in January of 2022. The cost of a Big Mac in each country was adjusted to US dollars based on current exchange rates. Is there evidence of a positive relationship between Big Mac cost (`dollar_price`) and the log GDP per person (`log_GDP`)?

- Upload and open the R script file for Week 13 lab.
- Upload the csv file, `big_mac_adjusted_index_S22.csv`.
- Enter the name of the data set for `datasetname` in the R script file in line 9.
- Highlight and run lines 1–9 to load the data.

```
# Read in data set
mac <- read.csv("datasetname")
```

Summarize and visualize the data

- To find the correlation between the variables, `log_GDP` and `dollar_price` highlight and run lines 13–16 in the R script file.

```
mac %>%
  select(c("log_GDP", "dollar_price")) %>%
  cor(use="pairwise.complete.obs") %>%
  round(3)
```

1. Report the value of correlation between the variables.
2. Calculate the value of the coefficient of determination between `log_GDP` and `dollar_price`.
3. Interpret the value of the coefficient of determination in context of the problem.

In the next part of the activity we will assess the linear model between Big Mac cost and log GDP.

- Enter the variable `log_GDP` for **response** and the variable `dollar_price` for **explanatory** in line 22.
- Highlight and run lines 22–23 to get the linear model output.

```
# Fit linear model: y ~ x
bigmacLM <- lm(response~explanatory, data=mac)
round(summary(bigmacLM)$coefficients,3) # Display coefficient summary
```

4. Give the value of the slope of the regression line. Interpret this value in context of the problem.

Conditions for the least squares line

5. Is there independence between the responses for the observational units? Justify your answer.

- Highlight and run lines 28–33 to create the scatterplot to check for linearity.

```
#Scatterplot
mac %>% # Pipe data set into...
  ggplot(aes(x = dollar_price, y = log_GDP))+ # Specify variables
  geom_point(alpha=0.5) + # Add scatterplot of points
  labs(x = "Big Mac Cost", # Label x-axis
       y = "log GDP", # Label y-axis
       title = "Scatterplot of Big Mac Cost vs. log GDP per person  
for Countries in 2022") + # Be sure to tile your plots
  geom_smooth(method = "lm", se = FALSE) # Add regression line
```

6. Is the linearity condition met to use regression methods to analyze the data? Justify your answer.

- Highlight and run lines 38–42 to produce the diagnostic plots needed to assess conditions to use theory-based methods.

```
#Diagnostic plots
bigmacLM <- lm(log_GDP~dollar_price, data = mac) # Fit linear regression model
par(mfrow=c(1,2)) # Set graphics parameters to plot 2 plots in 1 row
plot(bigmacLM, which=1) # Residual vs fitted values
hist(bigmacLM$resid, xlab="Residuals", ylab="Frequency",
     main = "Histogram of Residuals") # Histogram of residuals
```

7. Are the conditions met to use the t -distribution to approximate the sampling distribution of the standardized statistic? Justify your answer.

Ask a research question

8. Write out the null and alternative hypotheses in notation to test *correlation* between Big Mac cost and country GDP.

H_0 :

H_A :

Use statistical inferential methods to draw inferences from the data

Hypothesis test

Use the `regression_test()` function in R (in the `catstats` package) to simulate the null distribution of sample **correlations** and compute a p-value. We will need to enter the response variable name and the explanatory variable name for the formula, the data set name (identified above as `mac`), the summary measure used for the test, number of repetitions, the sample statistic (value of correlation), and the direction of the alternative hypothesis.

The response variable name is `log_GDP` and the explanatory variable name is `dollar_price`.

9. What inputs should be entered for each of the following to create the simulation to test correlation?

- Direction ("**greater**", "**less**", or "**two-sided**"):
- Summary measure (choose "**slope**" or "**correlation**"):
- As extreme as (enter the value for the sample correlation):
- Number of repetitions:

Using the R script file for this activity, enter your answers for question 9 in place of the `xx`'s to produce the null distribution with 1000 simulations.

- Highlight and run lines 47–53.
- Upload a copy of your plot showing the p-value to Gradescope for your group.

```
regression_test(log_GDP~dollar_price, # response ~ explanatory
  data = mac, # Name of data set
  direction = "xx", # Sign in alternative ("greater", "less", "two-sided")
  summary_measure = "xx", # "slope" or "correlation"
  as_extreme_as = xx, # Observed slope or correlation
  number_repetitions = 1000) # Number of simulated samples for null distribution
```

10. Report the p-value from the R output.

Simulation-based confidence interval

We will use the `regression_bootstrap_CI()` function in R (in the `catstats` package) to simulate the bootstrap distribution of sample **correlations** and calculate a confidence interval.

- Fill in the `xx`'s in the the provided R script file to find a 90% confidence interval.
- Highlight and run lines 58–62.

```
regression_bootstrap_CI(log_GDP~dollar_price, # response ~ explanatory
  data = mac, # Name of data set
  confidence_level = xx, # Confidence level as decimal
  summary_measure = "xx", # Slope or correlation
  number_repetitions = 1000) # Number of simulated samples for bootstrap distribution
```

11. Report the bootstrap 90% confidence interval in interval notation.

Communicate the results and answer the research question

12. Using a significance level of 0.1, what decision would you make?
13. What type of error is possible?
14. Interpret this error in context of the problem.
15. Write a paragraph summarizing the results of the study as if you are reporting these results in your local newspaper. **Upload a copy of your paragraph to Gradescope for your group.** Be sure to describe:
 - Summary statistic and interpretation
 - Summary measure (in context)
 - Value of the statistic
 - Order of subtraction when comparing two groups
 - P-value and interpretation
 - Statement about probability or proportion of samples
 - Statistic (summary measure and value)
 - Direction of the alternative
 - Null hypothesis (in context)
 - Confidence interval and interpretation
 - How confident you are (e.g., 90%, 95%, 98%, 99%)
 - Parameter of interest

- Calculated interval
 - Order of subtraction when comparing two groups
- Conclusion (written to answer the research question)
 - Amount of evidence
 - Parameter of interest
 - Direction of the alternative hypothesis
- Scope of inference
 - To what group of observational units do the results apply (target population or observational units similar to the sample)?
 - What type of inference is appropriate (causal or non-causal)?

Probability

12.1 Reading Guide: Special Topics

Textbook Chapter 23 Probability with tables

Vocabulary

Random process:

Probability:

Hypothetical two-way table:

Unconditional probability:

Notation:

Conditional probability:

Notation:

Event:

Notation:

Complement:

Notation:

Sensitivity:

Specificity:

Prevalence:

Notes

Method for creating a hypothetical two-way table:

1. Start with
2. Fill in the column or row totals using
3. Fill in the interior cells using
4. Add/Subtract to fill in the row/column totals not filled in at step 2.

To find unconditional probabilities from the table,

To find conditional probabilities from the table,

Example from section 23.4: Baby Jeff

1. Let D be the event a child has CPK. What does D^C represent?
2. Let T be the event a child tests positive for CPK. What does T^C represent?
3. Write each of the following values in proper probability notation:
 - a. $1/10000 = 0.0001 = P(\quad)$
 - b. $100\% = 1.0 = P(\quad)$
 - c. $99.98\% = 0.9998 = P(\quad)$
4. Write out the steps for creating the hypothetical two-way table in section 2.2.4 of your textbook, then copy the table below.

First,

Next,

After that,

Finally,

Hypothetical two-way table:

	Test Positive	Test Negative	Total
Has CPK			
Does not have CPK			
Total			100,000

5. What is the probability that a child who had a positive test result actually does have CPK? What probability notation should be used for this value?

6. Explain how the probability in #5 was calculated.

12.2 Video Notes: Probability

Read Chapters 22 in the course textbook. Use the following videos to complete the lecture notes for Module 12.

12.2.1 Course Videos

- Chapter23

Probability

Example: Two variables were collected on a random sample of people who had ever been married; whether a person had ever smoked and whether a person had ever been divorced. The data are displayed in the following table. This survey was based on a random sample in the United States in the early 1990s, so the data should be representative of the adult population who had ever been married at that time.

- Let event D be a person has gone through a divorce
- Let event S be a person smokes

	Has divorced	Has never divorced	Total
Smokes	238	247	485
Does not smoke	374	810	1184
Total	612	1057	1669

- What is the approximate probability that the person smoked?
- What is the approximate probability that the person had ever been divorced?
- Given that the person had been divorced, what is the probability that he or she smoked?
- Given that the person smoked, what is the probability that he or she had been divorced?
- Event: something that could occur, something we want to find the probability of
 - Getting a four when rolling a fair die
- Complement: opposite of the event

- Getting any value but a four when rolling a fair die
- The probability of an event is the _____ proportion of times the event would occur if the _____ process were repeated indefinitely.
 - For example, the probability of getting a four when rolling a fair die is _____.
- Unconditional probabilities
 - An _____ probability is calculated from the entire population not _____ on the occurrence of another event.
 - Examples:
 - * The probability of a single event
 - The probability a selected Stat 216 student is a computer science major.
 - * An “And” probability
 - The probability a selected Stat 216 student is a computer science major and a freshman.
- Conditional probabilities
 - A _____ probability is calculated _____ on the occurrence of another event.
 - Examples:
 - * The probability of event A given B
 - The probability a selected freshman Stat 216 student is a computer science major.
 - * The probability of event B given A
 - The probability a selected computer science Stat 216 student is a freshman
- Let event D be a person has gone through a divorce
- Let event S be a person smokes

	Has divorced	Has never divorced	Total
Smokes	238	247	485
Does not smoke	374	810	1184
Total	612	1057	1669

Calculate and interpret each of the following:

- $P(S^c) =$

- $P(D^c|S^c) =$

Creating a hypothetical two-way table

Steps:

- Start with a large number like 100000.
- Then use the unconditional probabilities to fill in the row or column totals.
- Now use the conditional probabilities to begin filling in the interior cells.
- Use subtraction to find the remaining interior cells.
- Add the column values together for each row to find the row totals.
- Add the row values together for each column to find the column totals.

Example: An airline has noticed that 30% of passengers pre-pay for checked bags at the time the ticket is purchased. The no-show rate among customers that pre-pay for checked bags is 5%, compared to 15% among customers that do not pre-pay for checked bags.

- Let event B = customer pre-pays for checked bag
- Let event N = customer no shows

Start by identifying the probability notation for each value given.

- $0.30 =$
- $0.05 =$
- $0.15 =$

	B	B^c	Total
N			
N^c			
Total			100,000

- What is the probability that a randomly selected customer who shows for the flight, pre-purchased checked bags?

Diagnostic tests

- Sensitivity:
- Specificity:
- Prevalence:

12.2.2 Lecture Notes Concept Checks

Be prepared for group discussion in the next class. One member from the table should write the answers to the following on the whiteboard.

1. Calculate and interpret the following: $P(D^c|S^c) =$.
2. What is the probability notation for 0.15 in the airline example?

12.3 Guided Activity 17: Probability Studies

12.3.1 Learning outcomes

- Recognize and simulate probabilities as long-run frequencies.
- Construct two-way tables to evaluate conditional probabilities.

12.3.2 Terminology review

In today's activity, we will cover two-way tables and probability. Some terms covered in this activity are:

- Proportions
- Probability
- Conditional probability
- Two-way tables

To review these concepts, see Chapter 23 in the textbook.

12.3.3 Calculating probabilities from a two-way table

In 2014, the website FiveThirtyEight examined the works of Bob Ross to see what trends could be found. They determined that of all the paintings he created, 95% of them contained at least one “happy tree.” Of those works with a happy tree, 43% contained at least one “almighty mountain.” Of the paintings that did not have at least one happy tree, only 10% contained at least one almighty mountain.

Let A = Bob Ross painting contains a happy tree, and B = Bob Ross painting contains an almighty mountain

	A	A^c	Total
B	40850	500	41350
B^c	54150	4500	58650
Total	95000	5000	100000

- What is the probability that a randomly selected Bob Ross painting contains both a “happy tree” and an “almighty mountain”? Use appropriate probability notation.
- What is the probability that a selected Bob Ross painting without an “almighty mountain” contains a “happy tree.” Use appropriate probability notation.
- What is the probability that a selected Bob Ross painting does not contain a “happy tree” given it does not contain an “almighty mountain”. Use appropriate probability notation.

2. A recent study of population decline of white-tailed deer in Wyoming due to chronic wasting disease (Edmunds 2016) (CWD) reported the prevalence of CWD to be 35.4%. The survival rate of CWD positive deer was 39.6% and the survival rate of CWD negative deer was 80.1%.

Let A = the event a deer is CWD positive, and B = the event the deer survived.

- a. Identify what each numerical value given in the problem represents in probability notation.

$$0.342 =$$

$$0.396 =$$

$$0.801 =$$

- b. Create a hypothetical two-way table to represent the situation.

	A	A^c	Total
B			
B^c			
Total			100,000

- c. Find $P(A \text{ and } B)$. What does this probability represent in the context of the problem?
- d. Find the probability that a deer that is positive for CWD did not survive. What is the notation used for this probability?
- e. What is the probability that a deer does not survive given they do not have CWD? What is the notation used for this probability?

12.3.4 Take home messages

1. Conditional probabilities are calculated dependent on a second variable. In probability notation, the variable following $|$ is the variable on which we are conditioning. The denominator used to calculate the probability will be the total for the variable on which we are conditioning.
2. When creating a two-way table we typically want to put the explanatory variable on the columns of the table and the response variable on the rows.
3. To fill in the two-way table, always start with the unconditional variable in the total row or column and then use the conditional probabilities to fill in the interior cells.

12.3.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

12.4 Activity 14: What's the probability?

12.4.1 Learning outcomes

- Recognize and simulate probabilities as long-run frequencies.
- Construct two-way tables to evaluate conditional probabilities.

12.4.2 Terminology review

In today's activity, we will cover two-way tables and probability. Some terms covered in this activity are:

- Proportions
- Probability
- Conditional probability
- Two-way tables

To review these concepts, see Chapter 23 in the textbook.

12.4.3 Probability

1. A dataset was collected on all NBA basketball players from inception of the league. The probability that an NBA player is above average height is 59.7%. Of NBA players that are above average height, 46.4% averaged at least four rebounds a game. The probability that an NBA player averages less than four rebounds in a game given they are below average height is 13.3%.

Let A = player is above average height, and B = player averages at least four rebounds a game.

	A	A^c	Total
B	27700.8	34940.1	62640.9
B^c	31999.2	5359.9	37359.1
Total	59700	40300	100000

- a. What is the probability that a randomly selected NBA player averages at least 4 rebounds a game? Use appropriate probability notation.
- b. What is the probability that a randomly selected NBA player is both above average height and averages at least 4 rebounds a game. Use appropriate probability notation.
- c. What is the probability that a randomly selected NBA player is not above average height given they do not average at least 4 rebounds a game. Use appropriate probability notation.

2. Since the early 1980s, the rapid antigen detection test (RADT) of group A *streptococci* has been used to detect strep throat. A recent study of the accuracy of this test shows that the **sensitivity**, the probability of a positive RADT given the person has strep throat, is 86% in children, while the **specificity**, the probability of a negative RADT given the person does not have strep throat, is 92% in children. The **prevalence**, the probability of having group A strep, is 37% in children. (Stewart et al. 2014)

Let A = the event the child has strep throat, and B = the event the child has a positive RADT.

- a. Identify what each numerical value given in the problem represents in probability notation.

$$0.86 =$$

$$0.92 =$$

$$0.37 =$$

- b. Create a hypothetical two-way table to represent the situation.

	A	A^c	Total
B			
B^c			
Total			100,000

- c. Find $P(A \text{ and } B)$. What does this probability represent in the context of the problem?
- d. Find the probability that a child with a positive RADT actually has strep throat. What is the notation used for this probability?
- e. What is the probability that a child does not have strep given that they have a positive RADT? What is the notation used for this probability?

12.4.4 Take home messages

1. Conditional probabilities are calculated dependent on a second variable. In probability notation, the variable following $|$ is the variable on which we are conditioning. The denominator used to calculate the probability will be the total for the variable on which we are conditioning.
2. When creating a two-way table we typically want to put the explanatory variable on the columns of the table and the response variable on the rows.
3. To fill in the two-way table, always start with the unconditional variable in the total row or column and then use the conditional probabilities to fill in the interior cells.

12.4.5 Additional notes

Use this space to summarize your thoughts and take additional notes on today's activity and material covered.

Semester Review

13.1 Group Final Exam Review

Use the provided data set from the Islands (Bulmer, n.d.) (FinalExamReviewData.csv) and the appropriate Exam 1 Review R script file to answer the following questions. Each adult (>21) islander was selected at random from all adult islanders. Note that some islanders choose not to participate in the study. These islanders that did not consent to be in the study are removed from the dataset before analysis. Variables and their descriptions are listed below. Here is some more information about some of the variables collected. Music type (classical or heavy metal) was randomly assigned to the Islanders. Time to complete the puzzle cube was measured after listening to music for each Islander. Heart rate and blood glucose levels were both measured before and then after drinking a caffeinated beverage.

Variable	Description
Island	Name of Island that the Islander resides on
City	Name of City in which the Islander resides
Population	Population of the City
Name	Name of Islander
Consent	Whether the Islander consented to be in the study (Declined, Consented)
Gender	Gender of Islander (M = male, F = Female)
Age	Age of Islander
Married	Marital status of Islander (yes, no)
Smoking_Status	Whether the Islander is a current smoker (nonsmoker, smoker)
Children	Whether the Islander has children (yes, no)
weight_kg	Weight measured in kg
height_cm	Height measured in cm
respiratory_rate	Breaths per minute
Type_of_Music	Music type Islander was randomly assigned to listen to (Classical, Heavy Metal)
After_PuzzleCube	Time to complete puzzle cube (minutes) after listening to assigned music
Education_Level	Highest level of education completed (highschool, university)
Balance_Test	Time balanced measured in seconds with eyes closed
Blood_Glucose_before	Level of blood glucose (mg/dL) before consuming assigned drink
Heart_Rate_before	Heart rate (bpm) before consuming assigned drink
Blood_Glucose_after	Level of blood glucose (mg/dL) after consuming assigned drink
Heart_Rate_after	Heart rate (bpm) after consuming assigned drink
Diff_Heart_Rate	Difference in heart rate (bpm) for Before - After consuming assigned drink
Diff_Blood_Glucose	Difference in blood glucose (mg/dL) for Before - After consuming assigned drink

1. Use the appropriate Exam 2 Review R script file and analyze the following research question: “Is there evidence that adult Islanders with a University degree are less likely to smoke than those with a High School degree?”

a. Parameter of Interest:

b. Null Hypothesis:

Notation:

Words:

c. Alternative Hypothesis:

Notation:

Words:

- d. Use the R script file to get the counts for each level and combination of variables. Fill in the following table with the variable names, levels of each variable, and counts using the values from the R output.

	Explanatory Variable		
Response variable	Group 1	Group 2	Total
Success			
Failure			
Total			

e. Calculate the value of the summary statistic to answer the research question. Give appropriate notation.

f. Interpret the value of the summary statistic in context of the problem:

g. Assess if the following conditions are met:

Independence (needed for both simulation and theory-based methods):

Success-Failure (must be met to use theory-based methods):

h. Use the provided R script file to find the simulation p-value to assess the research question. Report the p-value.

i. Interpret the p-value in the context of the problem.

j. Write a conclusion to the research question based on the p-value.

k. Using a significance level of $\alpha = 0.05$, what statistical decision will you make about the null hypothesis?

l. Use the provided R script file to find a 95% confidence interval.

m. Interpret the 95% confidence interval in context of the problem.

n. Regardless to your answer in part g, calculate the standardized statistic.

o. Interpret the value of the standardized statistic in context of the problem.

p. Use the provided R script file to find the theory-based p-value.

q. Use the provided R script file to find the appropriate z^* multiplier and calculate the theory-based confidence interval.

- r. Does the theory-based p-value and CI match those found using simulation methods? Explain why or why not.

- s. What is the scope of inference for this study?

2. Use the appropriate Final Exam Review R script file to analyze the following research question, “Is there evidence that adult Islanders that drink a caffeinated drink have an increase in blood glucose level, on average?” Use before – after as the order of subtraction.

a. Parameter of Interest:

b. Null Hypothesis:

Notation:

Words:

c. Alternative Hypothesis:

Notation:

Words:

- d. Use the R script file to get the summary statistics. Fill in the following table with the variable names, levels of each variable, and values from the R output.

Summary value	Variable
Mean	
Standard deviation	
Sample size	

e. Write the summary statistic to answer the research question with appropriate notation.

f. Interpret the value of the summary statistic in context of the problem:

g. Assess if the following conditions are met:

Independence (needed for both simulation and theory-based methods):

Normality:

- h. Use the provided R script file to find the simulation p-value to assess the research question. Report the p-value.
- i. Interpret the p-value in the context of the problem.
- j. Write a conclusion to the research question based on the p-value.
- k. Using a significance level of $\alpha = 0.1$, what statistical decision will you make about the null hypothesis?
- l. Use the provided R script file to find a 90% confidence interval.
- m. Interpret the 90% confidence interval in context of the problem.
- n. Regardless to your answer in part g, calculate the standardized statistic.
- o. Interpret the value of the standardized statistic in context of the problem.
- p. Use the provided R script file to find the theory-based p-value.
- q. Use the provided R script file to find the appropriate t^* multiplier and calculate the theory-based confidence interval.

- r. Does the theory-based p-value and CI match those found using simulation methods? Explain why or why not.
- s. What is the scope of inference for this study?

3. Use the appropriate Final Exam Review R script file to analyze the following research question: “Is there evidence that adult Islanders who listen to classical music take less time to complete the puzzle cube after listening to the music than for Islanders that listen to heavy metal music?” Use - classical - heavy metal as the order of subtraction.

a. Parameter of Interest:

b. Null Hypothesis:

Notation:

Words:

c. Alternative Hypothesis:

Notation:

Words:

- d. Use the R script file to get the summary statistics for each level of the explanatory variable. Fill in the following table with the variable names, levels of each variable, and values from the R output.

	Explanatory Variable	
Summary value	Group 1	Group 2
Mean		
Standard deviation		
Sample size		

e. Calculate the value of summary statistic to answer the research question. Give appropriate notation.

f. Interpret the value of the summary statistic in context of the problem:

g. Assess if the following conditions are met:

Independence (needed for both simulation and theory-based methods):

Normality:

- h. Use the provided R script file to find the simulation p-value to assess the research question. Report the p-value.
- i. Interpret the p-value in the context of the problem.
- j. Write a conclusion to the research question based on the p-value.
- k. Using a significance level of $\alpha = 0.05$, what statistical decision will you make about the null hypothesis?
- l. Use the provided R script file to find a 95% confidence interval.
- m. Interpret the 95% confidence interval in context of the problem.
- n. Regardless to your answer in part g, calculate the standardized statistic.
- o. Interpret the value of the standardized statistic in context of the problem.
- p. Use the provided R script file to find the theory-based p-value.
- q. Use the provided R script file to find the appropriate t^* multiplier and calculate the theory-based confidence interval.

- r. Does the theory-based p-value and CI match those found using simulation methods? Explain why or why not.
- s. What is the scope of inference for this study?

4. Use the appropriate Final Exam Review R script file to analyze the following research question: “Is there evidence that the height of adult Islanders is associated with the their balance time?”

a. Parameter of Interest:

b. Null Hypothesis:

Notation:

Words:

c. Alternative Hypothesis:

Notation:

Words:

- d. Use the R script file to get the summary statistics for this data. Fill in the following table using the values from the R output:

	y-intercept	slope	correlation
Summary value			

e. Interpret the value of slope in context of the problem.

f. Assess if the following conditions are met:

Independence (needed for both simulation and theory-based methods):

Linearity (needed for both simulation and theory-based methods):

Constant Variance:

Normality of Residuals:

- g. Use the provided R script file to find the simulation p-value to assess the research question. Report the p-value.
- h. Interpret the p-value in the context of the problem.
- i. Write a conclusion to the research question based on the p-value.
- j. Using a significance level of $\alpha = 0.01$, what statistical decision will you make about the null hypothesis?
- k. Use the provided R script file to find a 99% confidence interval.
- l. Interpret the 99% confidence interval in context of the problem.
- m. Regardless to your answer in part g, calculate the standardized statistic.
- n. Interpret the value of the standardized statistic in context of the problem.
- o. Use the provided R script file to find the theory-based p-value.
- p. Use the provided R script file to find the appropriate t^* multiplier and calculate the theory-based confidence interval.

- q. Does the theory-based p-value and CI match those found using simulation methods? Explain why or why not.
- r. What is the scope of inference for this study?

13.2 Golden Ticket to Descriptive and Inferential Statistical Methods

In this course, we have covered descriptive (summary statistics and plots) and inferential (hypothesis tests and confidence intervals) methods for five different scenarios:

- one categorical response variable
- two categorical variables
- one quantitative response variable or paired differences in a quantitative variable
- two quantitative variables
- one quantitative response variable and one categorical explanatory variable

The “golden ticket” shown on the next page presents a visual summary of the similarities and differences across these five scenarios.

Scenario	One Categorical Response	Two Categorical Variables	Paired Differences	Two Quantitative Variables	Quant. Response and Categ. Explanatory (independent samples)
Type of plot	Bar plot	Segmented bar plot, Mosaic plot	Dotplot, histogram, boxplot	Scatterplot	Side-by-sided boxplots, Stacked dotplots or histograms
Summary measure	Proportion	Difference in proportions	Mean Difference	Slope or correlation	Difference in means
Parameter notation	π	$\pi_1 - \pi_2$	μ_d	β_1 or ρ	$\mu_1 - \mu_2$
Statistic notation	\hat{p}	$\hat{p}_1 - \hat{p}_2$	\bar{x}_d	b_1 or r	$\bar{x}_1 - \bar{x}_2$
Null hypothesis	$H_0: \pi = \pi_0$	$H_0: \pi_1 - \pi_2 = 0$	$H_0: \mu_d = 0$	$H_0: \beta_1 = 0$ or $H_0: \rho = 0$	$H_0: \mu_1 - \mu_2 = 0$
Conditions for simulation methods	Independent cases;	Independence (within and between groups);	Independent cases;	Independent cases; Linear form;	Independence (within and between groups);
Simulation test (how to generate a null distn) p-value = proportion of null simulations at or beyond (H_A direction) the observed statistic	Spin spinner with probability equal to π_0 , n times or draw with replacement n times from a deck of cards created to reflect π_0 as probability of success. Plot the proportion of successes. Repeat 1000's of times. Centered at π_0	Label cards with response values from original data; mix cards together; shuffle into two new groups of sizes n_1 and n_2 . Plot difference in proportion of successes. Repeat 1000's of times. Centered at 0.	Shift the original data by adding $(0 - \bar{x}_d)$. Sample with replacement from the shifted data n times. Plot sample mean difference. Repeat 1000's of times. Centered at 0.	Hold the x values constant; shuffle new y 's to x 's. Find the regression line for shuffled data; plot the slope or the correlation for the shuffled data. Repeat 1000's of times. Centered at 0.	Label cards with response variable values from original data; mix cards together; shuffle into two new groups of sizes n_1 and n_2 . Plot difference in means. Repeat 1000's of times. Centered at 0.
Bootstrap CI (how to generate a boot. distn) X% CI: $\left(\frac{1-X}{2}\right)\%tile,$ $\left(X + \frac{1-X}{2}\right)\%tile$	Label n cards with the original responses. Randomly draw with replacement n times. Plot the resampled proportion of successes. Repeat 1000's of times. Centered at \hat{p} .	Label $n_1 + n_2$ cards with the original responses. Randomly draw with replacement n_1 times from group 1 and n_2 times from group 2. Plot the resampled difference in proportion of successes. Repeat 1000's of times. Centered at $\hat{p}_1 - \hat{p}_2$	Label n cards with the original responses. Randomly draw with replacement n times. Plot the resampled mean difference. Repeat 1000's of times. Centered at \bar{x}_d .	Label n cards with the original (response, explanatory) values. Randomly draw with replacement n times. Plot the resampled slope or correlation. Repeat 1000's of times. Centered at b_1 or r .	Label $n_1 + n_2$ cards with the original responses. Randomly draw with replacement n_1 times from group 1 and n_2 times from group 2. Plot the resampled difference in means. Repeat 1000's of times. Centered at $\bar{x}_1 - \bar{x}_2$.
Theory-based distribution	Standard Normal	Standard Normal	t - distribution with $n - 1$ df	t - distribution with $n - 2$ df	t - distribution with min of n_1-1 or n_2-1 df
Conditions for theory-based hypothesis tests and confidence intervals	Independent cases; Number of successes and number of failures in the sample both at least 10.	Independence (within and between groups); Number of successes and number of failures in EACH sample all at least 10. (All four cell counts at least 10.)	Independent cases; $n < 30$ with no clear outliers OR $30 \leq n < 100$ with no extreme outliers OR $n \geq 100$	Linear form; Independent cases; Nearly normal residuals; Variability around the regression line is roughly constant.	Independent cases (within and between groups); In each sample, $n < 30$ with no clear outliers OR $30 \leq n < 100$ with no extreme outliers OR $n \geq 100$
Theory-based standardized statistic (test statistic)	$z = \frac{\hat{p} - \pi_0}{SE_0(\hat{p})}$ $SE_0(\hat{p}) = \sqrt{\frac{\pi_0 \times (1 - \pi_0)}{n}}$	$z = \frac{\hat{p}_1 - \hat{p}_2 - 0}{SE_0(\hat{p}_1 - \hat{p}_2)}$ $SE_0(\hat{p}_1 - \hat{p}_2) = \sqrt{\widehat{p_{pool}} \times (1 - \widehat{p_{pool}}) \times \left(\frac{1}{n_1} + \frac{1}{n_2}\right)}$	$t = \frac{\bar{x}_d - 0}{SE(\bar{x}_d)}$ $SE(\bar{x}_d) = \frac{s_d}{\sqrt{n}}$	$t = \frac{b_1 - 0}{SE(b_1)}$ $SE(b_1)$ is the reported standard error (std. error) of the slope term in the lm() output from R.	$t = \frac{\bar{x}_1 - \bar{x}_2 - 0}{SE(\bar{x}_1 - \bar{x}_2)}$ $SE(\bar{x}_1 - \bar{x}_2) = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}$
Theory-based confidence interval	$\hat{p} \pm z^* \times SE(\hat{p})$ $SE(\hat{p}) = \sqrt{\frac{\hat{p} \times (1 - \hat{p})}{n}}$	$\hat{p}_1 - \hat{p}_2 \pm z^* \times SE(\hat{p}_1 - \hat{p}_2)$ $SE(\hat{p}_1 - \hat{p}_2) = \sqrt{\frac{\hat{p}_1 \times (1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2 \times (1 - \hat{p}_2)}{n_2}}$	$\bar{x}_d \pm t^* \times SE(\bar{x}_d)$ $SE(\bar{x}_d) = \frac{s_d}{\sqrt{n}}$	$b_1 \pm t^* \times SE(b_1)$ $SE(b_1)$ is the reported standard error (std. error) of the slope term in the lm() output from R.	$\bar{x}_1 - \bar{x}_2 \pm t^* \times SE(\bar{x}_1 - \bar{x}_2)$ $SE(\bar{x}_1 - \bar{x}_2) = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}$

References

- “Average Driving Distance and Fairway Accuracy.” 2008. <https://www.pga.com/> and <https://www.lpga.com/>.
- Banton, et al, S. 2022. “Jog with Your Dog: Dog Owner Exercise Routines Predict Dog Exercise Routines and Perception of Ideal Body Weight.” *PLoS ONE* 17(8).
- Bhavsar, et al, A. 2022. “Increased Risk of Herpes Zoster in Adults 50 Years Old Diagnosed with COVID-19 in the United States.” *Open Forum Infectious Diseases* 9(5).
- Bulmer, M. n.d. “Islands in Schools Project.” <https://sites.google.com/site/islandsinschoolsprojectwebsite/home>.
- “Bureau of Transportation Statistics.” 2019. <https://www.bts.gov/>.
- “Child Health and Development Studies.” n.d. <https://www.chdstudies.org/>.
- Darley, J. M., and C. D. Batson. 1973. “From Jerusalem to Jericho”: A Study of Situational and Dispositional Variables in Helping Behavior.” *Journal of Personality and Social Psychology* 27: 100–108.
- Davis, Smith, A. K. 2020. “A Poor Substitute for the Real Thing: Captive-Reared Monarch Butterflies Are Weaker, Paler and Have Less Elongated Wings Than Wild Migrants.” *Biology Letters* 16.
- Du Toit, et al, G. 2015. “Randomized Trial of Peanut Consumption in Infants at Risk for Peanut Allergy.” *New England Journal of Medicine* 372.
- Edmunds, et al, D. 2016. “Chronic Wasting Disease Drives Population Decline of White-Tailed Deer.” *PLoS ONE* 11(8).
- Education Statistics, National Center for. 2018. “IPEDS.” <https://nces.ed.gov/ipeds/>.
- “Great Britain Married Couples: Great Britain Office of Population Census and Surveys.” n.d. <https://discovery.nationalarchives.gov.uk/details/r/C13351>.
- Group, TODAY Study. 2012. “A Clinical Trial to Maintain Glycemic Control in Youth with Type 2 Diabetes.” *New England Journal of Medicine* 366: 2247–56.
- Hamblin, J. K., K. Wynn, and P. Bloom. 2007. “Social Evaluation by Preverbal Infants.” *Nature* 450 (6288): 557–59.
- Hirschfelder, A., and P. F. Molin. 2018. “I Is for Ignoble: Stereotyping Native Americans.” Retrieved from <https://www.ferris.edu/HTMLS/news/jimcrow/native/homepage.htm>.
- Hutchison, R. L., and M. A. Hirthler. 2013. “Upper Extremity Injuries in Homer’s Iliad.” *Journal of Hand Surgery (American Volume)* 38: 1790–93.
- “IMDb Movies Extensive Dataset.” 2016. <https://kaggle.com/stefanoleone992/imdb-extensive-dataset>.
- Kalra, et al., D. I. 2022. “Trustworthiness of Indian Youtubers.” Kaggle. <https://doi.org/10.34740/KAGGLE/DSV/4426566>.
- Keating, D., N. Ahmed, F. Nirappil, Stanley-Becker I., and L. Bernstein. 2021. “Coronavirus Infections Dropping Where People Are Vaccinated, Rising Where They Are Not, Post Analysis Finds.” *Washington Post*. <https://www.washingtonpost.com/health/2021/06/14/covid-cases-vaccination-rates/>.
- Laeng, Mathisen, B. 2007. “Why Do Blue-Eyed Men Prefer Women with the Same Eye Color?” *Behavioral Ecology and Sociobiology* 61(3).
- Levin, D. T. 2000. “Race as a Visual Feature: Using Visual Search and Perceptual Discrimination Tasks to Understand Face Categories and the Cross-Race Recognition Deficit.” *Journal of Experimental Psychology* 129(4).
- Madden, et al, J. 2020. “Ready Student One: Exploring the Predictors of Student Learning in Virtual Reality.” *PLoS ONE* 15(3).
- Miller, G. A. 1956. “The Magical Number Seven, Plus or Minus Two: Some Limits on Our Capacity for Processing Information.” *Psychological Review* 63(2).
- Moquin, W., and C. Van Doren. 1973. “Great Documents in American Indian History.” Praeger.
- “More Americans Are Joining the ‘Cashless’ Economy.” 2022. <https://www.pewresearch.org/short-reads/2022/10/05/more-americans-are-joining-the-cashless-economy/>.
- National Weather Service Corporate Image Web Team. n.d. “National Weather Service – NWS Billings.” <https://w2.weather.gov/climate/xmacis.php?wfo=byz>.
- O’Brien, Lynch, H. D. 2019. “Crocodylian Head Width Allometry and Phylogenetic Prediction of Body Size in Extinct Crocodyliforms.” *Integrative Organismal Biology* 1.

- “Ocean Temperature and Salinity Study.” n.d. <https://calcofi.org/>.
- “Older People Who Get Covid Are at Increased Risk of Getting Shingles.” 2022. <https://www.washingtonpost.com/health/2022/04/19/shingles-and-covid-over-50/>.
- “Physician’s Health Study.” n.d. <https://phs.bwh.harvard.edu/>.
- Porath, Erez, C. 2017. “Does Rudeness Really Matter? The Effects of Rudeness on Task Performance and Helpfulness.” *Academy of Management Journal* 50.
- Quinn, G. E., C. H. Shin, M. G. Maguire, and R. A. Stone. 1999. “Myopia and Ambient Lighting at Night.” *Nature* 399 (6732): 113–14. <https://doi.org/10.1038/20094>.
- Ramachandran, V. 2007. “3 Clues to Understanding Your Brain.” https://www.ted.com/talks/vs_ramachandran_3_clues_to_understanding_your_brain.
- “Rates of Laboratory-Confirmed COVID-19 Hospitalizations by Vaccination Status.” 2021. CDC. <https://covid.cdc.gov/covid-data-tracker/#covidnet-hospitalizations-vaccination>.
- Richardson, T., and R. T. Gilman. 2019. “Left-Handedness Is Associated with Greater Fighting Success in Humans.” *Scientific Reports* 9 (1): 15402. <https://doi.org/10.1038/s41598-019-51975-3>.
- Stephens, R., and O. Robertson. 2020. “Swearing as a Response to Pain: Assessing Hypoalgesic Effects of Novel “Swear” Words.” *Frontiers in Psychology* 11: 643–62.
- Stewart, E. H., B. Davis, B. L. Clemans-Taylor, B. Littenberg, C. A. Estrada, and R. M. Centor. 2014. “Rapid Antigen Group a Streptococcus Test to Diagnose Pharyngitis: A Systematic Review and Meta-Analysis” 9 (11). <https://doi.org/10.1371/journal.pone.0111727>.
- Stroop, J. R. 1935. “Studies of Interference in Serial Verbal Reactions.” *Journal of Experimental Psychology* 18: 643–62.
- Subach, et al, A. 2022. “Foraging Behaviour, Habitat Use and Population Size of the Desert Horned Viper in the Negev Desert.” *Soc. Open Sci* 9.
- Sulheim, S., A. Ekeland, I. Holme, and R. Bahr. 2017. “Helmet Use and Risk of Head Injuries in Alpine Skiers and Snowboarders: Changes After an Interval of One Decade” 51 (1): 44–50. <https://doi.org/10.1136/bjsports-2015-095798>.
- “Titanic.” n.d. <http://www.encyclopedia-titanica.org>.
- “US COVID-19 Vaccine Tracker: See Your State’s Progress.” 2021. Mayo Clinic. <https://www.mayoclinic.org/coronavirus-covid-19/vaccine-tracker>.
- US Environmental Protection Agency. n.d. “Air Data – Daily Air Quality Tracker.” <https://www.epa.gov/outdoor-air-quality-data/air-data-daily-air-quality-tracker>.
- Wahlstrom, et al, K. 2014. “Examining the Impact of Later School Start Times on the Health and Academic Performance of High School Students: A Multi-Site Study.” *Center for Applied Research and Educational Improvement*.
- Weiss, R. D. 1988. “Relapse to Cocaine Abuse After Initiating Desipramine Treatment.” *JAMA* 260(17).
- “Welcome to the Navajo Nation Government: Official Site of the Navajo Nation.” 2011. Retrieved from <https://www.navajo-nsn.gov/>.
- Wilson, Woodruff, J. P. 2016. “Vertebral Adaptations to Large Body Size in Theropod Dinosaurs.” *PLoS ONE* 11(7).